



# Insight into

# DBMS:

## Design and Implementation

## Overview of the projects



[mlinking@126.com](mailto:mlinking@126.com)

+86 15010255486

# SAP Labs Waterloo

Established Database R&D site going back 35+ years

~220 people

- Majority SAP Database Research and Development
- Work on HANA, Cloud-native databases, distributed databases, Edge Computing

Co-op & Grad Student intern programs

Strong academic collaborations

## We are hiring!

- Full-time positions
- Student jobs - master/bachelor theses, graduate internships



# Outline of the chapters covered

---

- Introduction
- Overview of the projects
- Demonstration of Development environment
  - Watch and practice by yourself
- My understanding about (R)DBMS
  - History and D&I
- SQL translation with 2 conversions
  - SQL → RA (Relational Algebra)
  - RA → Sequence of File operations
- Transaction control
- Deeper
  - File, (R)DBMS, ERP, DW, Big Data (No SQL, SQL again)
  - SQL on MPP and Hadoop (Greenplum, **HAWQ**)



# Overview of the projects

## □ Overview of the projects

### ● Basic (70)

1. Index File - 10
2. Simulate multiple users - 10
3. Try a Syntax parsing example - 20
4. Lock table management - 20
5. Debug PostgreSQL/Greenplum/HAWQ - 10

### ● 1 Web-based Application project (20)

- Java/JSP/PHP/Python/R + PostgreSQL/Greenplum/HAWQ - R.D.I.P

### ● Advanced (10) PostgreSQL/Greenplum/HAWQ

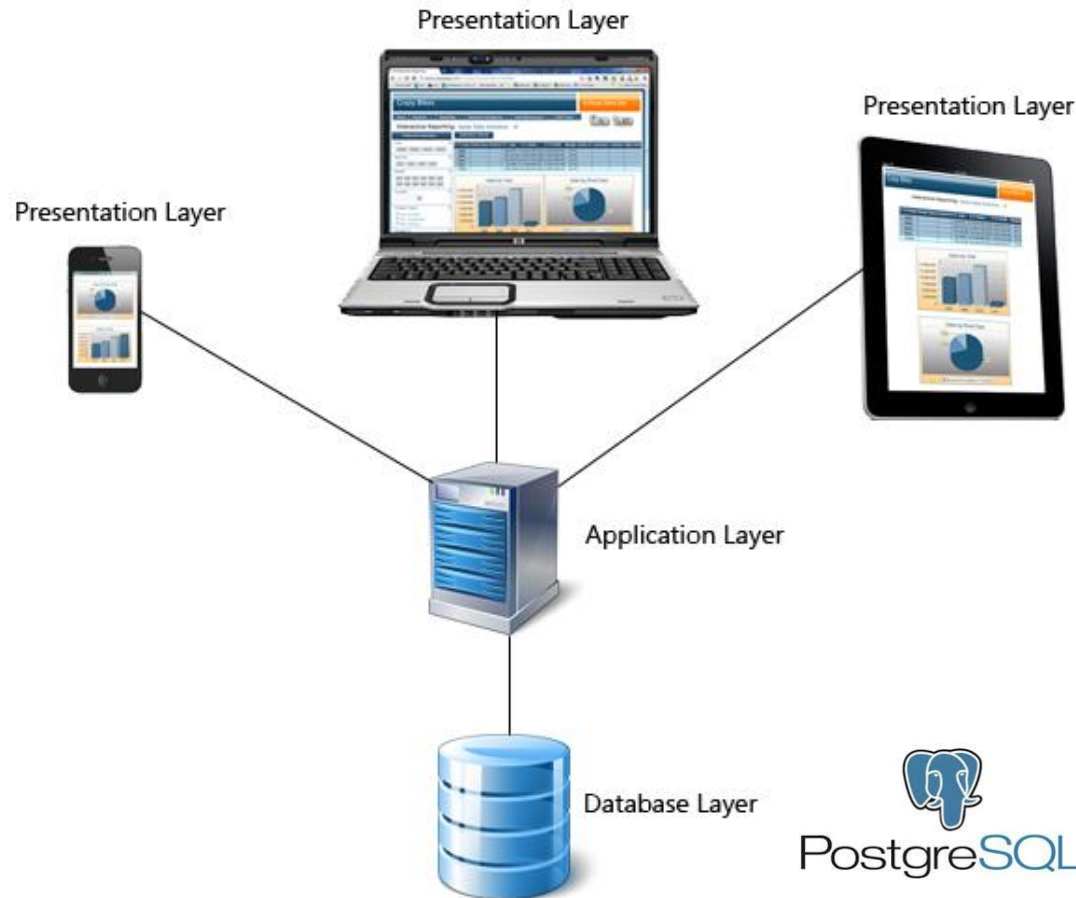
- [transaction mechanism][SQL execution]



## Modern DBMS is important for IT age

### As backend service provider

➤ 3 tier architecture is popular in many applications



# By Modern DBMS? - Properties

## □ Core is definitely

- Data Management (IDUS)

## □ Top goal

- Concurrent data management for many users in a server or a distributed system

## □ With 2 other properties

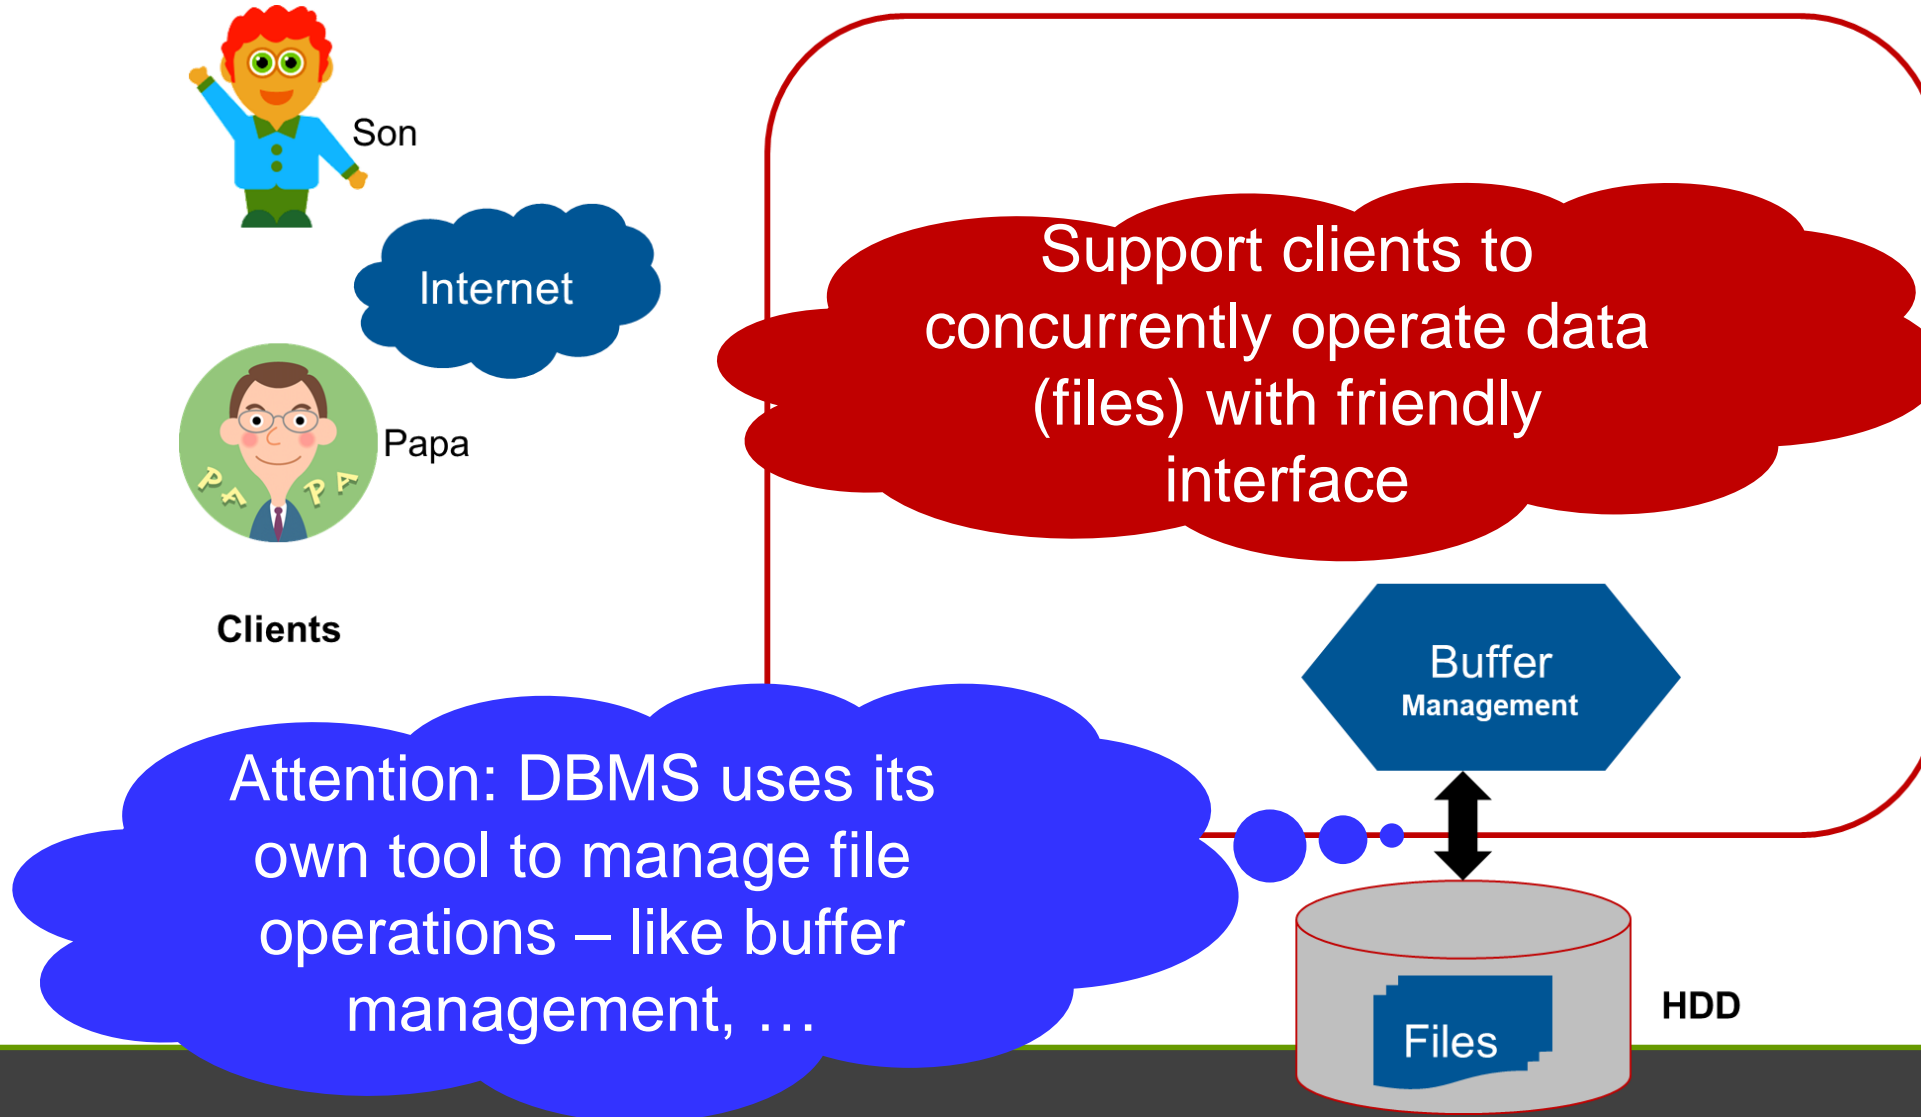
1. **Concurrent Data Management** by using files (which are supported by OS)
  - Deadlock and Data inconsistency

2. **Friendly/Flexible data operations** for

- Define specific data files
- Insert, Delete, Update, Select
- Support some statistics
- Provide flexible way for users to access
- ...

```
Select fname, lname  
From employee, department  
Where dno=dnumber  
and  
dname = 'Research'
```

# Sketch of modern DBMS (More details later)



# You can start some projects ASAP because you have learned so far

Because related skills  
are easy for you to  
follow

## ■ Basic (70)

1. Index File
2. Simulate multiple users
3. Try a Syntax parsing example
4. Lock table management
5. Debug PostgreSQL/Greenplum

My responsibility is to  
focus on the D&I of the  
core, and checking  
your work

## ■ 1 Web-based Application project (20)

➤ Java/JSP/PHP/Python/R + PostgreSQL/Greenplum/HAWQ -  
R.D.I.P

## ■ Advanced (10) in PostgreSQL/Greenplum/HAWQ

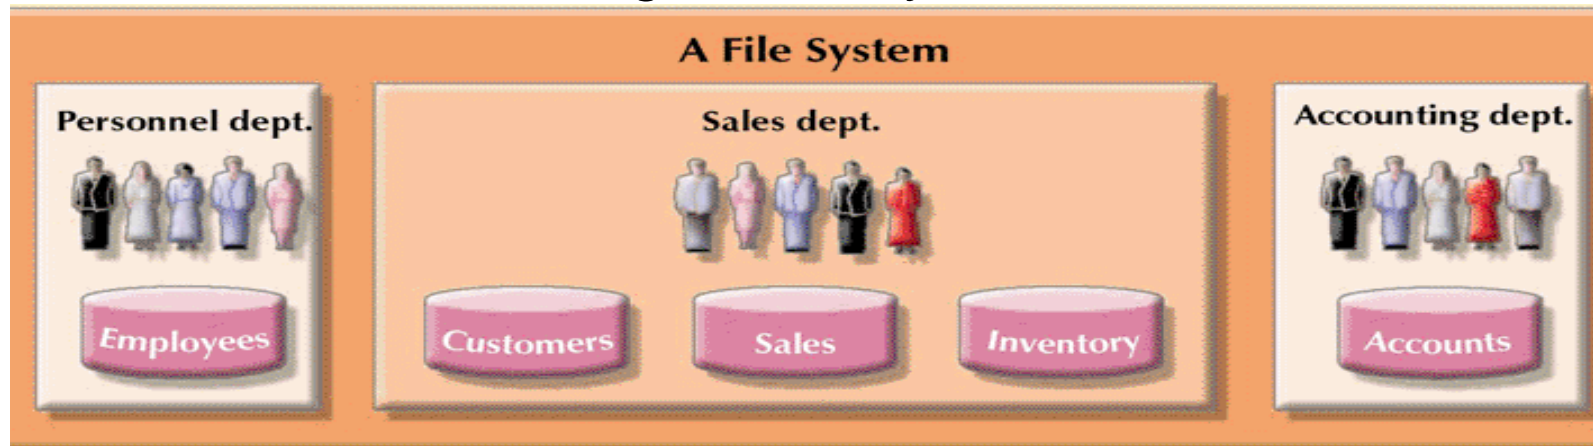
➤ [transaction mechanism][SQL execution]



# Project requirements

## □ Project 1: Index File [10 pts]

- You all have learned how to use file to store and access data – in fact this is the early version of Data Management System



- If the tuples/records are not so huge – 10 thousands, the performance seems OK. However millions of thousands of xxx? – Index File is needed!
- By index, you have learned many index data structures – AVL tree (named after inventors Adelson-Velsky and Landis), **B tree** (Balanced), KD-tree (K-dimensional) etc. **Your practice is MM (Main Memory) version!**

## □ You are given

■ On-line retails – 541909 records

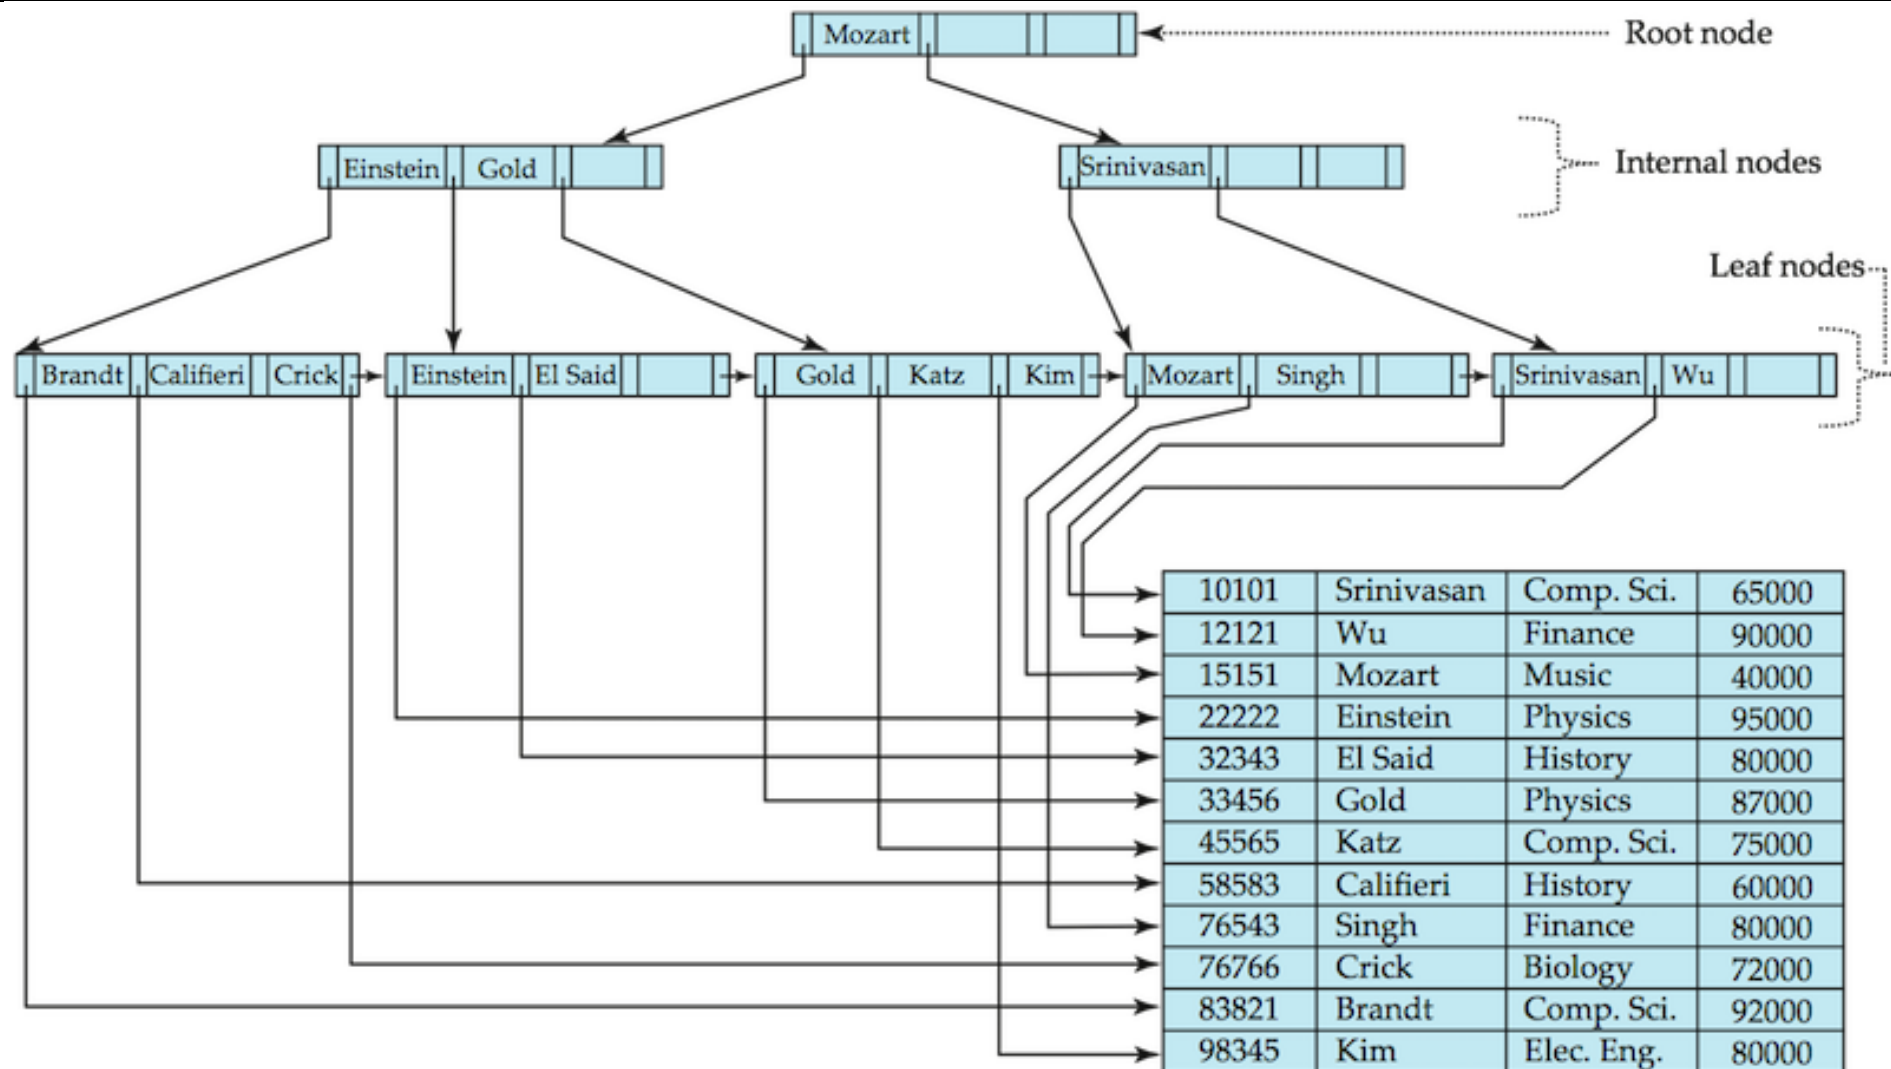
➤ <https://archive.ics.uci.edu/ml/datasets/online+retail>

|    | A         | B         | C                                   | D        | E              | F         | G          | H              |
|----|-----------|-----------|-------------------------------------|----------|----------------|-----------|------------|----------------|
| 1  | InvoiceNo | StockCode | Description                         | Quantity | InvoiceDate    | UnitPrice | CustomerID | Country        |
| 2  | 536365    | 85123A    | WHITE HANGING HEART T-LIGHT HOLDER  | 6        | 2010/12/1 8:26 | 2.55      | 17850      | United Kingdom |
| 3  | 536365    | 71053     | WHITE METAL LANTERN                 | 6        | 2010/12/1 8:26 | 3.39      | 17850      | United Kingdom |
| 4  | 536365    | 84406B    | CREAM CUPID HEARTS COAT HANGER      | 8        | 2010/12/1 8:26 | 2.75      | 17850      | United Kingdom |
| 5  | 536365    | 84029G    | KNITTED UNION FLAG HOT WATER BOTTLE | 6        | 2010/12/1 8:26 | 3.39      | 17850      | United Kingdom |
| 6  | 536365    | 84029E    | RED WOOLLY HOTTIE WHITE HEART.      | 6        | 2010/12/1 8:26 | 3.39      | 17850      | United Kingdom |
| 7  | 536365    | 22752     | SET 7 BABUSHKA NESTING BOXES        | 2        | 2010/12/1 8:26 | 7.65      | 17850      | United Kingdom |
| 8  | 536365    | 21730     | GLASS STAR FROSTED T-LIGHT HOLDER   | 6        | 2010/12/1 8:26 | 4.25      | 17850      | United Kingdom |
| 9  | 536366    | 22633     | HAND WARMER UNION JACK              | 6        | 2010/12/1 8:28 | 1.85      | 17850      | United Kingdom |
| 10 | 536366    | 22632     | HAND WARMER RED POLKA DOT           | 6        | 2010/12/1 8:28 | 1.85      | 17850      | United Kingdom |
| 11 | 536367    | 84879     | ASSORTED COLOUR BIRD ORNAMENT       | 32       | 2010/12/1 8:34 | 1.69      | 13047      | United Kingdom |
| 12 | 536367    | 22745     | POPPY'S PLAYHOUSE BEDROOM           | 6        | 2010/12/1 8:34 | 2.1       | 13047      | United Kingdom |
| 13 | 536367    | 22748     | POPPY'S PLAYHOUSE KITCHEN           | 6        | 2010/12/1 8:34 | 2.1       | 13047      | United Kingdom |

## □ We are interested to know

- The transactions which buy some specific product
  - White Metal Lantern
- How to find the invoice number? – “[White Metal Lantern](#)”
  - Location of the record should be connected with/embedded into the index
- You're required to
  1. Direct traverse the file to count the number of transactions containing “White Metal Lantern”
  2. **B+-tree** index (MM version) to finish 1
    - ✓ Finish the performance comparison of 1 and 2
  3. Design and Implement an Index File s.t.
    - ✓ Efficient to **reconstruct** from the Index File



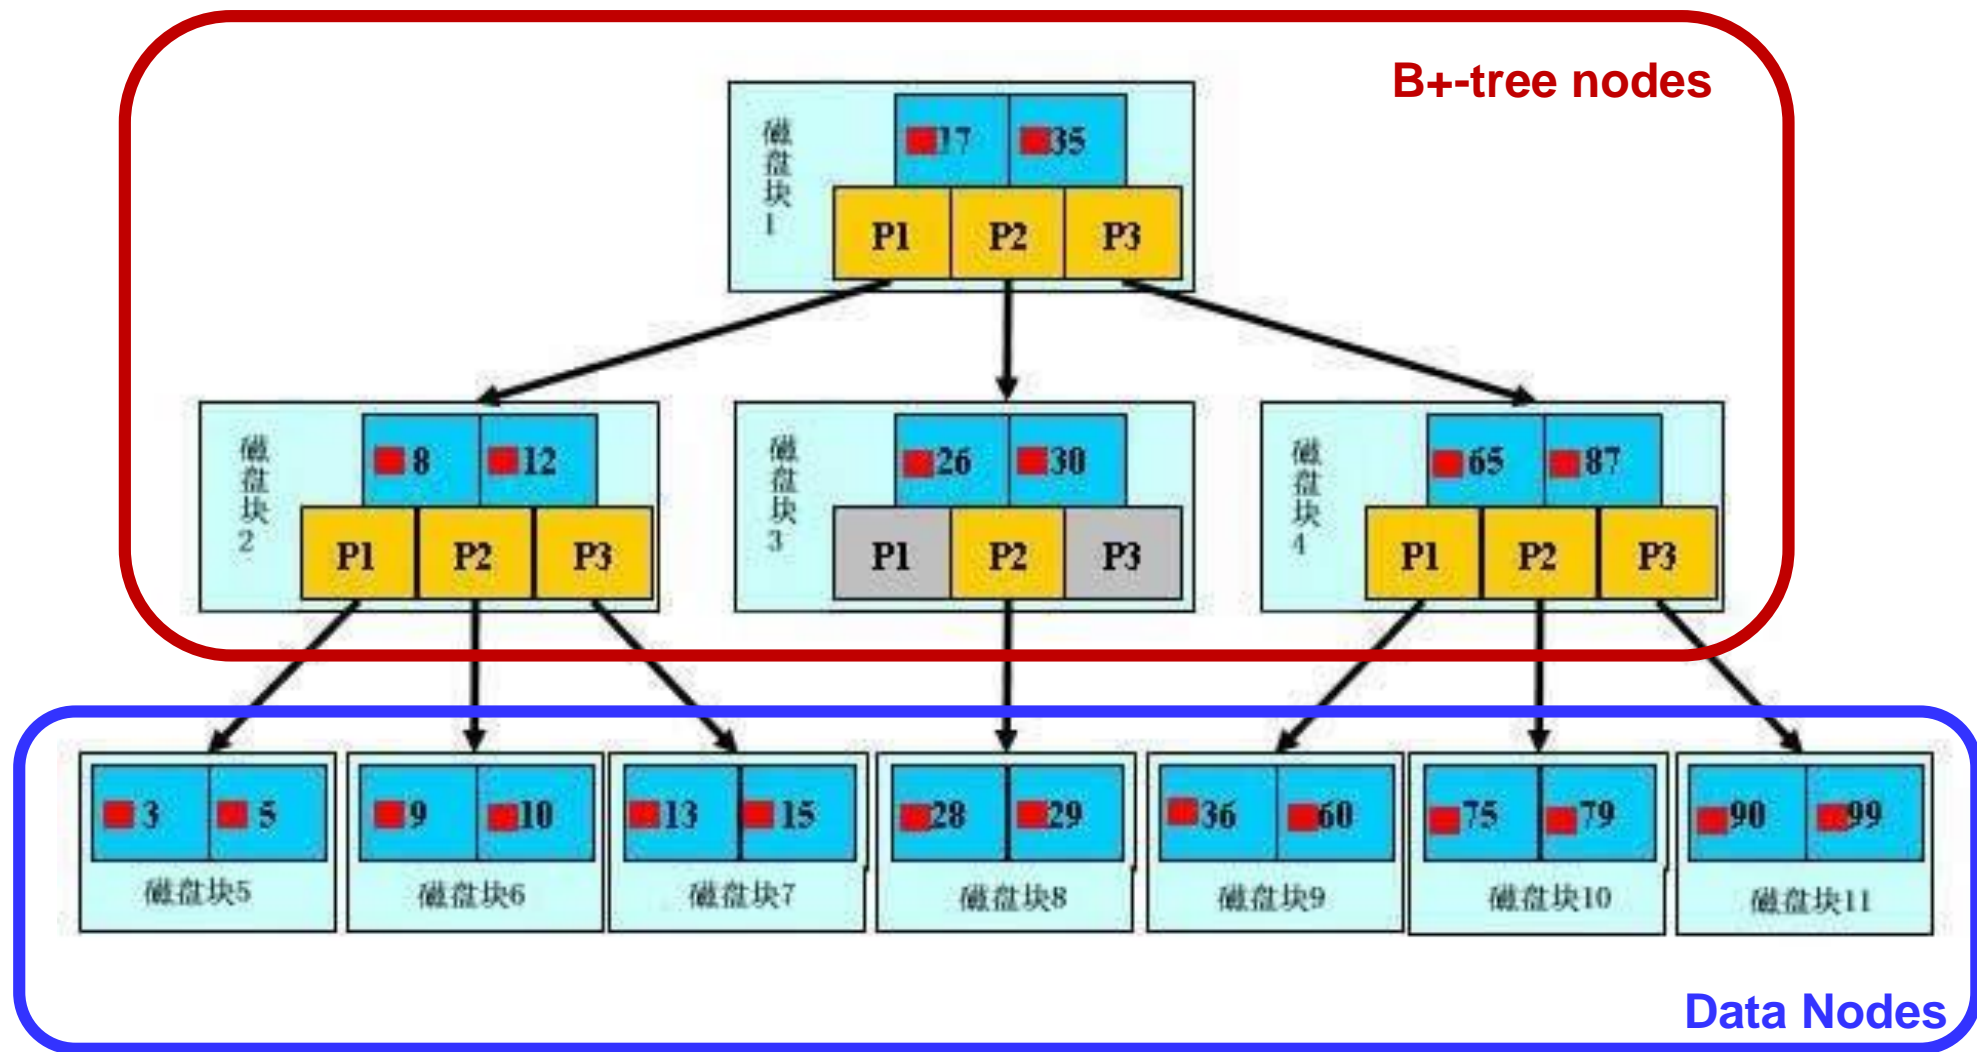


# BLOCK?

## □ In modern OS, **PAGING** based main memory management is **popularly taken**

- Your (machine code version) program is split into fixed size pages
  - 4KByte is the default size of page in Windows
- The swapping of your program into/out of MM is based on **PAGEs**
- To ensure the consistency of **FILE** management on HDD, **BLOCK** is presented for file organization in HDD
  - Its role is just like **PAGE** in MM management
  - All files are mapped to **BLOCKs**
  - Therefore, **BLOCK** size is usually same as **PAGE** size
    - ✓ Not MUST – maybe  $P=4KB$ , while  $B=8KB$  as in Windows NT





# Project requirements

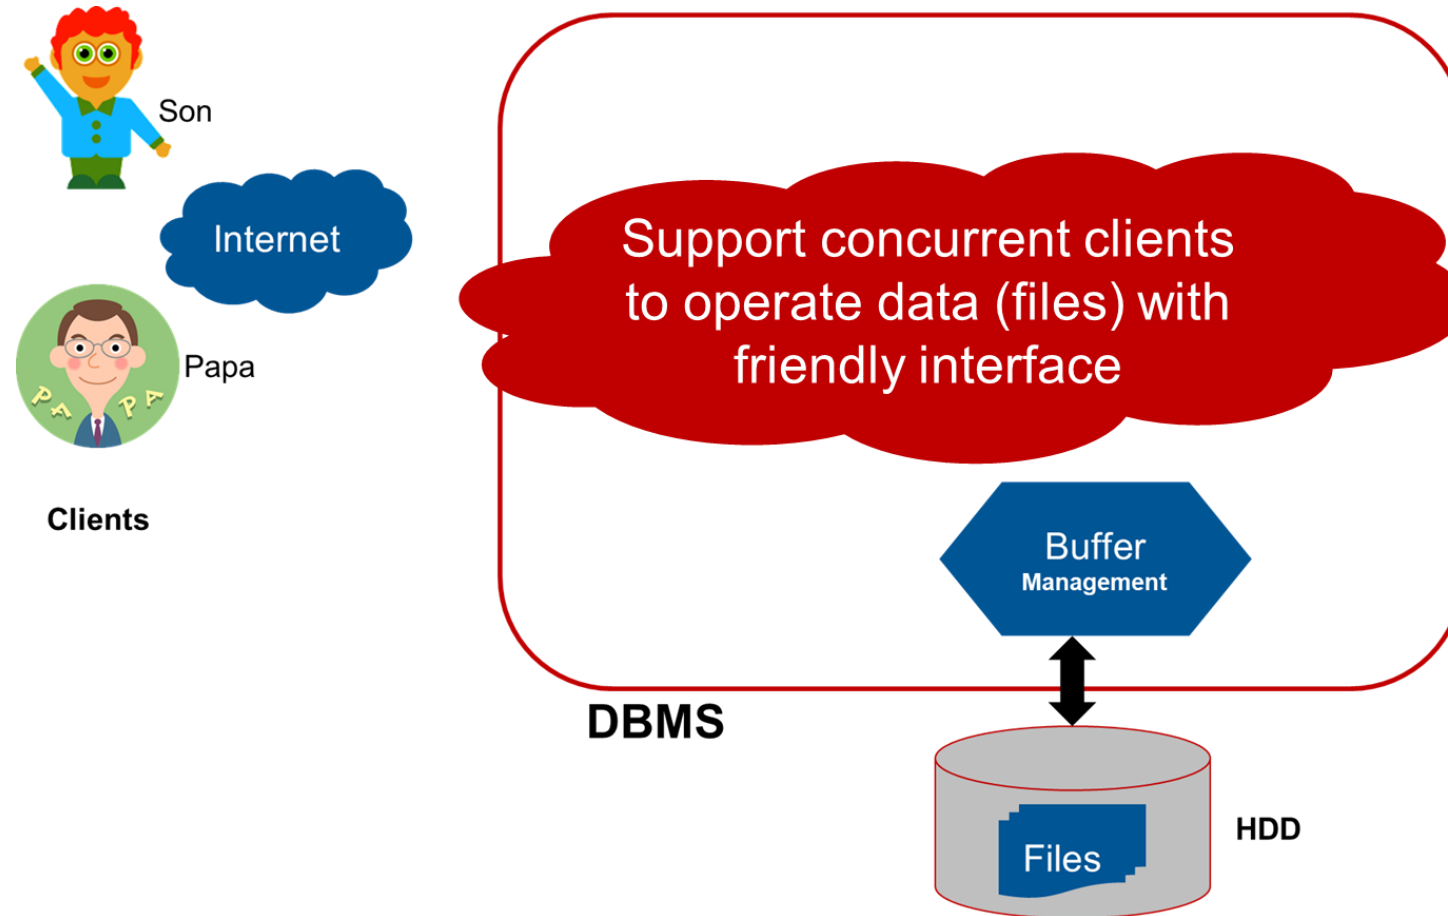
## □ Project 2 [10 pts]: Simulate Concurrent Users

- By using multi-**process** programming skills, you're required to construct a simulation system
  - + C/S structure (Networking)
  - + **Process** pool (Server side)
  - + Support hundreds of clients
- Logic of the simulation
  - Each client has same logic
    - ✓ Send words "Access some character" – random letter from A to Z.
      - » For example: "Access X"
  - Server based on **Process** Pool
    - ✓ Echo "Client x" + received words after random sleep
      - » After 5 ms echo "Client 12 for X"

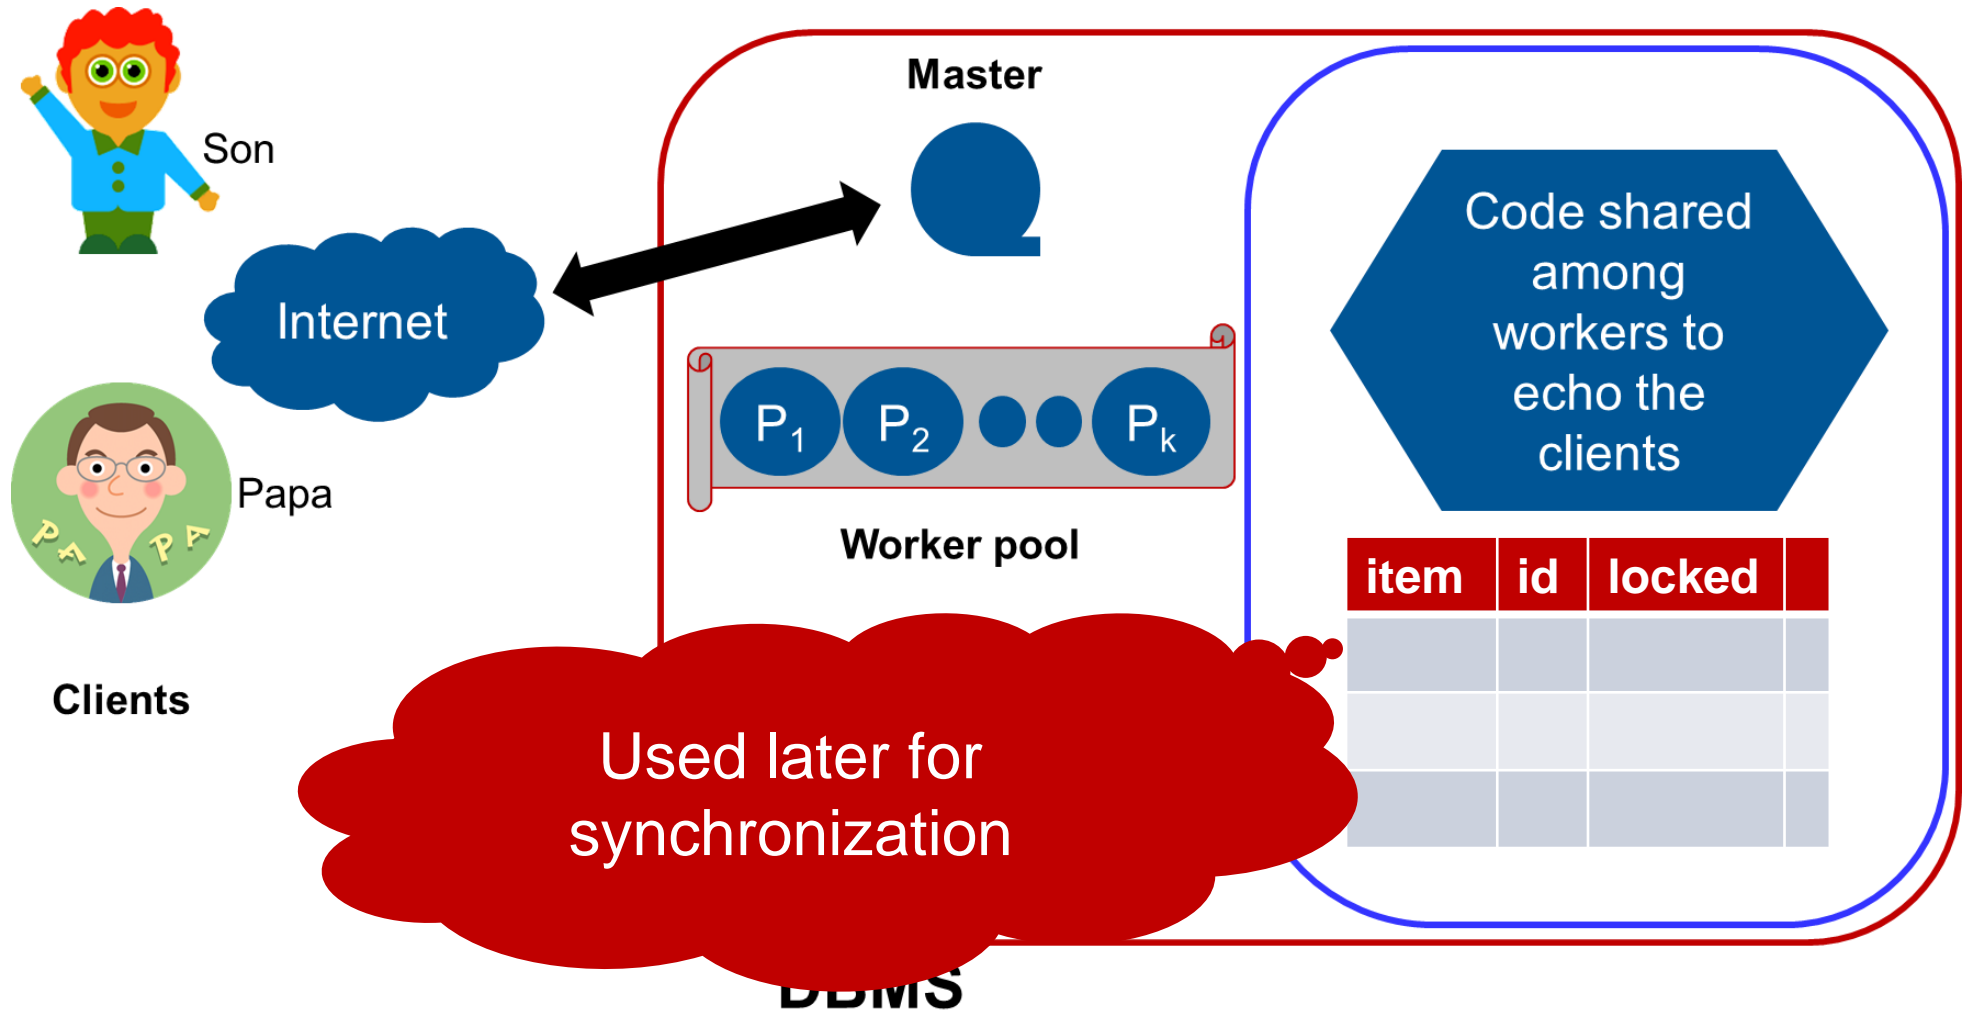




- If there is no available thread in the server pool, client should wait for a random sleep







---

□ You are required further to compare the performance (time to finish all the clients' requests) with different number of **processes** in the pool

- 10, 30, 50, 70, 100
- Draw the comparison

# Multi-Processed Programming with Python

- 如果你打算编写多进程的服务程序，Unix/Linux无疑是正确的选择。由于Windows没有fork调用，难道在Windows上无法用Python编写多进程的程序？
- 由于Python是跨平台的，自然也应该提供一个跨平台的多进程支持。multiprocessing模块就是跨平台版本的多进程模块。
  - multiprocessing模块提供了一个Process类来代表一个进程对象

```
from multiprocessing import Process
import os

# 子进程要执行的代码
def run_proc(name):
    print('Run child process %s (%s)...' % (name, os.getpid()))

if __name__ == '__main__':
    print('Parent process %s.' % os.getpid())
    p = Process(target=run_proc, args=('test',))
    print('Child process will start.')
    p.start()
    p.join()
    print('Child process end.')
```



# MultiProcess Pool

- 如果要启动大量的子进程，可以用进程池的方式批量创建子进程：

```
from multiprocessing import Pool
import os, time, random

def long_time_task(name):
    print('Run task %s (%s)...' % (name, os.getpid()))
    start = time.time()
    time.sleep(random.random() * 3)
    end = time.time()
    print('Task %s runs %0.2f seconds.' % (name, (end - start)))

if __name__ == '__main__':
    print('Parent process %s.' % os.getpid())
    p = Pool(4)
    for i in range(5):
        p.apply_async(long_time_task, args=(i,))
    print('Waiting for all subprocesses done...')
    p.close()
    p.join()
    print('All subprocesses done.')
```



# 子进程

- 很多时候，子进程并不是自身，而是一个外部进程。我们创建了子进程后，还需要控制子进程的输入和输出。
- subprocess模块可以让我们非常方便地启动一个子进程，然后控制其输入和输出。
- 例子演示了如何在Python代码中运行命令nslookup www.python.org，这和命令行直接运行的效果是一样的：

```
import subprocess

print('$ nslookup www.python.org')
r = subprocess.call(['nslookup', 'www.python.org'])
print('Exit code:', r)
```



## □ 如果子进程还需要输入，则可以通过communicate()方法输入：

```
import subprocess

print('$ nslookup')
p = subprocess.Popen(['nslookup'], stdin=subprocess.PIPE, stdout=subprocess.PIPE, stderr=subprocess.PIPE)
output, err = p.communicate(b'set q=mx\npython.org\nexit\n')
print(output.decode('utf-8'))
print('Exit code:', p.returncode)
```

# 进程间通信

- Process之间肯定是需要通信的，操作系统提供了很多机制来实现进程间的通信。Python的multiprocessing模块包装了底层的机制，提供了Queue、Pipes等多种方式来交换数据。
- 以Queue为例，在父进程中创建两个子进程，一个往Queue里写数据，一个从Queue里读数据：

[ask.csdn.net > questions](#) ▾

[python-PyCharm-执行程序PermissionError: \[WinError 5\] 拒绝 ...](#)

2019年1月9日 - ... in put if not self.\_sem.acquire(block, timeout): PermissionError: [WinError 5] 拒绝访问。 Process finished with exit code 0. 而直接在cmd中, ...

[python-多进程在运行的时候只有一个子进程会运行, 怎么解决 ...](#) 2020年5月13日

[图片-安装tensorflow的问题? ——CSDN问答频道](#)

2018年3月28日

[ask.csdn.net站内的其它相关信息](#)

```
from multiprocessing import Process, Queue
import os, time, random
```

# 写数据进程执行的代码:

```
def write(q):
    print('Process to write: %s' % os.getpid())
    for value in ['A', 'B', 'C']:
        print('Put %s to queue...' % value)
        q.put(value)
        time.sleep(random.random())
```

# 读数据进程执行的代码:

```
def read(q):
    print('Process to read: %s' % os.getpid())
    while True:
        value = q.get(True)
        print('Get %s from queue.' % value)
```

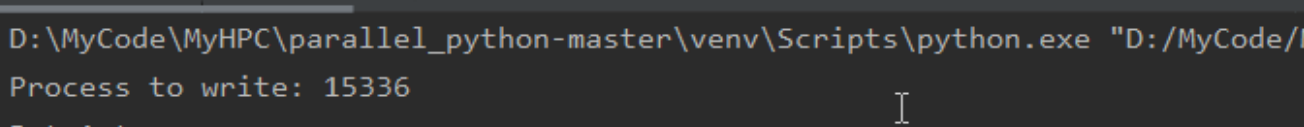
```
if __name__ == '__main__':
    # 父进程创建Queue, 并传给各个子进程:
    q = Queue()
    pw = Process(target=write, args=(q,))
    pr = Process(target=read, args=(q,))
    # 启动子进程pw, 写入:
    pw.start()
    # 启动子进程pr, 读取:
    pr.start()
    # 等待pw结束:
    pw.join()
    # pr进程里是死循环, 无法等待其结束, 只能强行终止:
    pr.terminate()
```

## 导入Manager

```
from multiprocessing import Process, Manager
```

将 `q = Queue()` 改为 `q = Manager().Queue()`

编辑于: 2019.06.16 11:16



```
Run: MultiProcessComm x MultiProcessPool x
D:\MyCode\MyHPC\parallel_python-master\venv\Scripts\python.exe "D:/MyCode/MyHPC/parallel
Process to write: 15336
Put A to queue...
Process to read: 11064
Get A from queue.
Put B to queue...
Get B from queue.
Put C to queue...
Get C from queue.
```



# Project requirements

---

## □ **Project 5 [10 pts]: Debug PostgreSQL**

- <http://hsqldb.org/>

- Very helpful to follow source code to understand the D&I of many softwares
  - Minix/Linux for OS
  - HyperSQL, H2, **PostgreSQL** , MySQL for DBMS

---

❑ I prepared some slides to demonstrate how to use VS+Perl etc. to build/compile the source code of **PostgreSQL**

❑ **Tasks of Project 5**

1. Use PostgreSQL

➤ With the given dataset – “Online Retail.xlsx” – 541909 records

✓ <https://archive.ics.uci.edu/ml/datasets/Online+Retail>

✓ Try to load it into PostgreSQL

» Define the table same as xlsx? Redundant or not – you’ll learn later in DBMS Lecture part?

✓ Try SQL sentences

» I.D.U.S

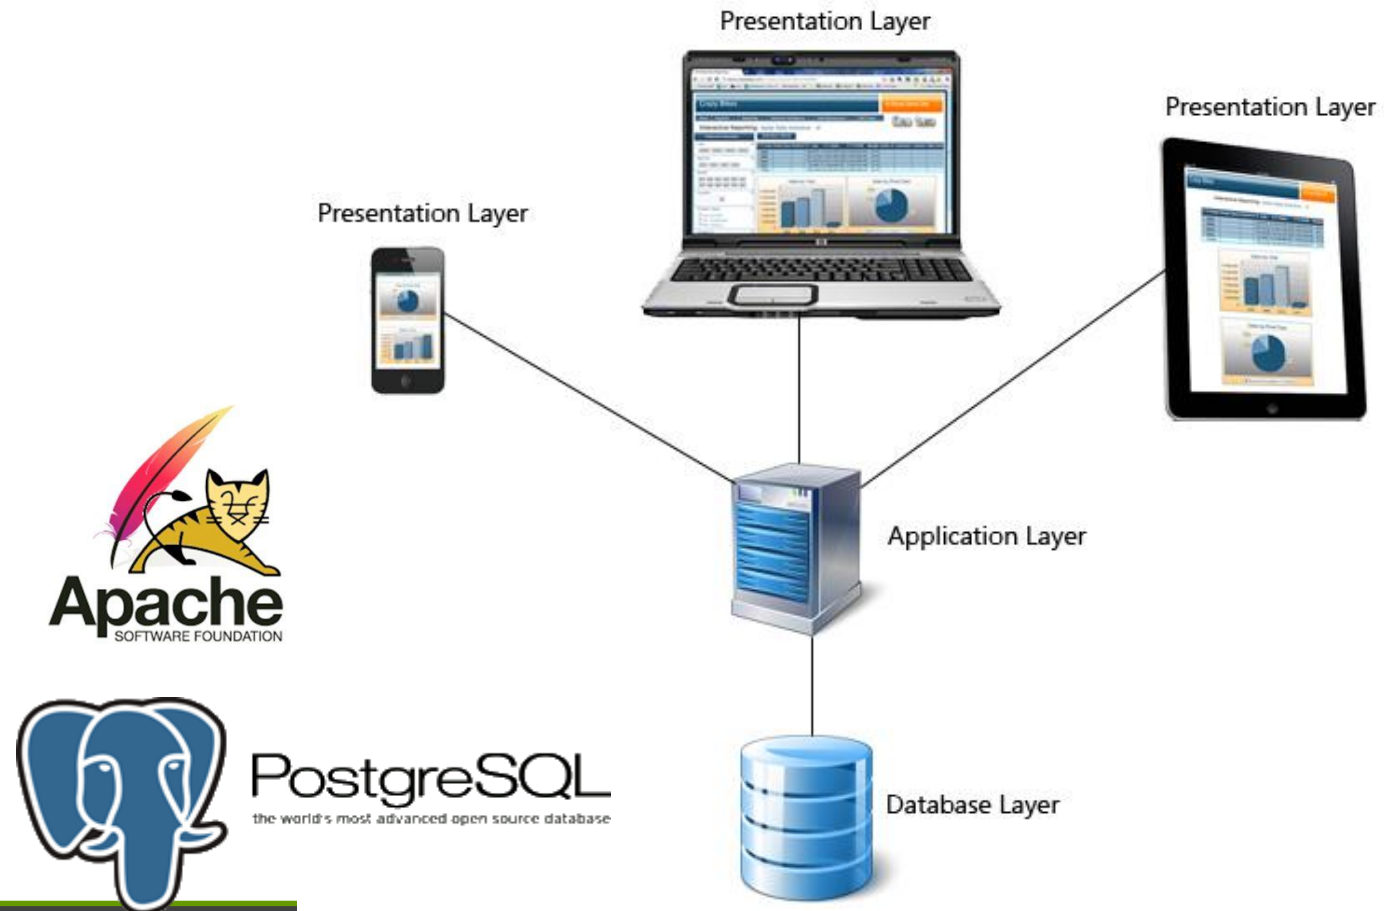
## □ Tasks of Project 5

### 2. Debug PostgreSQL

- a) Change server prompt to **your name**
  - ✓ lbkong>>
- b) Cut the SQL processing in PostgreSQL (Client and Server), and just **echo** the inputted SQL
  - ✓ lbkong>> your input is: SELECT \* From students
- c) Try to **print out the parse tree of a simple SQL** supported by PostgreSQL, like “SELECT \* from students”
  - ✓ Of course you should create a “students” table first, and input some students records
  - ✓ Name, studentID, Memo (for 40 words)

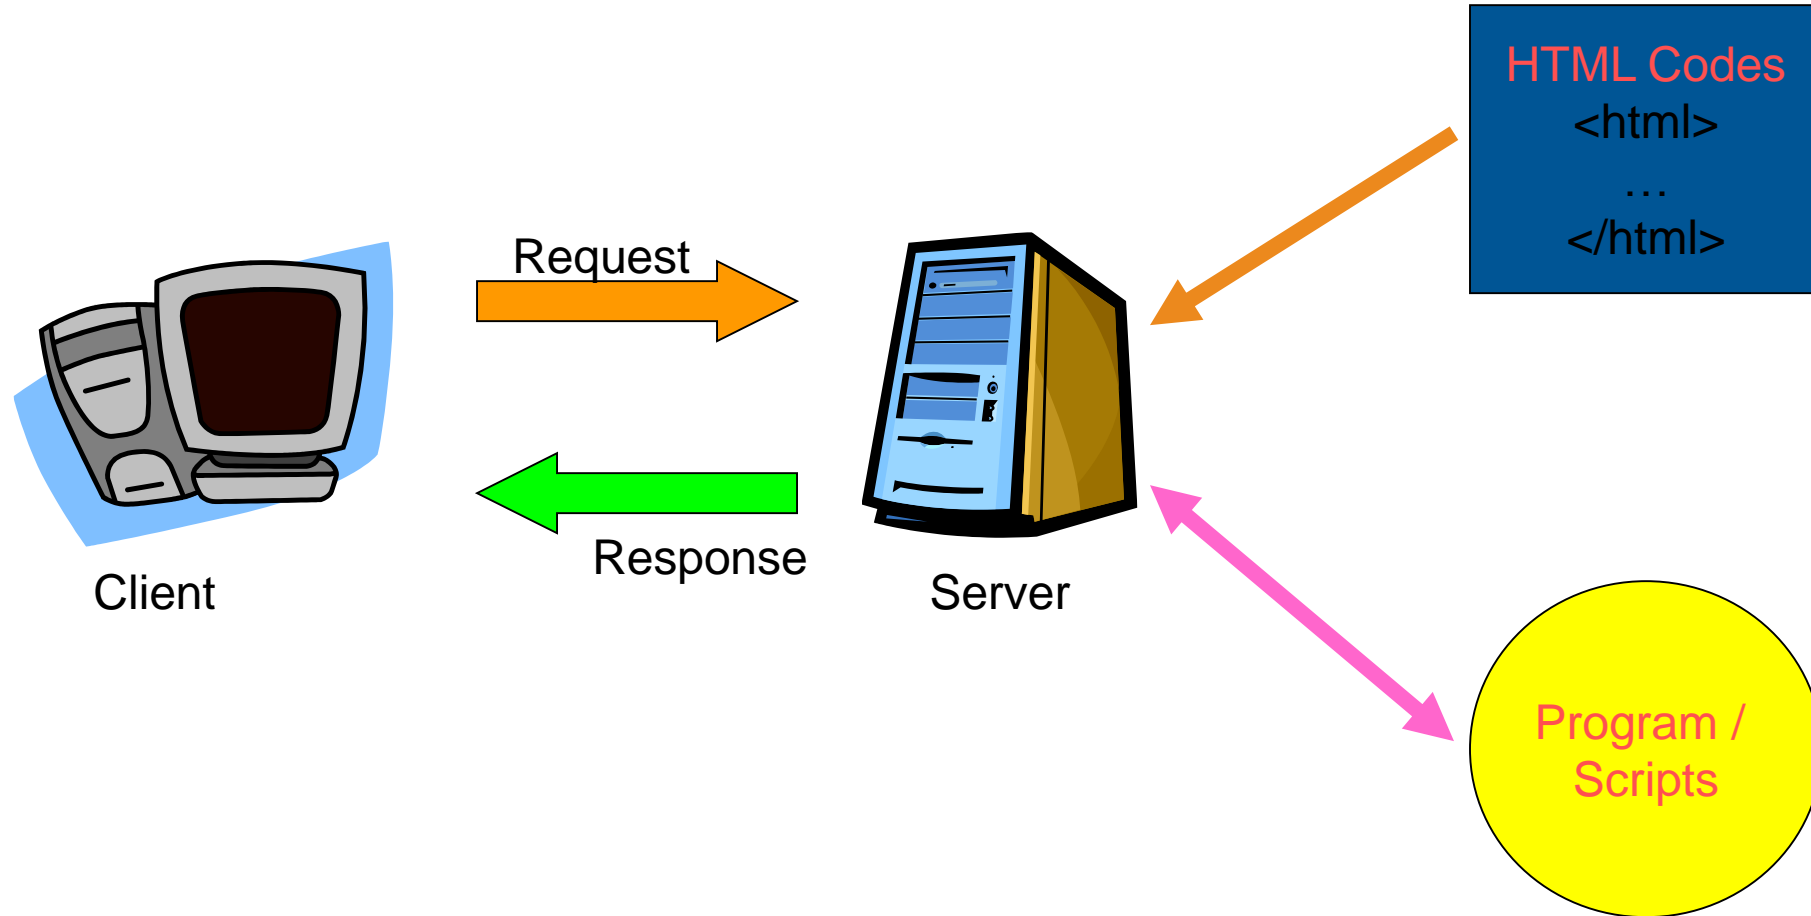
# Java/JSP/PHP/Python/R + PostgreSQL/Greenplum/HAWQ - R.D.I.P

- **Web Application[20 pts]:** such as JSP/Tomcat+ PostgreSQL
  - DBMS is the critical software supporting our modern society



# Browser is the popular frontend in 3-tier

## □ The HTTP Request/Response Model



# Dynamic page idea is powerful – improve more interaction through Web pages!

---

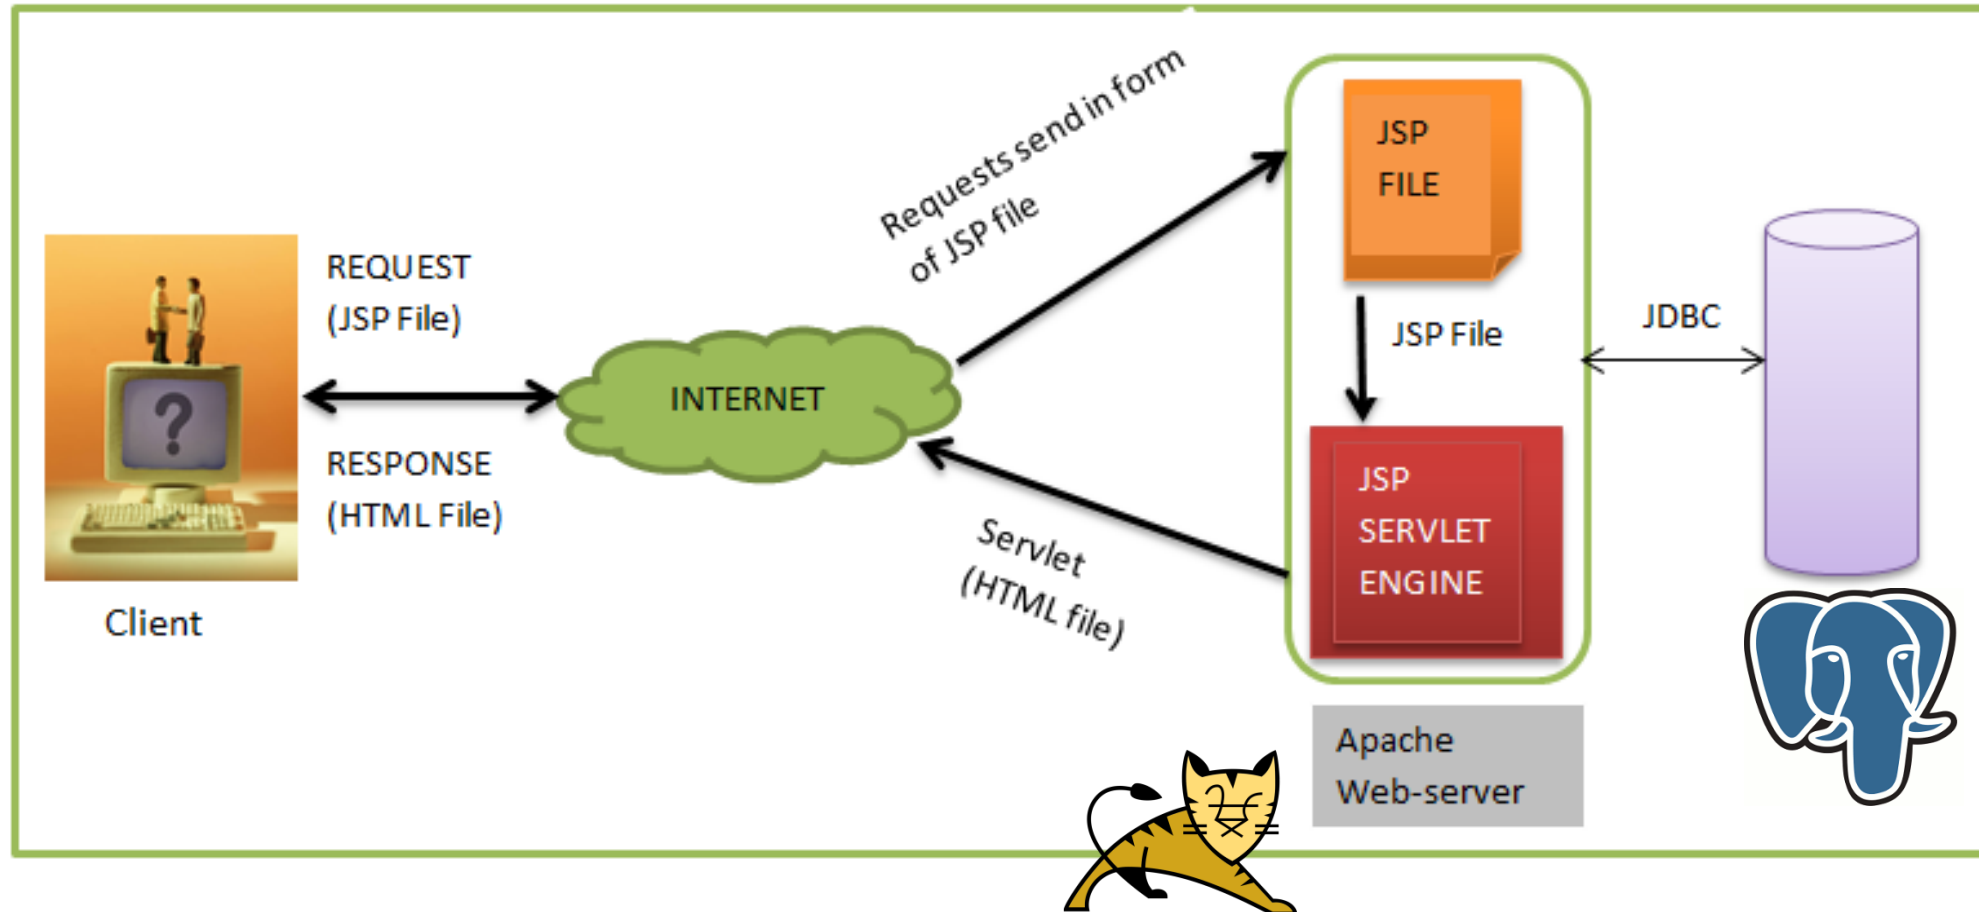
## □ Dynamic? – [Web pages call functions for powerful service]

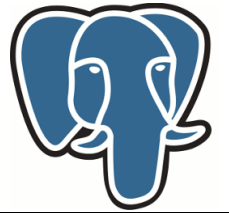
- The server could compose a webpage **on the fly** according to the client's request
  - The string returned to client is a complete webpage

## □ Before xSPs (ASP, JSP) or PHP/Python..., the traditional way to produce dynamic webpages is CGI – Common Gateway Interface

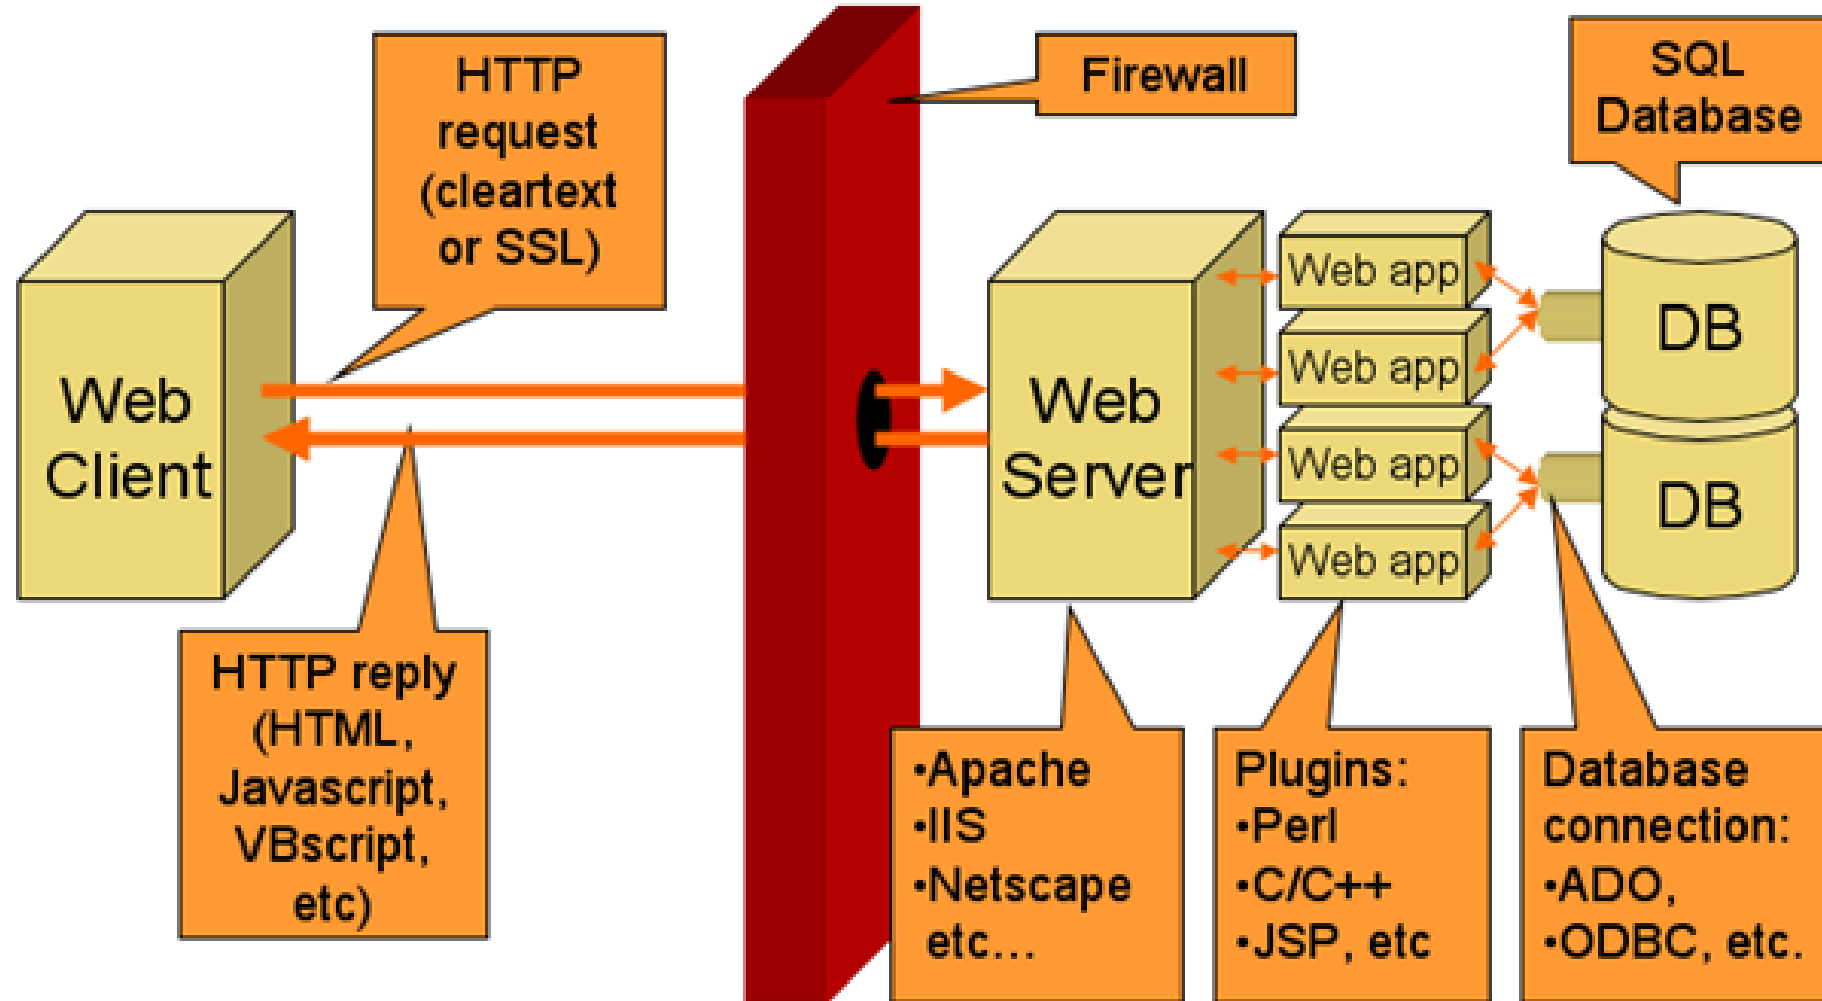
- a program which could produce HTML file
  - It follows some standard by which World Wide Web servers may access external programs so that data is returned automatically in the form of a web page
- You can use C to implement a **CGI** [Common Gate Interface] program which is called by web server

□ JSP (JavaServer Pages ) is one of those to support dynamic





## □ Now many techniques for dynamic – JSP is one of them





# Clarify some concepts

## □ Taxonomy

I hope you could understand the difference by yourself

|        |                                 |  |
|--------|---------------------------------|--|
|        | embedded in HTML                | separate   |
| server | SSI, ASP, PHP, <b>JSP</b> , CFM | server API (NSAPI), <b>cgi</b> , <b>servlets</b> |
| client | <b>JavaScript</b>               | Java applets, plug-in                            |

# CGI example in C

- ❑ CGI is the server-side program to produce HTML to client.
- ❑ So, we can also use CGI to produce a web page with FORM

```
# include <stdio.h>
# include <stdlib.h>
void main(){
    printf( "Content-type:text/html;charset=utf-8\n\n");
    printf( " <html> \n ");
    printf( " <meta http-equiv=\"Content-Type\" content=\"text/html; charset=utf-8\" /> \n ");
    printf( " <title> Your first web page </title> \n ");
    printf( " <body> \n ");

    printf(" <form action=\"http://localhost/cgi-bin/Untitled6.cgi\">");
    printf(" <p>Pls input two numbers for multiplication.</p>");
    printf(" <input name=\"m\" size=\"5\">");
    printf(" <input name=\"n\" size=\"5\"> <br/>");
    printf(" <input type=submit value=\"OK\" ></form> ");

    printf( " </body> ");
    printf( " </html> ");
}
```



# Sample Servlet

## □ Servlet works like a CGI

```
import java.io.*;
import javax.servlet.*;
import javax.servlet.http.*;

public class HelloWorld extends
    HttpServlet {

    public void doGet(HttpServletRequest
        request, HttpServletResponse
        response)
        throws IOException, ServletException
    {
        response.setContentType("text/html");
```

## Hello World!

```
        PrintWriter out = response.getWriter();
        out.println("<html>");
        out.println("<body>");
        out.println("<head>");
        out.println("<title>Hello World!</title>");
        out.println("</head>");
        out.println("<body>");
        out.println("<h1>Hello World!</h1>");
        out.println("</body>");
        out.println("</html>");
    }
```



## □ JavaScript example – executed in client

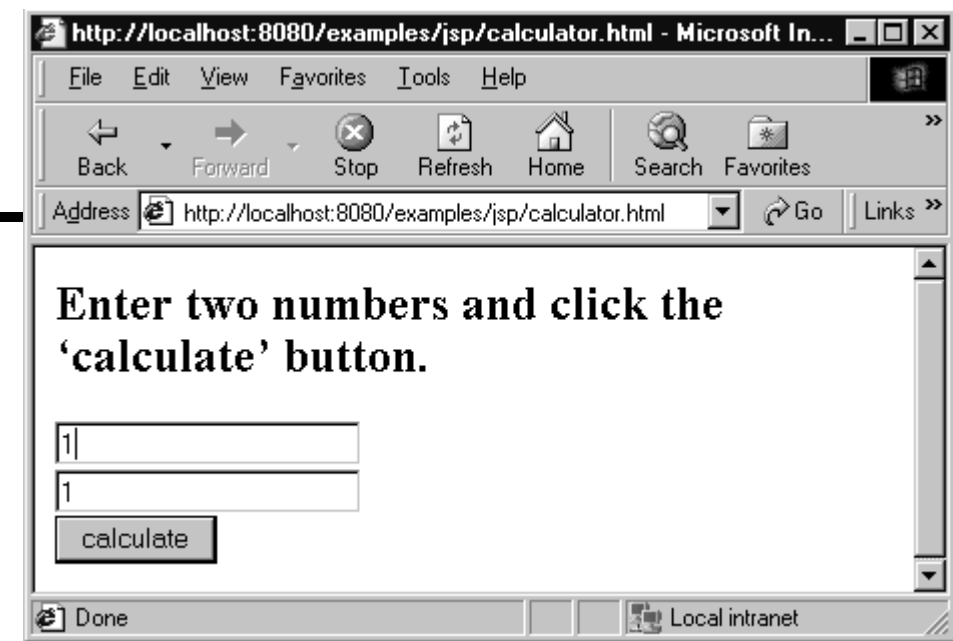
```
<!DOCTYPE html>␣  
<!-- JSExAlertWrite.html -->␣  
<html lang="en">␣  
<head>␣  
  <meta charset="utf-8">␣  
  <title>JavaScript Example: Functions alert() and document.write()</title>␣  
  <script>␣  
    alert("Hello, world!");␣  
  </script>␣  
</head>␣  
<body>␣  
  <h1>My first JavaScript says:</h1>␣  
  <script>␣  
    document.write("<h2><em>Hello world, again!</em></h2>");␣  
    document.write("<p>This document was last modified on "␣  
      + document.lastModified + ".</p>");␣  
  </script>␣  
</body>␣  
</html>
```



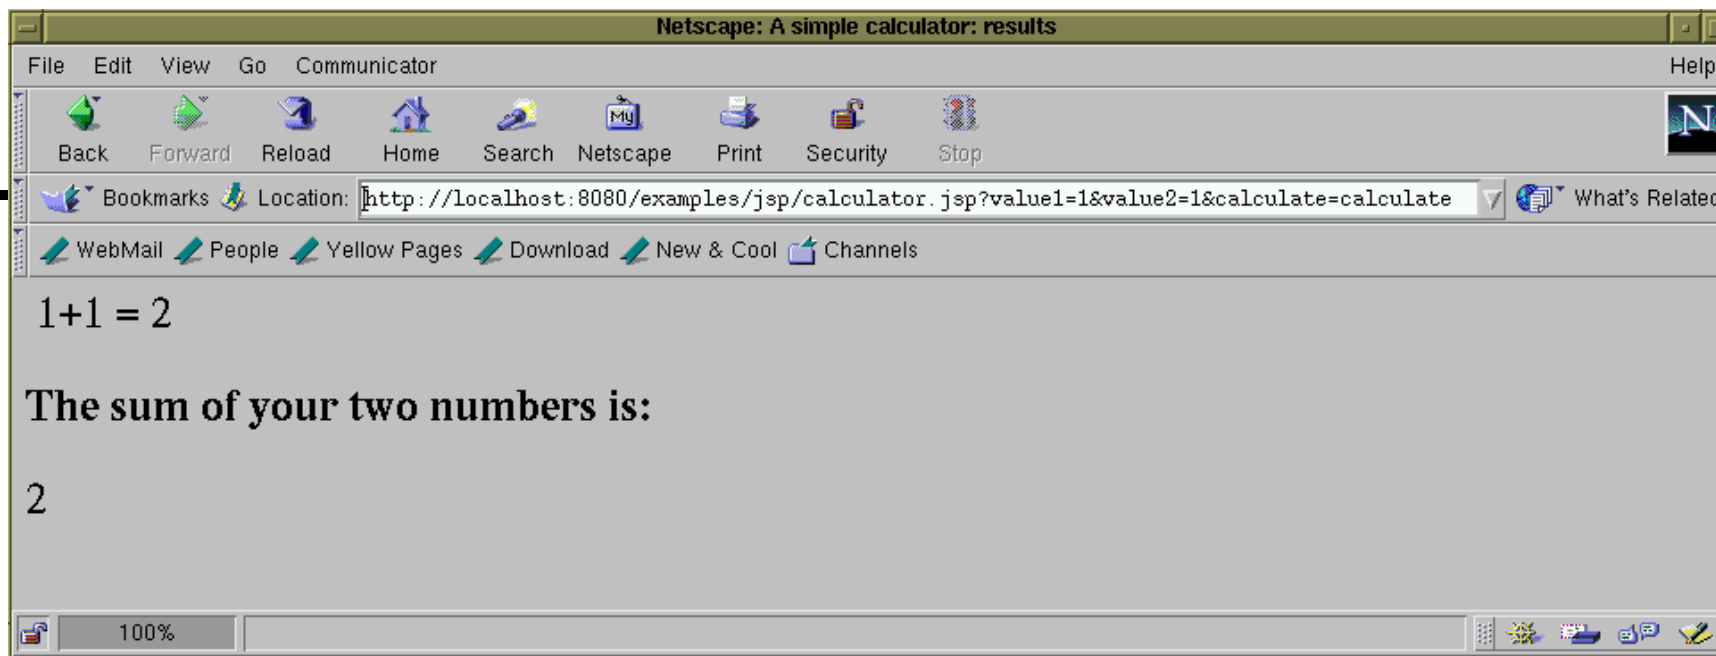
# JSP example – Calculator

```
<html>
<head></head>
<body>
<p>Enter two numbers and click the
'calculate' button.</p>
```

```
<form action="calculator.jsp" method="get">
<input type="text" name="value1"><br>
<input type="text" name="value2 "><br>
<input type="submit" name="calculate" value="calculate">
</form>
</body>
</html>
```



Calculator.html



```
<html>
<head><title>A simple calculator: results</title></head>
<body>
<!-- A simpler example 1+1=2 -->
1+1 = <%= 1+1 %>
<!-- A simple calculator -->
<h2>The sum of your two numbers is:</h2>
<%= Integer.parseInt(request.getParameter("value1")) +
      Integer.parseInt(request.getParameter("value2")) %>
</body>
</html>
```

Calculator.jsp

# CLI → GUI (C/S) → B/S (3 tiers)

Last login: Fri Jan 10 12:38:06 2014 from 10.10.0.106

```
Interface eth0 IP: 10.10.0.251
Interface eth1 IP: 172.24.140.1
```

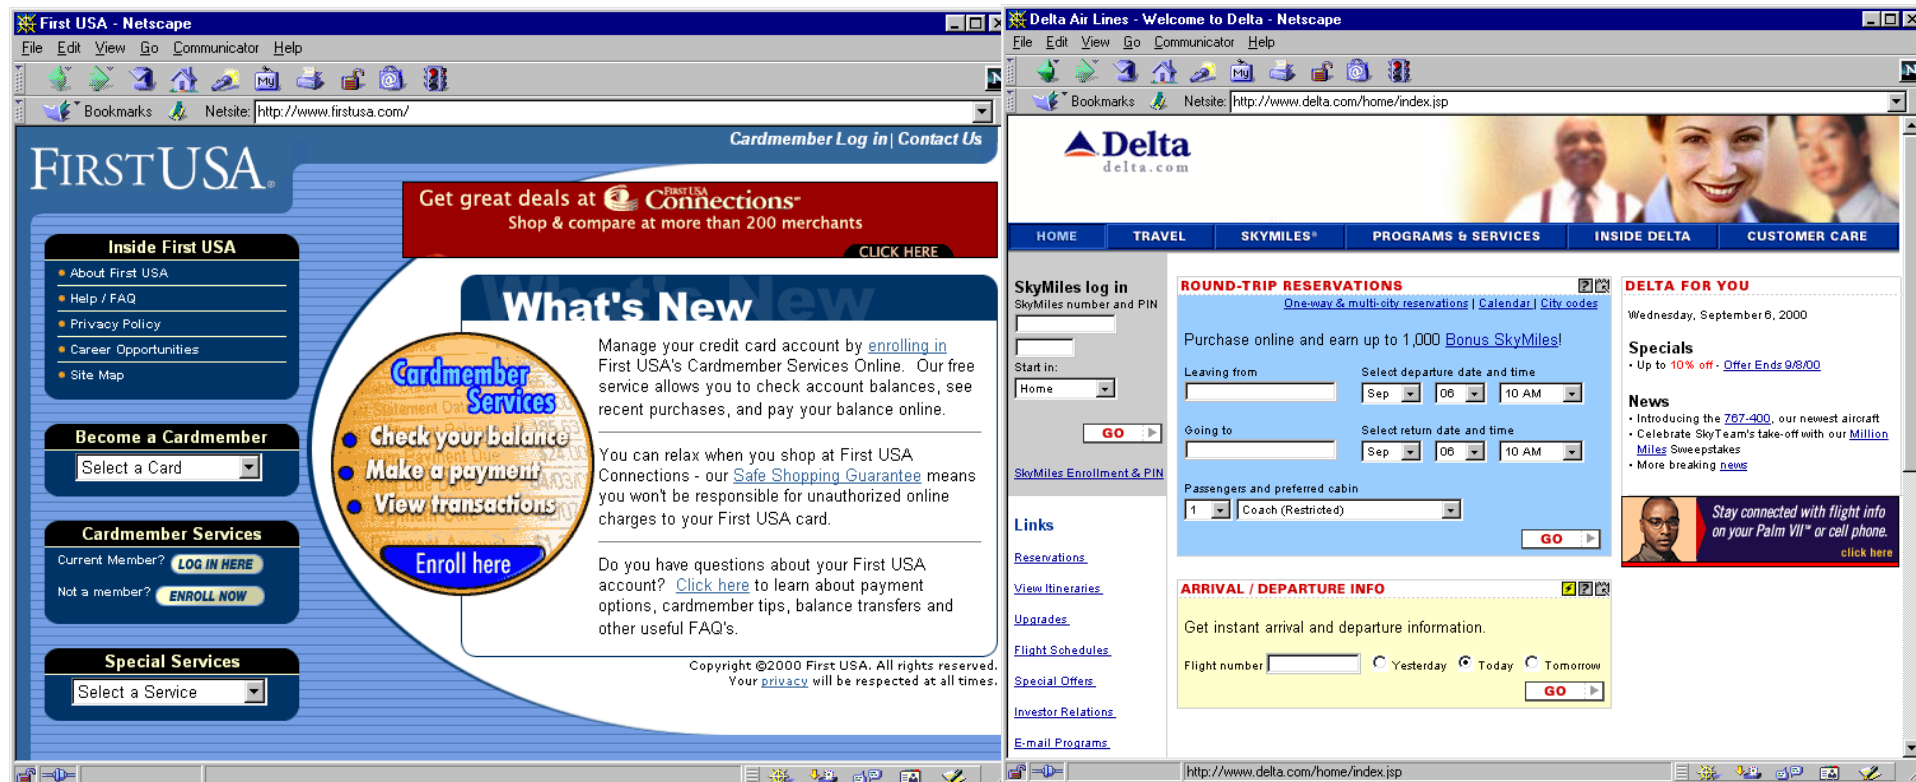
Please note most tasks should be handled via the FreePBX GUI.  
You can access the FreePBX GUI by typing `http://10.10.10.10:8080`  
For support please visit <http://www.freepbx.org>

```
[root@freepbx-a ~]#
```



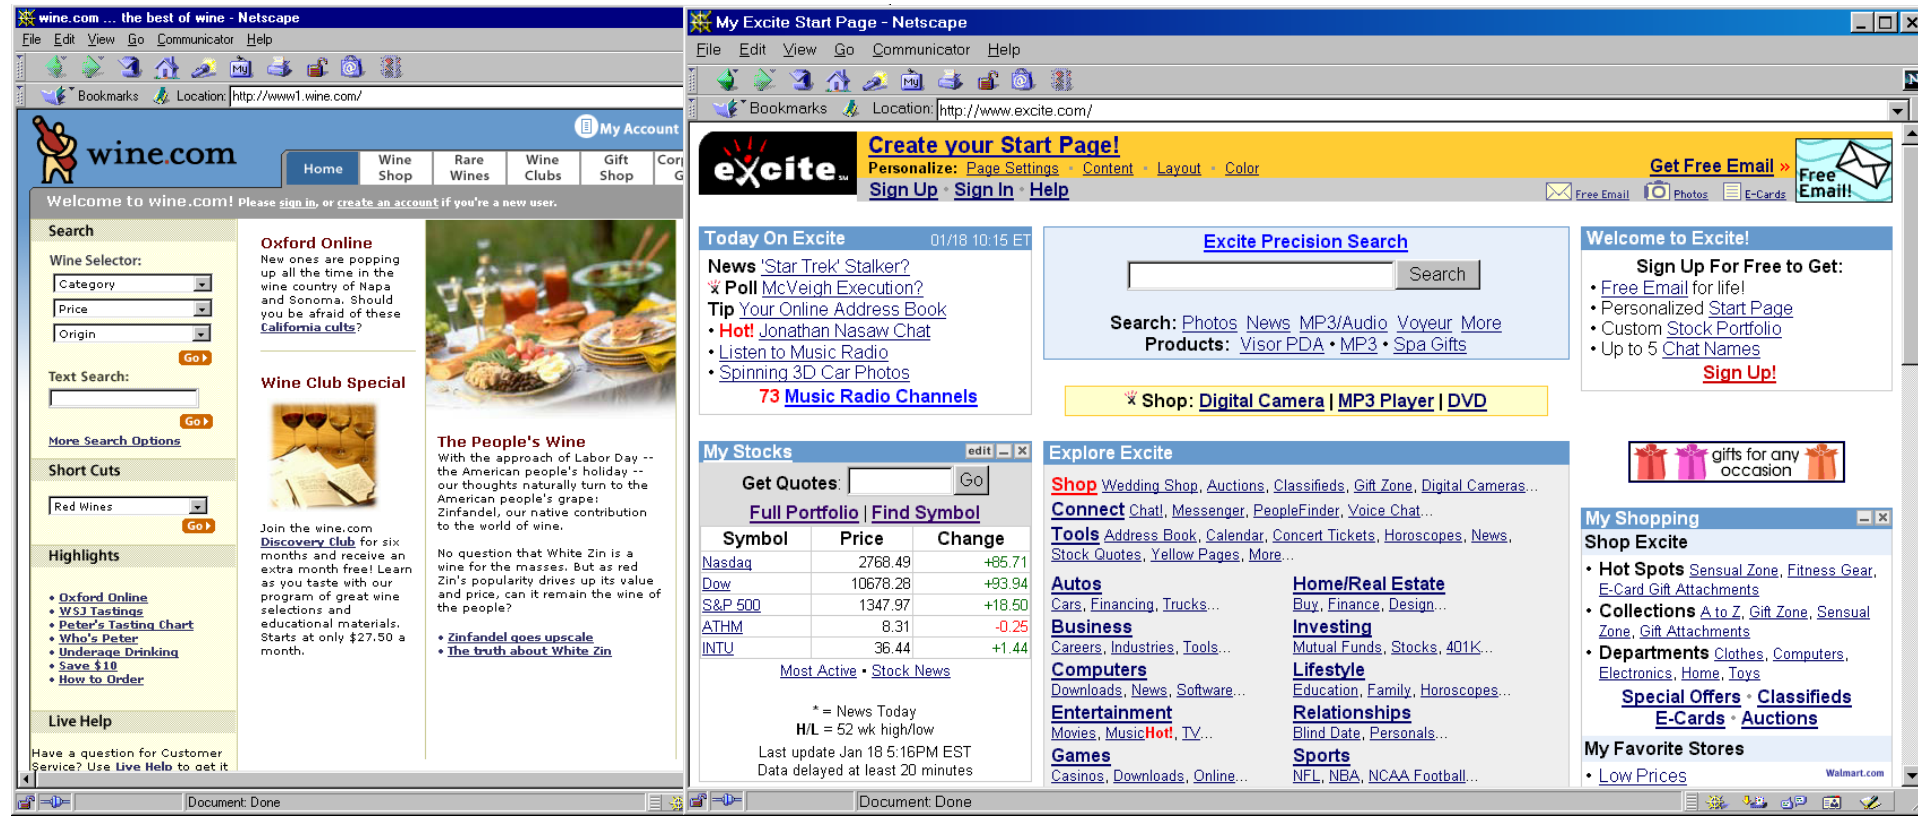
## □ JSP/Servlets in the Real World

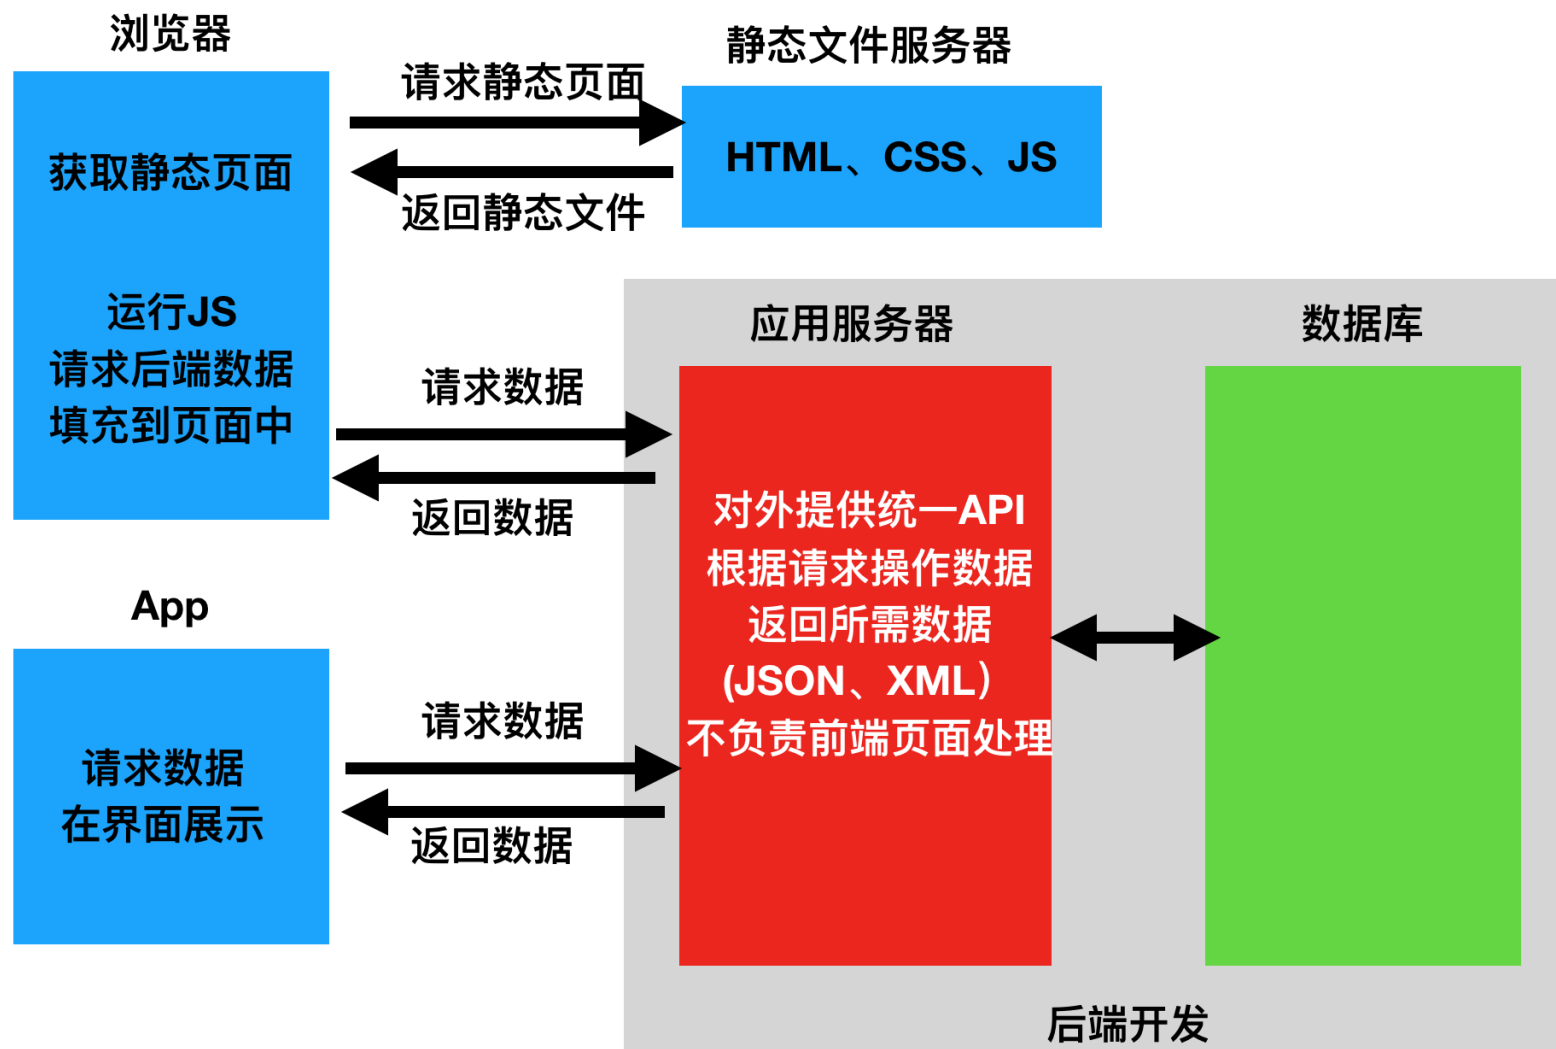
- **Delta Airlines:** entire Web site, including real-time schedule info
- **First USA Bank:** largest credit card issuer in the world; most on-line banking customers

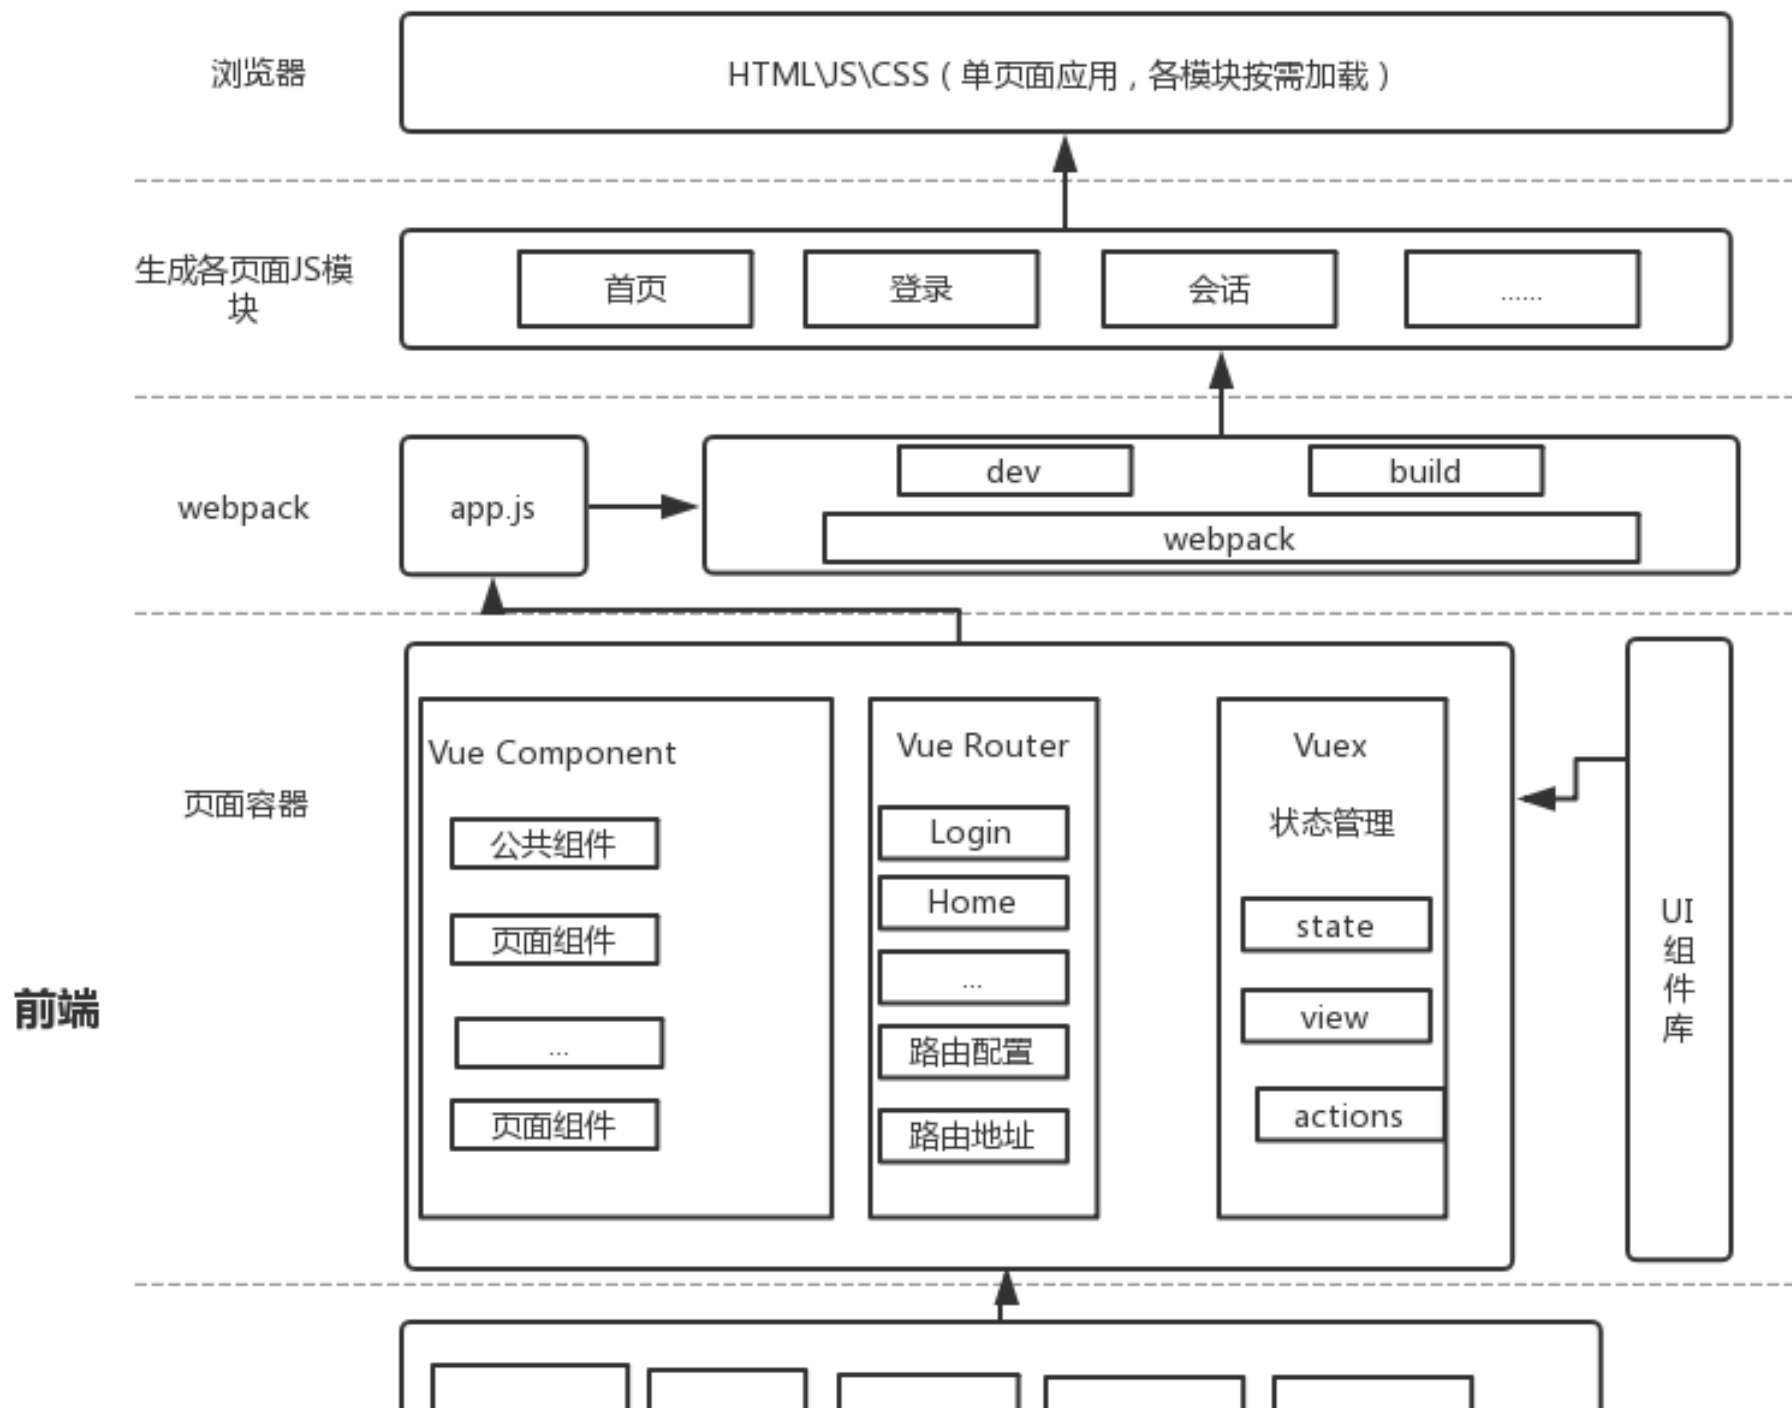


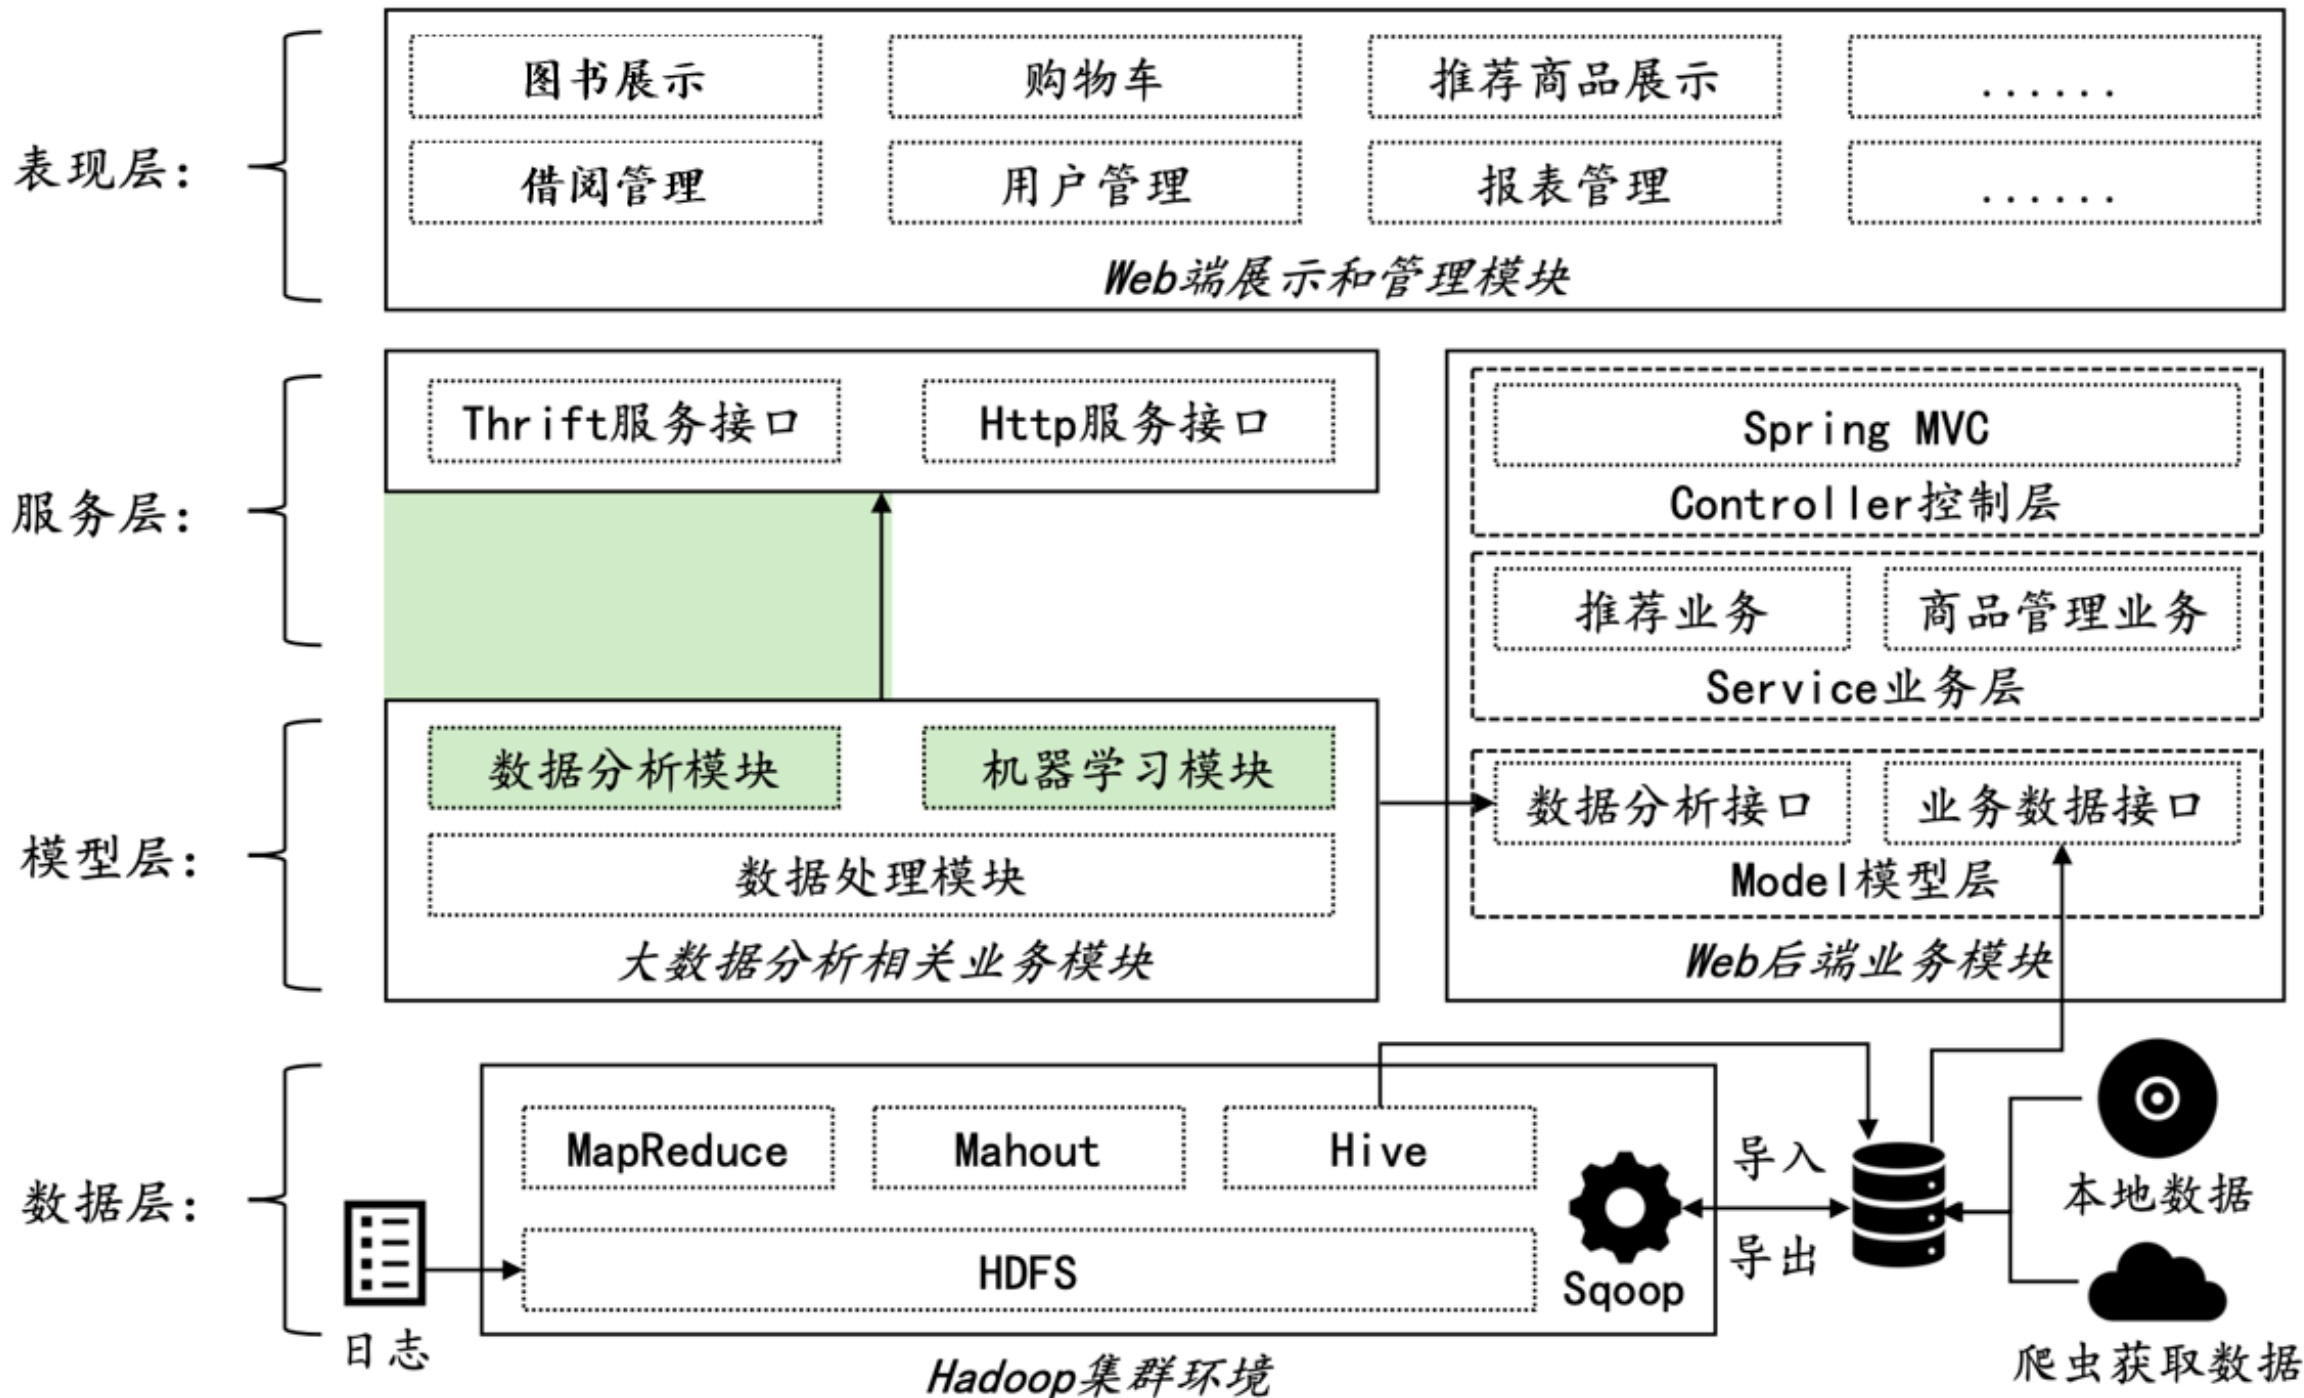


- **Excite:** one of the top five Internet portals; one of the ten busiest sites on the Web









## □ Requirement for Project 5

### ■ The JSP based Web app should support

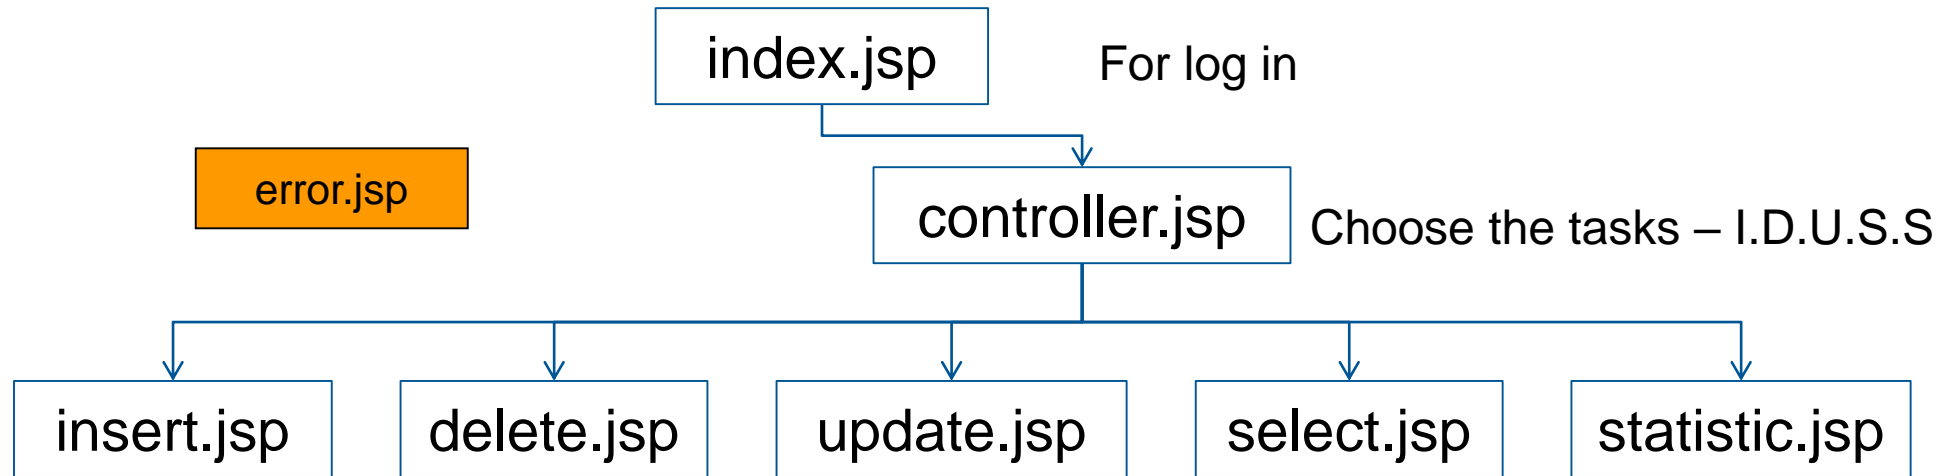
- Log in module
- I.D.U.S
- Some simple statistics like AVG, Top 100, ...
- Try transaction commit

It's your responsibility  
to learn JSP and finish  
this “Web Application”

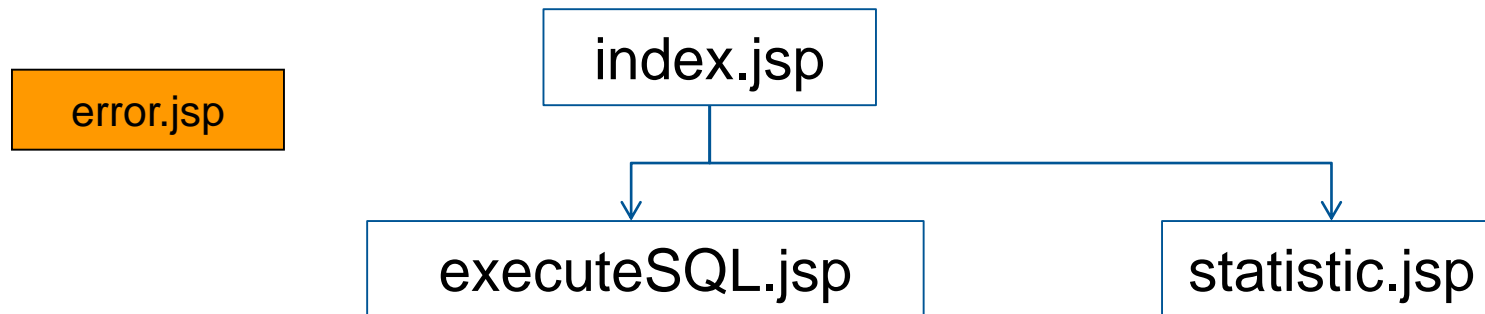
### ■ You can choose your own data (or “on-line retail”) to finish your Web app

- Data should be kept in PostgreSQL
  - ✓ Show the db design in document
    - » You can use PowerDesigner to design and try database first
  - ✓ How to connect PostgreSQL is your responsibility

# Hints



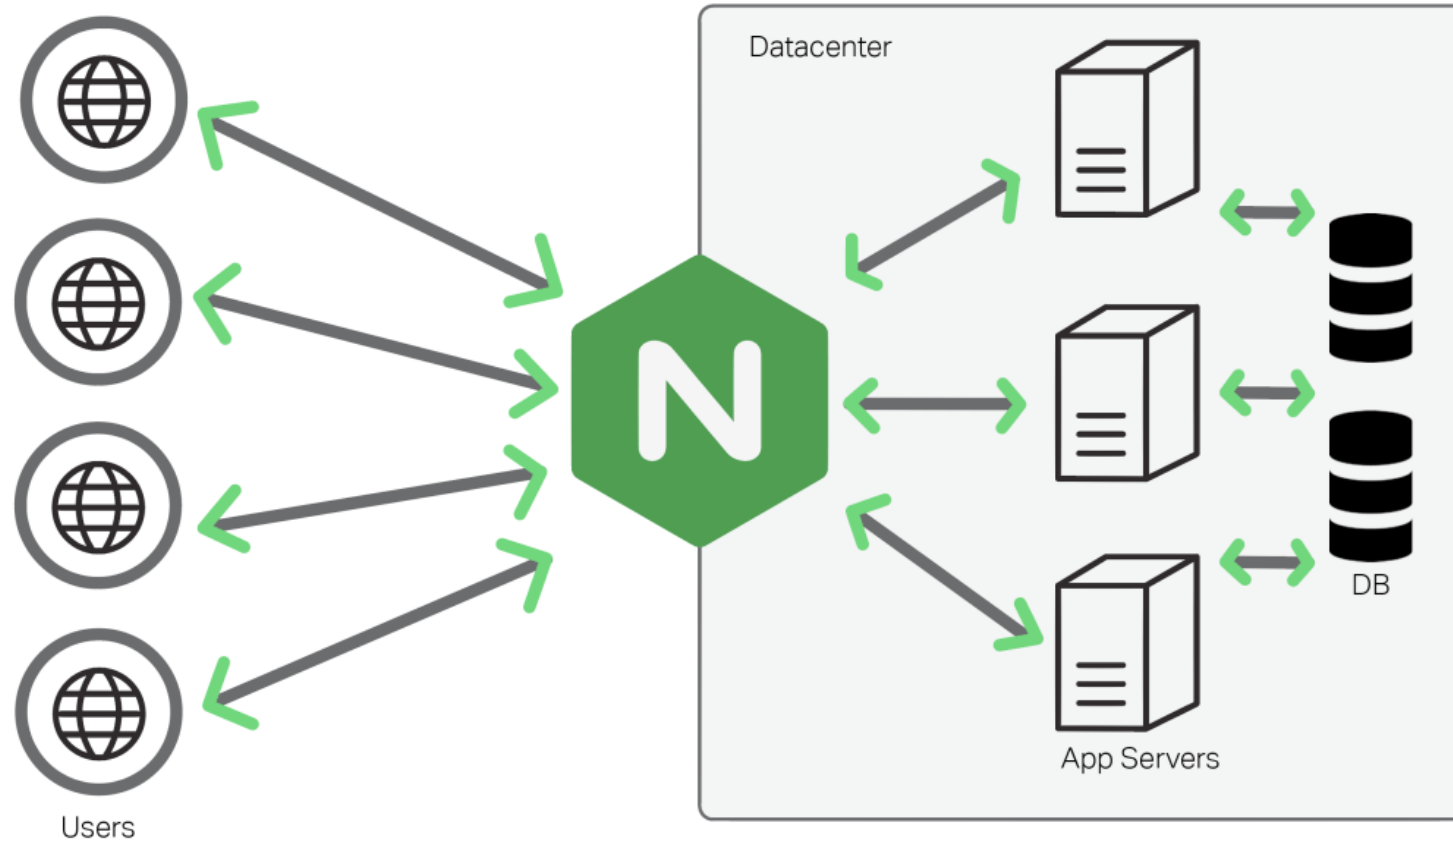
OR



You input SQL, and execute it  
Show the result

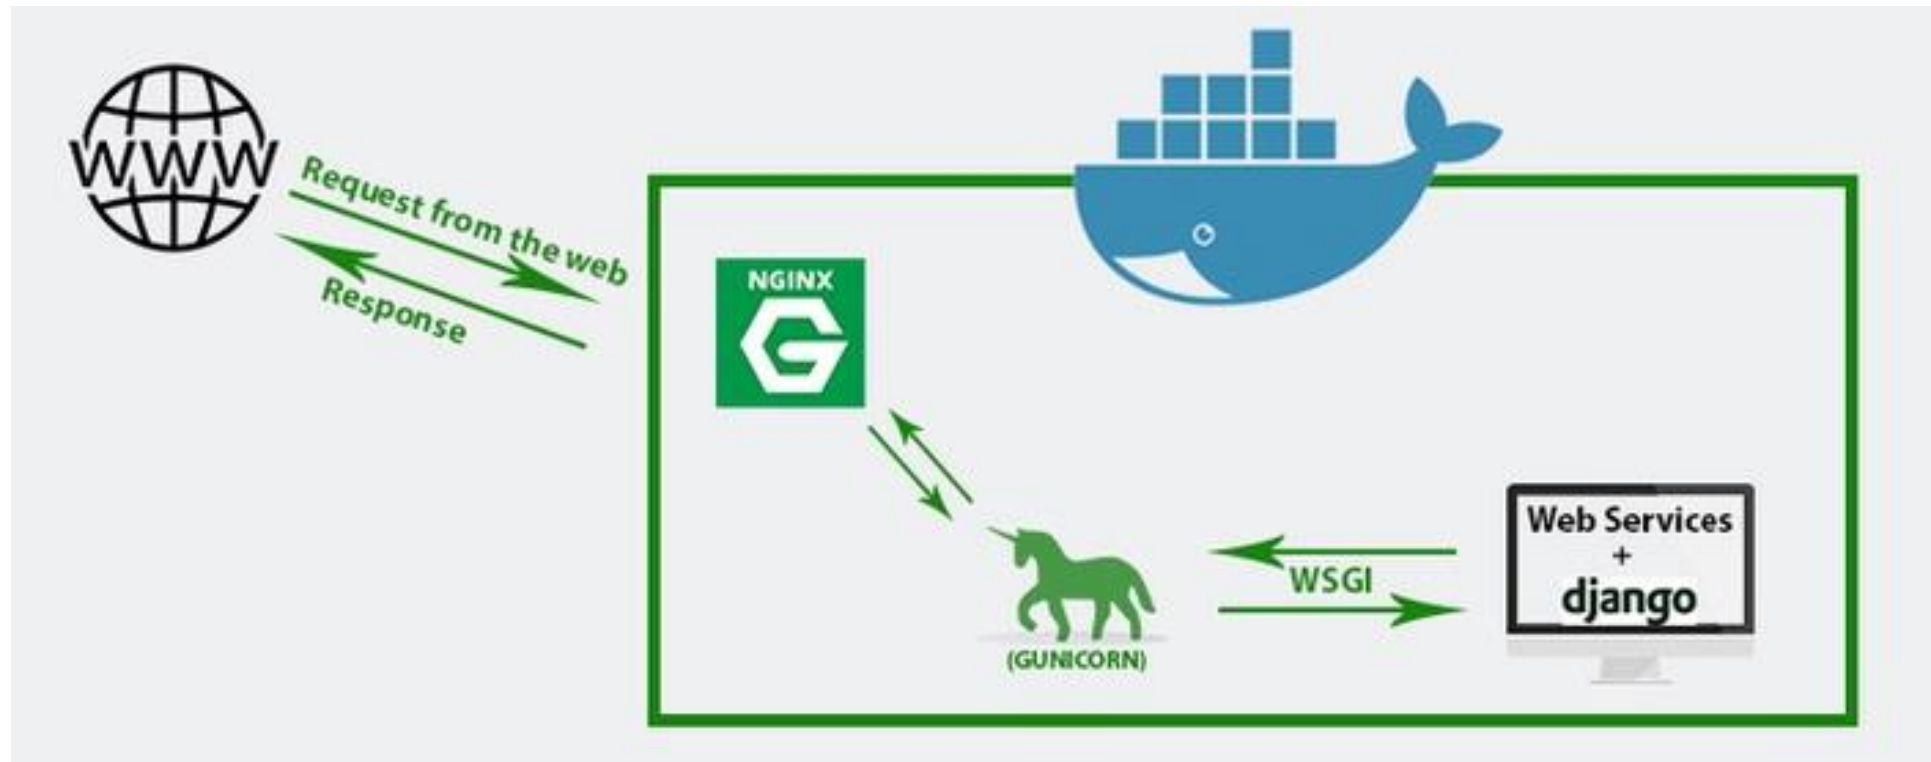
# More complex architecture In Big Data age

□ + Nginx



# More complex architecture In Big Data age

□ **Linux** + Docker, Nginx, Django etc.



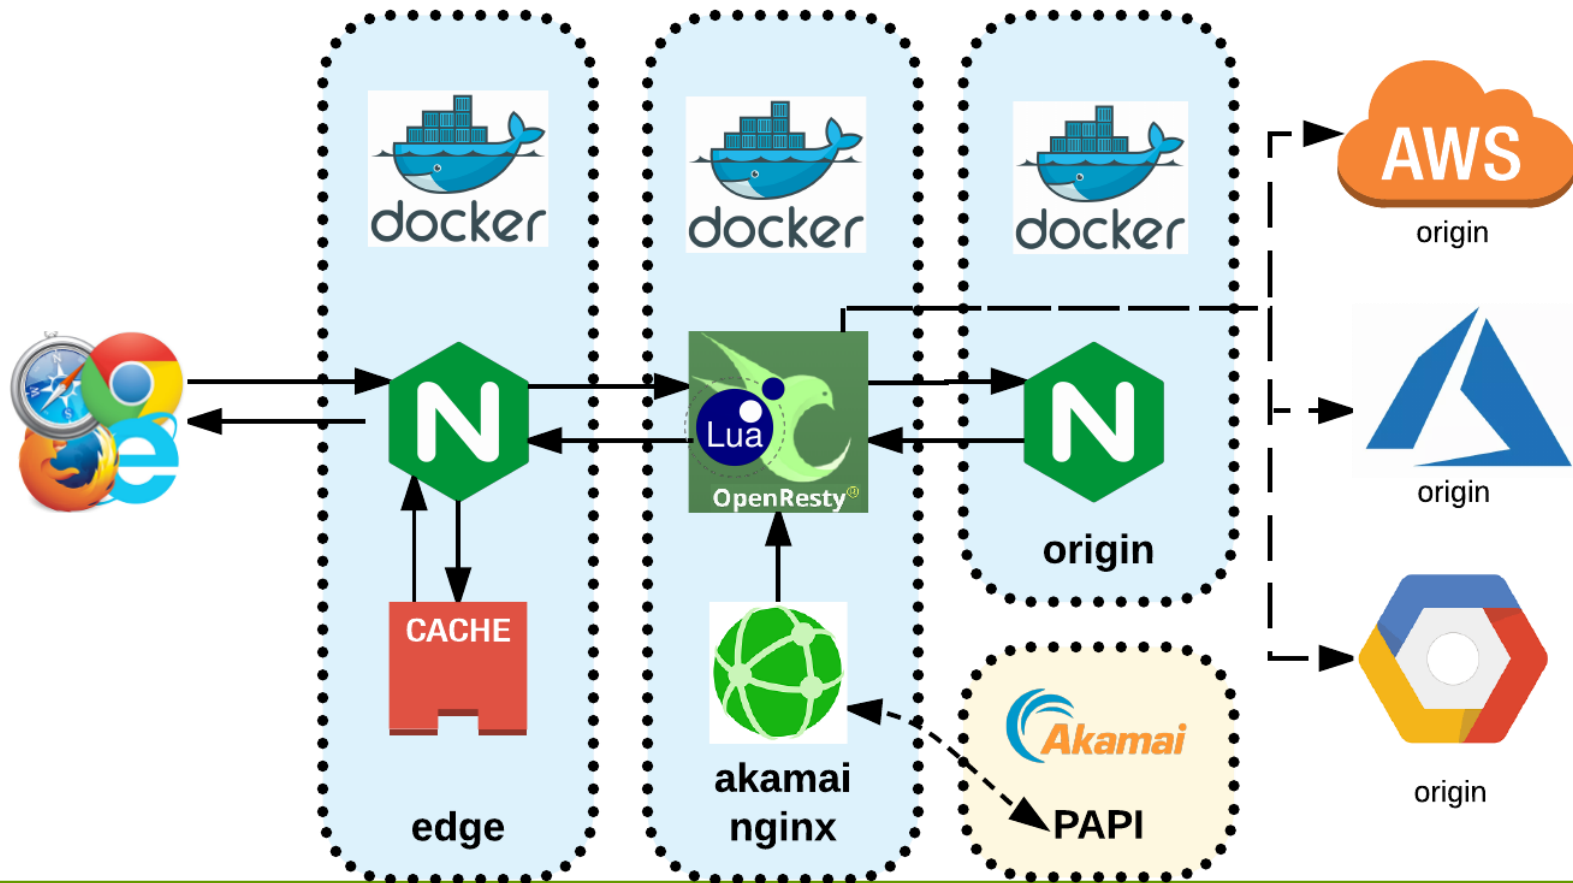
<https://www.codementor.io/samueljames/nginx-setting-up-a-simple-proxy-server-using-docker-and-python-django-f7hy4e6jv>

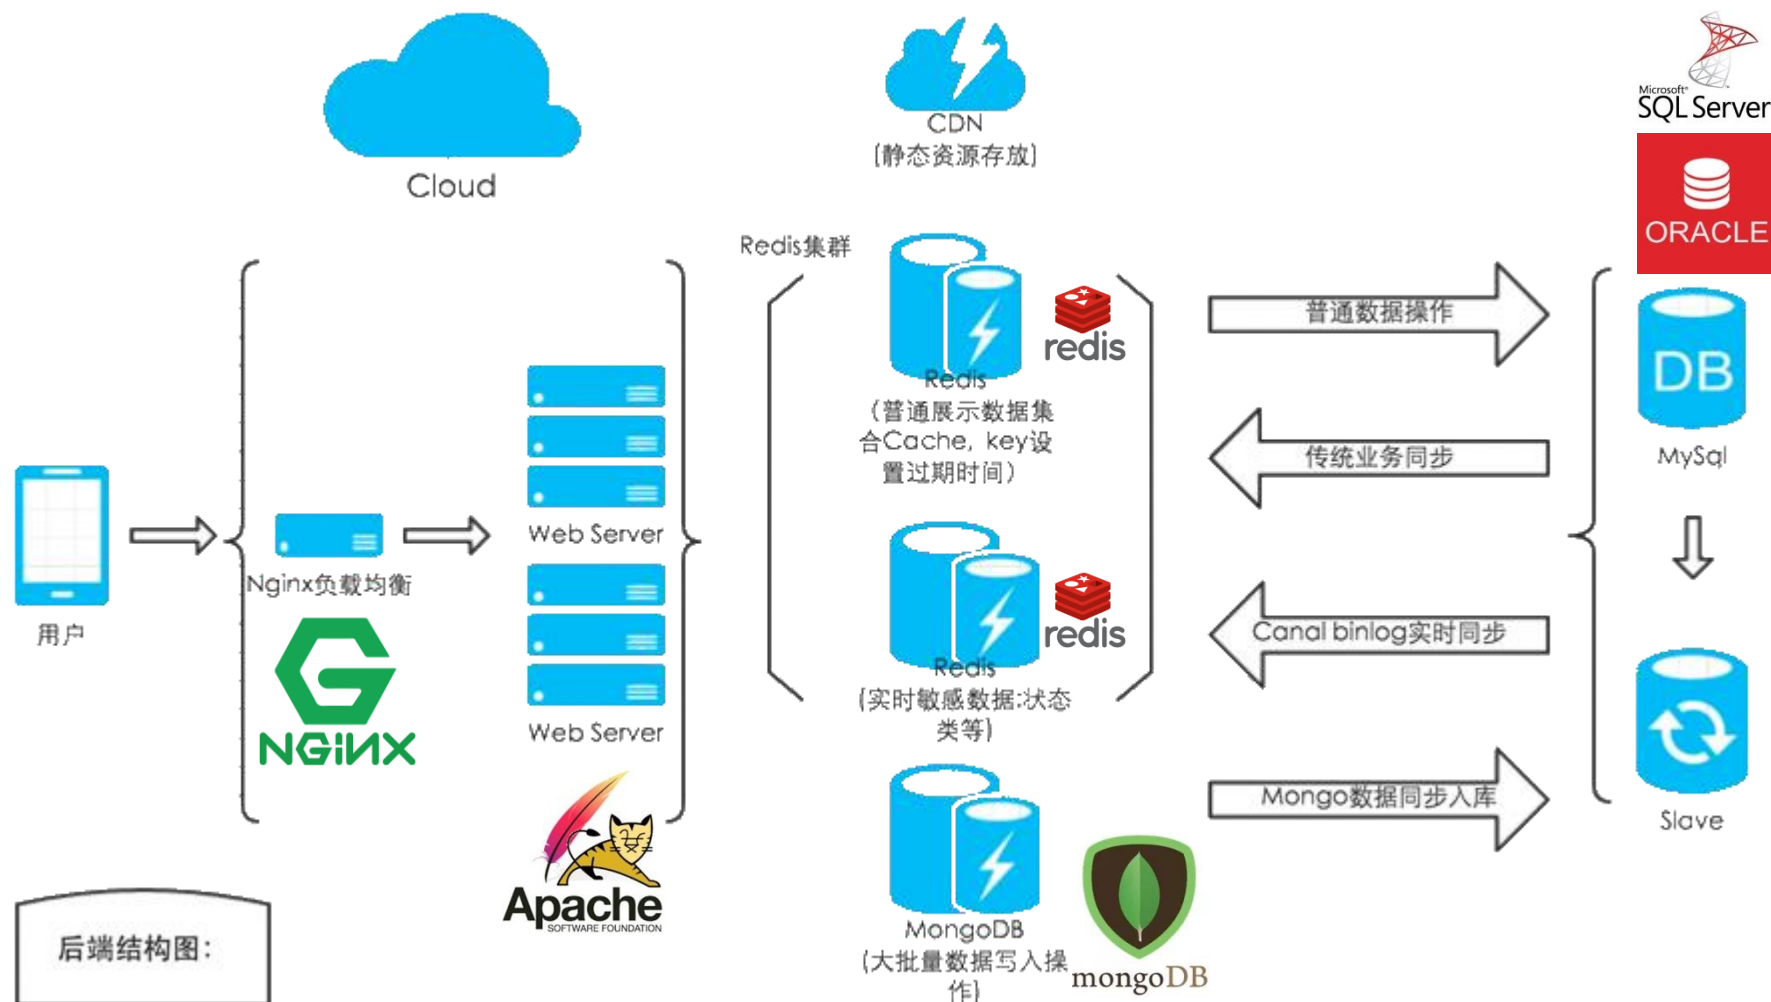


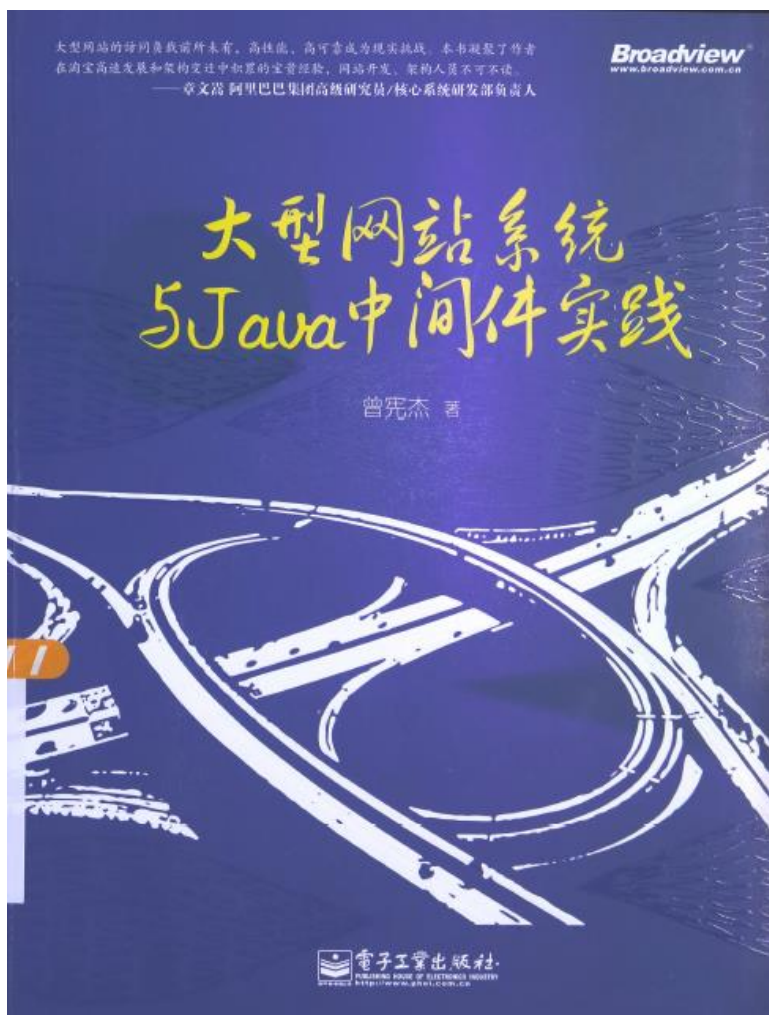
# More complex architecture In Big Data age

□ **Linux** + Docker, Nginx, Lua etc.

akamai-nginx request flow







## □ 大型网站系统与Java 中间件实践

## □ 然后再看看下面这本书—— 更详细的技术细节

### ■ 亿级流量网站架构核心技术



Broadview®  
www.broadview.com.cn

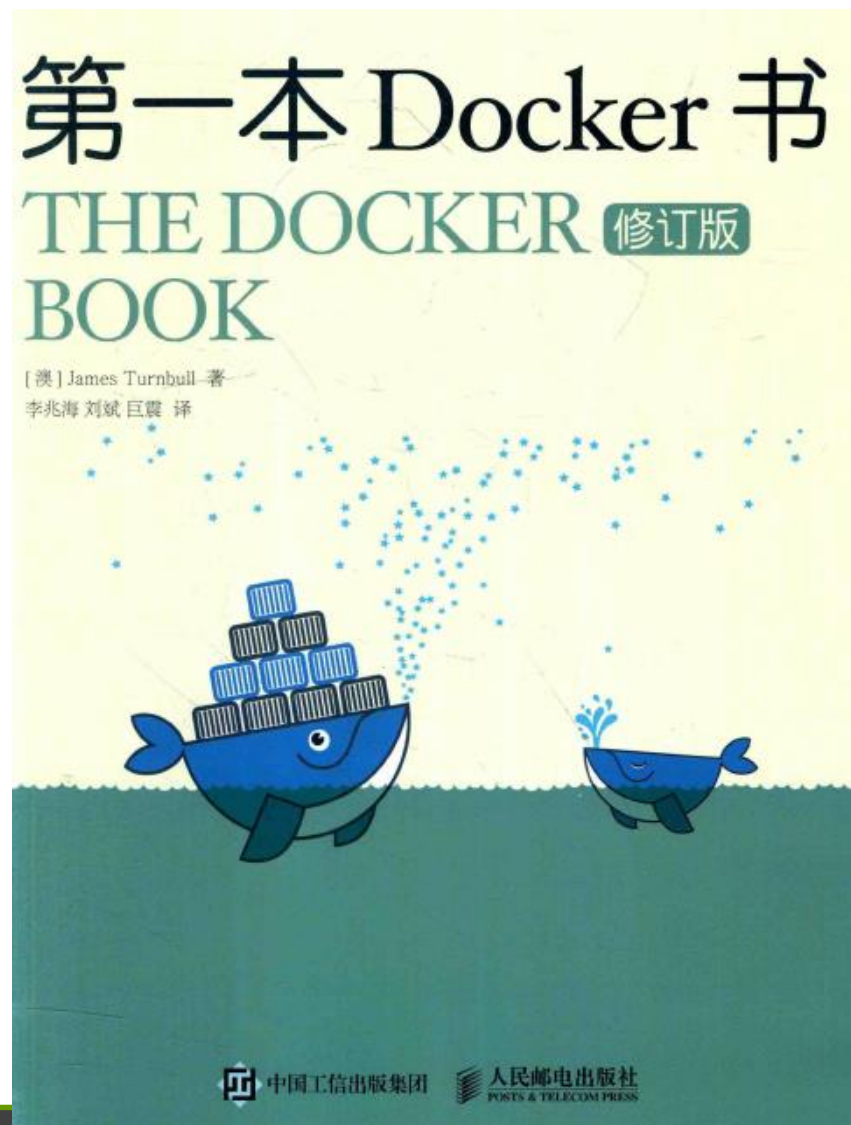
# 大型网站技术架构

核心原理与案例分析

李智慧 著



电子工业出版社  
PUBLISHING HOUSE OF ELECTRONICS INDUSTRY  
www.phei.com.cn



- 第一本Docker 书
- The Docker Book

# The way this semester to evaluate **100 pts** based on

---

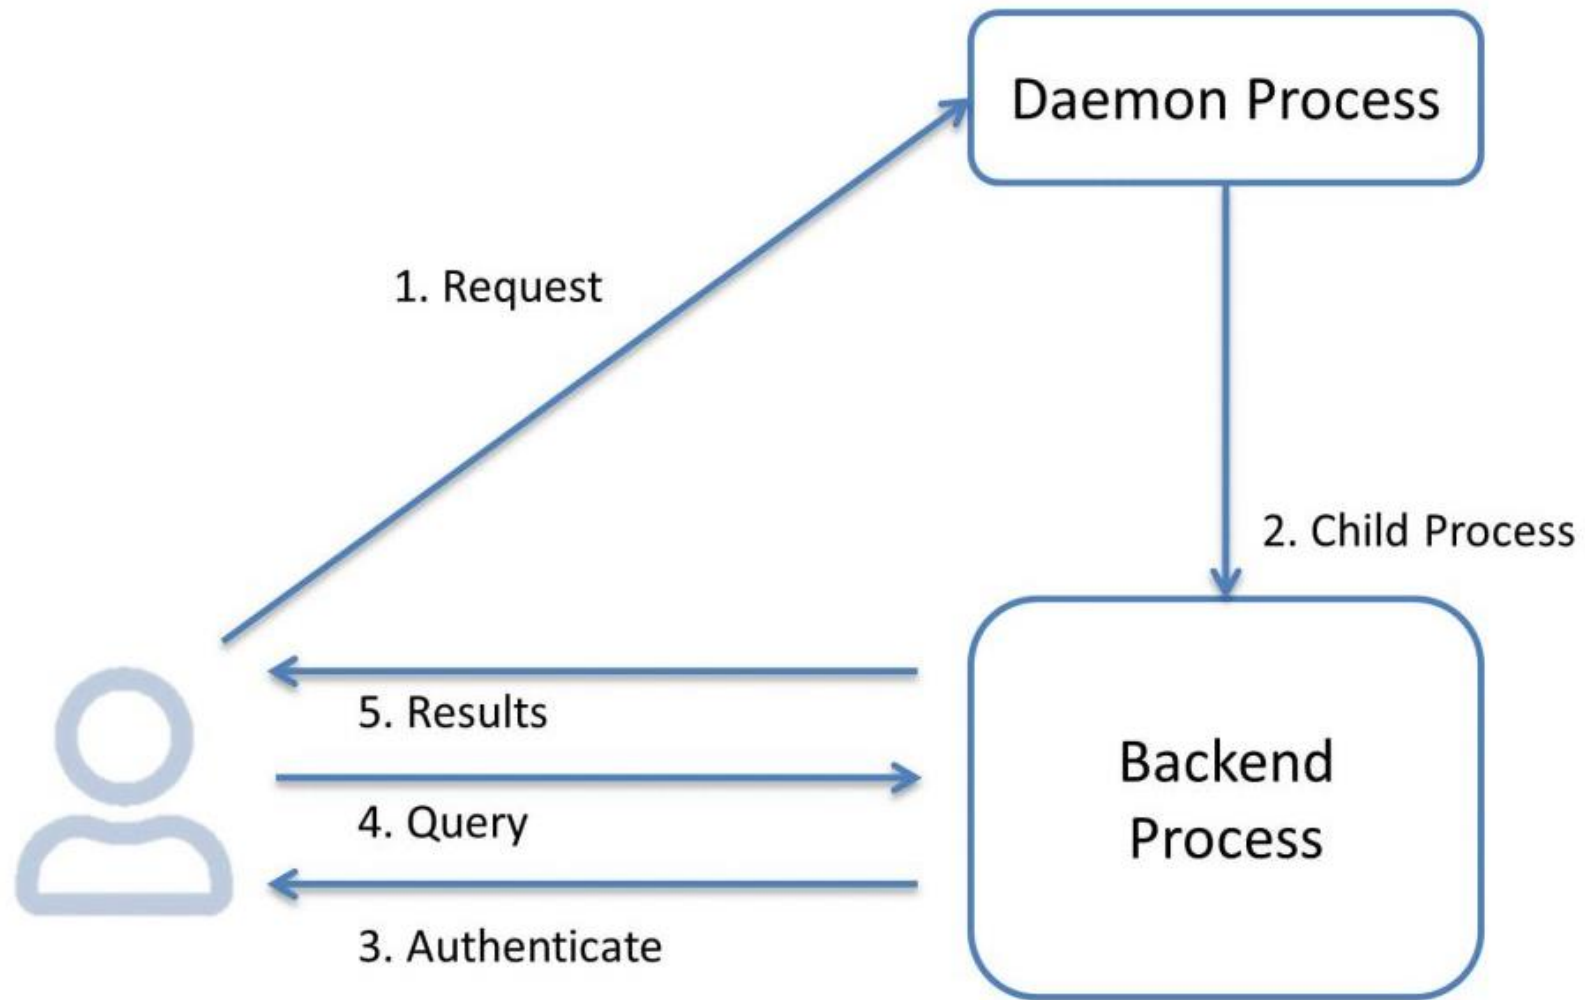
## ■ Basic (70)

- 1.
- 2.
3. Try a Syntax parsing example
4. Lock table management
- 5.

- Understanding [multi-process]

## ■ Advanced (10) in PostgreSQL/Greenplum/HAWQ

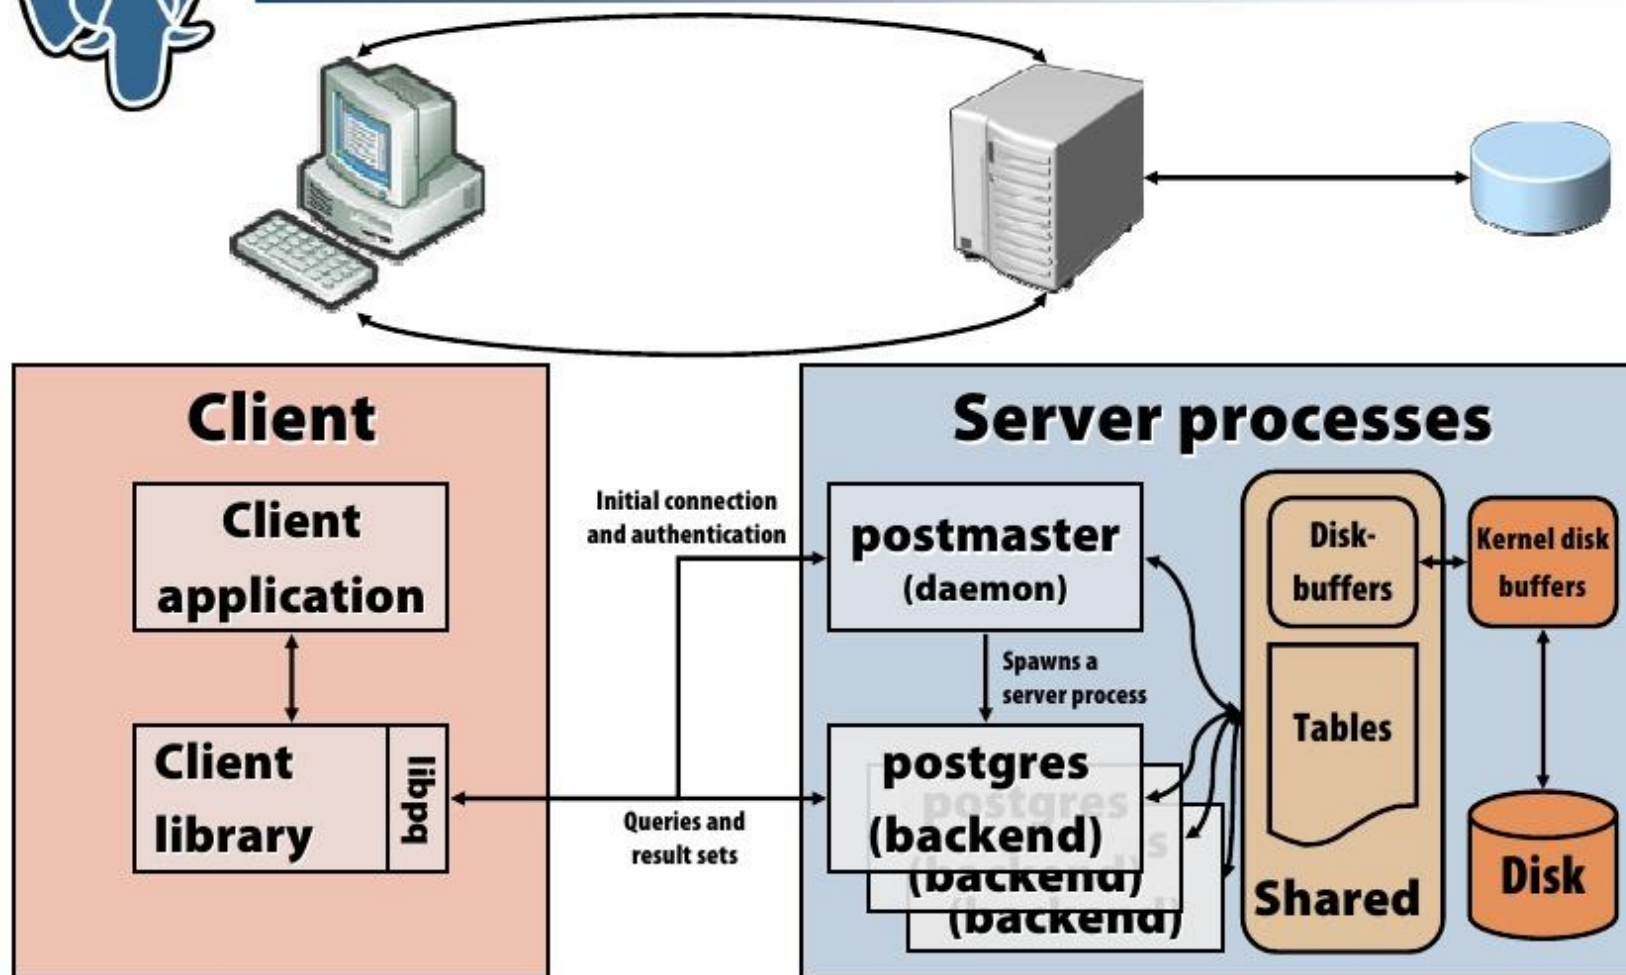
➤ [transaction mechanism][SQL execution]







# Overall architecture



```
23961 Ss 0:05.64 /usr/local/bin/postmaster (postgres)
23963 S 0:01.13 postmaster: stats buffer process (postgres)
23966 S 0:03.24 postmaster: stats collector process (postgres)
36324 I 0:00.43 postmaster: oddbjorn testdb [local] idle (postgres)
36428 I 0:00.23 postmaster: oddbjorn testdb [local] idle (postgres)
```



# The way this semester to evaluate **100 pts** based on

---

## ■ Basic (70)

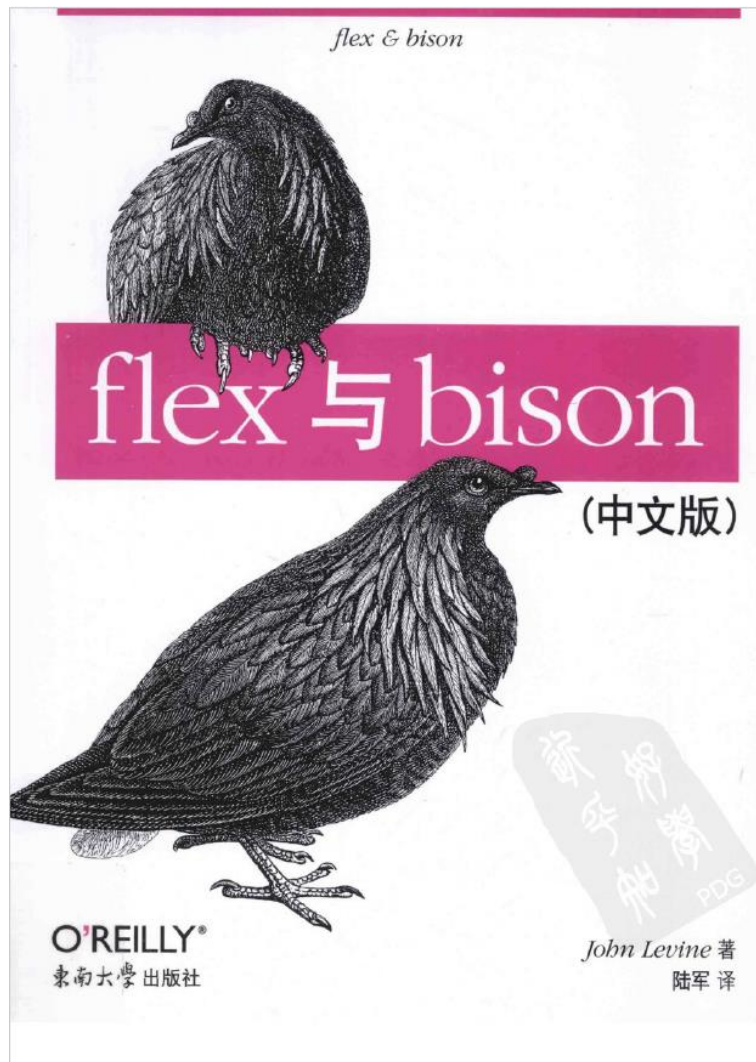
- 1.
- 2.
3. Try a Syntax parsing example
4. Lock table management
- 5.



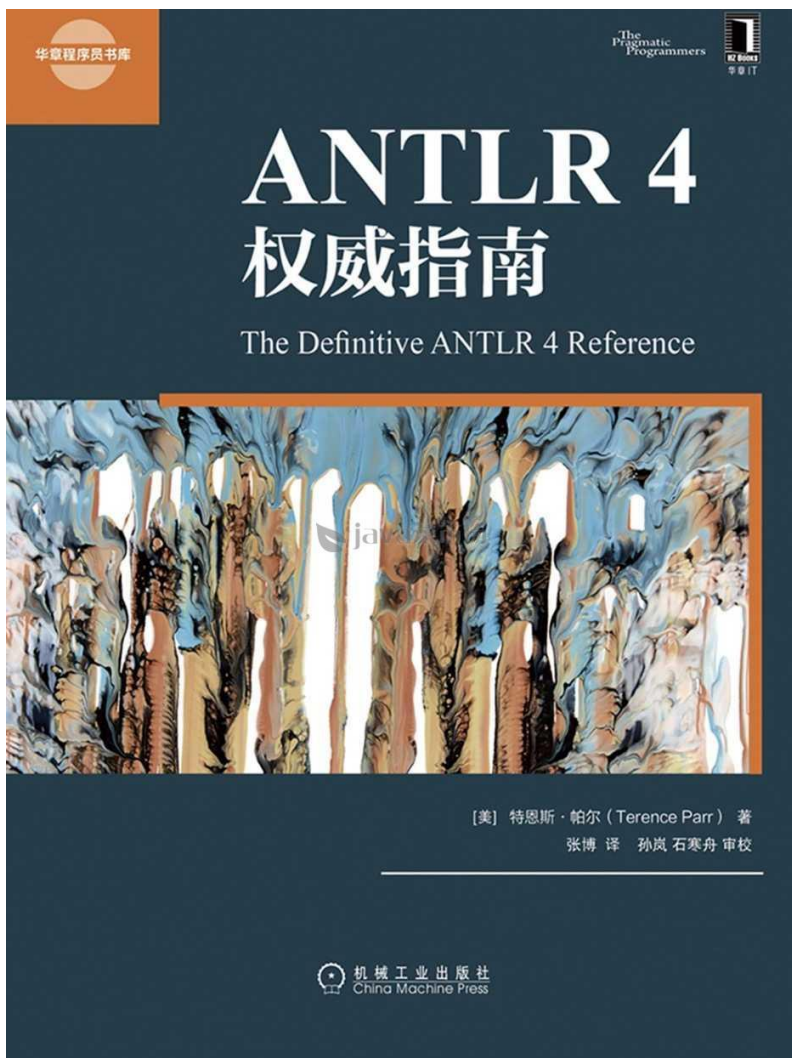
Later

## ■ Advanced (10) ] in PostgreSQL/Greenplum/HAWQ

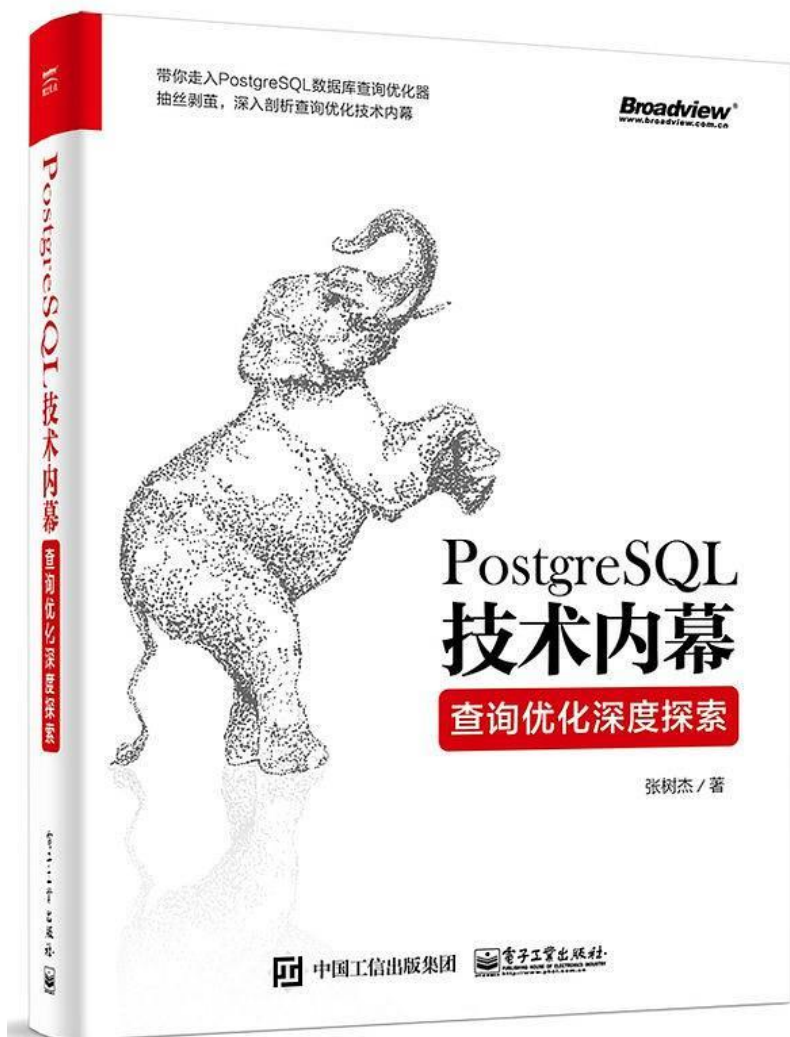
➤ [transaction mechanism][SQL execution]



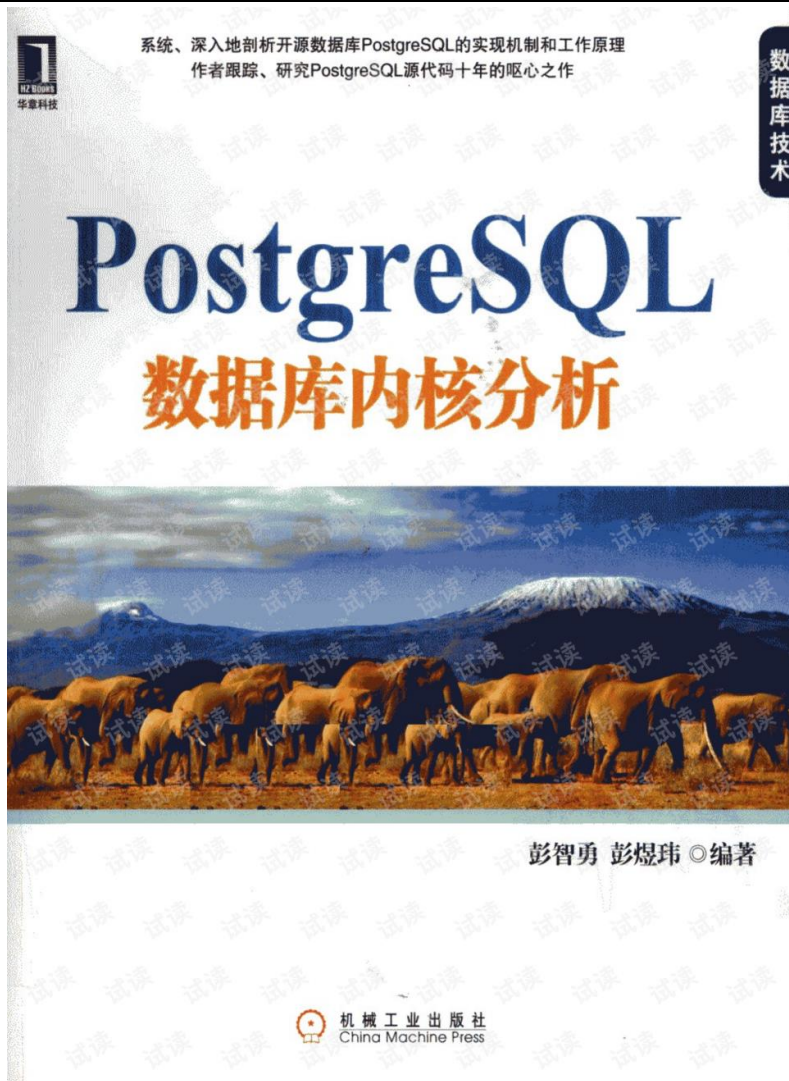
- ❑ flex & bison - Unix Text Processing Tools
- ❑ John Levine
- ❑ O'Reilly Media
- ❑ 2009



- ❑ ANTLR 4权威指南
- ❑ 作者: Terence Parr
- ❑ 出版社: 机械工业出版社
- ❑ 译者: 张博
- ❑ 出版年: 2017-5-1
- ❑ 页数: 262
- ❑ 定价: 69元
- ❑ 丛书: 华章程序员书库
- ❑ ISBN: 9787111566489



- ❑ PostgreSQL技术内幕：查询优化深度探索
- ❑ 作者：张树杰
- ❑ 出版社：电子工业出版社
- ❑ 出版年：2018-6
- ❑ 定价：79
- ❑ ISBN: 9787121341489



- PostgreSQL 数据库内核分析
- 作者: 彭智勇 / 彭煜玮
- 出版社: 机械工业出版社华章公司
- 出版年: 2012-1