



Insight into

DBMS: Design and Implementation

Chapter 4 B: SQL to File Operations

RA → File Operations



孔令波

minking@126.com

+86 15010255486

Outline of the chapters covered

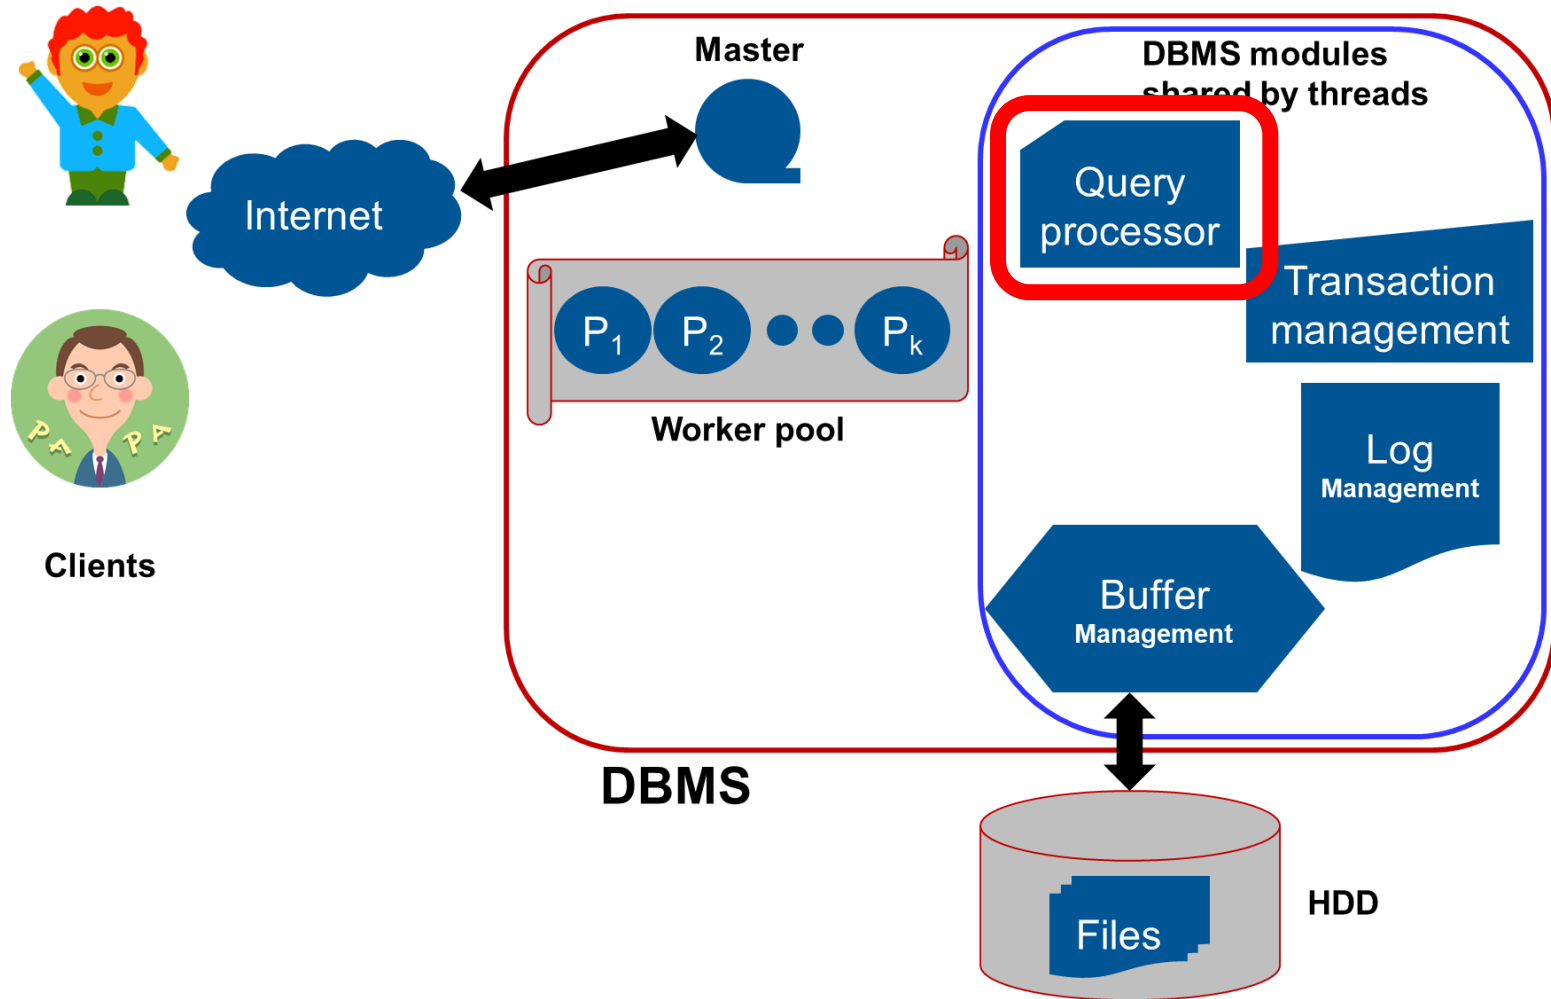
- Introduction
- Overview of the projects
- Demonstration of Development environment
 - Watch and practice by yourself
- My understanding about (R)DBMS
 - History and D&I
- SQL translation with 2 conversions
 - SQL → RA (Relational Algebra)
 - RA → Sequence of File operations
- Transaction control
- Deeper
 - File, (R)DBMS, ERP, DW, Big Data (No SQL, SQL again)
 - SQL on MPP and Hadoop (Greenplum, **HAWQ**)



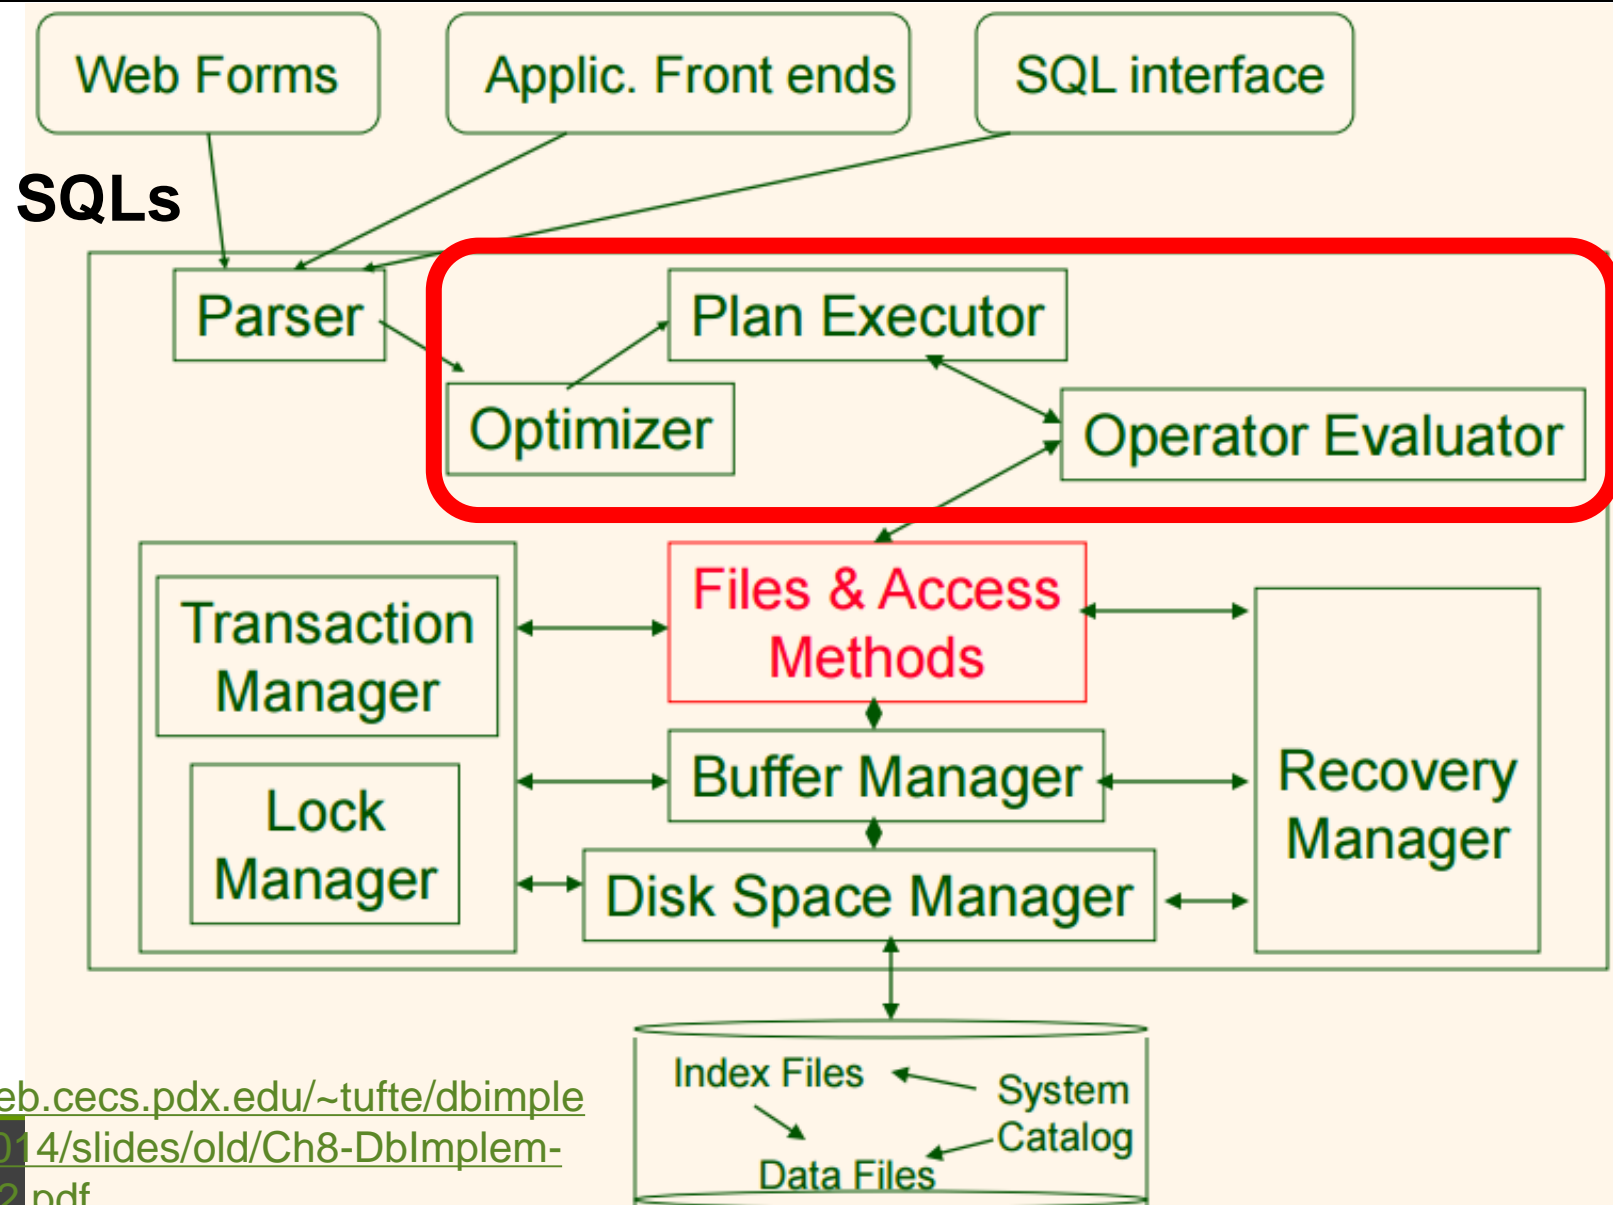
□ Sit back, but **not relax!**

- We'll start the most exciting yet challenging part – SQL translator !





Database Architecture



<http://web.cecs.pdx.edu/~tuft/dbimplem-spr2014/slides/old/Ch8-DbImplem-Fall2012.pdf>



Chapter 3 B: How to get File operations for SQL really

□ My understanding about DBMS

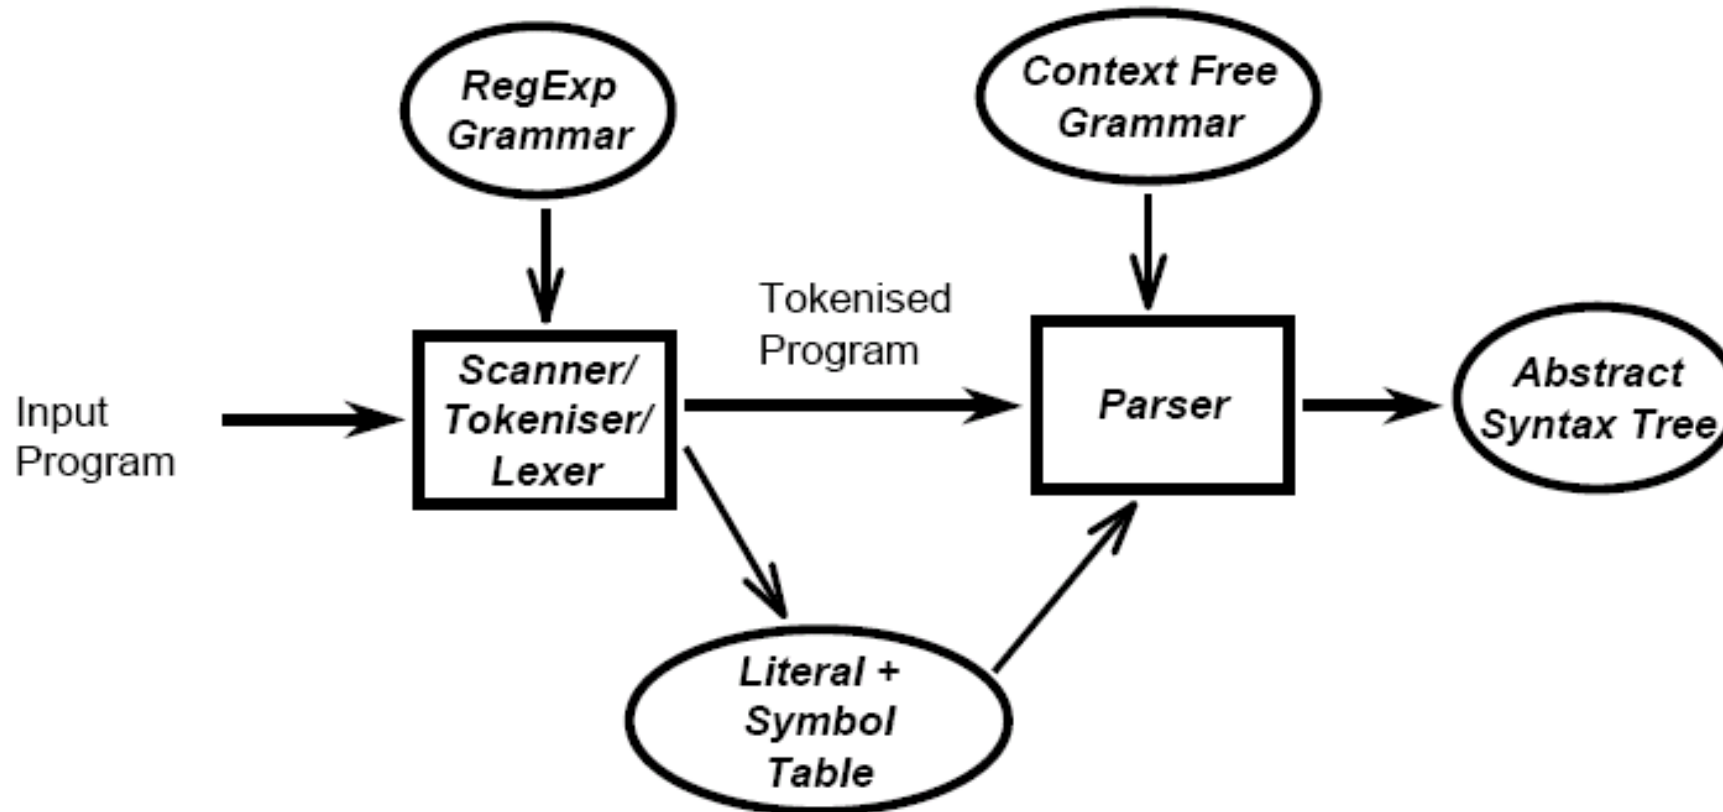
- 3 stages before SQL is really executed
 - Parse tree → Logical Query Plan (LQP)/RA tree
 - Optimize LQP: Select optimal plan
 - ✓ Algebraic Optimization
 - LQP → Physical Query Plan (PQP)
 - ✓ Physical optimization
 - Execution of PQP
 - ✓ Materialize or pipeline

□ Advanced project in reading PostgreSQL

- Parse and Execute SQL in PostgreSQL

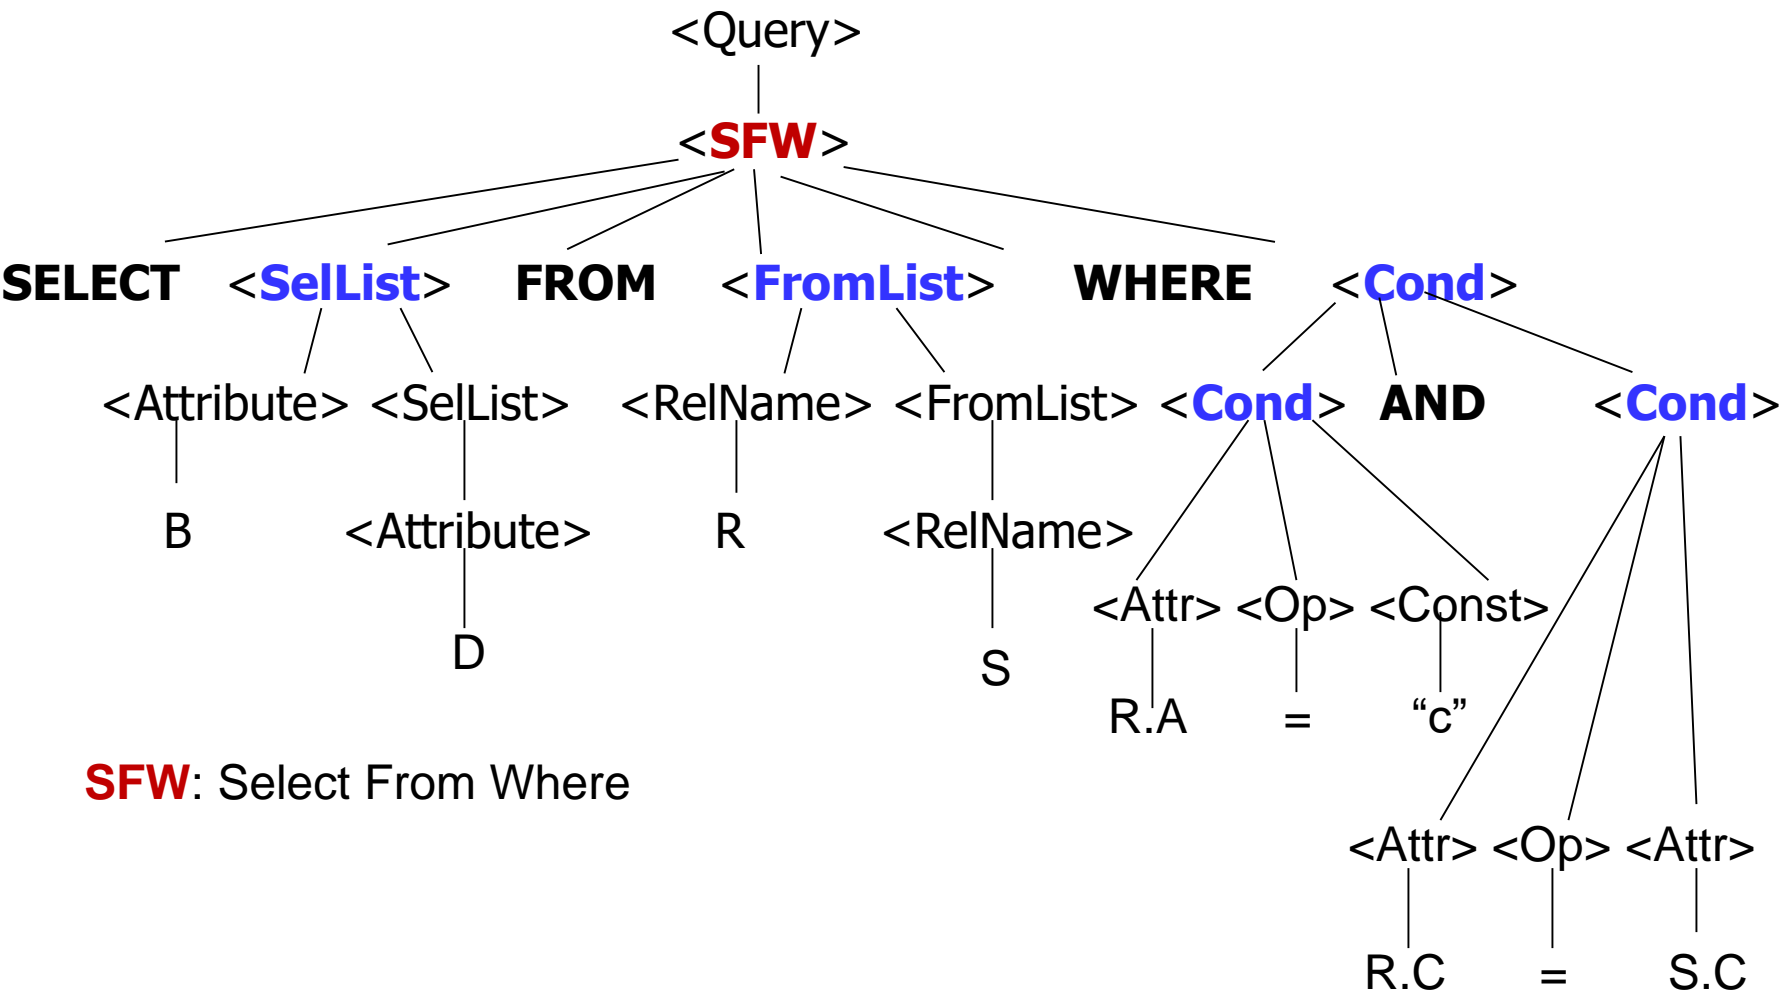
Parse Tree

Select B,D
From R,S
Where $R.A = "c" \wedge R.C = S.C$



Parse Tree

Select B,D
From R,S
Where R.A = "c" \wedge R.C=S.C



SFW: Select From Where



Relational Algebra [关系代数] is needed to carry out SQL's translation



The most important motivation for the research work that resulted in the relational model was the objective of providing a sharp and clear boundary between the logical and physical aspects of database management.

(E. F. Codd)

Enrolled

cid	grade	sid
Carnatic101	C	53666
Reggae203	B	53666
Topology112	A	53650
History105	B	53666

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

❑ Pre-Relational:

write a program

❑ Relational SQL

Select name, cid from students s, enrolled e where s.sid = e.sid

❑ Relational Algebra

$$\pi_{name, cid} (Students \bowtie_{sid} Enrolled)$$

❑ Relational Calculus

$$\{R \mid R.name = S.name \wedge R.cid = S.cid \wedge \exists S(S \in Students \wedge \exists E(E \in Enrolled \wedge E.sid = S.sid))\}$$



With RA (Relational Algebra)

□ Three issues are concerned to carry out SQL's translation into File operations

■ SQL → Parse Tree

➤ We have learned before

■ Parse Tree → Relational Algebra expression (RA)/Logical Query plan (LQP)

➤ Optimizing will be sketched

■ LQP → Physical Query Plan (PQP)

➤ Map RA into corresponding functions [File Operations] – iterators?

✓ TableScan(a), JoinScan(a, b), ...

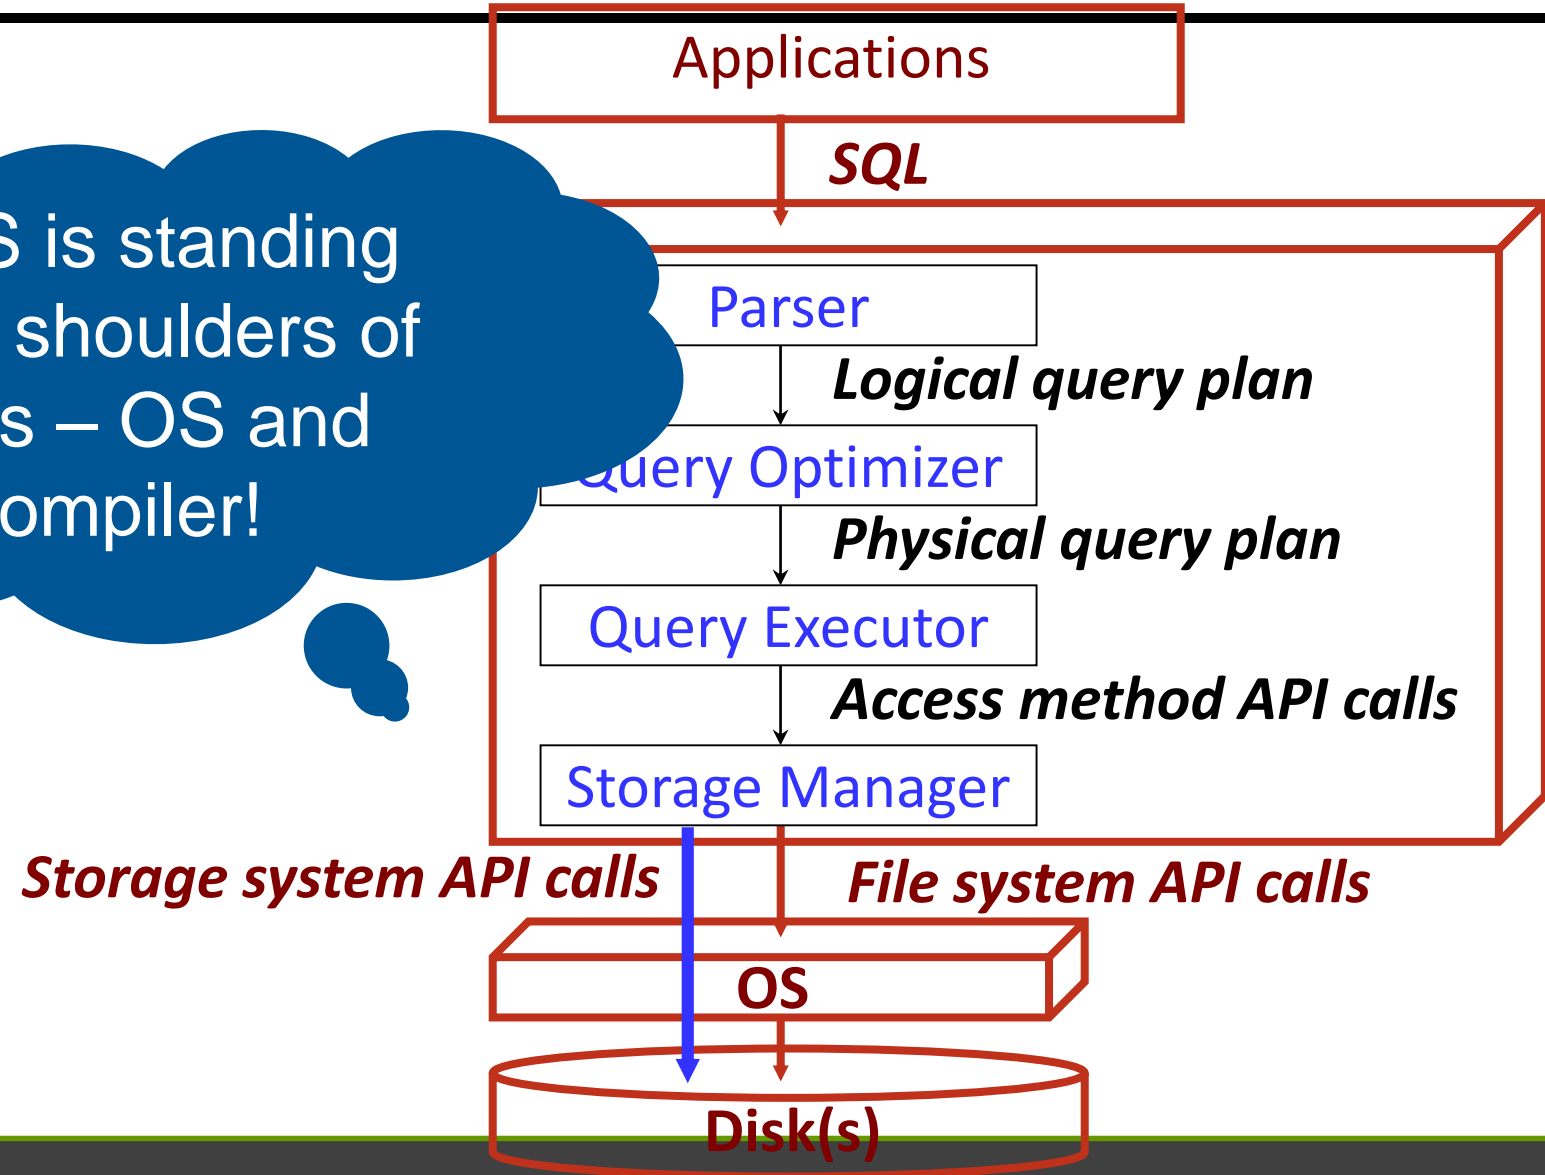
■ Executing

➤ Derive operation sequence and execute



Leave SQL execution later

DBMS is standing
on the shoulders of
giants – OS and
Compiler!



Chapter 3 B: How to get File operations for SQL really

❑ My understanding about DBMS

- 3 stages before SQL is really executed
 - Parse tree → Logical Query Plan (LQP)/RA tree
 - Optimize LQP: Select optimal plan
 - ✓ Algebraic Optimization
 - LQP → Physical Query Plan (PQP)
 - ✓ Physical optimization
 - Execution of PQP
 - ✓ Materialize or pipeline

❑ Advanced project in reading PostgreSQL

- Parse and Execute SQL in PostgreSQL

□ Mapping from an SQL query to an RA expression requires:

- a collection of **templates** for particular kinds of queries
- a matching process to ...
 - determine what kind of query we have (i.e. choose a template)
 - bind components of actual query to slots in the template
- mapping rules to ...
 - convert the matched query into relational algebra
 - filling slots in RA expression from matched components

□ May need to use multiple templates to map whole SQL statement.

Parse Tree to Relational Algebra

- This first rule allows a simple parse tree to be transformed into relational algebra
 - A **<Query>** with a **<Condition>** without sub-queries is replaced by a relational algebra expression
- The relational algebra expression consists of
 - The **product** (\times) of all the relations in the **<FromList>**, which is an argument to
 - A **selection** σ_C where C is the **<Condition>**, which is an argument to
 - A **projection** π_L where L consists of the attributes in the **<SelList>**
- There are still many other rules ... your turn!



□ How the initial tree is built:

- Apply Cartesian cross for relations in the **FROM** clause
- Then the conditions (join & selections) from **WHERE** clause are added
- Then projections from the **SELECT** clause

□ Canonical Query Translation

Canonical translation of SQL queries into algebra expressions.

Structure:

```
select distinct  $a_1, \dots, a_n$   
from            $R_1, \dots, R_k$   
where           $p$ 
```



(Simplified) Canonical translation of SQL to algebra

1. Let $R_1 \dots R_k$ be the entries in the **from** clause of the query. Construct the expression

$$F := \begin{cases} R_1 & \text{if } k = 1 \\ ((\dots (R_1 \times R_2) \times \dots) \times R_k) & \text{else} \end{cases}$$

2. The **where** clause is optional in SQL. Therefore, we distinguish the cases that there is no **where** clause and that the **where** clause exists and contains a predicate p . Construct the expression

$$W := \begin{cases} F & \text{if there is no **where** clause} \\ \sigma_p(F) & \text{if the **where** clause contains } p \end{cases}$$


3. Let s be the content of the **select distinct** clause. For the canonical translation it must be of either '*' or a list a_1, \dots, a_n of attribute names. Construct the expression

$$S := \begin{cases} W & \text{if } s = '*' \\ \Pi_{a_1, \dots, a_n}(W) & \text{if } s = a_1, \dots, a_n \end{cases}$$

4. Return S .




Initial Logical Plan



Cartesian product of R and S

Select B,D
From R,S
Where R.A = "c" \wedge
R.C=S.C



Relational Algebra: $\Pi_{B,D} [\sigma_{R.A=\text{"c"} \wedge R.C = S.C} (R \times S)]$

Chapter 3 B: How to get File operations for SQL really

❑ My understanding about DBMS

- 3 stages before SQL is really executed
 - Parse tree → Logical Query Plan (LQP)/RA tree
 - Optimize LQP: Select optimal plan
 - ✓ Algebraic Optimization
 - LQP → Physical Query Plan (PQP)
 - ✓ Physical optimization

❑ Advanced project in reading PostgreSQL

- Parse and Execute SQL in PostgreSQL

One SQL could be represented as different logical plans

□ Optimization of Query trees

■ Making the initial tree more efficient

- Many different relational algebra expressions correspond to the same query
- The trick is to take the initial tree and convert it into an equivalent, more efficient tree

■ There are rules to help do this

- Basically, we want to move select and project as far down the tree as possible and still maintain equivalence
 - ✓ This is because they limit the size of the relation

□ Expression Rewriting Rules

- Since RA is a **well-defined formal system**, there exist many algebraic **laws** on RA expressions, which can be used as a basis for expression rewriting in order to produce **equivalent** (more-efficient) **expressions**

□ Expression transformation based on such rules can be used

- to simplify/improve SQL→RA mapping results
- to generate new plan variations during query optimization

Relational Algebra Laws

Commutative and Associative Laws:

- $R \times S \leftrightarrow S \times R$
- $(R \times S) \times T \leftrightarrow R \times (S \times T)$
- $R \bowtie S \leftrightarrow S \bowtie R$
- $(R \bowtie S) \bowtie T \leftrightarrow R \bowtie (S \bowtie T)$ (natural join)
- $R \cup S \leftrightarrow S \cup R$
- $(R \cup S) \cup T \leftrightarrow R \cup (S \cup T)$
- $R \cap S \leftrightarrow S \cap R$
- $(R \cap S) \cap T \leftrightarrow R \cap (S \cap T)$
- $R \bowtie_{\text{Cond}} S \leftrightarrow S \bowtie_{\text{Cond}} R$ (theta join)

But it is **not** true in general that

- $(R \bowtie_{\text{Cond1}} S) \bowtie_{\text{Cond2}} T \leftrightarrow R \bowtie_{\text{Cond1}} (S \bowtie_{\text{Cond2}} T)$

Example: $R(a, b), \quad S(b, c), \quad T(c, d), \quad (R \text{ Join}_{[R.b=S.b]} S) \text{ Join}_{[a=d]} T$

Cannot rewrite as $R \text{ Join}_{[R.b=S.b]} (S \text{ Join}_{[a=d]} T)$ because neither $S.a$ nor $T.a$ exists.



... Relational Algebra Laws

Selection commutativity (where c is a condition):

$$\sigma_c (\sigma_d (R)) \iff \sigma_d (\sigma_c (R))$$

Selection splitting (where c and d are conditions):

$$\begin{aligned} \sigma_c \wedge d (R) &\iff \sigma_c (\sigma_d (R)) \\ \sigma_c \vee d (R) &\iff \sigma_c (R) \cup \sigma_d (R) \quad (\text{but only if } R \text{ is a set}) \end{aligned}$$

Selection pushing ($\sigma_c (R \text{ op } S)$):

$$\begin{aligned} \sigma_c (R \cup S) &\iff \sigma_c R \cup \sigma_c S \\ &\quad (\text{must be pushed into both arguments of union}) \\ \sigma_c (R - S) &\iff \sigma_c R - S, \quad \sigma_c (R - S) \iff \sigma_c R - \sigma_c S \\ &\quad (\text{must be pushed into left branch of difference, may be pushed to right branch}) \end{aligned}$$



□ Algebraic/Logical Optimization [代数/逻辑优化]

■ Make use of algebraic equivalences:

- examine query expression
- search for applicable transformation rules (heuristics)
- generate equivalent (and "better") algebraic expressions

■ Most commonly used heuristic:

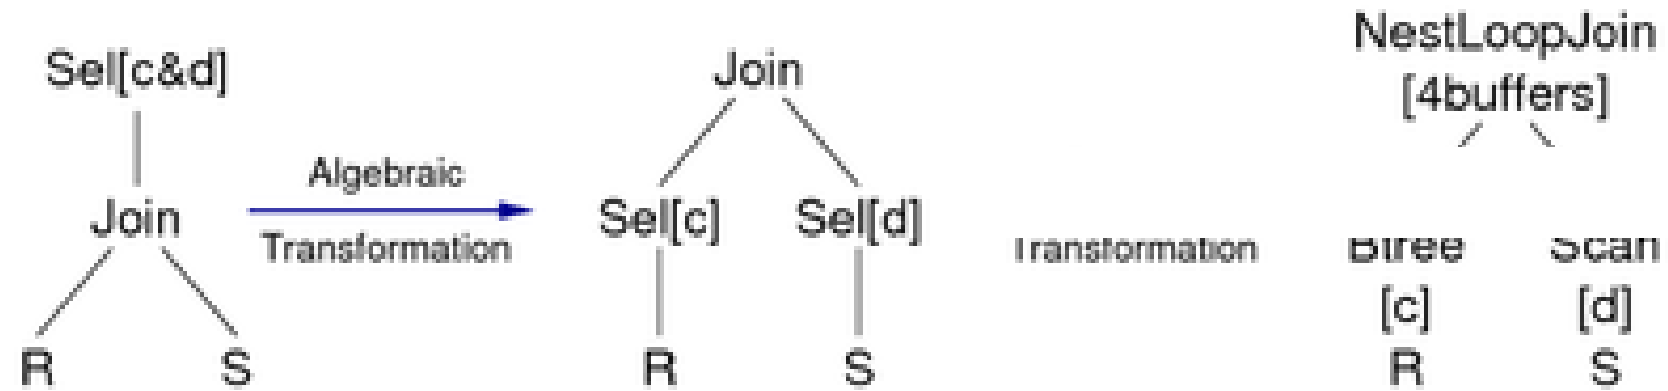
- Apply **Select and Project before Join**

■ Rationale: minimizes size of intermediate relations.

- Can potentially be done during the SQL→RA mapping phase.

■ Algebraic optimization cannot assist with finding good join order.

Example of optimisation transformations:



For join, may also consider sort/merge join and hash join.

HEURISTIC-BASED OPTIMIZATION


Define static rules that transform logical operators to a physical plan.

- Perform most restrictive selection early
- Perform all selections before joins
- Predicate/Limit/Projection pushdowns
- Join ordering based on cardinality



Stonebraker

Example: Original versions of INGRES
(until mid 1990s)



HEURISTICS + COST-BASED JOIN SEARCH

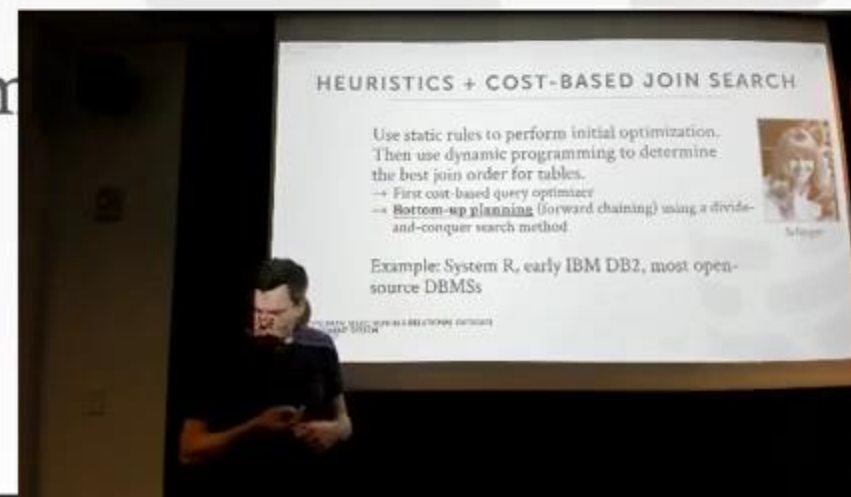
Use static rules to perform initial optimization.
Then use dynamic programming to determine the best join order for tables.

- First cost-based query optimizer
- **Bottom-up planning** (forward chaining) using a divide-and-conquer search method



Selinger

Example: System R, early IBM DB2, most open-source DBMSs



Chapter 3 B: How to get File operations for SQL really

□ My understanding about DBMS

- 3 stages before SQL is really executed
 - Parse tree → Logical Query Plan (LQP)/RA tree
 - Optimize LQP: Select optimal plan
 - ✓ Algebraic Optimization
 - LQP → Physical Query Plan (PQP)
 - ✓ Physical optimization
 - Execution of PQP
 - ✓ Materialize or pipeline

□ Advanced project in reading PostgreSQL

- Parse and Execute SQL in PostgreSQL


LQP→PQP

□ Here we have another conversion/translation from **RA/LQP** to PQP

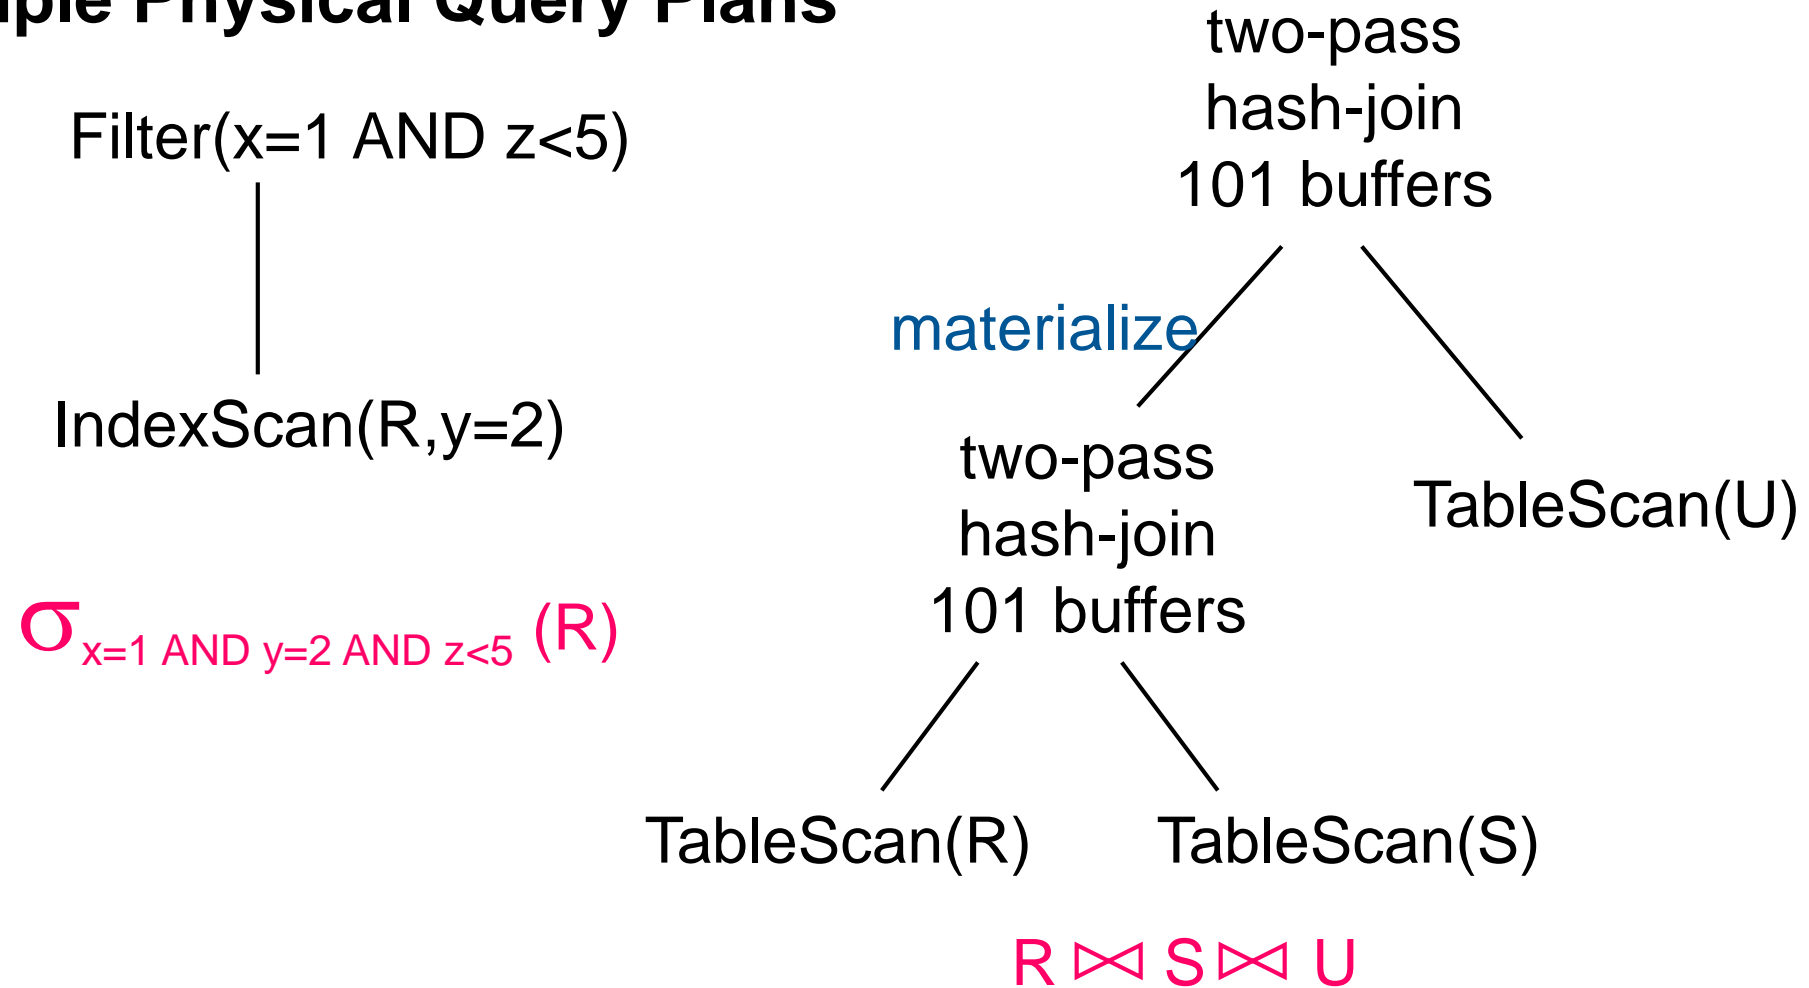
- Mapping Relational operators (Select, Product, Project, ...) to **file functions** (TableScan(...), Join(...), Btree(...), ...)

■ Rules

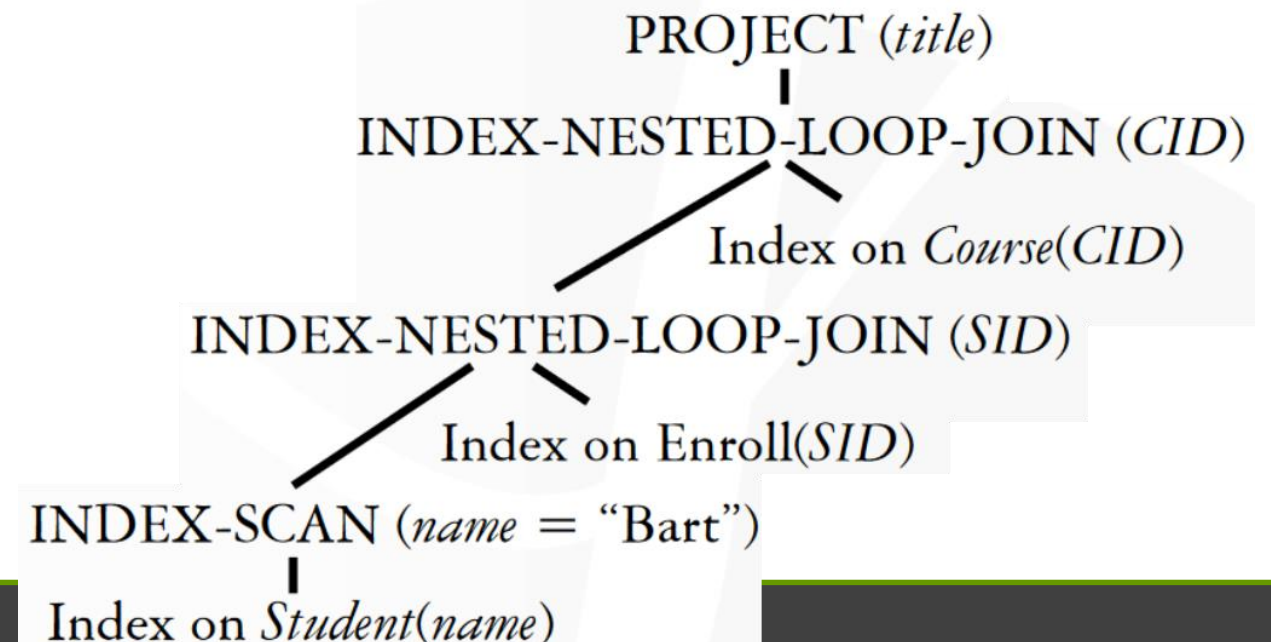
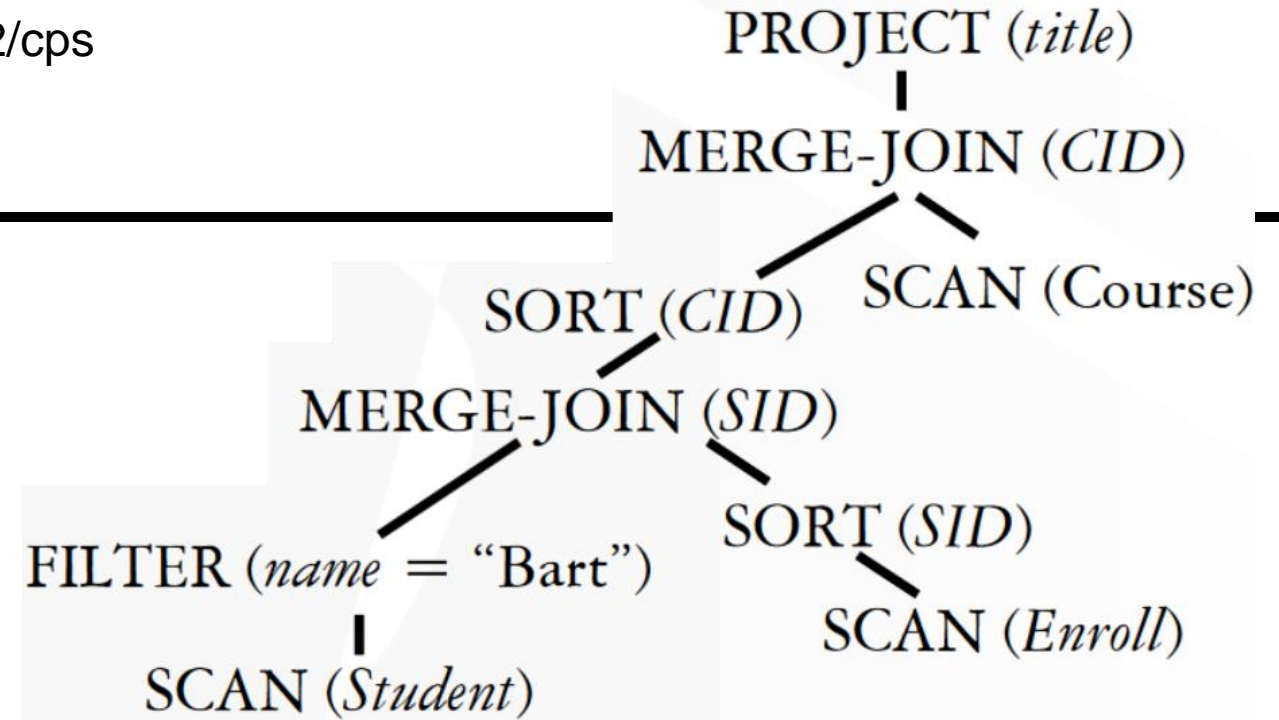
- leaves of LQP (stored relations) become scan operators
- internal nodes of LQP (operators) become one or more physical operations (algorithms)
- edges of LQP are marked as "pipeline" or "materialize"



□ Example Physical Query Plans



- Node are physical operators that implement particular algorithms (e.g., scanning, sorting, hashing, ...)
- There are even more equivalent physical plans
 - Even a single logical plan can have different physical plans



□ Physical Optimization [物理优化]

- Makes use of **execution cost analysis [based on statistics]**:
 - examine query evaluation plan
 - determine efficient join sequences
 - select access method for each operation (e.g. index for select)
 - for distributed DB, select best sites (closest, best bandwidth)
 - determine total cost for evaluation plan
 - repeat for all possible plans and choose best
- Physical optimization is also called query evaluation plan generation.

□ Then physical optimization is carried out

- Among many pre-optimal PQPs
- Based on catalog information

Tables

name	file	#tuples	size
Tables	...	6	1
Attributes	...	23	1
Views	...	1	1
Indexes	...	0	1
Sailors	...	40,000	500
Reserves	...	100,000	1000

Attributes

name	table	type	pos
name	Tables	string	1
file	Tables	string	2
#tuples	Tables	integer	3
size	Tables	integer	4
name	Attributes	string	1
table	Attributes	string	2
type	Attributes	string	3
pos	Attributes	integer	4
⋮	⋮	⋮	⋮
sid	Sailors	integer	1
sname	Sailors	string	2
rating	Sailors	integer	3
age	Sailors	real	4
sid	Reserves	integer	1
bid	Reserves	integer	2
day	Reserves	date	3
rname	Reserves	string	4

Views

name	text
Captains	SELECT * FROM Sailors WHERE...

Indexes

name	file	type	#keys	size
Boats	...	B+Tree	100	1

Query Evaluation Example

Example database schema (multi-campus university):

Employee(**eid**, ename, status, city)

Department(**dname**, city, address)

Subject(**sid**, sname, syllabus)

Lecture(**subj**, **dept**, **empl**, time)

Query: Which depts in Sydney offer database subjects?

Example database instance statistics:

```
select dname
from   Department D, Subject S, Lecture L
where  D.city = 'Sydney' and S.sname like '%Database%'
       and D.dname = L.dept and S.sid = L.subj
```

Relation	r	R	C	b
Employee	1000	100	10	100
Department	100	200	5	50
Subject	500	95	10	100
Lecture	2000	100	10	200

Additional information (needed to determine query costs):

- 5 departments are located in Sydney
- 80 subjects are about databases
- 300 lectures are on databases
- 100 lectures are in Sydney
- 3 of these are on databases

... Query Evaluation Example

Strategy #1 for answering query:

1. *TMP1* \leftarrow *Subject* \times *Lecture*
2. *TMP2* \leftarrow *TMP1* \times *Department*
3. *TMP3* \leftarrow *Select[check](TMP2)*

check = (*city*='Sydney' & *sname*='Databases' & *dname*=*dept* & *sid*=*subj*)

4. *RESULT* \leftarrow *Project[dname](TMP3)*

Costs involved in using strategy #1:

Reln	<i>r</i>	<i>R</i>	<i>C</i>	<i>b</i>
TMP1	1000000	195	5	20000
TMP2	100000000	395	2	50000000
TMP3	3	395	2	2
RESULT	3	100	10	1

Total I/Os

= Cost_{Step1} + Cost_{Step2} + Cost_{Step3} + Cost_{Step4}

= (100+100*200+ 20000 + (20+20*20000+ 5*10⁷ + 5*10⁷)+2) + 2

= 100,440,124

Relation	<i>r</i>	<i>R</i>	<i>C</i>	<i>b</i>
Employee	1000	100	10	100
Department	100	200	5	20
Subject	500	95	10	100
Lecture	2000	100	10	200



Strategy #2 for answering query:

-
1. *TMP1* \leftarrow *Join[sid=subj](Subject, Lecture)*
 2. *TMP2* \leftarrow *Join[dept=dname](TMP1, Department)*
 3. *TMP3* \leftarrow *Select[city='Sydney' & sname='Databases'](TMP2)*
 4. *RESULT* \leftarrow *Project[dname](TMP3)*

Costs involved in using strategy #2:

Reln	<i>r</i>	<i>R</i>	<i>C</i>	<i>b</i>
TMP1	2000	195	5	400
TMP2	2000	395	2	1000
TMP3	3	395	2	2
RESULT	3	100	10	1

Total I/Os

$$\begin{aligned} &= \text{CostStep1} + \text{CostStep2} + \text{CostStep3} + \text{CostStep4} \\ &= (100 + 100 * 200 + 400) + (20 + 20 * 400 + 1000) + (1000 + 2) + 2 \\ &= 30,524 \end{aligned}$$



Strategy #3 for answering query:

-
1. *TMP1* \leftarrow *Join*[*sid=subj*](*Subject*, *Lecture*)
 2. *TMP2* \leftarrow *Select*[*sname='Databases'*](*TMP1*)
 3. *TMP3* \leftarrow *Join*[*dept=dname*](*TMP2*, *Department*)
 4. *TMP4* \leftarrow *Select*[*city='Sydney'*](*TMP3*)
 5. *RESULT* \leftarrow *Project*[*dname*](*TMP4*)

Costs involved in using strategy #3:

Reln	<i>r</i>	<i>R</i>	<i>C</i>	<i>b</i>
TMP1	2000	195	5	400
TMP2	20	195	5	4
TMP3	20	395	2	10
TMP4	3	395	2	2
RESULT	3	100	10	1

Total I/Os

$$\begin{aligned} &= \text{CostStep1} + \text{CostStep2} + \text{CostStep3} + \text{CostStep4} + \text{CostStep5} \\ &= (100+100*200+400) + (400+4) + (4+4*20+10) + (10+2) + 1 \\ &= 21,011 \end{aligned}$$



Strategy #4 for answering query:

-
- | | | | |
|----|---------------|--------------|---|
| 1. | <i>TMP1</i> | \leftarrow | <i>Select[sname='Databases'](Subject)</i> |
| 2. | <i>TMP2</i> | \leftarrow | <i>Select[city='Sydney'](Department)</i> |
| 3. | <i>TMP3</i> | \leftarrow | <i>Join[sid=subj](TMP1, Lecture)</i> |
| 4. | <i>TMP4</i> | \leftarrow | <i>Join[dept=dname](TMP3, TMP2)</i> |
| 5. | <i>RESULT</i> | \leftarrow | <i>project[dname](TMP4)</i> |
-

Costs involved in using strategy #4:

Reln	<i>r</i>	<i>R</i>	<i>C</i>	<i>b</i>
TMP1	80	95	5	16
TMP2	5	200	5	1
TMP3	300	195	5	60
TMP4	3	395	2	2
RESULT	3	100	10	1

Total I/Os

$$\begin{aligned} &= \text{Cost}_{\text{Step1}} + \text{Cost}_{\text{Step2}} + \text{Cost}_{\text{Step3}} + \text{Cost}_{\text{Step4}} + \text{Cost}_{\text{Step5}} \\ &= (100+16) + (20+1) + (16+16*200+60) + (1+1*60+2) + 2 \\ &= 3478 \end{aligned}$$



Strategy #5 for answering query:

1. $TMP1 \leftarrow \text{Select}[sname='Databases'](\text{Subject})$
2. $TMP2 \leftarrow \text{Select}[city='Sydney'](\text{Department})$
3. $TMP3 \leftarrow \text{Join}[dept=dname](TMP2, \text{Lecture})$
4. $TMP4 \leftarrow \text{Join}[sid=subj](TMP3, TMP1)$
5. $RESULT \leftarrow \text{project}[dname](TMP4)$

... Query Evaluation Example

Costs involved in using strategy #5:

Reln	r	R	C	b
TMP1	80	95	5	16
TMP2	5	200	5	1
TMP3	100	295	3	34
TMP4	3	395	2	2
RESULT	3	100	10	1

Total I/Os

$$\begin{aligned} &= \text{Cost}_{\text{Step1}} + \text{Cost}_{\text{Step2}} + \text{Cost}_{\text{Step3}} + \text{Cost}_{\text{Step4}} + \text{Cost}_{\text{Step5}} \\ &= (100+16) + (20+1) + (1+1*200+34) + (16+16*34+2) + 2 \\ &= 936 \end{aligned}$$



Chapter 3 B: How to get File operations for SQL really

❑ My understanding about DBMS

- 3 stages before SQL is really executed
 - Parse tree → Logical Query Plan (LQP)/RA tree
 - Optimize LQP: Select optimal plan
 - ✓ Algebraic Optimization
 - LQP → Physical Query Plan (PQP)
 - ✓ Physical optimization
 - Execution of PQP
 - ✓ Materialize or pipeline

❑ Advanced project in reading PostgreSQL

- Parse and Execute SQL in PostgreSQL

Query Execution: Just like the execution of math expression

- ❑ The query execution engine is like a **virtual machine** that runs physical query plans
 - Either **materialize** or **pipeline**

- ❑ **Two architectural considerations:**
 - Granularity of processing
 - Control flow:
 - Sequential
 - ✓ In-order Traversal, and one after another
 - Or
 - Concurrent
 - ✓ Cut into segments, and execute concurrently

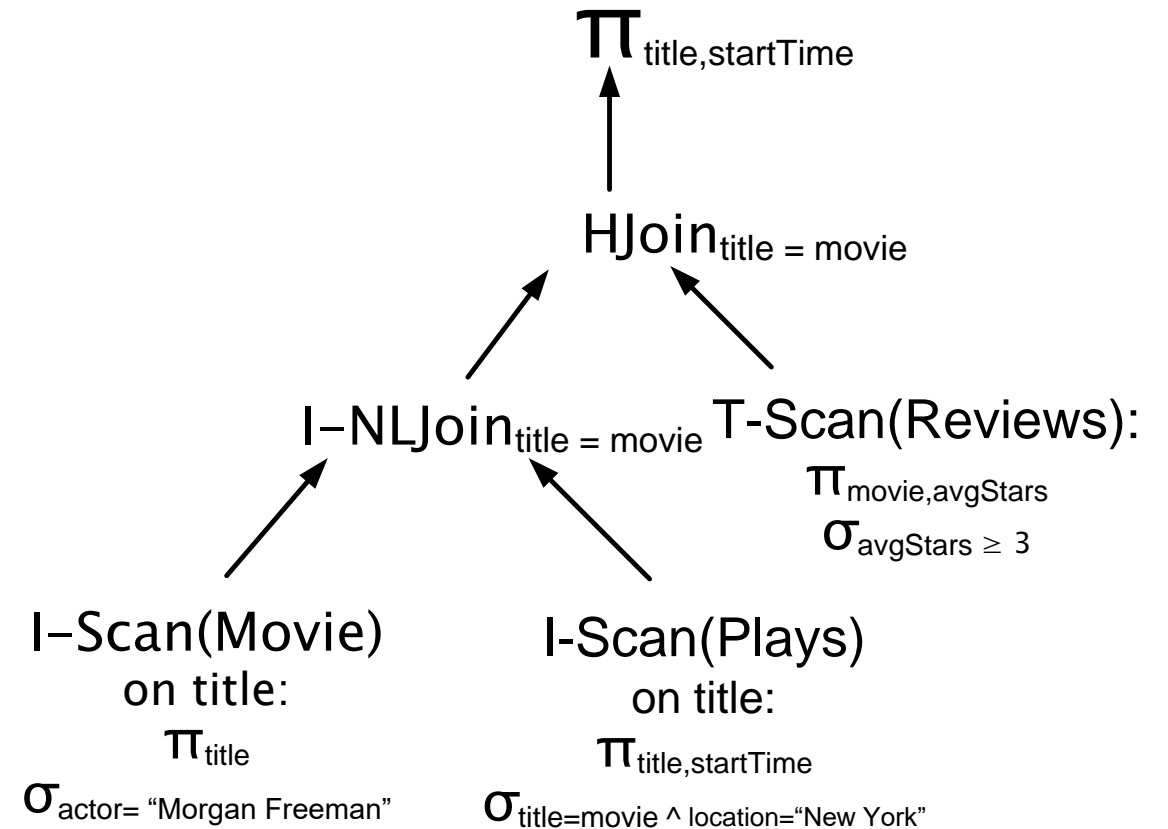
Granularity of Processing

Options for granularity:


- Operators process one **relation** at a time
- Operators process one **tuple** at a time (**pipelining**) – better responsiveness and parallelism

Blocking algorithms can't pipeline output, e.g., GROUP BY over unsorted data

Some operators are partly blocking: e.g., hash join (HJoin) must buffer one input before joining



Operator Plumbing [水管工 – 组装...]



□ Materialization:

- output of one operator written to disk, next operator reads from the disk
- perform operations in series and write intermediate results to disk


□ Pipelining:

- output of one operator directly fed to next operator

Two ways to transfer
intermediate results
among the operations

Materialization

□ Produce intermediate result [中间结果]



https://www.cs.duke.edu/courses/fall08/cps216/Lectures/03_iterators_and_rewrites.ppt



□ Materialization Example:

- select s.name, s.id, e.course, e.mark
- from Student s, Enrolment e
- where e.student = s.id and
- e.semester = '05s2' and s.name = 'John';

□ might be executed as

Temp1 = BtreeSelect[semester=05s2](Enrolment)

Temp2 = **BtreeSelect**[name=John](Student)

-- indexes on name and semester

-- produce sorted Temp1 and Temp2

Temp3 = **SortMergeJoin**[e.student=s.id](Temp1,Temp2)

-- SMJoin especially effective, since

-- Temp1 and Temp2 are already sorted

Result = **Project**[name,id,course,mark](Temp3)

□ Pipelining

■ How pipelining is organized between two operators:

- blocks execute "concurrently" **as producer/consumer** pairs
 - ✓ first operator acts as producer; second as consumer
- structured as interacting **iterators** (open; while(next); close)


■ Advantage:

- no requirement for disk access (results passed via memory buffers)

■ Disadvantage:

- each operator accesses inputs via linear scan

Iterators: Pipelining



→ Each operator supports:

- **Open()**
- **GetNext()**
- **Close()**

→ Parameters are those “List”s

Implementation of single-relation selection iterator:

```
typedef struct {  
    File    inf; // input file  
    Cond    cond; // selection condition  
    Buffer   buf; // buffer holding current page  
    int     curp; // current page during scan  
    int     curr; // index of current record in page  
} Iterator;
```

Iterator structure contains information:

- related to operation being performed (e.g. `cond`)
- information giving current execution state (e.g. `curp`, `curr`)

Implementation of single-relation selection iterator (cont):

```
Iterator *open(char *relName, Condition cond) {
    Iterator *iter = malloc(sizeof(Iterator));
    iter->inf = openFile(fileName(relName), READ);
    iter->cond = cond;
    iter->curp = 0;
    iter->curr = -1;
    readBlock(iter->inf, iter->curp, iter->buf);
    return iter;
}

void close(Iterator *iter) {
    closeFile(iter->inf);
    free(iter);
}
```



Implementation of single-relation selection iterator (cont):

```
Tuple next(Iterator *iter) {
    Tuple rec;
    do {
        // check if reached end of current page
        if (iter->curr == nRecs(iter->buf)-1) {
            // check if reached end of data file
            if (iter->curp == nBlocks(iter->inf)-1)
                return NULL;
            iter->curp++;
            iter->buf = readBlock(iter->inf, iter->curp);
            iter->curr = -1;
        }
        iter->curr++;
        rec = getRec(iter->buf, iter->curr);
    } while (!matches(rec, iter->cond));
    return rec;
}
// curp and curr hold indexes of most recently read page/record
```




Consider the query:

```
select s.id, e.course, e.mark
from   Student s, Enrolment e
where  e.student = s.id and
       e.semester = '05s2' and s.name = 'John' ;
```

which maps to the RA expression

Proj[id, course, mark](Join[student=id](Sel[05s2](Enr), Sel[John](Stu)))

Modelled as communication between RA tree nodes:



□ This query might be executed as

System:

```
iter0 = open(Result)
while (Tup = next(iter0)) { display Tup }
close(iter0)
```

Result:

```
iter1 = open(Join)
while (T = next(iter1))
    { T' = project(T); return T' }
close(iter1)
```

Sel1:

```
iter4 = open(Btree(Enrolment, 'semester=05s2'))
while (A = next(iter4)) { return A }
close(iter4)
```

Join: -- nested-loop join

```
iter2 = open(Sel1)
while (R = next(iter2)) {
    iter3 = open(Sel2)
    while (S = next(iter3))
        { if (matches(R,S) return (RS) }
    close(iter3) // better to reset(iter3)
}
close(iter2)
```

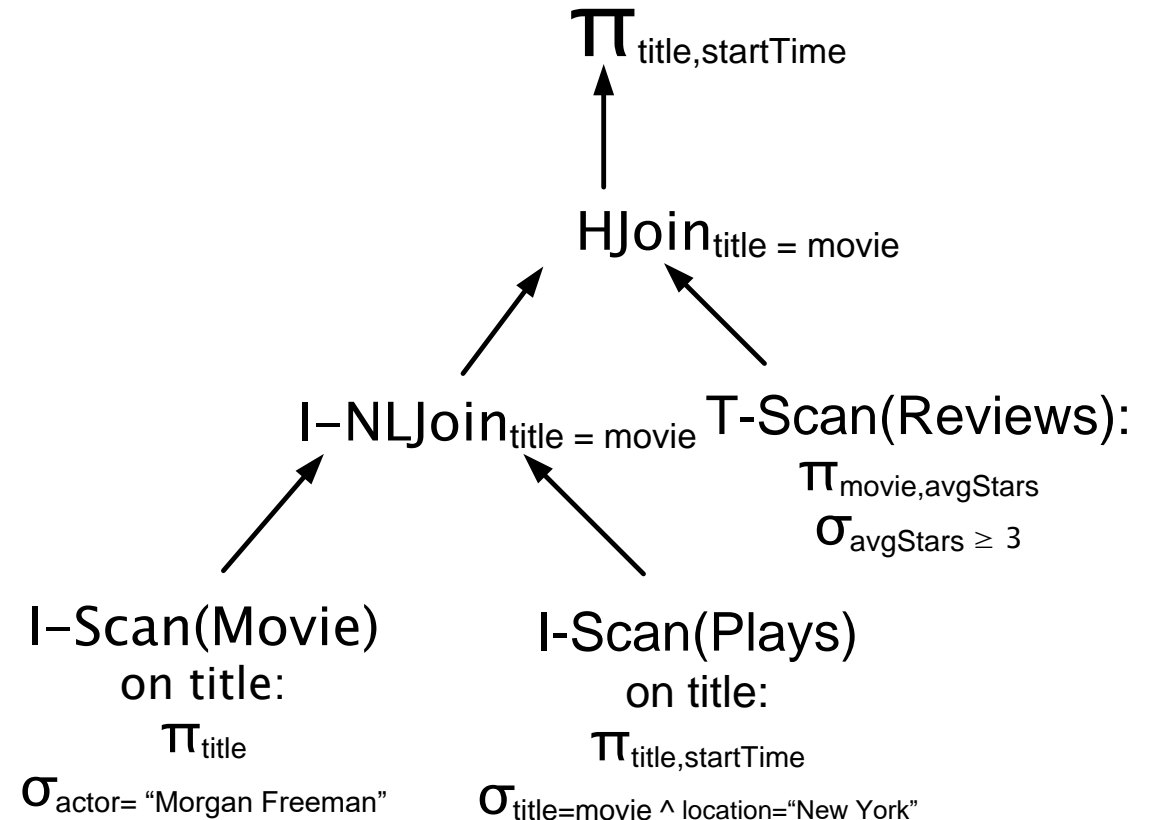
Sel2:

```
iter5 = open(Btree(Student, 'name=John'))
while (B = next(iter5)) { return B }
close(iter5)
```

Control Flow


Options for control flow:

- ❑ **Iterator-based** – execution begins at root
 - *open, next, close*
 - Propagate calls to children
 - May call multiple child *nexts*
- ❑ **Dataflow-driven** – execution driven by tuple arrival at leaves
- ❑ **Hybrids** of the two



Iterator for Select

$\sigma_{R.A = \text{"c"}}$




```
Open() {  
    /** initialize child */  
    Child.Open();  
}  
  
Close() {  
    /** inform child */  
    Child.Close();  
}
```

```
GetNext() {  
    LOOP:  
    t = Child.GetNext();  
    IF (t == EOT) {  
        /** no more tuples */  
        RETURN EOT;  
    }  
    ELSE IF (t.A == "c")  
        RETURN t;  
    ENDLOOP:  
}
```



Iterator for Tuple Nested Loop Join (TNLJ)



□ TNLJ (conceptually)

for each $r \in \text{Lexp}$ do

for each $s \in \text{Rexp}$ do

if $\text{Lexp.C} = \text{Rexp.C}$, output r, s

https://www.cs.duke.edu/courses/fall08/cps216/Lectures/03_iterators_and_rewrites.ppt



In short

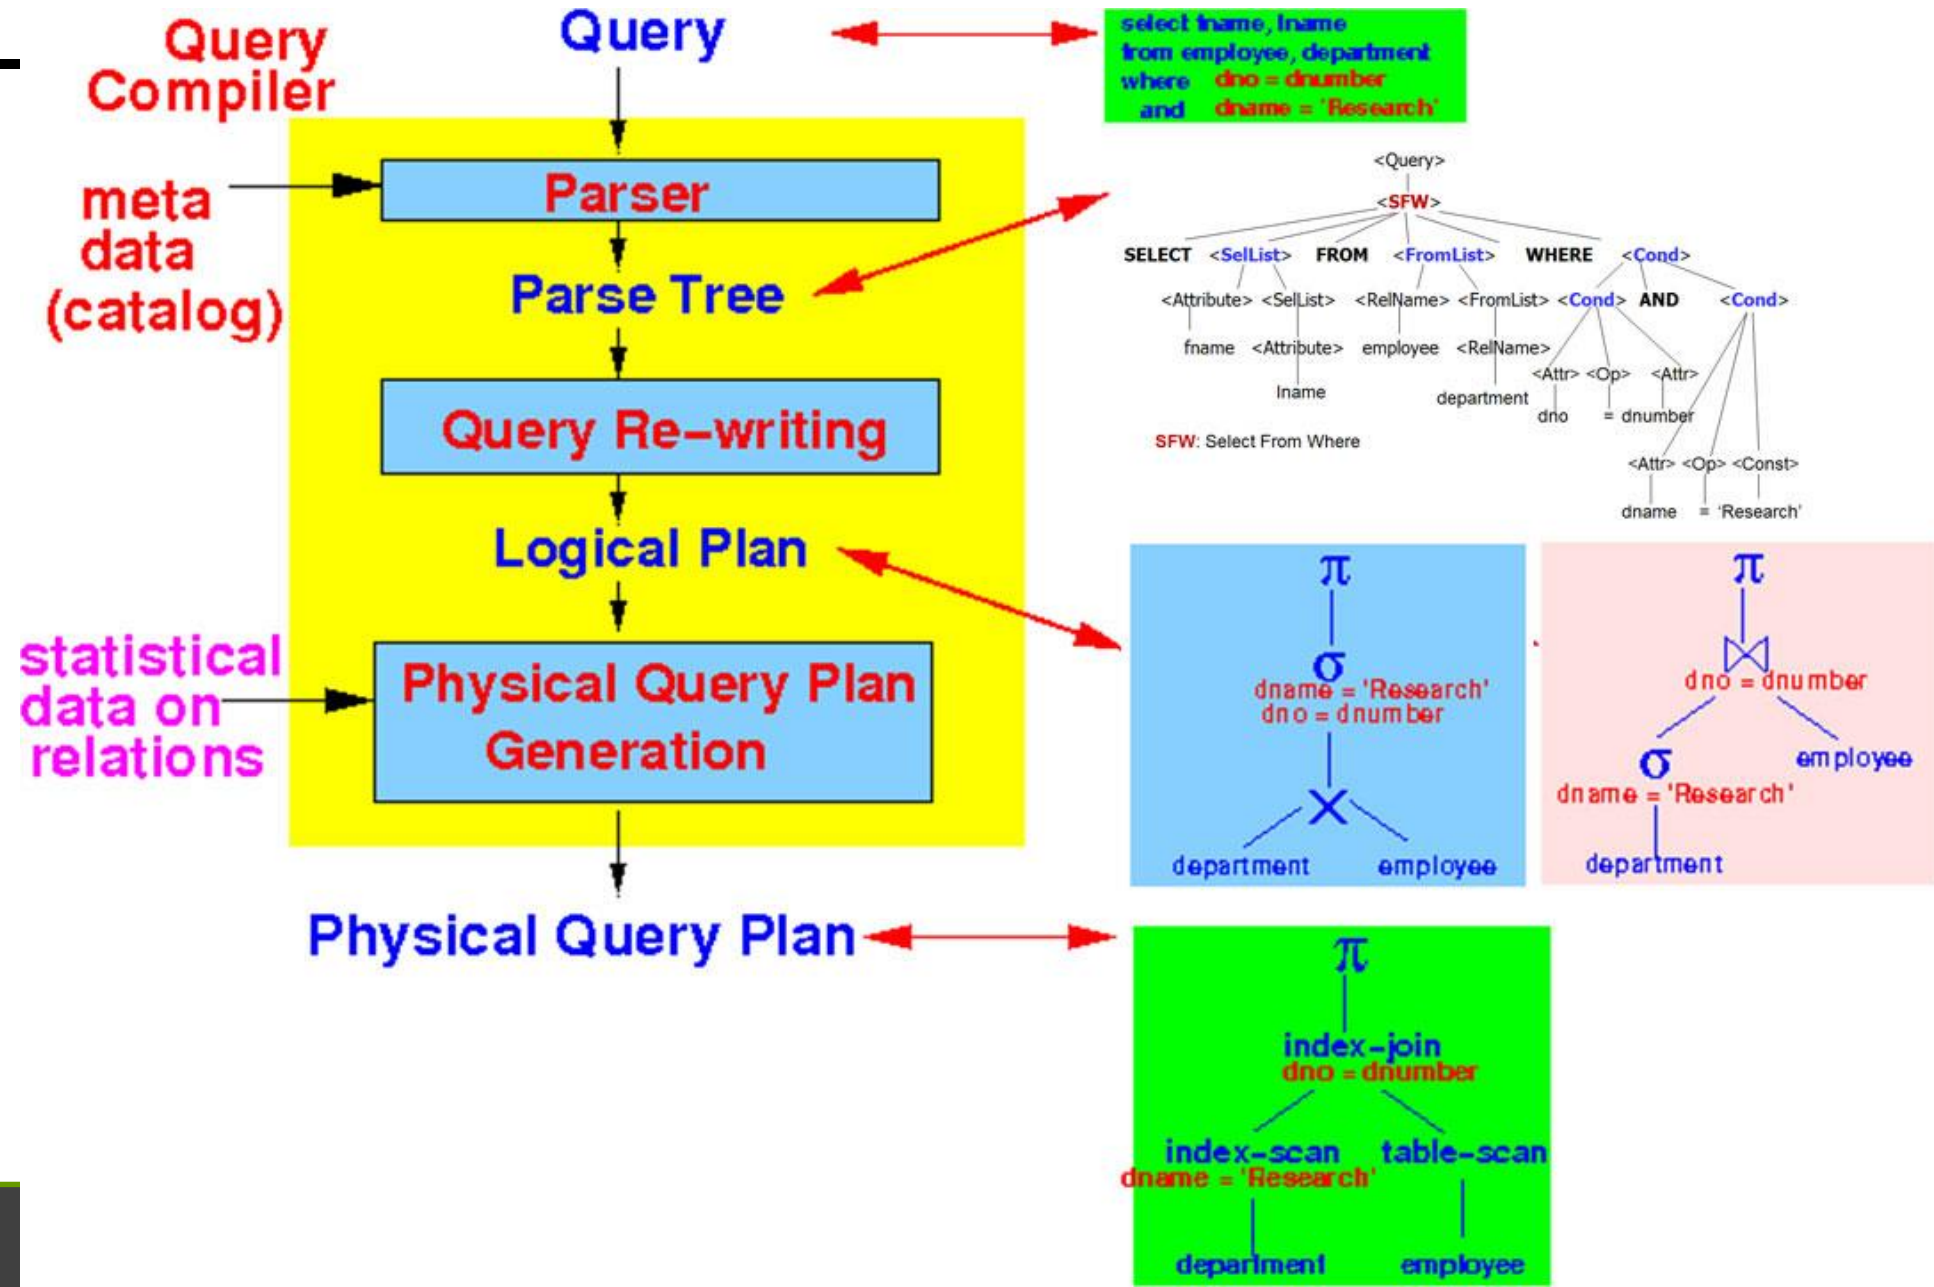
□ Math Expression → AST

- Helpful to understand SQL processing

□ SQL Processing

- SQL → Parse Tree [Math → AST]
- Parse Tree → Canonical RA [Tree Traverse]
 - Canonical RA is converted into equivalent RAs (for Algebraic optimization)
- RA → PQP (Physical Query Plan) [Math → AST]
 - RA could be mapped to many PQPs
 - With help of catalog information, physical optimization
- The optimal PQP is executed [Tree Traverse][Iterator]
 - Materialize or Pipeline

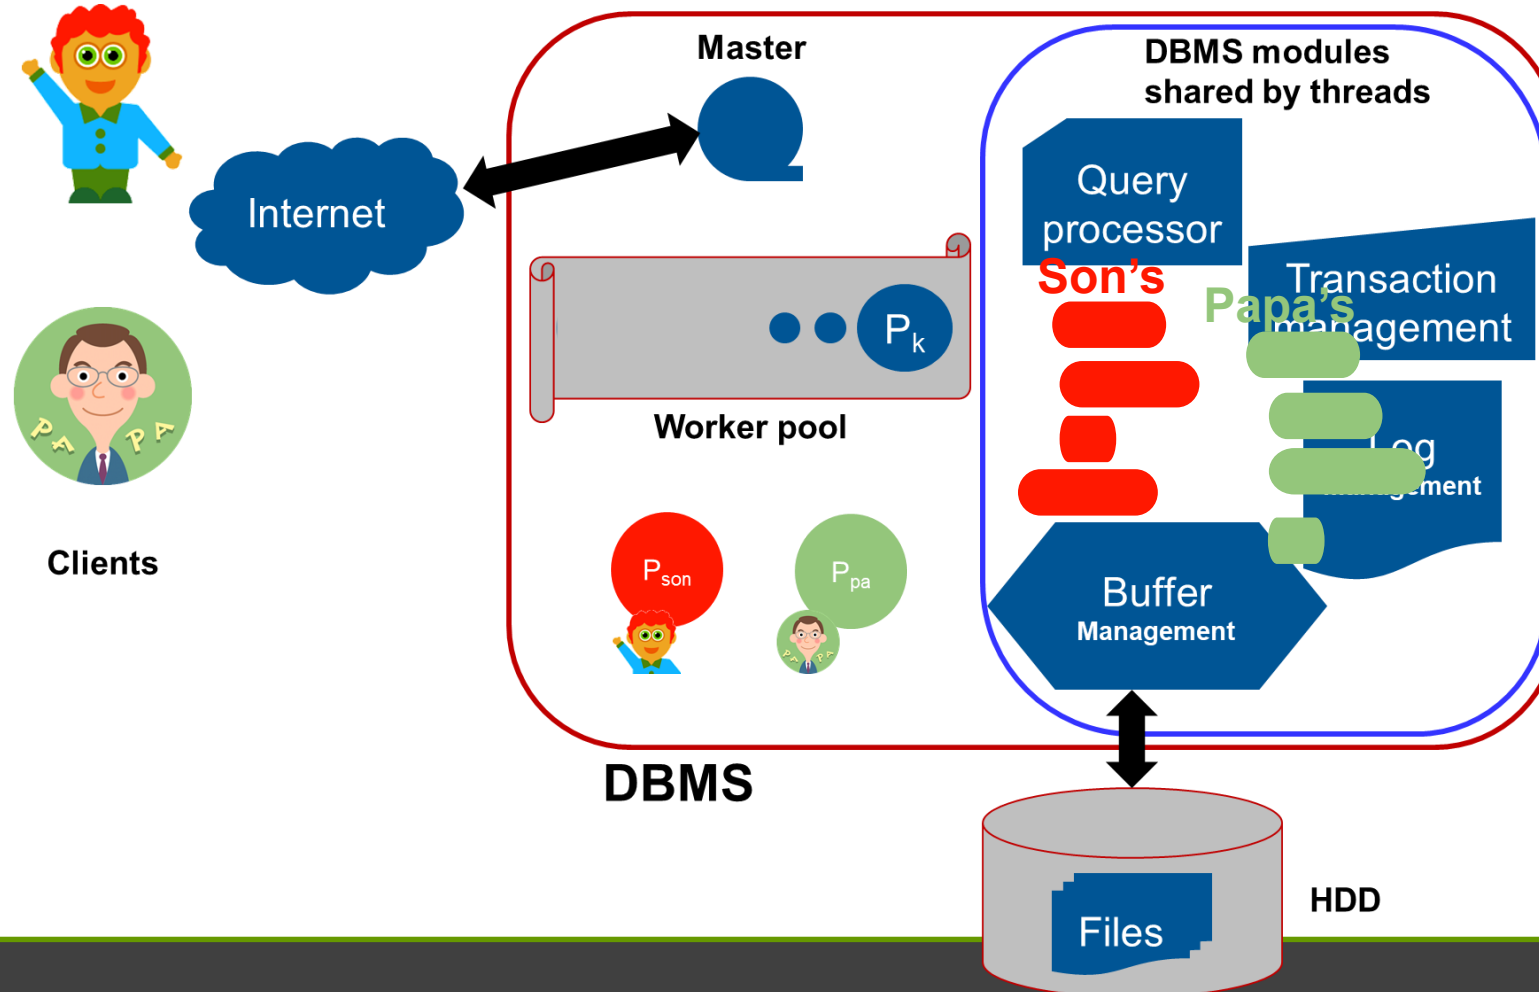
In short



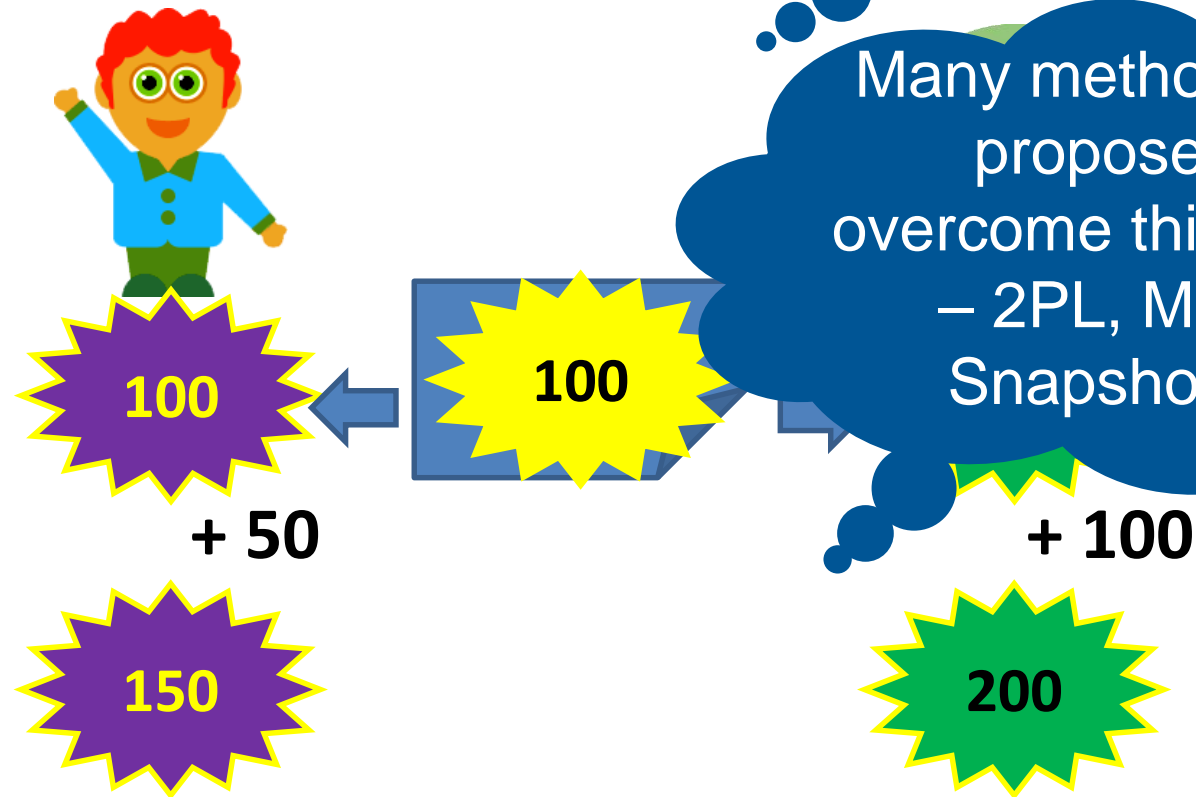
Now we get the corresponding file operations

SQLs → File Operation sequences

- But file operations of SQLs from concurrent clients/users could be interleaved



□ Data Inconsistency!



Chapter 3 B:

How to get File operations for SQL really

□ My understanding about DBMS

- 3 stages before SQL is really executed
 - Parse tree → Logical Query Plan (LQP)/RA tree
 - Optimize LQP: Select optimal plan
 - ✓ Algebraic Optimization
 - LQP → Physical Query Plan (PQP)
 - ✓ Physical optimization
 - Execution of PQP
 - ✓ Materialize or pipeline

□ Advanced project in reading HyperSQL

- Parse and Execute SQL in HyperSQL



□ Distill the workflow of SQL processing with an example

- Major classes/functions
- + Data Structure
- + Relationships of function calls
- + your understanding about the mechanism

Reading codes – Understand

scitools.com <https://scitools.com/>

126网易免费邮--你的 北京交通大学·校园 Ebook-dl | Free Dow 西林街 :: 全能型网盘 难以忽视的真相 2,资 VGG Practical https://libgen.pw

Understand™
Legacy Code Maintenance Visual Static Analysis C++ Editor

GitAhead™
Revision Control

Free Trial Update Buy Free Trial Buy



□ Chapter 60

CHAPMAN & HALL/CRC COMPUTER and INFORMATION SCIENCE SERIES

Handbook of DATA STRUCTURES and APPLICATIONS

