



# Insight into

## DBMS: Design and Implementation

**Chapter 6 A: Go deeper and wider  
File, DBMS, ERP/MRP, DW/DM, Big Data**



孔令波

[mlinking@126.com](mailto:mlinking@126.com)

+86 15010255486

# Outline of the chapters covered

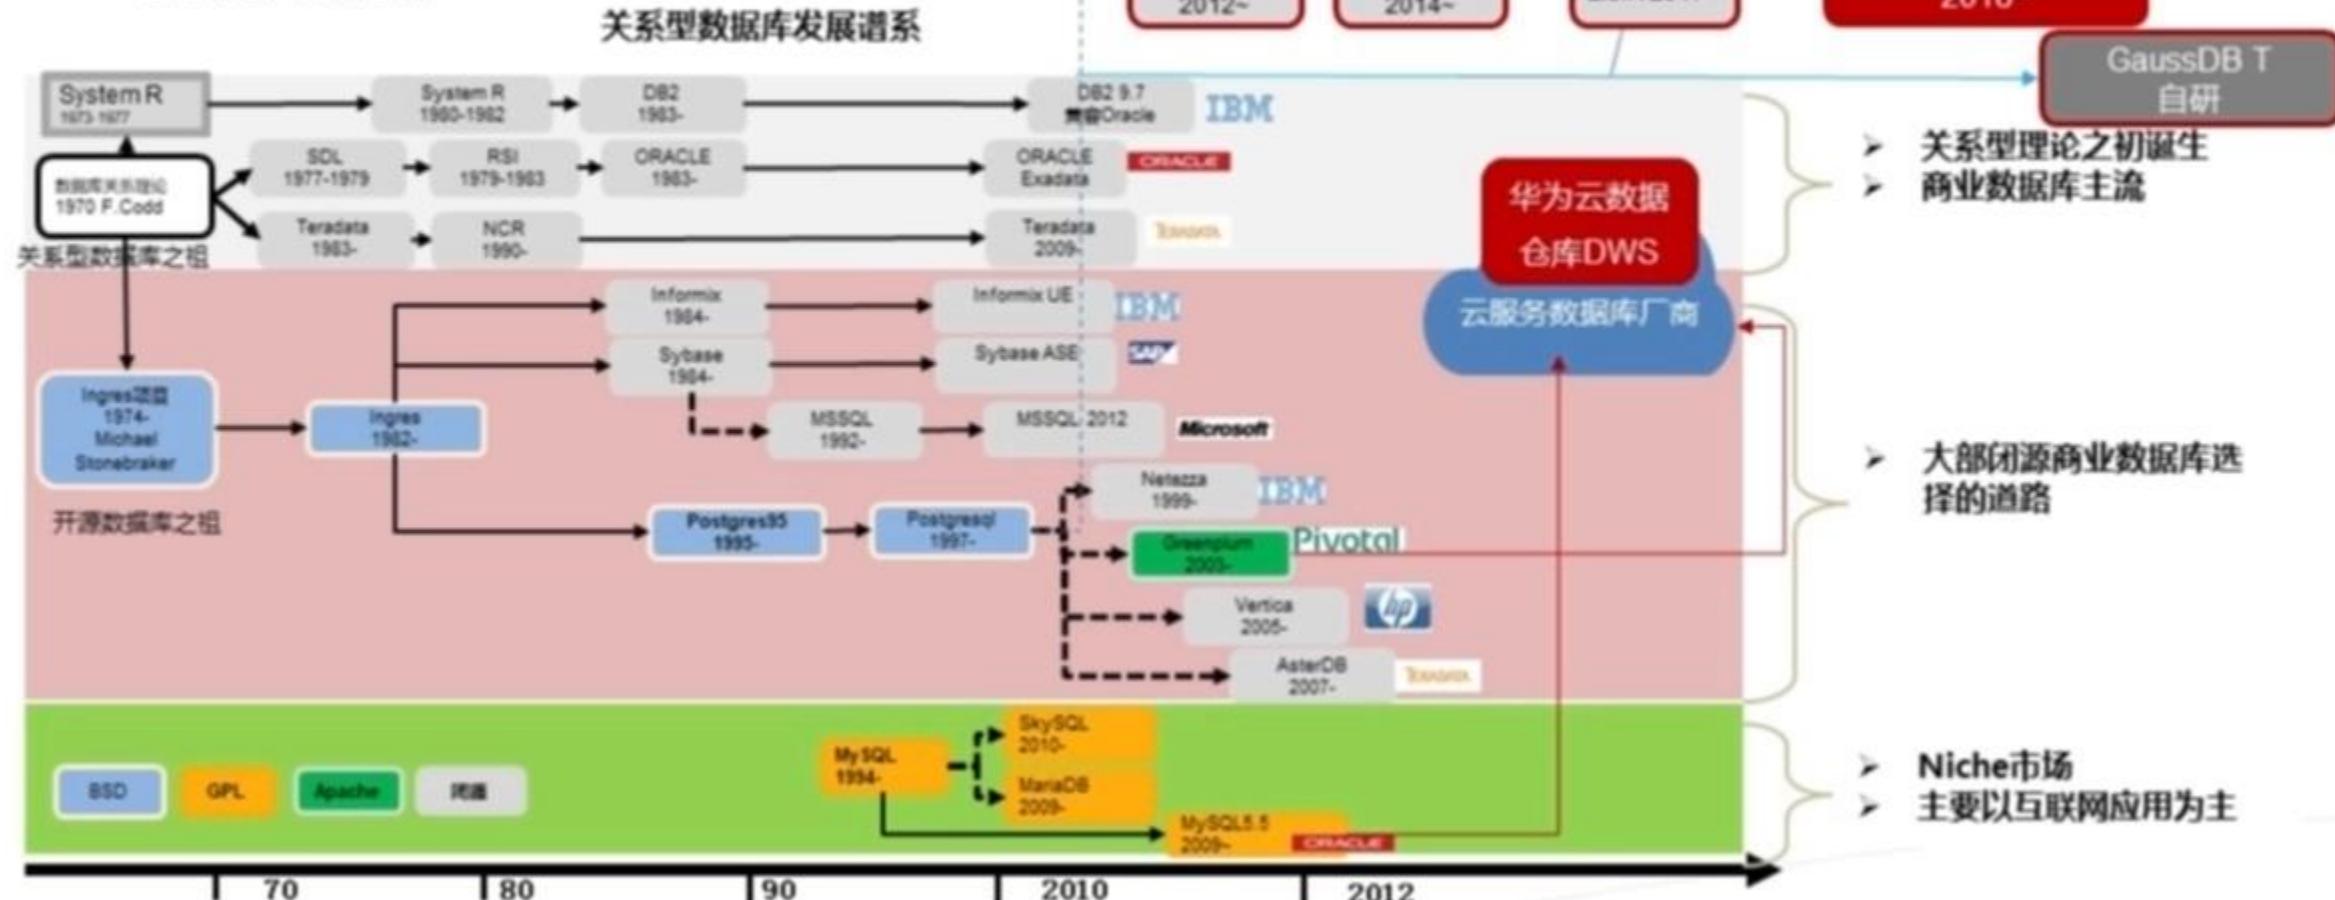
---

- **Introduction**
- **Overview of the projects**
- **Demonstration of Development environment**
  - Watch and practice by yourself
- **My understanding about (R)DBMS**
  - History and D&I
- **SQL translation with 2 conversions**
  - SQL → RA (Relational Algebra)
  - RA → Sequence of File operations
- **Transaction control**
- **Deeper**
  - File, (R)DBMS, ERP, DW, Big Data (No SQL, SQL again)
  - SQL on MPP and Hadoop (Greenplum, **HAWQ**)



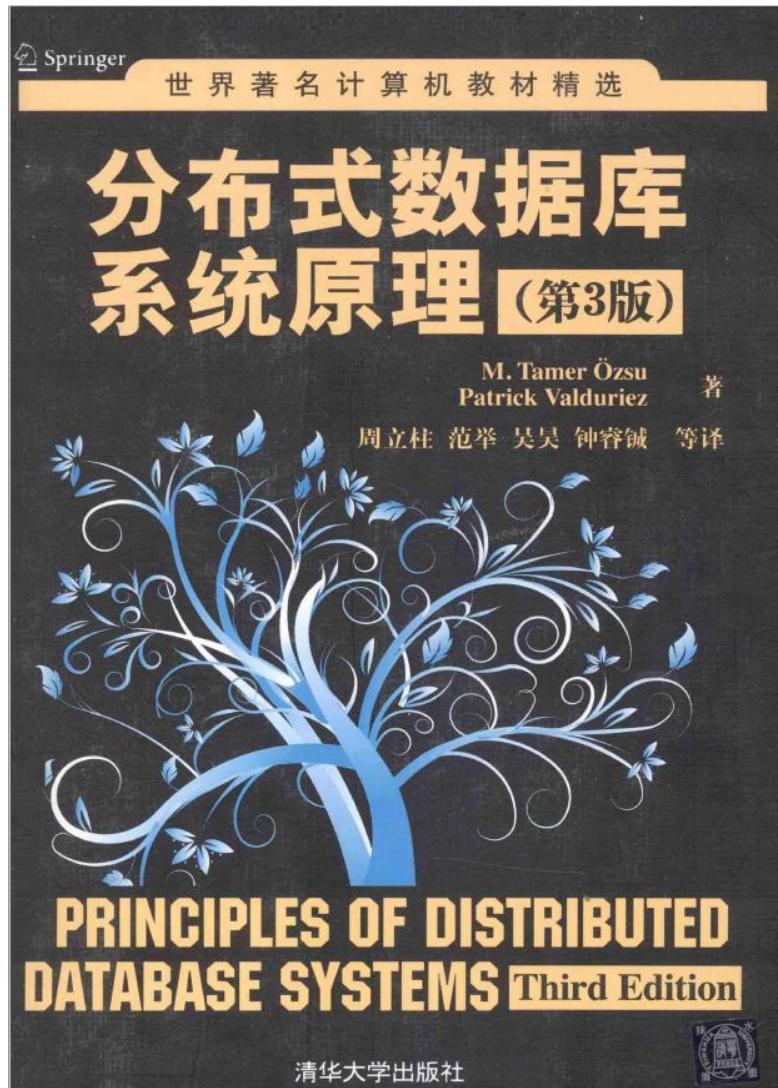
# 数据库发展历史

→ 一脉相承  
- - - 生态兼容，内核自研



① 注：改名原因为华为自研数据库成为独立品牌GaussDB, 其中T表示OLTP, A表示OLAP数据库

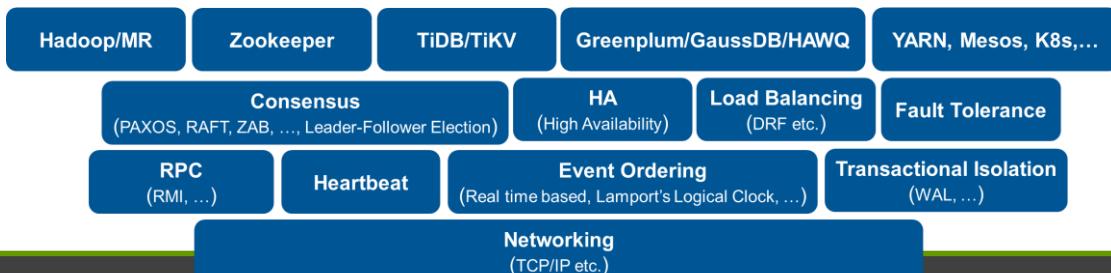
# 2021年7月7日21:49:57 以后再翻翻吧



- 分布式数据库系统原理
- 厄兹叙 (M.Tamer Ozsu) / Patrick Valduriez 著  
, 周立柱 / 范举 译
- 2014

觉得还是应该从 那个层次结构的方式来介绍 Distributed DBMS  
2021年7月7日 21:50:50

在草稿之上还有 Reading Torlerance, SAUCR(?)的信息，课后查一下吧  
2021年5月26日 12:57:28



# Chapter 5: Go deeper and wider

## □ Overview about the evolution of Data Management & Analytics

- Evolution of DM
  - Before IT
  - In IT : File system, RDBMS, ERP, DW, Big Data
- Hints –
  - Integrate DM with DA (Data Analytics)
    - ✓ Data – diverse (from structured to unstructured) and huge (from single to distributed)

## □ SQL again

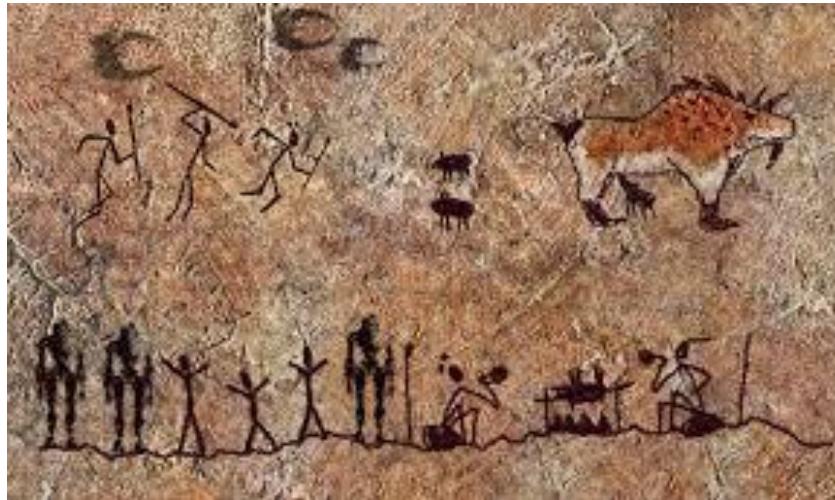
## □ Final (if you want to share)

# Kernel is Data management

## – storage + processing

### □ By that, a long history

- Early human records-clay tablets, hieroglyphics, cave paintings, paper records of family histories, treaties, inventories, and so on



cave painting



Hieroglyphics [haɪərə'glifiks]  
hieroglyphist [haɪərə'glifist]  
n. 象形文字研究者, 书写象形文字者

# Piles of paper files

---



All previous data processing is **MANKIND**  
😊 - Read, Write, Create,  
Compose, Print, ...

# Chapter 8: Go deeper and wider

## □ Overview about the evolution of Data Management & Analytics

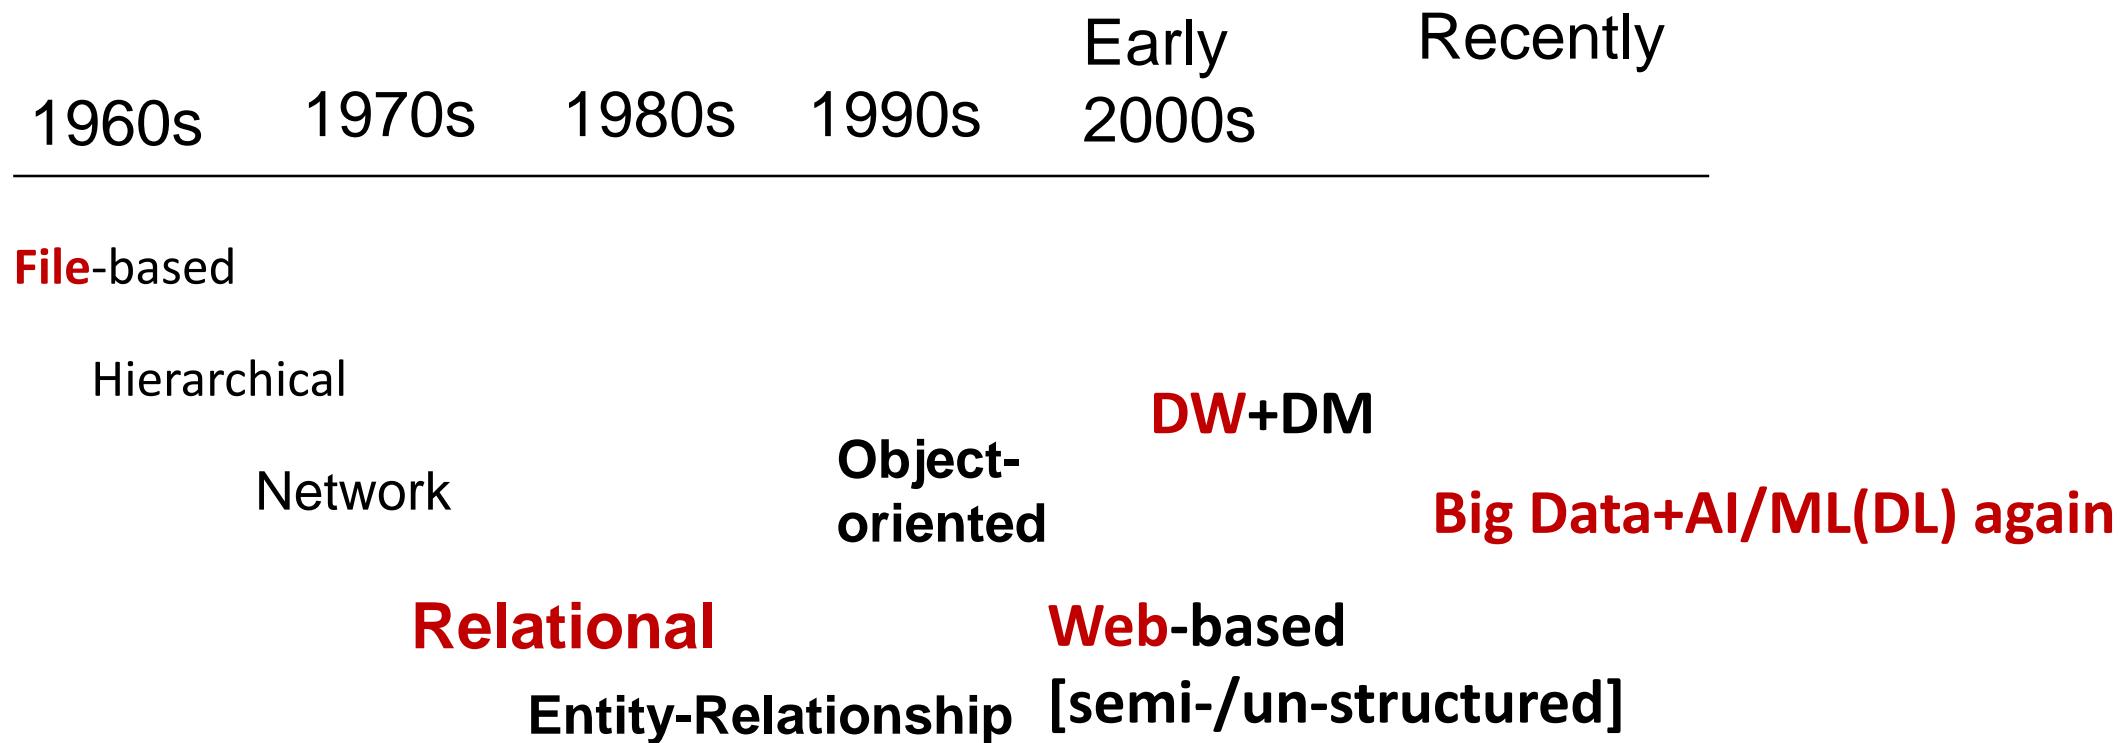
- Evolution of DM
  - Before IT
  - In IT : File system, RDBMS, ERP, DW, Big Data
- Hints –
  - Integrate DM with DA (Data Analytics)
  - Integrate DM with DA (Data Analytics)
    - ✓ Data – diverse (from structured to unstructured) and huge (from single to distributed)

## □ SQL again

## □ Final (if you want to share)

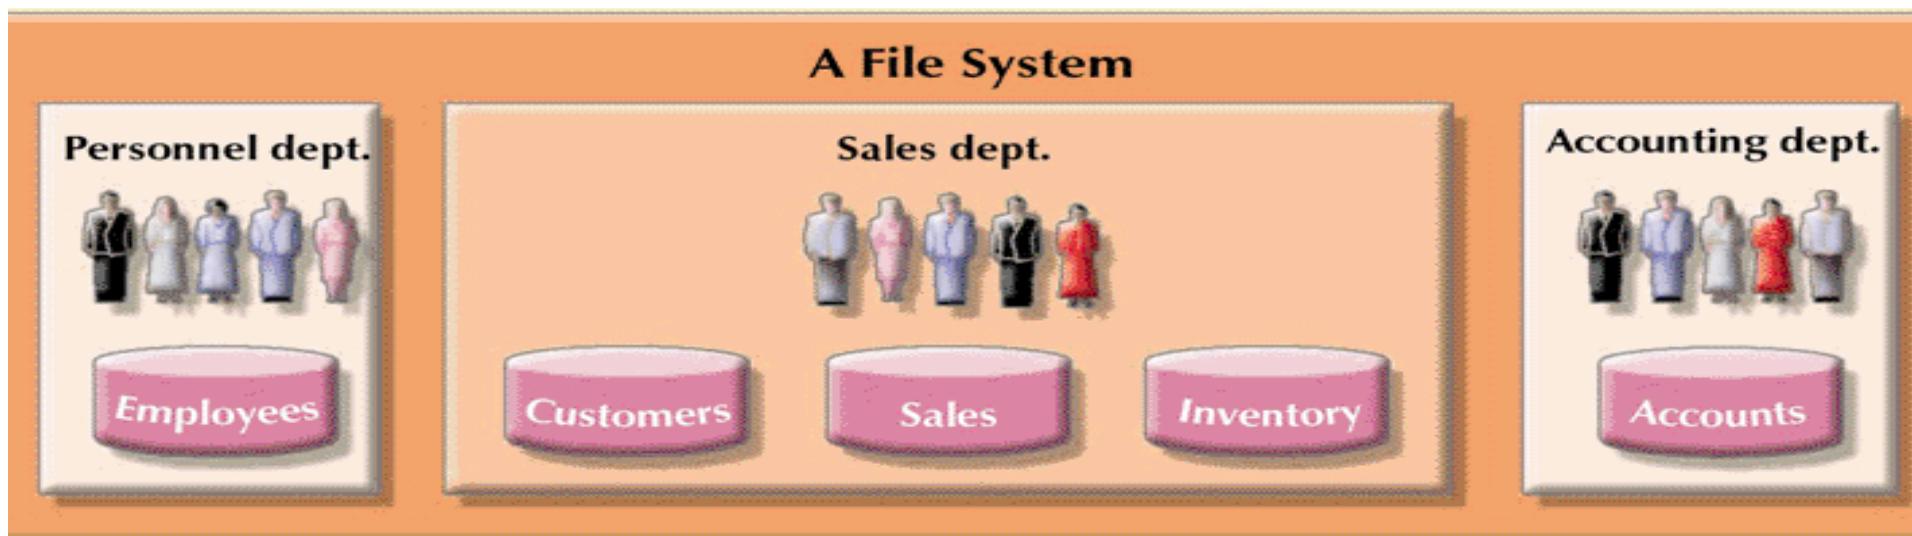
# Evolution of Data Models

## □ Timeline



## □ Old way to manage data – Simple File System

- Each file was used by its own application programs
- Each file was owned by individual or department who commissioned its creation



# Historical Perspective (1960-)

- Companies began automating their **back-office bookkeeping** in the 1960s.
- **COBOL** and its record-oriented file model were the work-horses of this effort.
- Typical work-cycle:
  1. a batch of transactions was applied to the old-tape-master
  2. a new-tape-master produced
  3. printout for the next business day.
- **COmmon Business-Oriented Language** (COBOL 2002 standard)

**COBOL**  
THE NEW MEANING...





# COBOL

A quote by Prof. dr. E.W. Dijkstra,  
18 June 1975:

cripple ['kripəl]  
vt.使跛, 受伤致残

"The use of COBOL cripples the mind; its teaching should, therefore, be regarded as a criminal offence."

**COBOL**  
THE NEW MEANING...

COMpletely  
OUTDATED  
BADLY  
OVERUSED  
LanguagE



But: In 2012 still vacancies available for COBOL programmers.

# COBOL Code (just an example!)

```
01 LOAN-WORK-AREA.  
  03 LW-LOAN-ERROR-FLAG      PIC 9(01)    COMP.  
  03 LW-LOAN-AMT            PIC 9(06)V9(02) COMP.  
  03 LW-INT-RATE             PIC 9(02)V9(02) COMP.  
  03 LW-NBR-PMTS            PIC 9(03)    COMP.  
  03 LW-PMT-AMT              PIC 9(06)V9(02) COMP.  
  03 LW-INT-PMT              PIC 9(01)V9(12) COMP.  
  03 LW-TOTAL-PMTS          PIC 9(06)V9(02) COMP.  
  03 LW-TOTAL-INT            PIC 9(06)V9(02) COMP.  
  
*  
 004000-COMPUTE-PAYMENT.  
  
*  
  MOVE 0 TO LW-LOAN-ERROR-FLAG.  
  
  IF (LW-LOAN-AMT ZERO)  
    OR  
    (LW-INT-RATE ZERO)  
    OR  
    (LW-NBR-PMTS ZERO)  
    MOVE 1 TO LW-LOAN-ERROR-FLAG  
    GO TO 004000-EXIT.  
  
  COMPUTE LW-INT-PMT = LW-INT-RATE / 1200  
  ON SIZE ERROR  
    MOVE 1 TO LW-LOAN-ERROR-FLAG  
    GO TO 004000-EXIT.
```

# Before RA – network, hierarchical model, ...

## Historical Perspective (1970's)

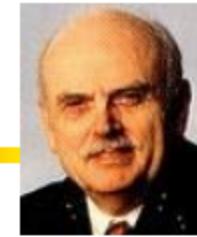
- Transition from handling transactions in **daily batches** to systems that managed an **on-line database** that could capture transactions as they happened.
- At first these systems were **ad hoc**
- Late in the **60's**, "**network**" and "**hierarchical**" database products emerged.

### Data Base Task Group

- A **network data model** standard (DBTG) was defined, which formed the basis for most commercial systems during the **1970's**.
- In **1980** DBTG-based **Cullinet** was the leading software company.



# The "relational" data model



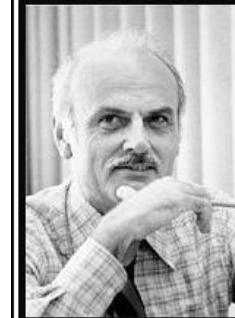
The "relational" data model, by Ted Codd in his landmark 1970 article "*A Relational Model of Data for Large Shared Data Banks*", was a major advance over DBTG.

- The relational model unified data and metadata => only one form of data representation.
- A non-procedural data access language based on algebra or logic.
- The data model is easier to visualize and understand than the pointers-and-records-based DBTG model.
- Programs written in terms of the "abstract model" of the data, rather than the actual database design => programs insensitive to changes in the database design.

# For more friendly management

Born: August 23, 1923  
Died: April 18, 2003 (aged 79)

## □ 1970 - E.F. Codd and the Relational Model



The most important motivation for the research work that resulted in the relational model was the objective of providing a sharp and clear boundary between the logical and physical aspects of database management.

(E. F. Codd)

## ■ I prefer to use to the top goal and 2 roles to indicate the benefit of DBMS

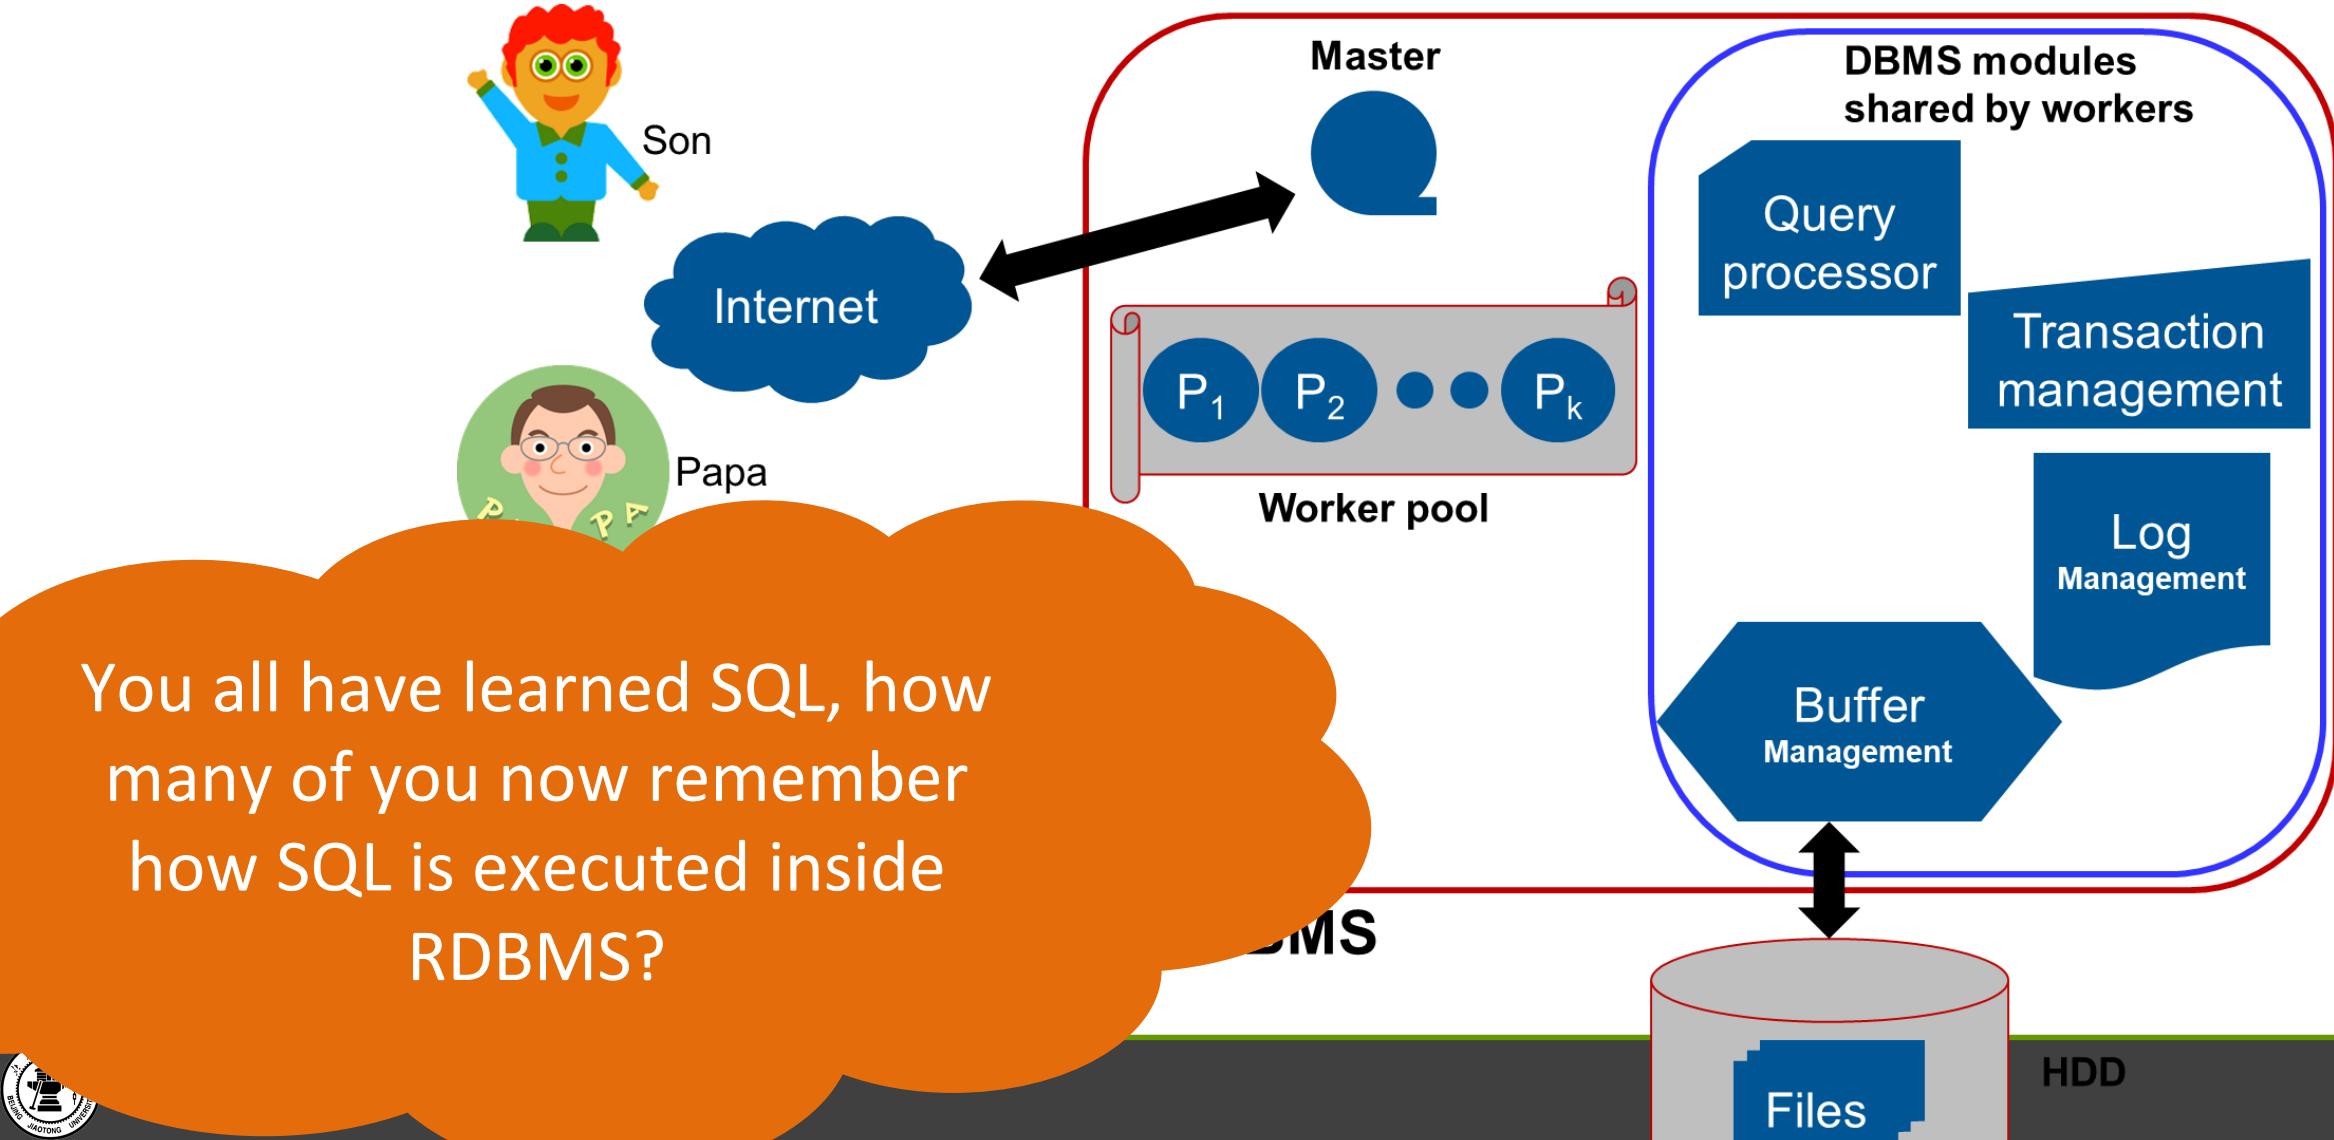
### ➤ Top goal:

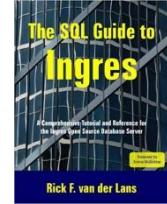
- ✓ Support the **data access of many users concurrently**

### ➤ 2 roles:

1. Concurrent data management by many users
2. Provide friendly/flexible interaction for common users

# Sketch of internal modules of modern RDBMS





# Old RDBMSs – Ingres, System R ...

## The "relational" data model success

- Both industry and university research communities embraced the relational data model and extended it during the 1970s.
- It was shown that a high-level relational database query language could give performance comparable to the best record-oriented database systems. (!)
- This research produced a generation of systems and people that formed the basis for IBM's DB2, Ingres, Sybase, Oracle, Informix and others.



### Ingres at UC Berkeley in 1972 (1/2)

---

Inspired by Codd's work on the relational database model, (Stonebraker, Rowe, Wong, and others) a project that resulted in:

- the design and build of a relational database system
- the query language (QUEL)
- relational optimization techniques
- a language binding technique
- storage strategies
- pioneering work on distributed databases

## Ingres at UC Berkeley in 1972 (2/2)

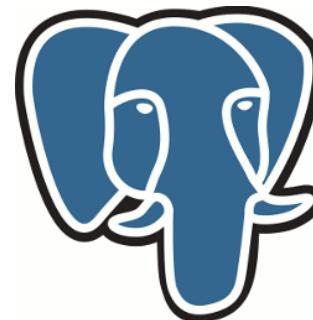
---

The academic system evolved into Ingres from Computer Associates.

Nowadays: PostgreSQL; also the basis for a new object-relational system.

Further work on:

- distributed databases
- database inference
- active databases (automatic responding)
- extensible databases.



**PostgreSQL**  
the world's most advanced open source database

## IBM: System R (1/2)

- Codd's ideas were inspired by the problems with the DBTG network data model and with IBM's product based on this model ([IMS](#)). **controversial** [kɔntrə've:ʃəl]  
adj. 有争议的, 引起争议的
- Codd's relational model was very [controversial](#):
  - too simplistic
  - could never give good performance.
- IBM Research chartered a [10-person effort](#) to [prototype a relational system](#) -> a prototype, System R ([evolved into the \[DB2\]\(#\) product](#))

[http://liacs.leidenuniv.nl/~bakkerem2/dbdm/  
Leiden-DBDM-02-Databases-1x1.pdf](http://liacs.leidenuniv.nl/~bakkerem2/dbdm/Leiden-DBDM-02-Databases-1x1.pdf)



## IBM: System R (2/2)

---

- Defined the fundamentals on:
  - query optimization,
  - data independence (views),
  - transactions (logging and locking), and
  - security (the grant-revoke model).
- SQL from System R became more or less the standard.
- The System R group further research:
  - distributed databases ([project R\\*](#)) and
  - object-oriented extensible databases ([project Starburst](#)).  
[http://iacs.leidenuniv.nl/~bakkerem2/dbdm/  
Leiden-DBDM-02-Databases-1x1.pdf](http://iacs.leidenuniv.nl/~bakkerem2/dbdm/Leiden-DBDM-02-Databases-1x1.pdf)

## USA funded database research period 1970 - 1996:

- Projects at UCLA => Teradata

**TERADATA.**

THE BEST DECISION POSSIBLE™

- Projects at CCA (SDD-1, Daplex, Multibase, and HiPAC):
  - distributed database technology
  - object-oriented database technology
- Projects at Stanford:
  - deductive database technology
  - data integration technology
  - query optimization technology.
- Projects at CMU:
  - general transaction models
  - => Transarc Corporation.



[http://liacs.leidenuniv.nl/~bakkerem2/dbdm/  
Leiden-DBDM-02-Databases-1x1.pdf](http://liacs.leidenuniv.nl/~bakkerem2/dbdm/Leiden-DBDM-02-Databases-1x1.pdf)

# Extending RDBMS – distributed, OO, 3-tier applications, GIS, Temporal DBMS, ...

## The database research agenda of the 1980's Extending Relational Databases

- geographically distributed databases
- parallel data access.
- Theoretical work on distributed databases led to prototypes which in turn led to products.
  - *Note: Today, all the major database systems offer the ability to distribute and replicate data among nodes of a computer network.*
- Execution of each of the relational data operators in parallel => hundred-fold and thousand-fold speedups.
  - *Note: The results of this research appear nowadays in the products of several major database companies. Especially beneficial for data warehousing and decision support systems: <http://dacs.leidenuniv.nl/~bakkerem2/dbdm/> effective application in the area of OLTP is challenging.*

## □ Application architecture

### ■ C/S access to DBs

- Client/Server architecture is a general model for systems where a service is provided by one system (the server) to another (the client)

### ■ Web-based access to DBs

- Browser/Server architecture is a general model, in which Web server serves pages to browsers (clients) and can access database(s) ← You've tried JSP/Tomcat or Flask/Mango...

## □ Many types of data are adopted into DBMS

### ■ Object Oriented DBs

### ■ Semi-structured Data and XML: The research topic during my Ph.D at PKU

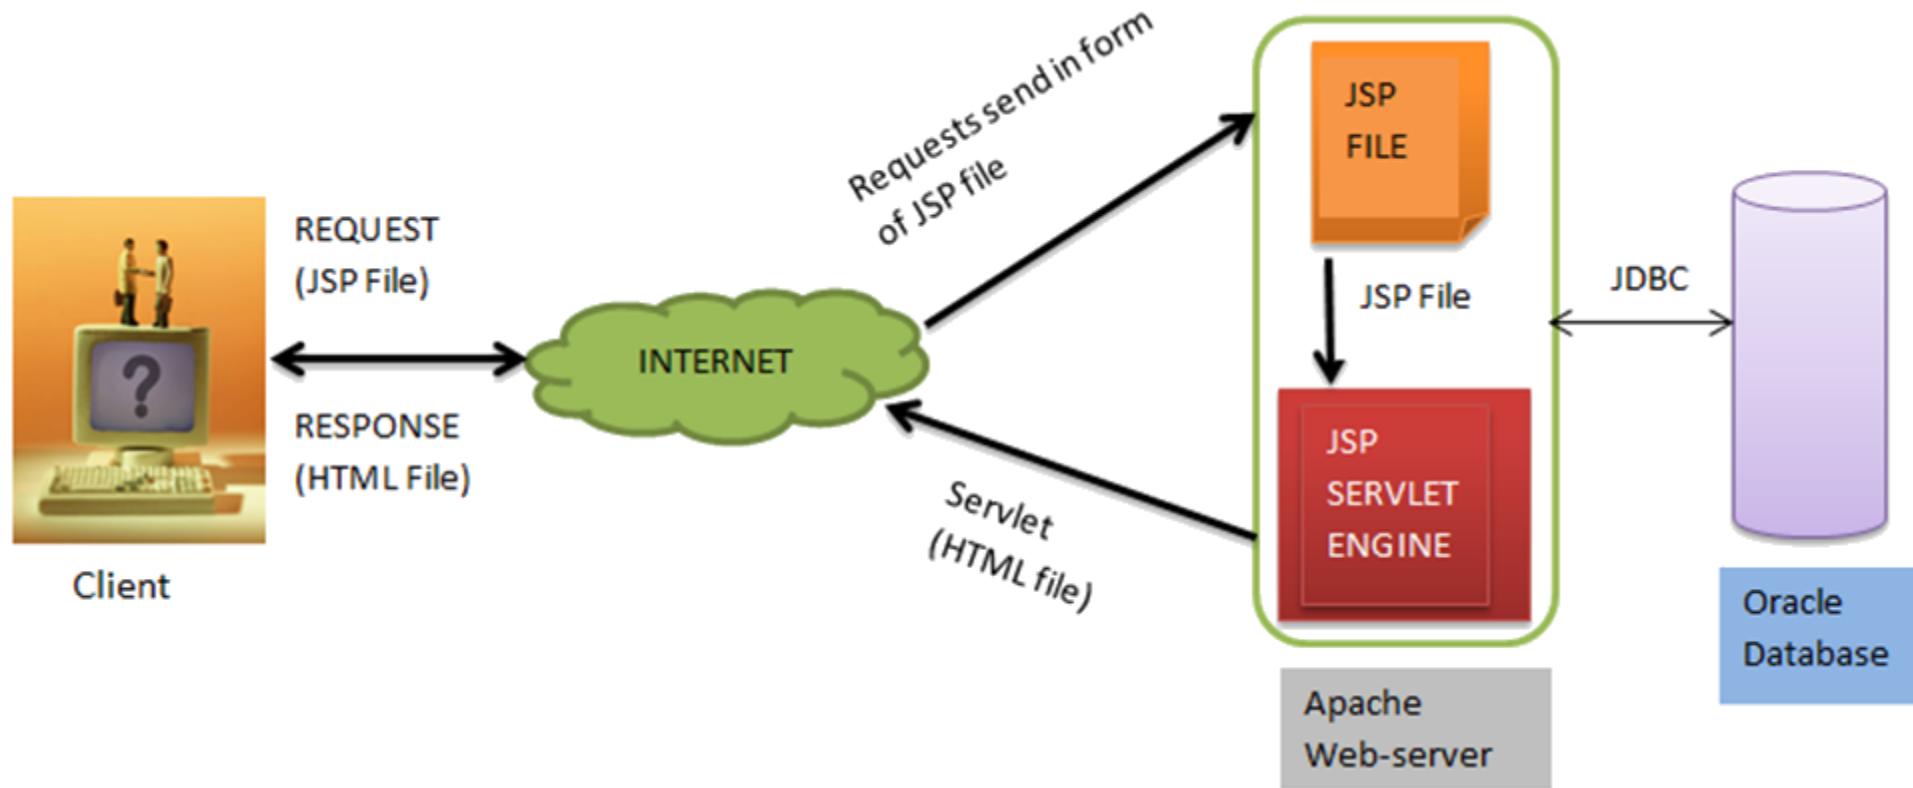
### ■ Multimedia DBs

### ■ Temporal DBs

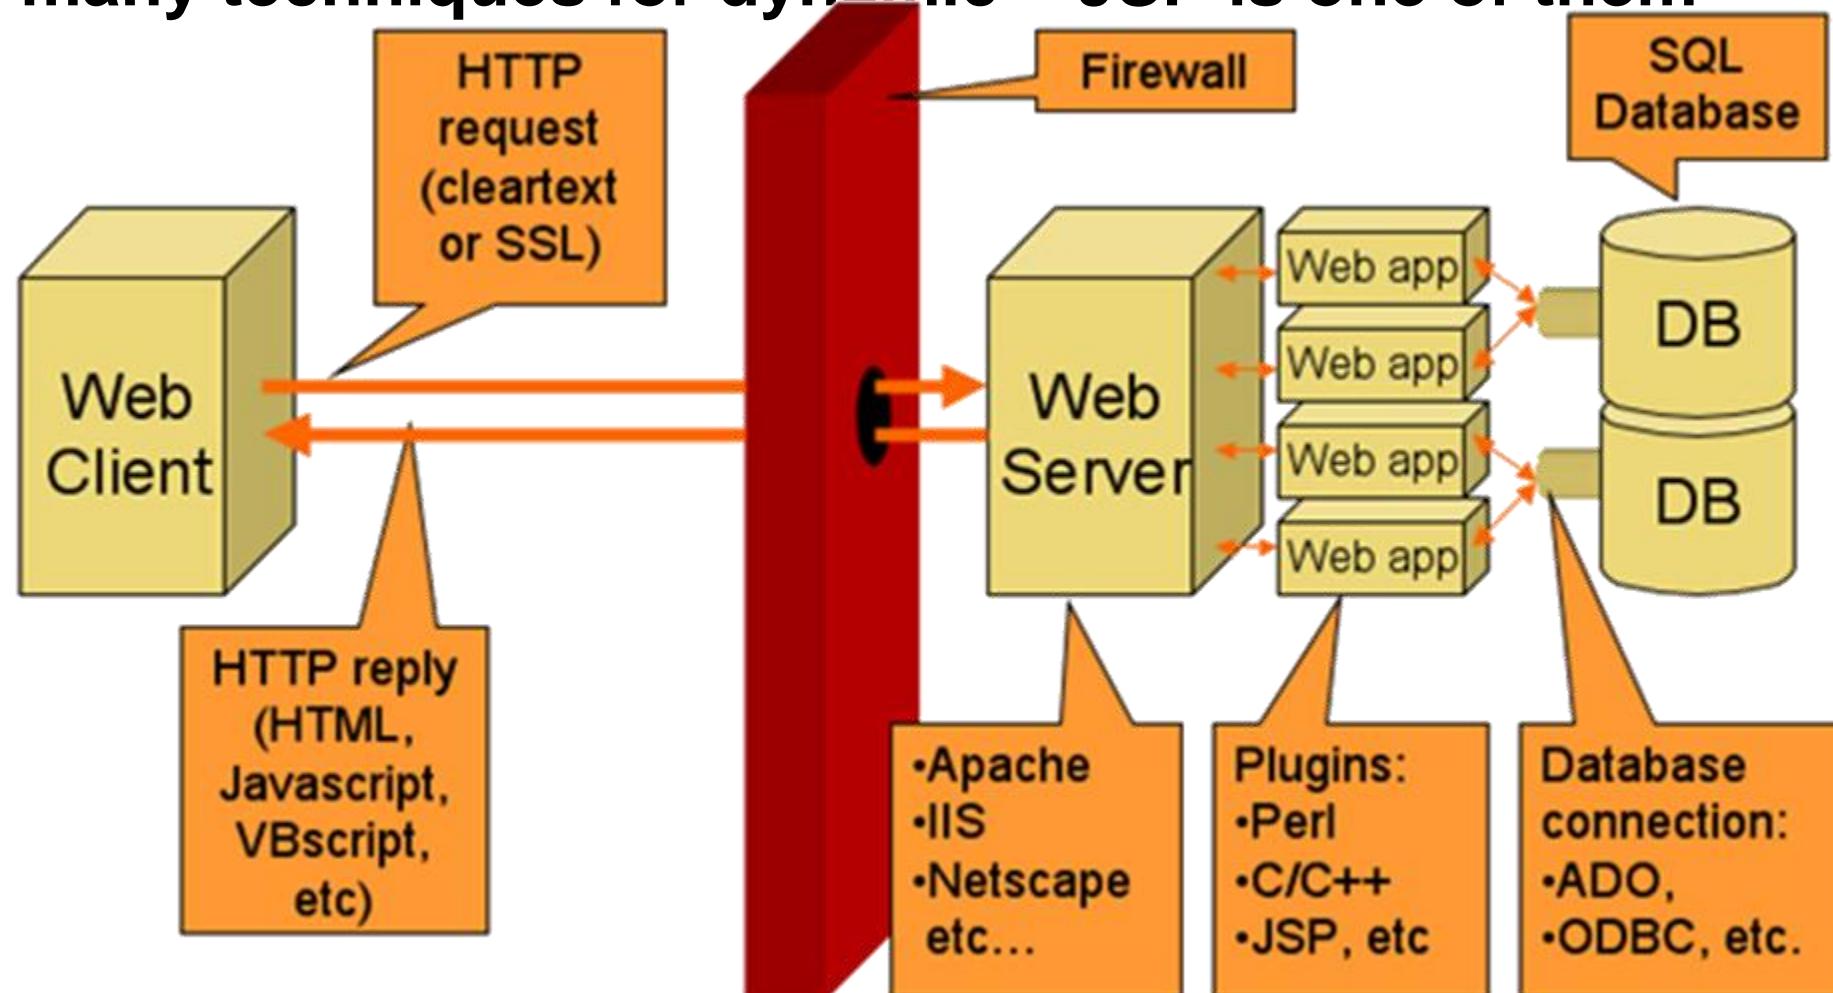
### ■ Logic DBs .....

Want to manage  
different data types as  
many as possible~~~

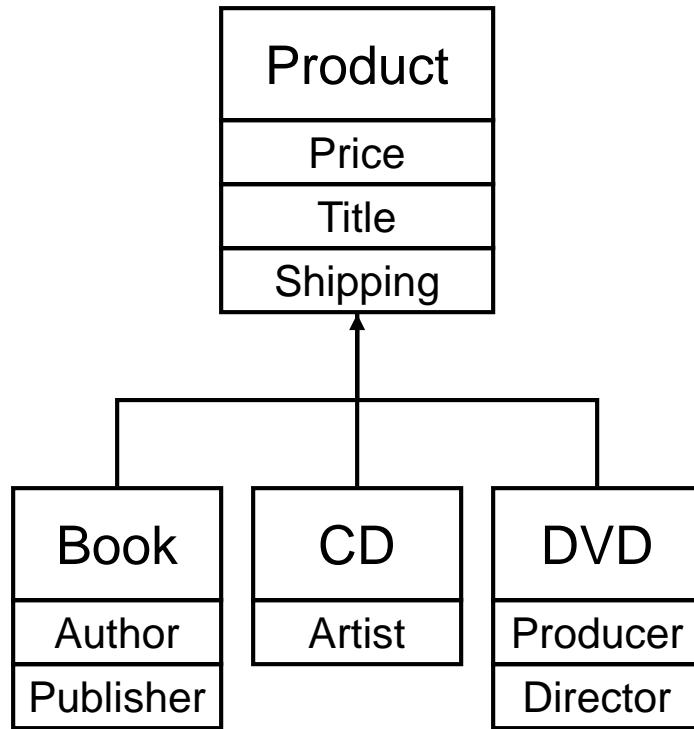
# Web-based Databases



## □ Now many techniques for dynamic – JSP is one of them



# OODB Example



## □ Product is abstract

- You cannot make a Product directly
- You can, however, make a Book, CD, or DVD, and these are Products

---

## □ Advantages

- Good integration with Java, C++, etc.
- Can store complex information
- Fast to recover whole objects
- Has the advantages of the (familiar) object paradigm

## □ Disadvantages

- There is no underlying theory to match the relational model
- Can be more complex and less efficient
- OODB queries tend to be procedural, unlike SQL

## ■ Multimedia DBs

- Store Images, Music and audio, Video and animation, Full texts of books, Web pages

## ■ Querying Multimedia DBs

### ➤ Metadata searches

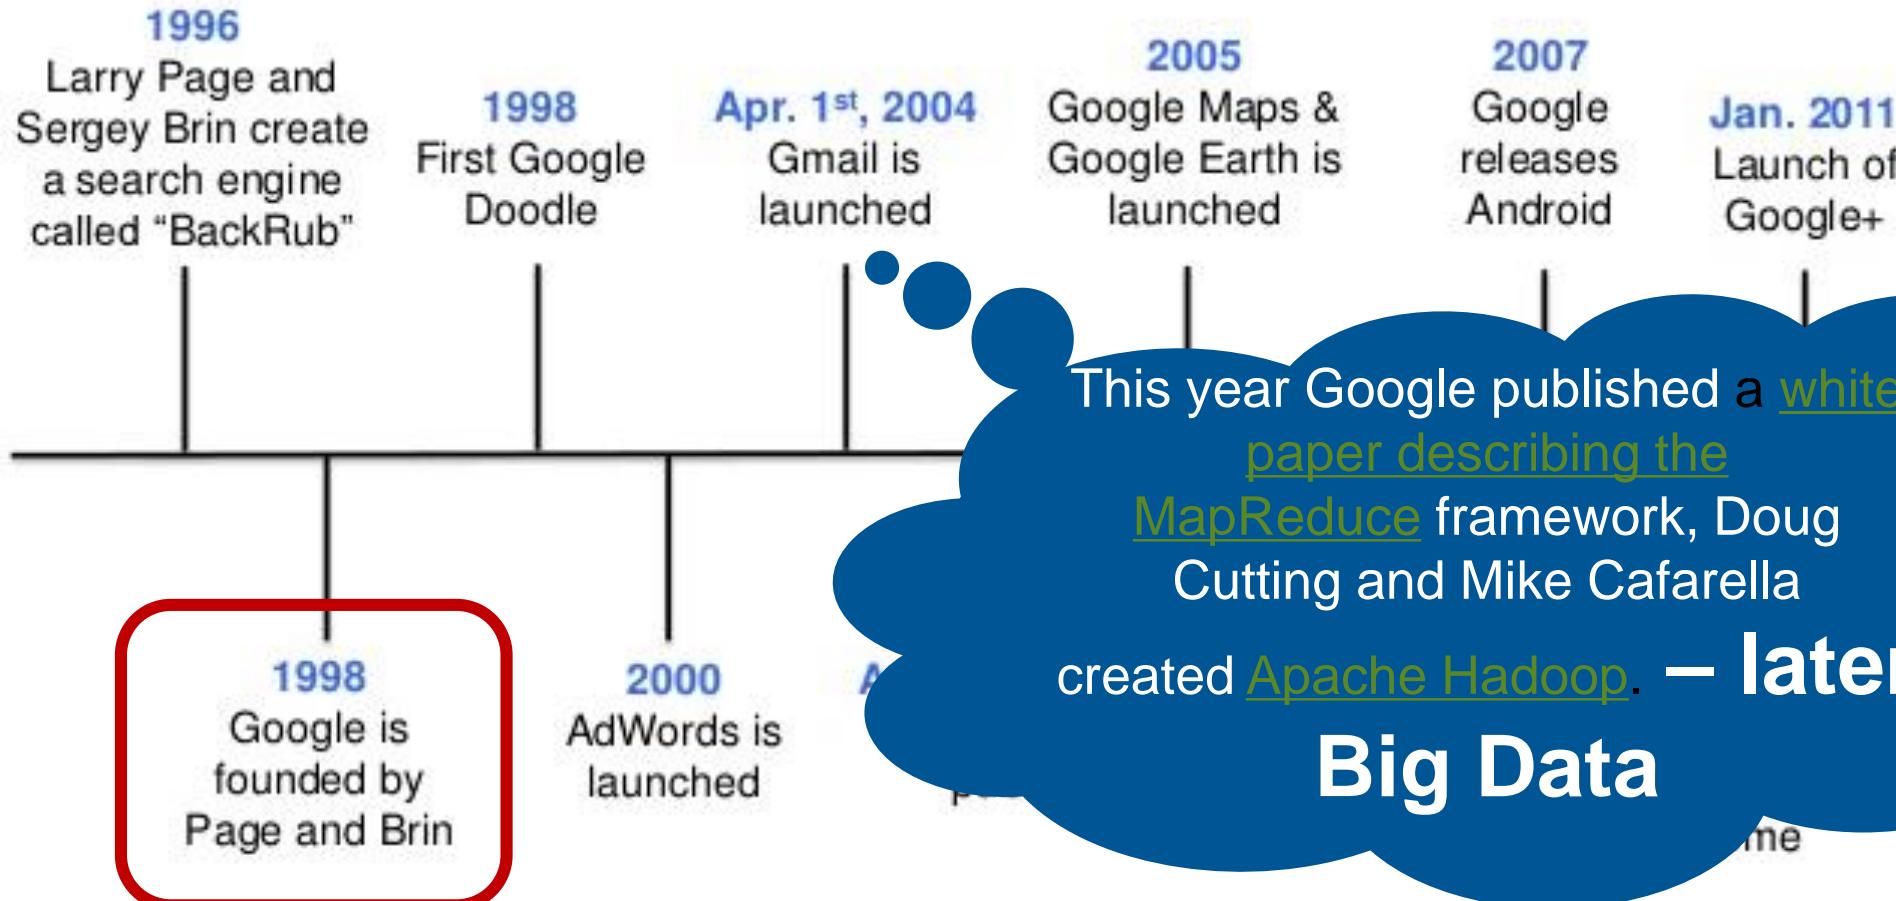
- ✓ Information about the multimedia data (metadata) is stored
- ✓ This can be kept in a standard relational database and queried normally
- ✓ Limited by the amount of metadata available

### ➤ Content searches

- ✓ The multimedia data is searched directly
- ✓ Potential for much more flexible search
- ✓ Depends on the type of data being used
- ✓ Often difficult to determine what the ‘correct’ results are

# Google triggers other researches ...

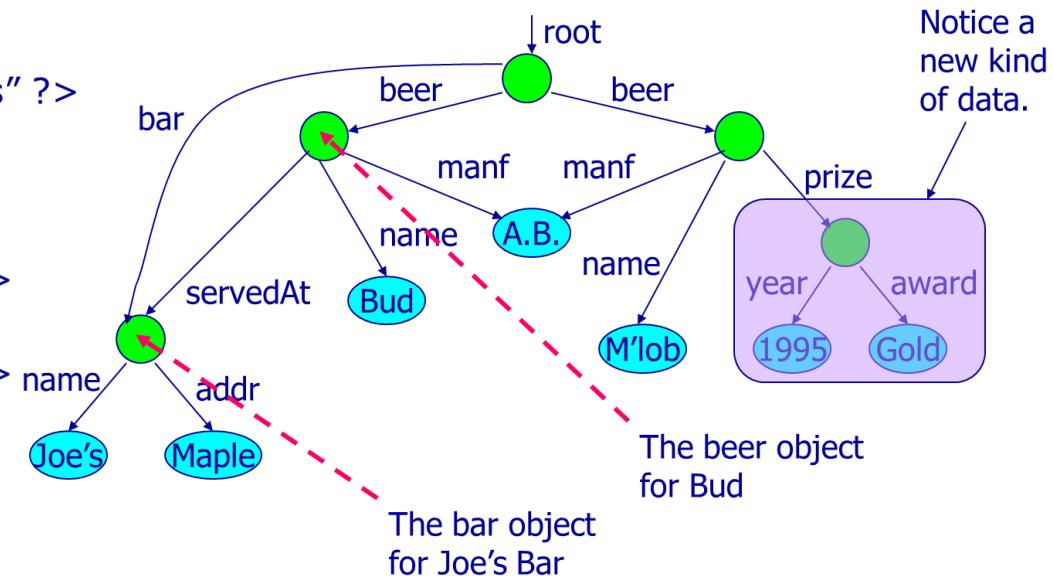
## □ CBIR, XML, ...



## Semi-structured Data and XML

- XML oriented keyword search – The research topic during my Ph.D at PKU
- And Graphs of Semi-structured Data

```
<? XML VERSION = "1.0" STANDALONE = "yes" ?>
<BARS>
  <BAR><NAME>Joe's Bar</NAME>
    <BEER><NAME>Bud</NAME>
      <PRICE>2.50</PRICE></BEER>
    <BEER><NAME>Miller</NAME>
      <PRICE>3.00</PRICE></BEER>
  </BAR>
  <BAR> ...
</BARS>
```



---

## □ XPATH and XQUERY

### □ XPATH is a language for describing paths in XML documents.

- Really think of the semistructured data graph and *its* paths.
- Why do we need path description language: can't get at the data using just Relation.Attribute expressions.

### □ XQUERY is a full query language for XML documents with power similar to OQL (Object Query Language, query language for object-oriented databases).

# Content-Based Retrieval



QBIC™ (Query By Image Content)  
from IBM - searches for images  
having similar colour or layout



'Modern' Databases

<http://wwwqbic.almaden.ibm.com/cgi-bin/stamps-demo>

# Content-Based Retrieval

---

## □ Image retrieval is hard

- It is often not clear when two images are 'similar'
- Image interpretation is unsolved and expensive
- Different people expect different things



## □ Do we look for?

- Images of roses
- Images of red things?
- Images of flowers?
- Images of red flowers?
- Images of red roses?

# Chapter 5: Go deeper and wider

## □ Overview about the evolution of Data Management & Analytics

- Evolution of DM

- Before IT
- In IT : File system, RDBMS, **ERP**, **DW**, Big Data

- Hints –

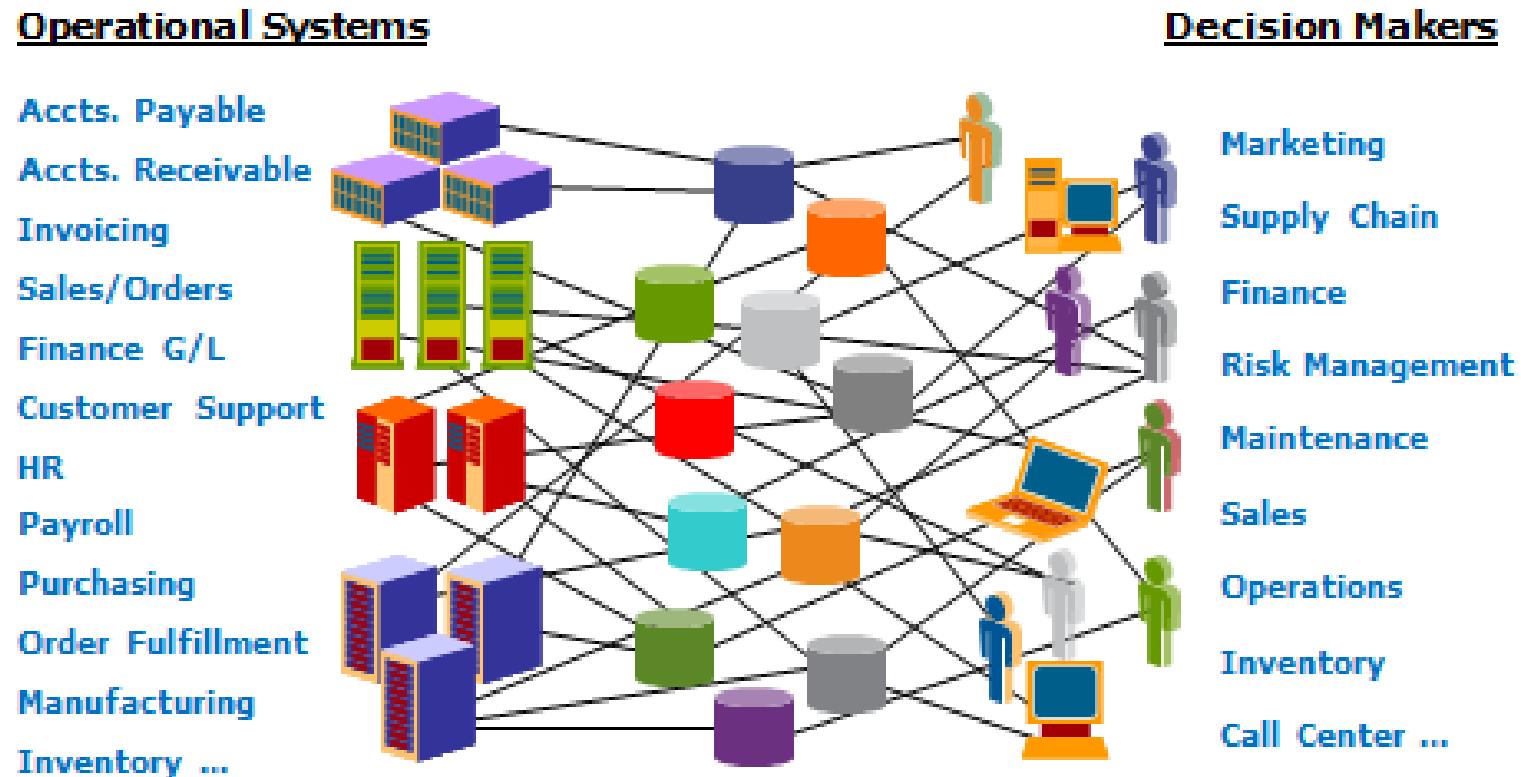
- Integrate DM with DA (Data Analytics)
  - ✓ Data – diverse (from structured to unstructured) and huge (from single to distributed)

## □ SQL again

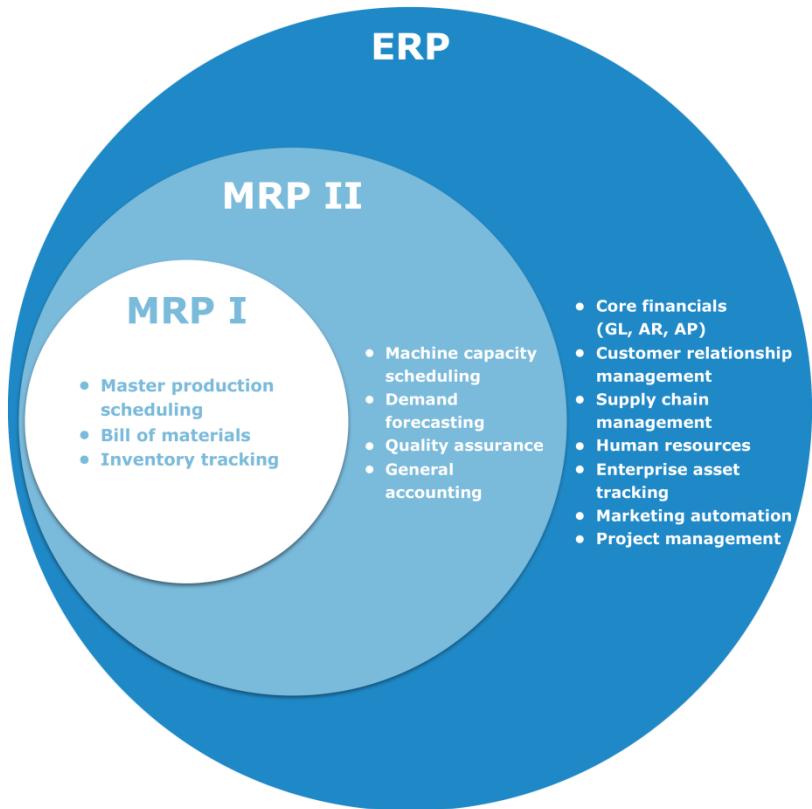
## □ Final (if you want to share)

# 1 Integration of the Data and Analytics

- DBMS: Transaction data are stored in many scattered databases
  - It's easy to do IDUQ (insert, delete, update, query) – OLTP style



# From MRP to ERP



## □ MRP I/II: Manufacturing Resource Planning

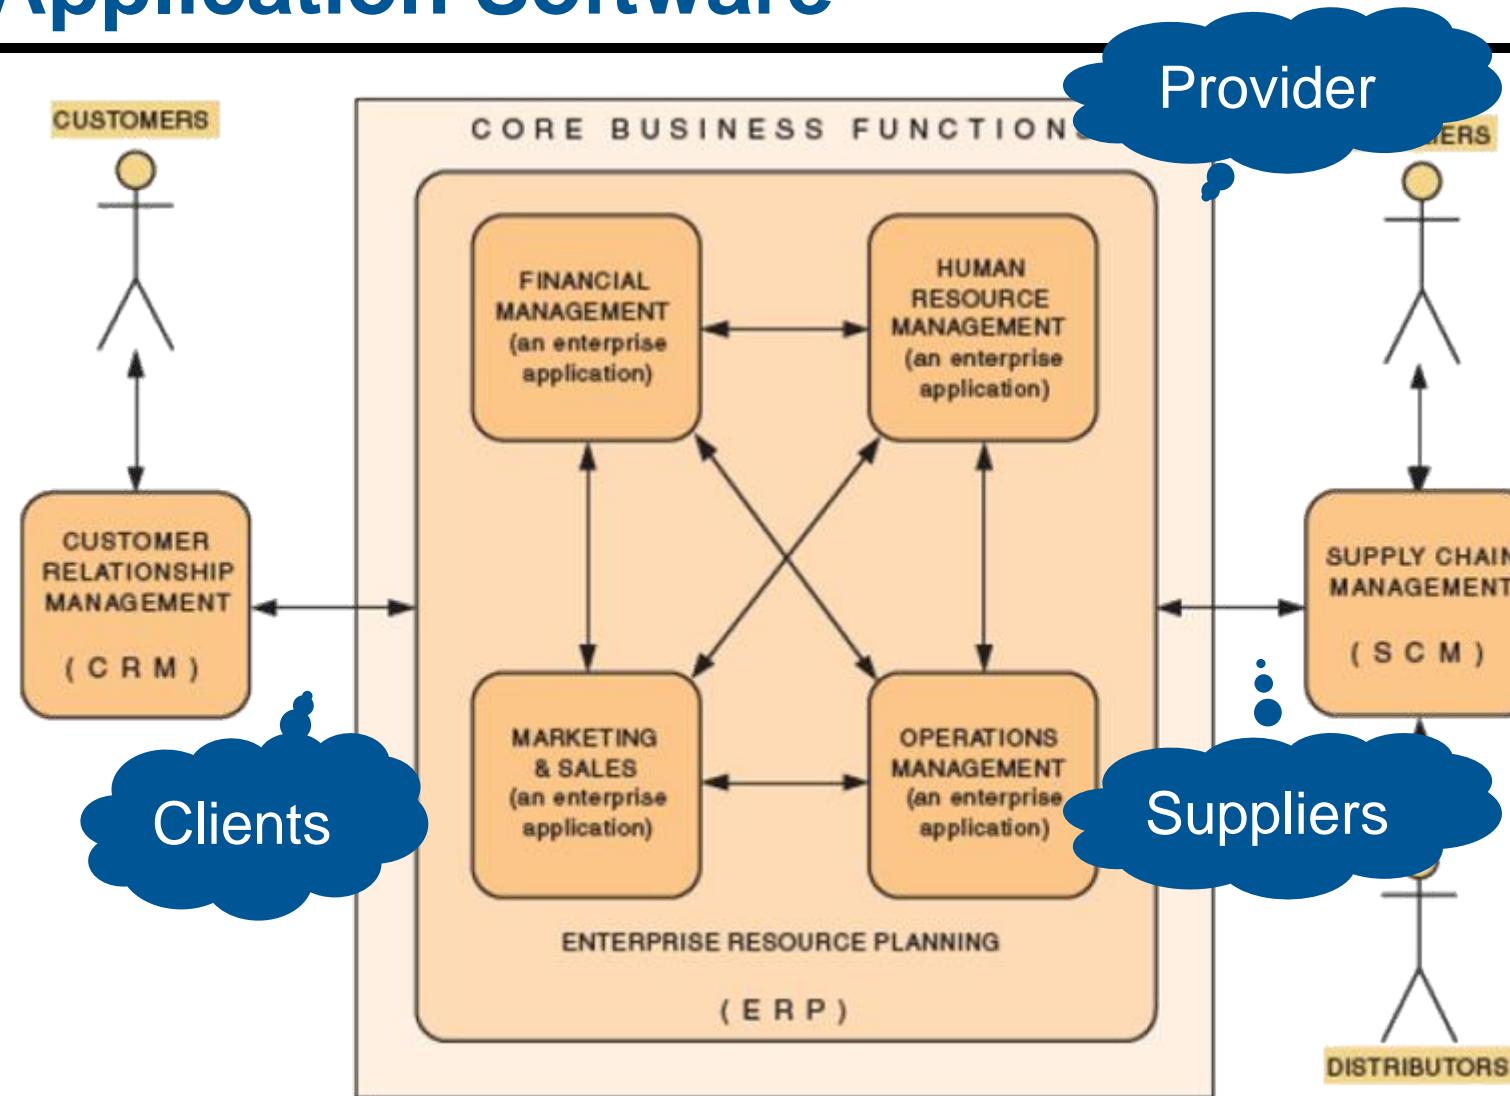
- MRP I was some of the first business software to be widely adopted during the 1970s.
  - Manufacturers sought these systems in order to improve efficiency and accuracy when it came to basic processes such as production scheduling and inventory management.
- By the 1980s, manufacturers realized they needed software that could also tie into their accounting systems and forecast inventory requirements - MRP II

## □ ERP: Enterprise resource planning

- Business analysis
- ERP systems are regarded as the successors to MRP II software. ERP suites include applications that are well outside the scope of manufacturing.

<http://www.softwareadvice.com/resources/mrp-vs-mrp-ii/>

# Great applications are constructed – Enterprise Application Software





<https://www.voortman.net/en/company/news/448-mrp-erp-or-mis>

- SAP, TeraData etc.

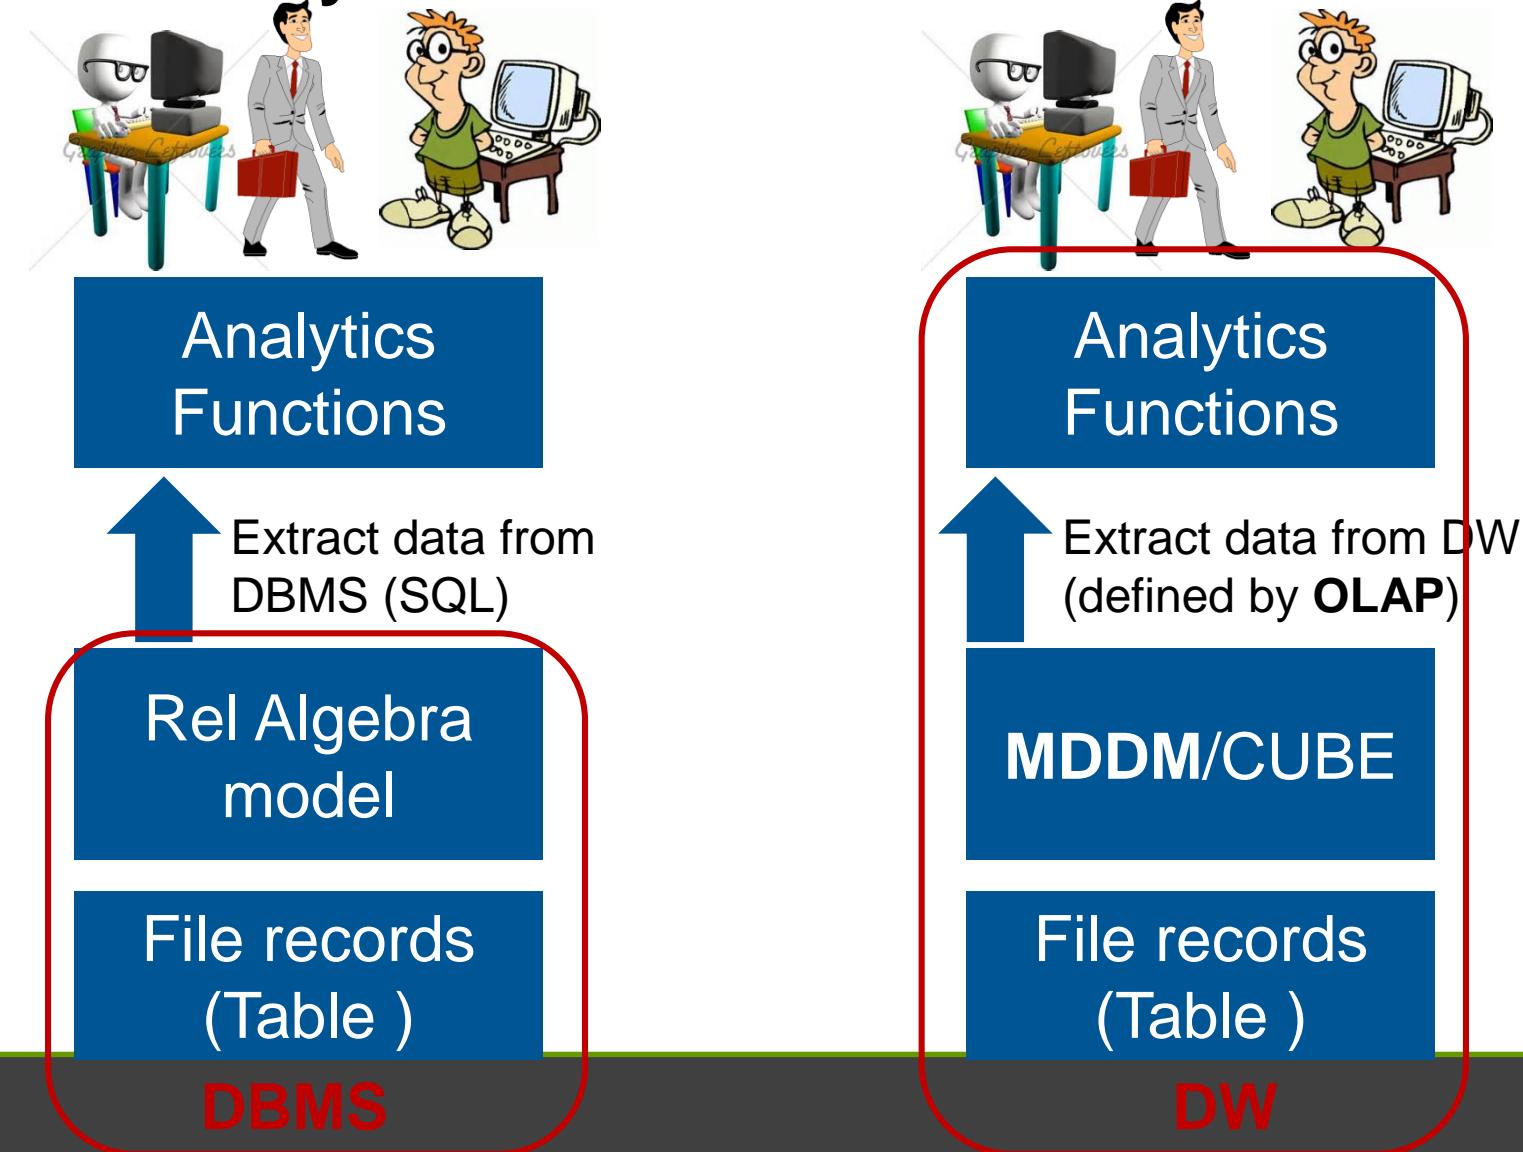
**TERADATA**  
THE BEST DECISION POSSIBLE™



# But, analytics is still separated

- + **Integrate** Data analytics into RDBMS was tried in  
1990s

- + **diverse**  
data



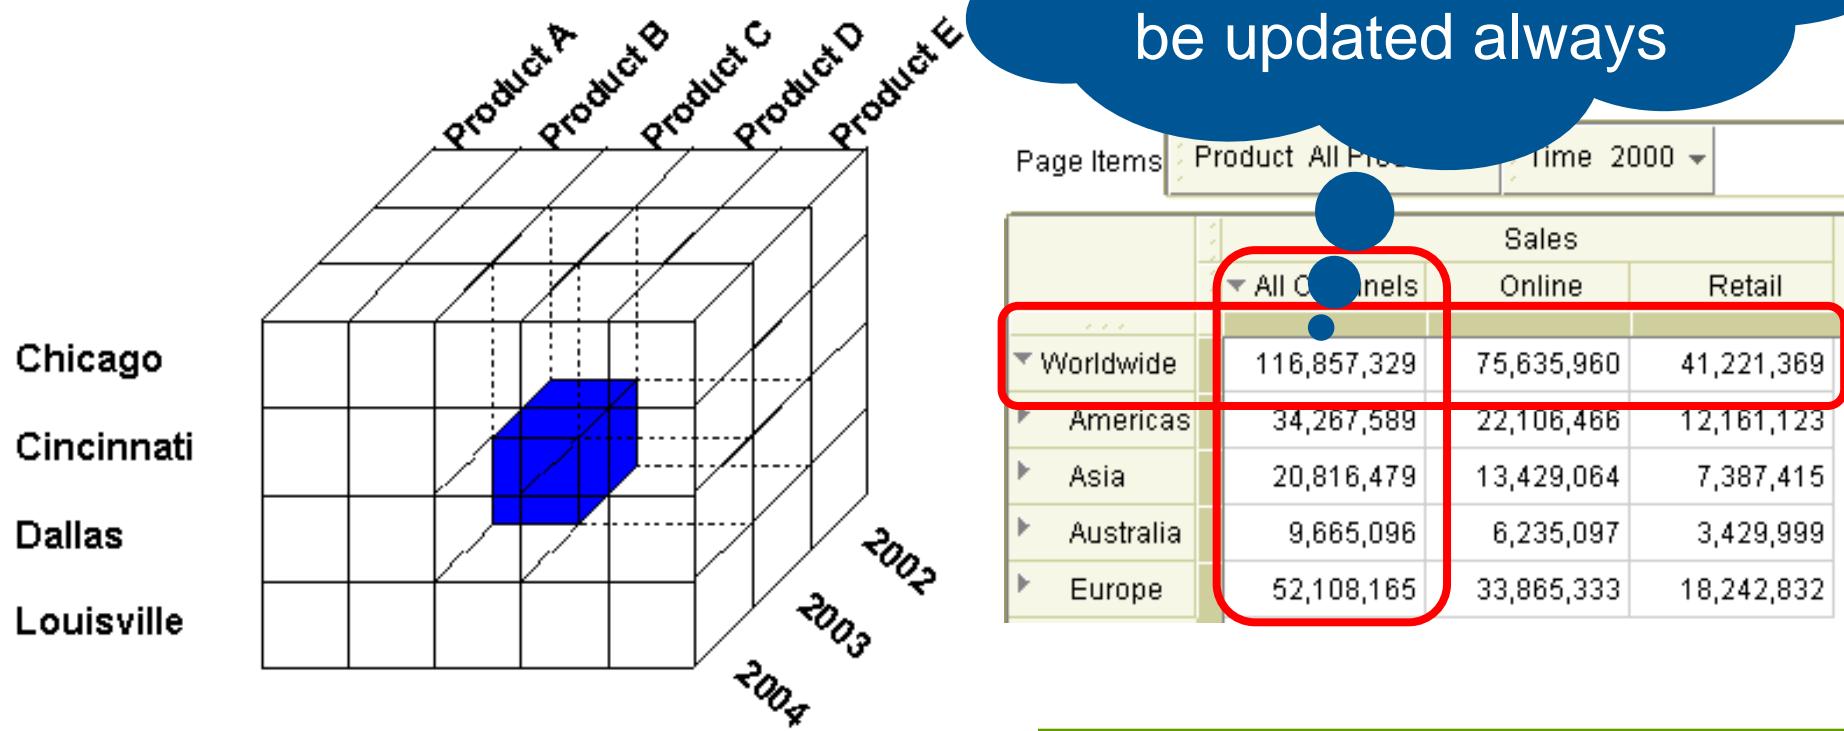
# Statistic result is stored in advance

[http://docs.oracle.com/html/B13915\\_04/i\\_olap\\_chapter.htm](http://docs.oracle.com/html/B13915_04/i_olap_chapter.htm)

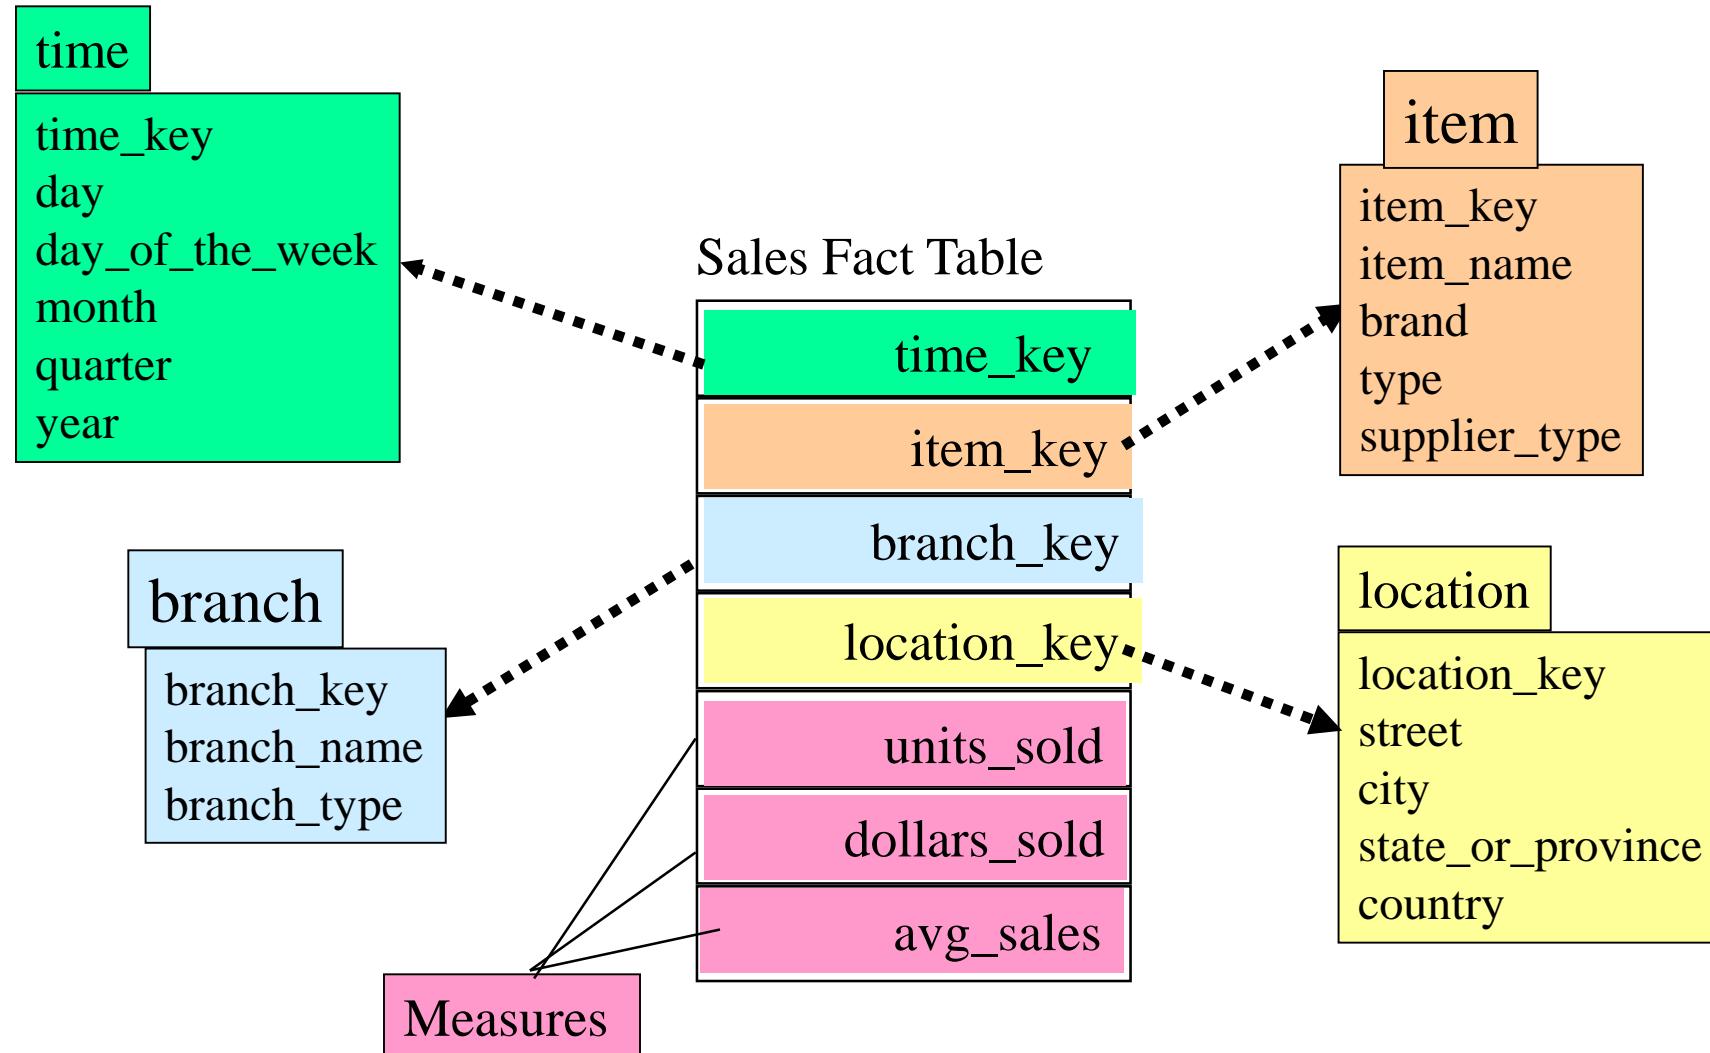
## □ Define the data with the multi-dimensional view

- A multidimensional data is optimized for analytical processing. Such data sources are sometimes called analytical processing (OLAP) data sources.

Some cumulative information is stored in advanced and will be updated always

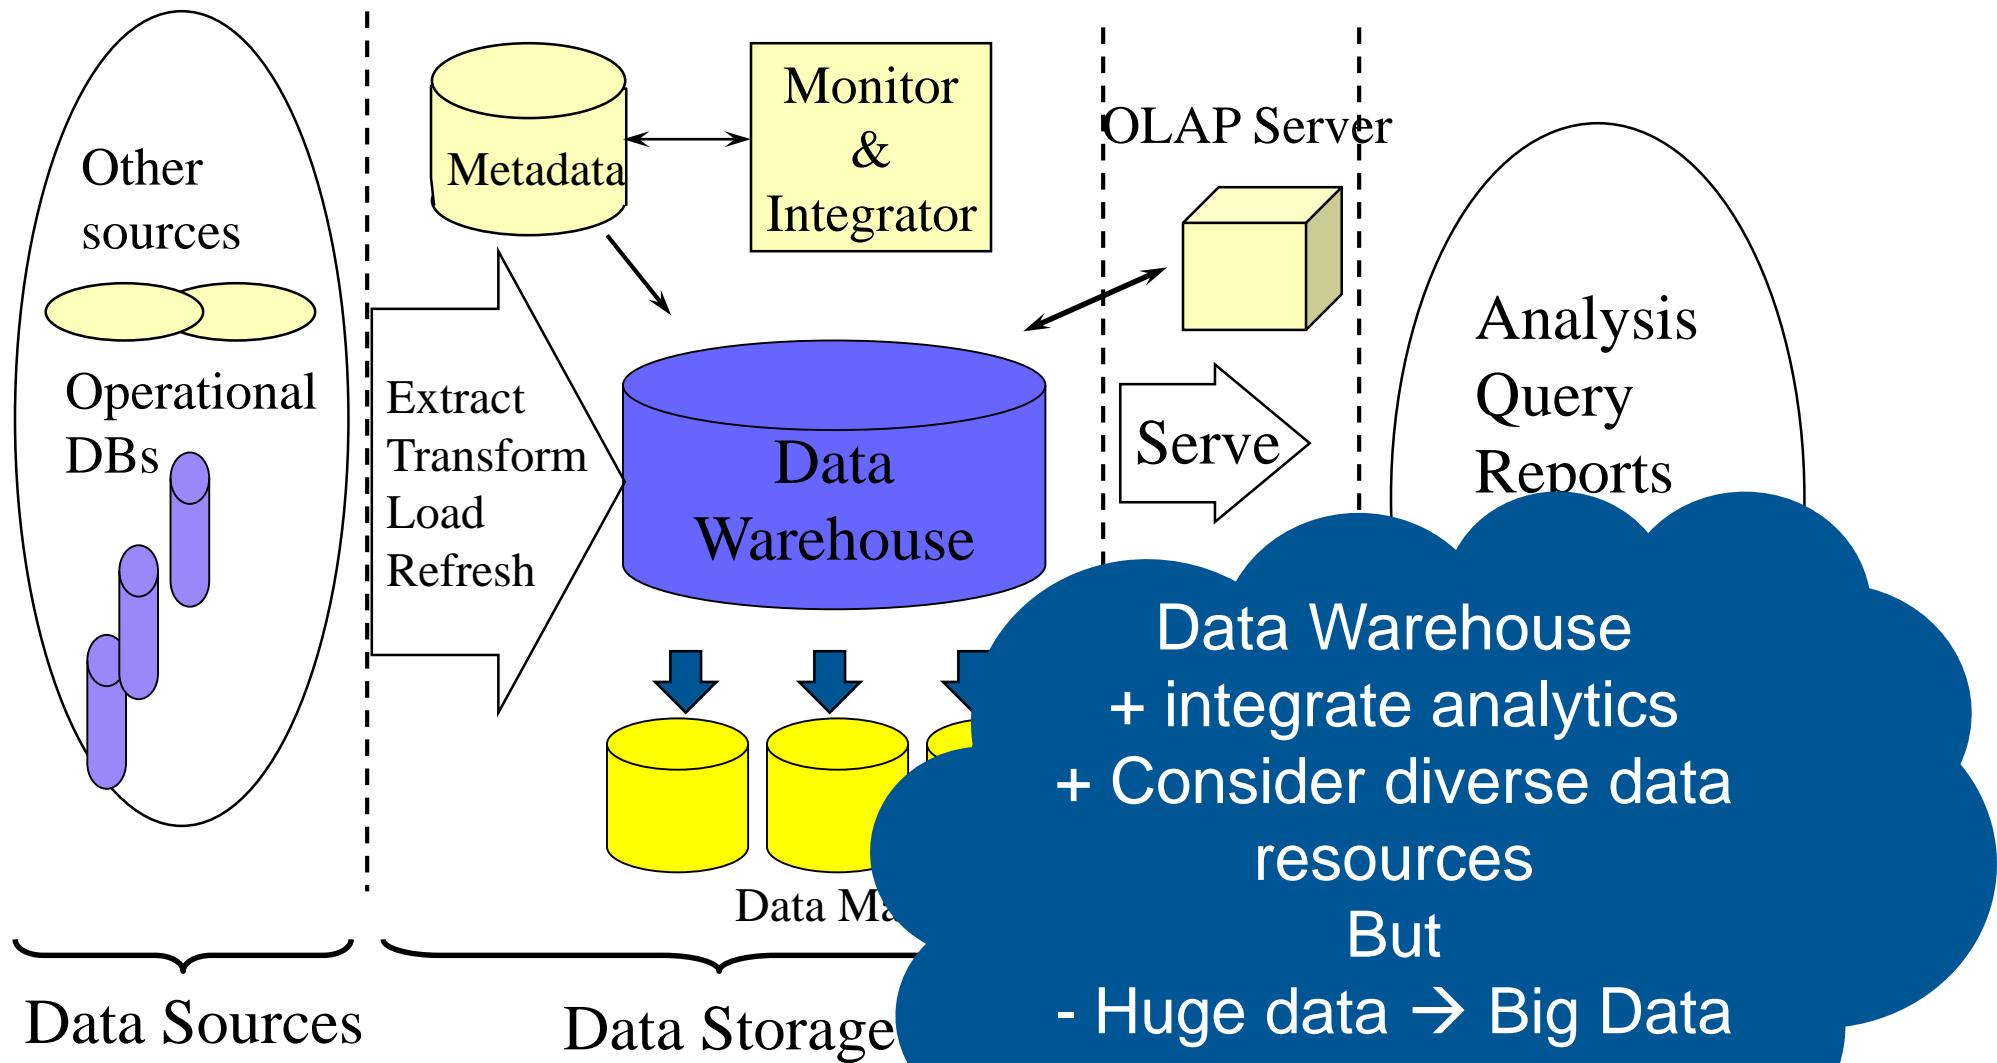


# MDDM of Star Schema (based on RDBMS)

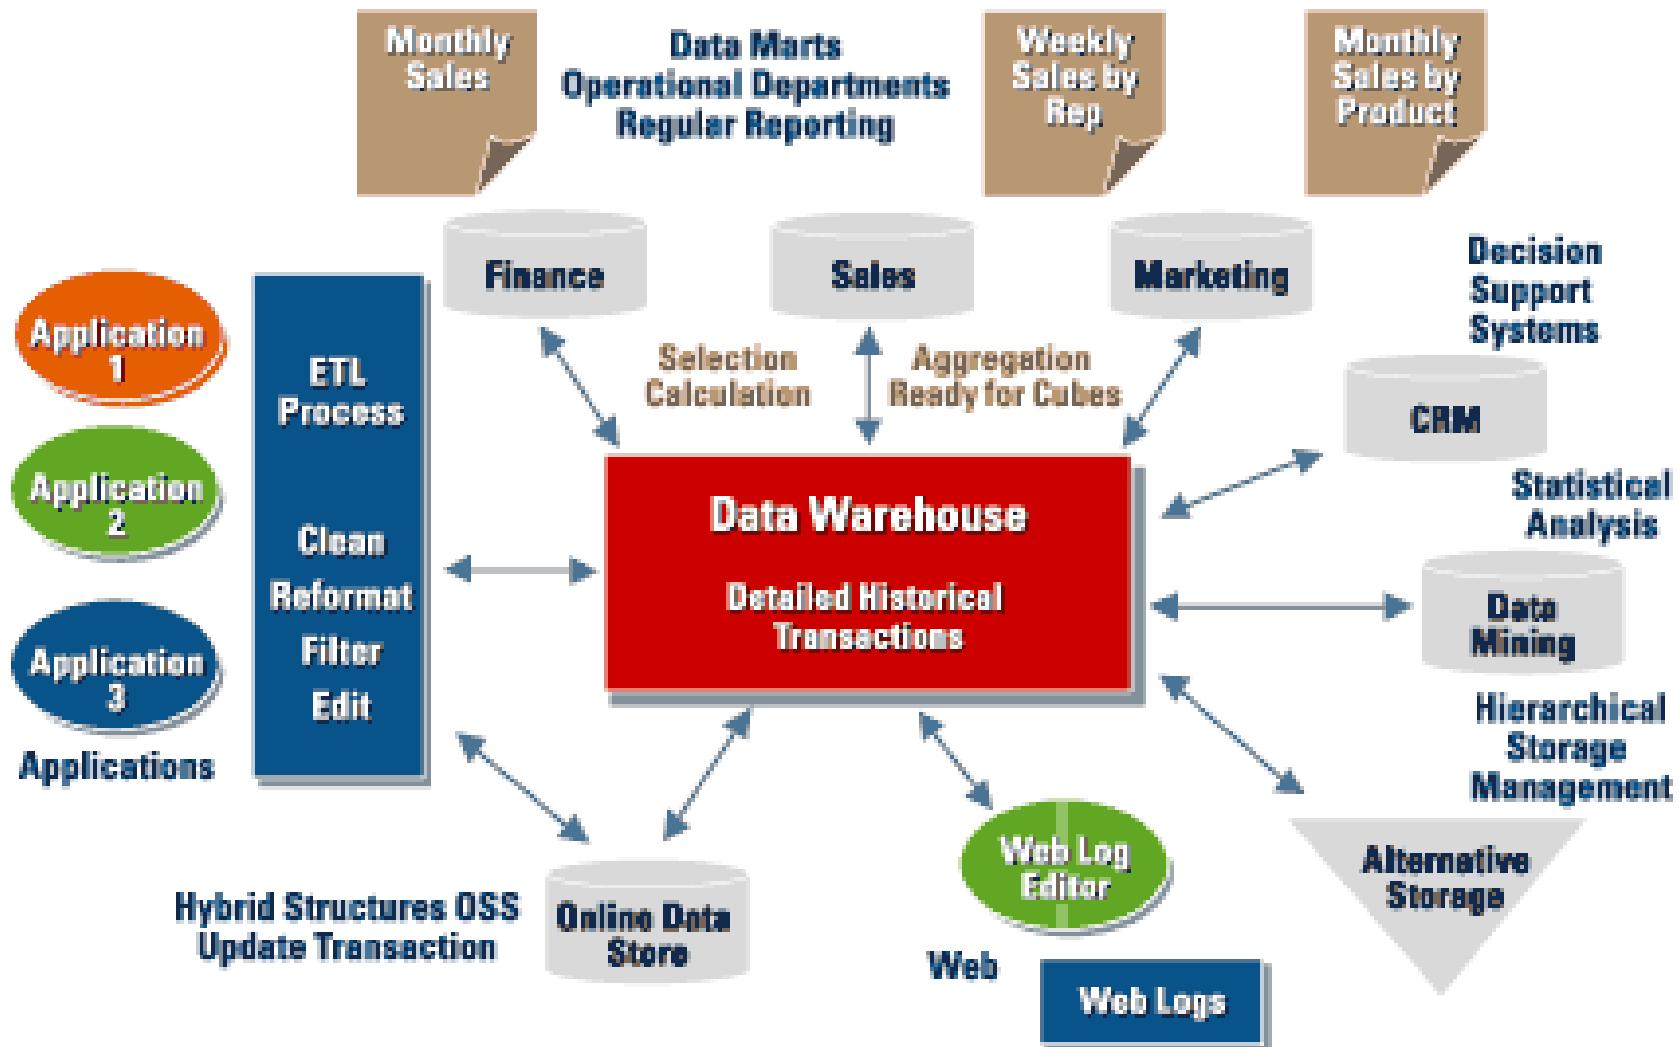


# Data Warehouse

## integrate analytics with data management



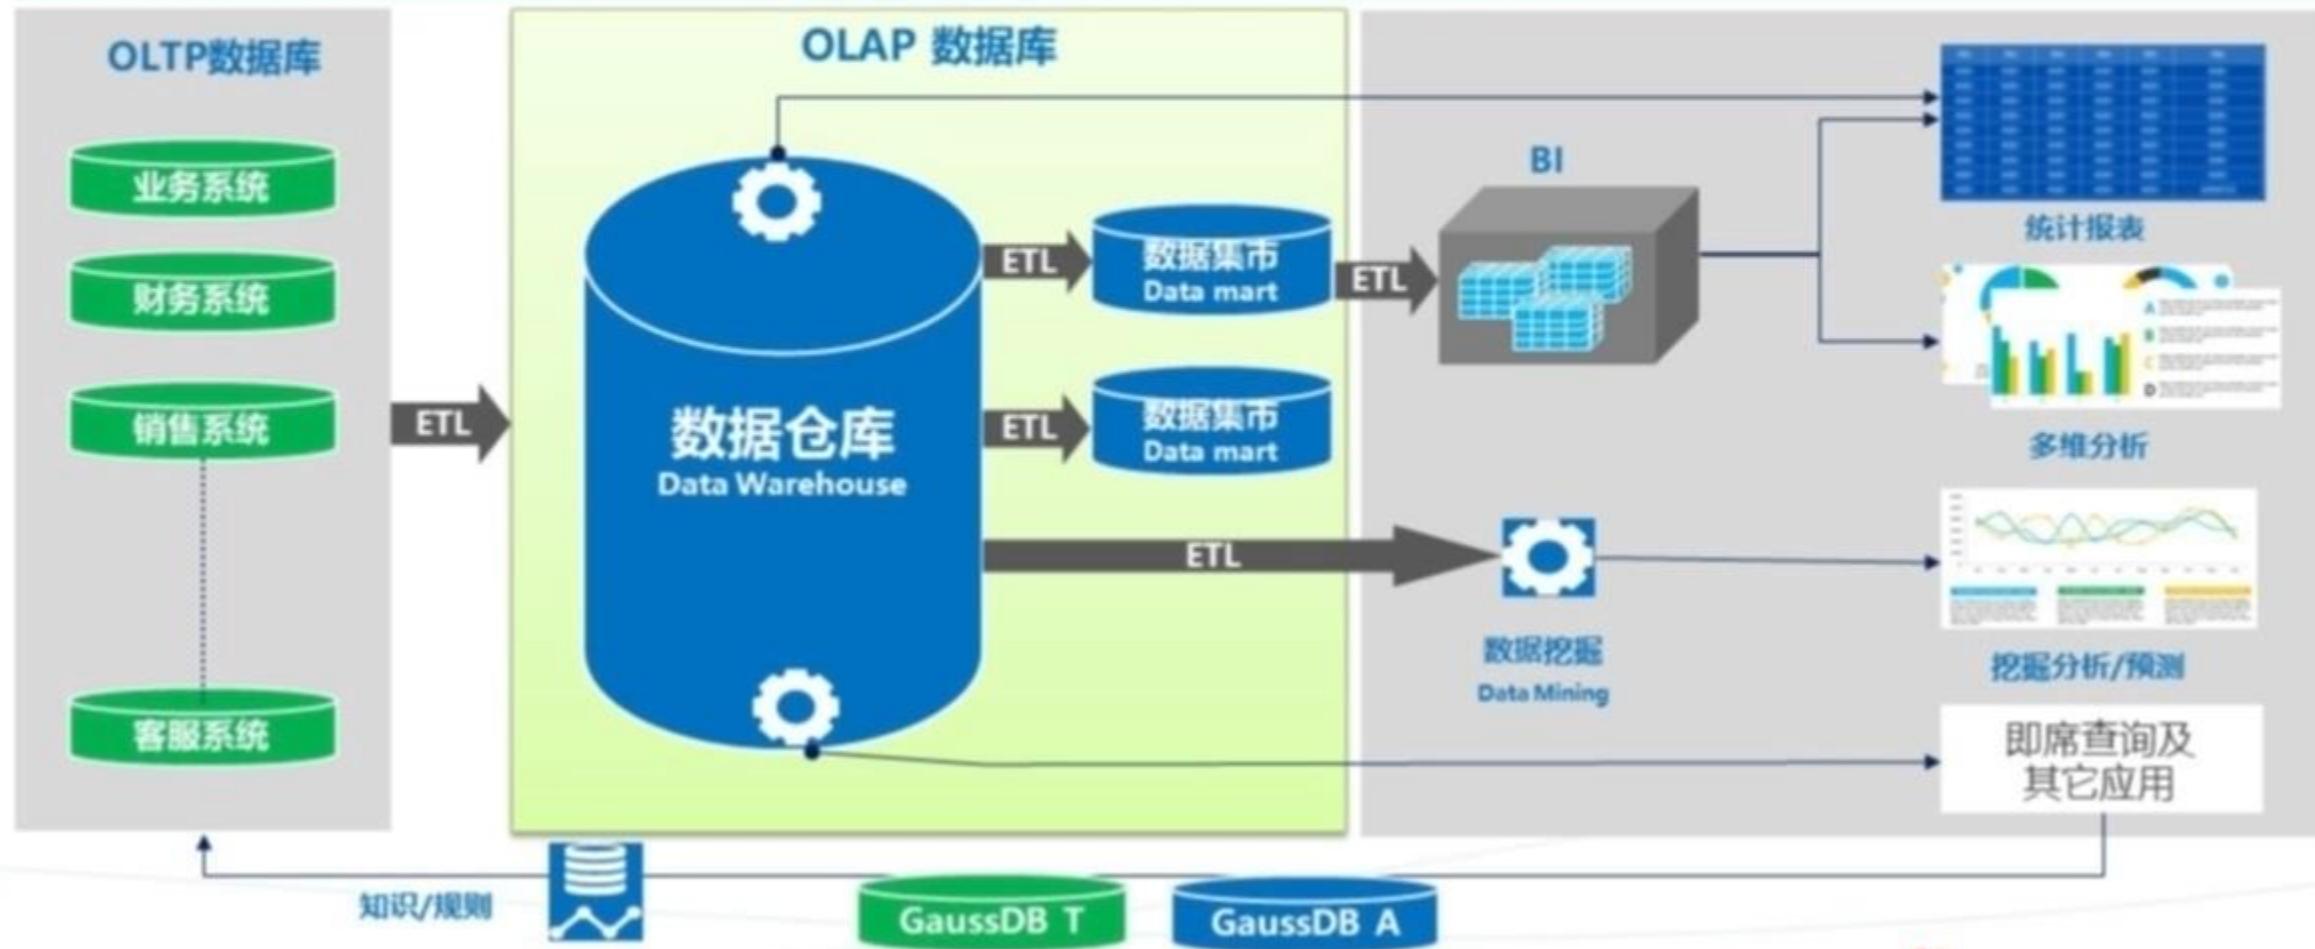
# Data Warehouse tries to integrate (Later)



Based on a graphic developed by Bill Inmon

# 典型的企业OLTP和OLAP数据库

- 联机事务处理(OLTP): 存储/查询业务应用中活动的数据以支撑日常的业务活动;
- 联机分析处理(OLAP): 存储历史数据以支撑复杂的分析操作，侧重决策支持;



# Data Mining

---

## □ Tools that:

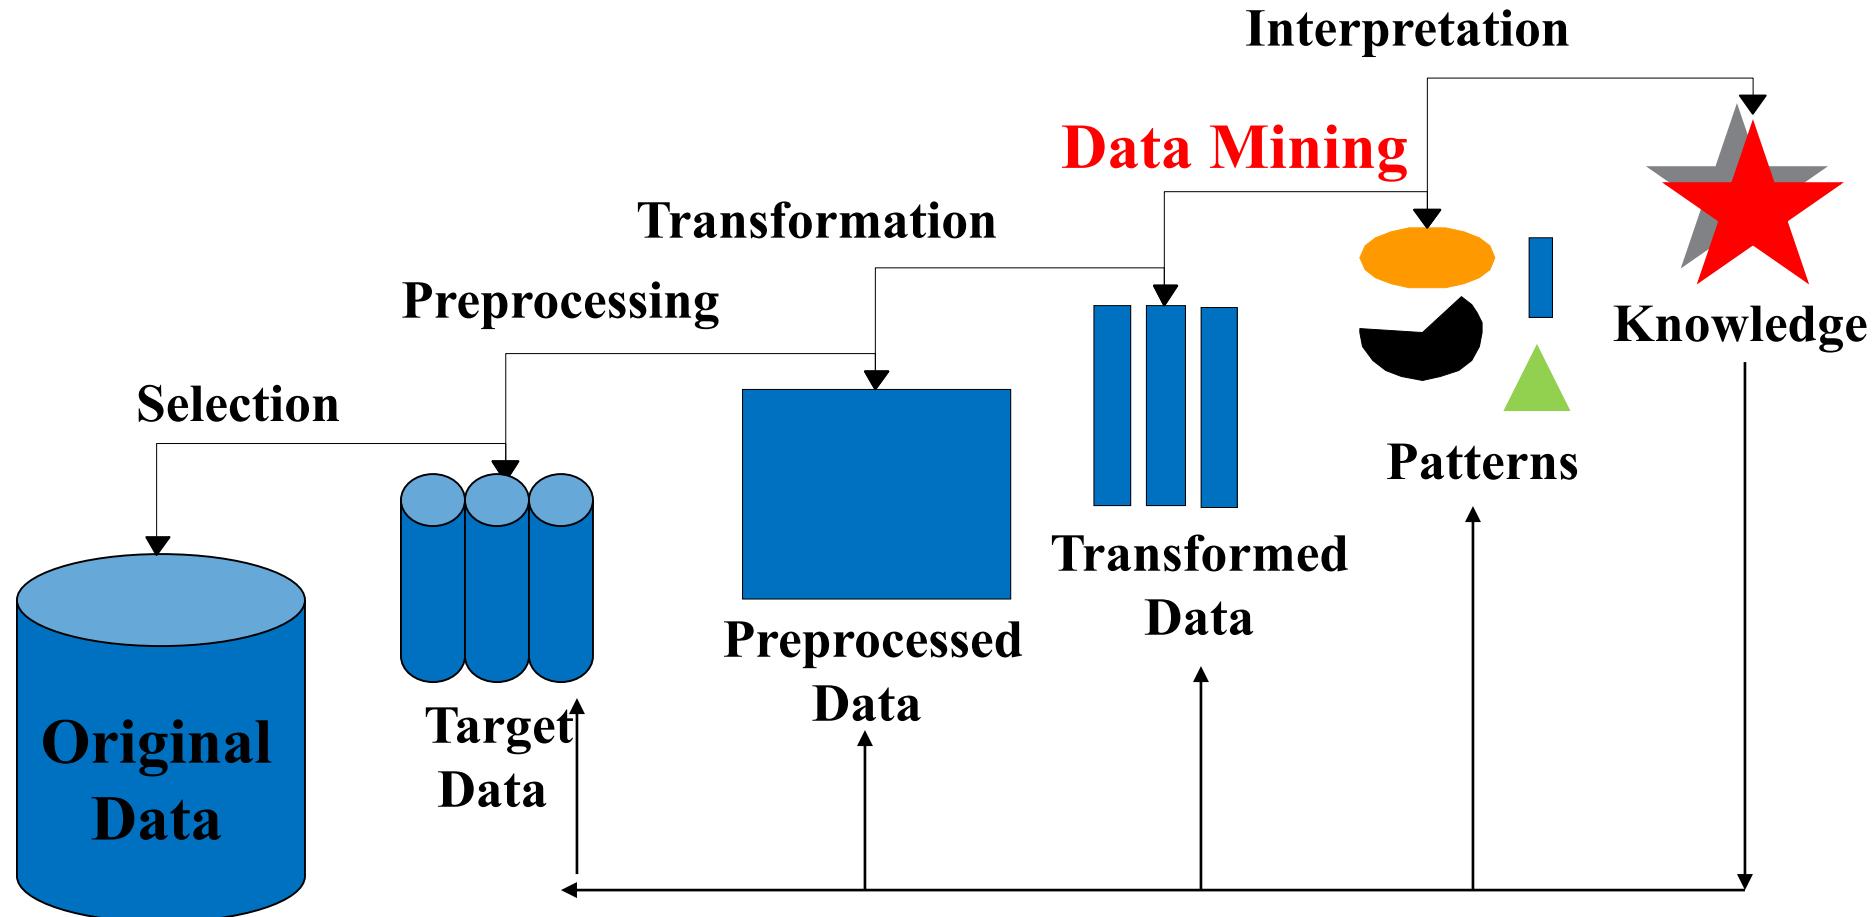
- Proactively and automatically search the data
- uncover problems or opportunities hidden in data relationships
- form computer models based on their findings, and then
- use the models to predict business behavior

## □ A methodology designed to perform knowledge discovery expeditions over the database data with only minimal end-user intervention during the discovery phase



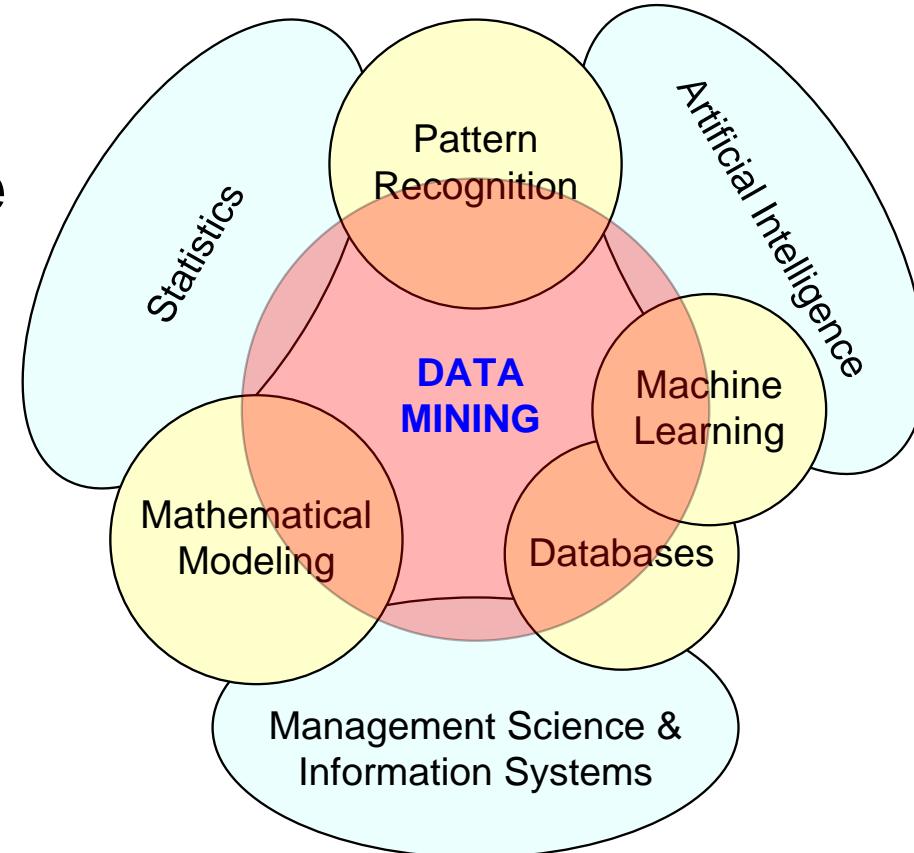
# More popular model for Data Mining- KDD

## □ KDD: Knowledge Discovery from Data

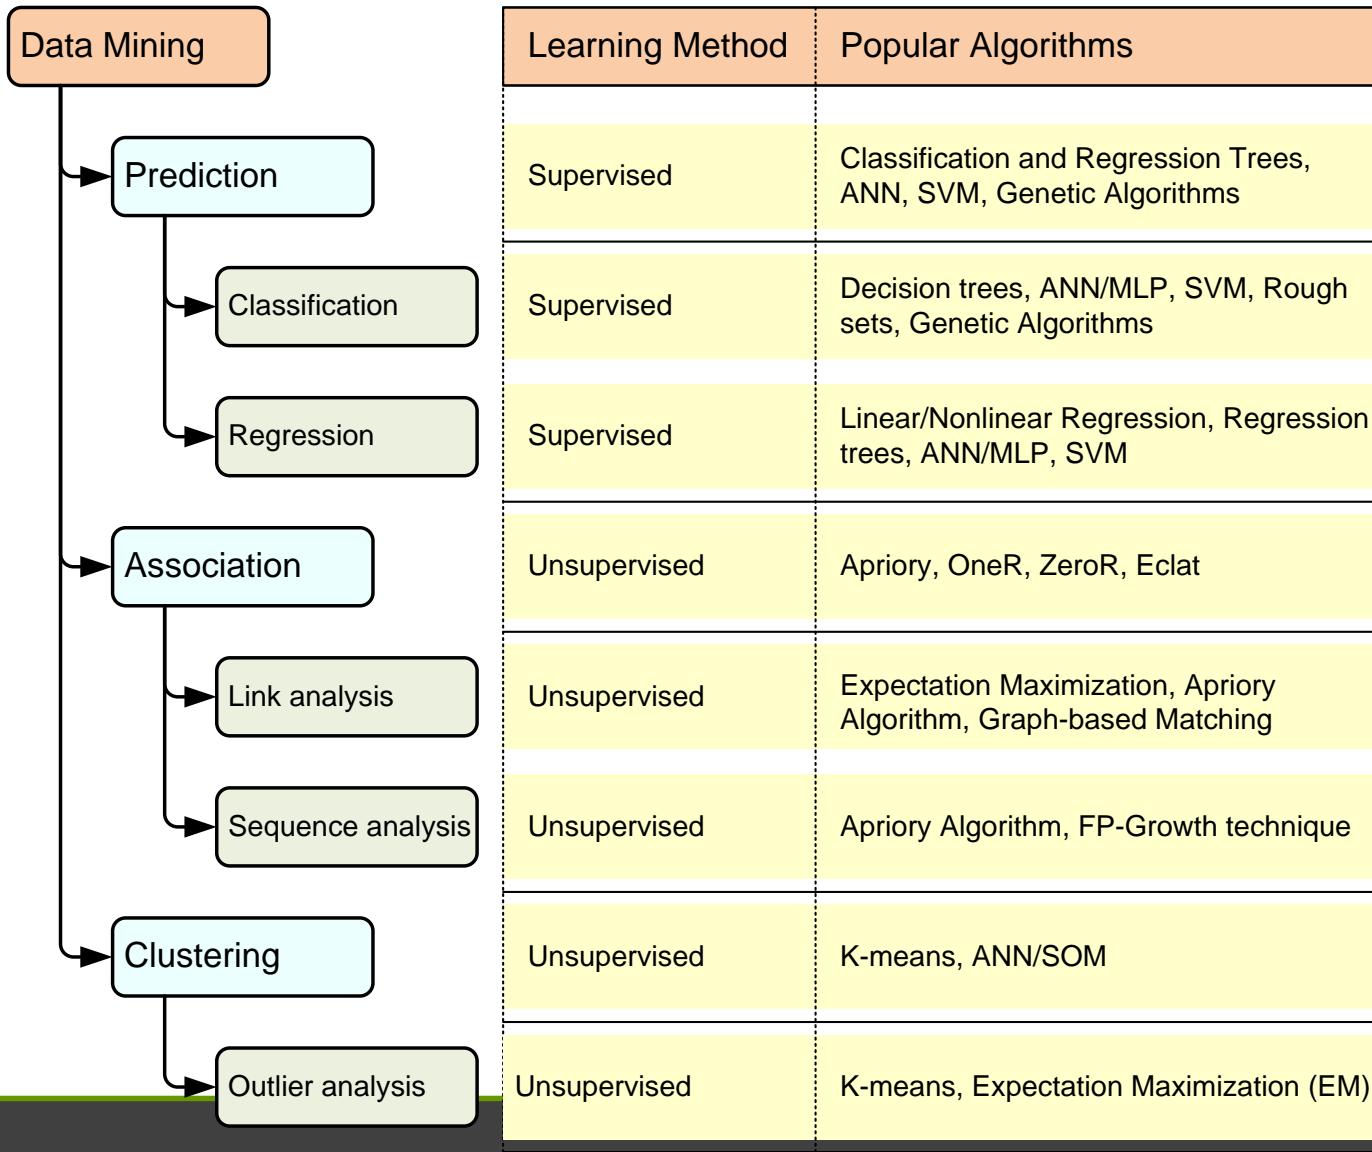


# Origins of Data Mining

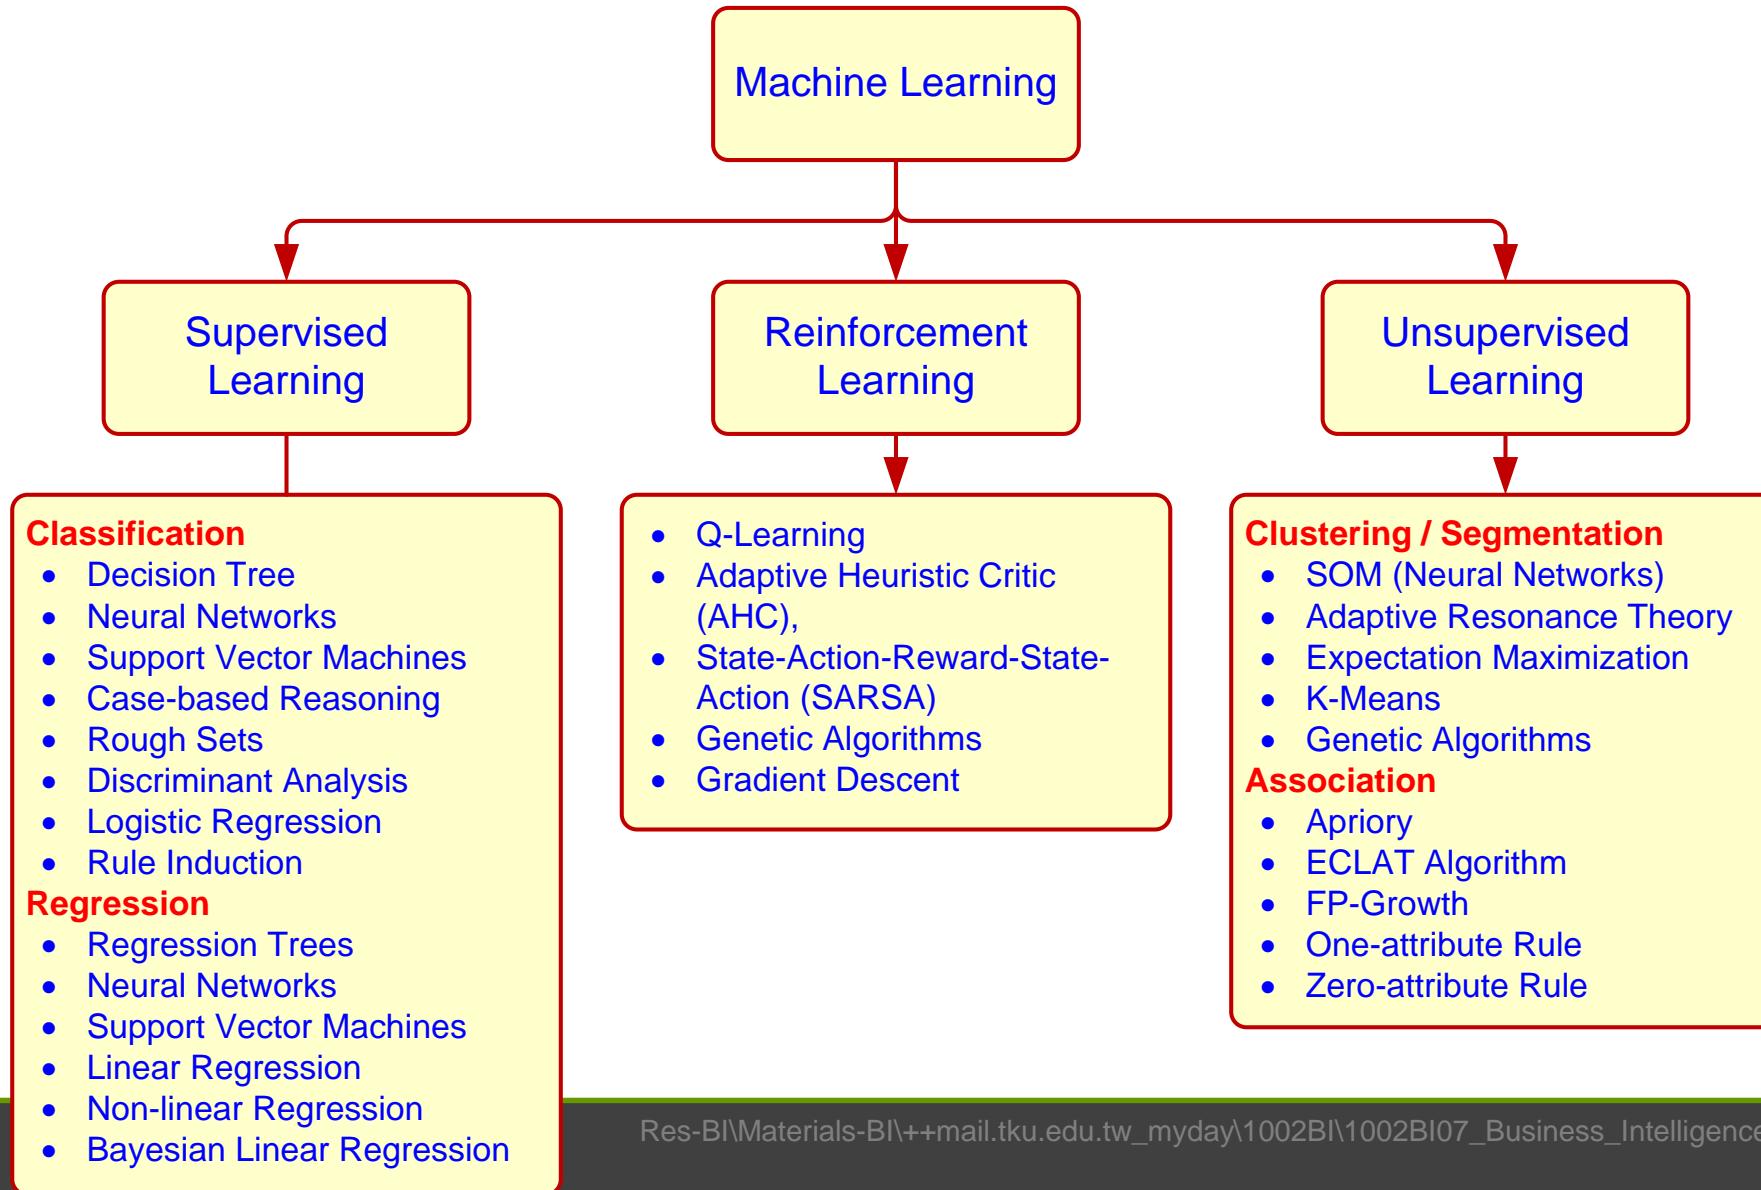
- Draws ideas from AI, machine learning, pattern recognition, statistics, and database systems
- Traditional Techniques may be unsuitable due to
  - Enormity of data
  - High dimensionality of data
  - Heterogeneous, distributed nature of data



# A Taxonomy for Data Mining Tasks



# Machine Learning Methods

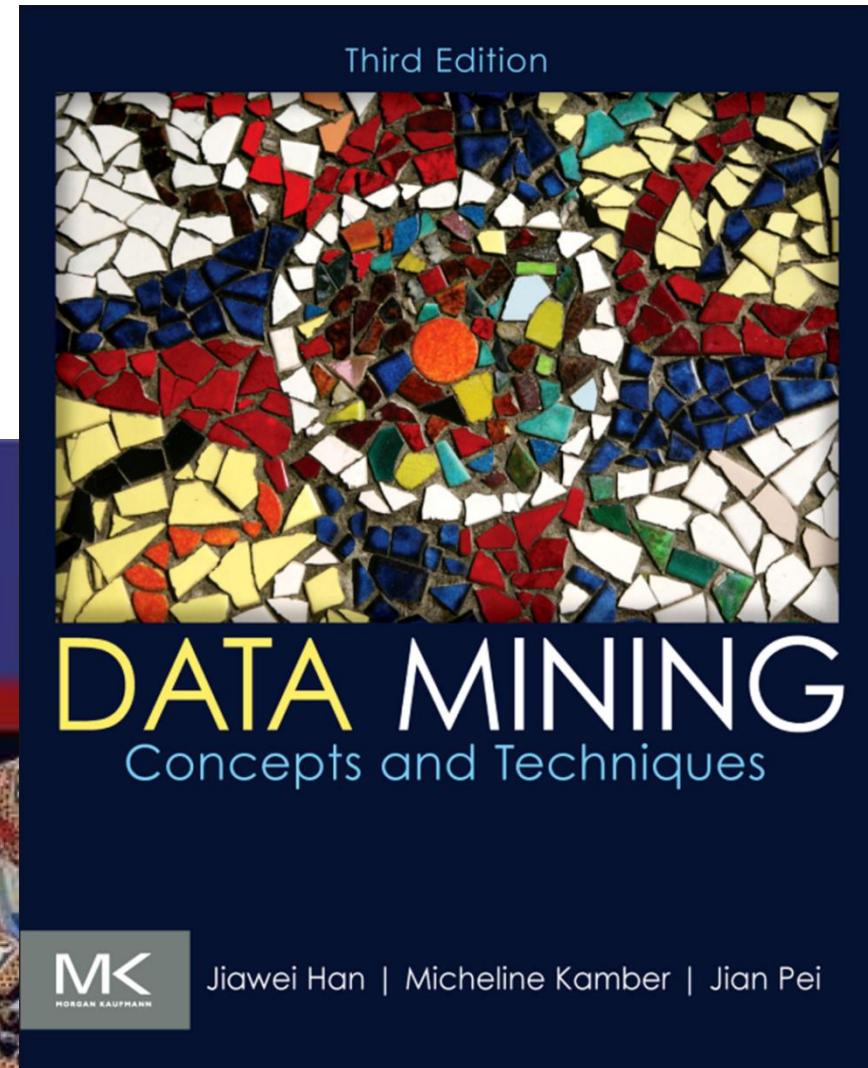
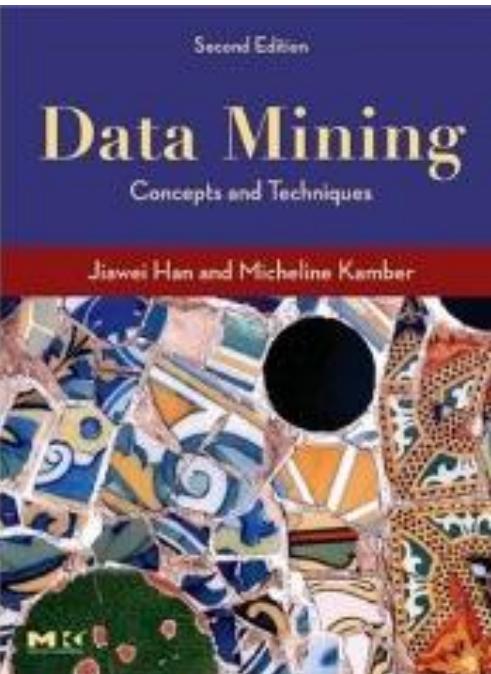
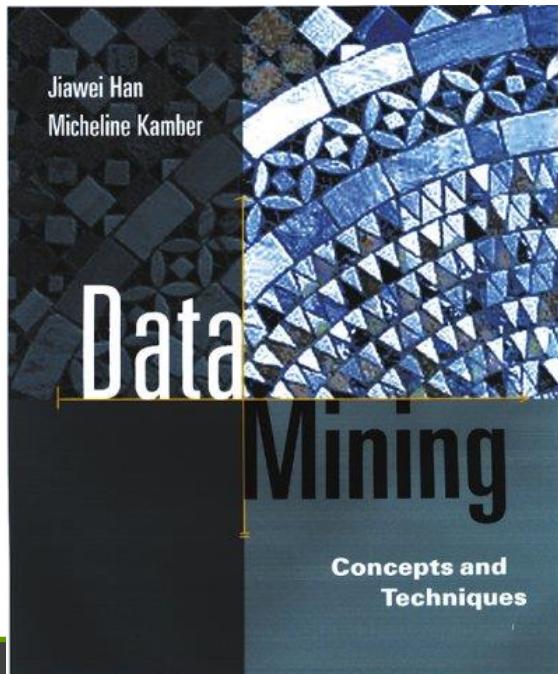


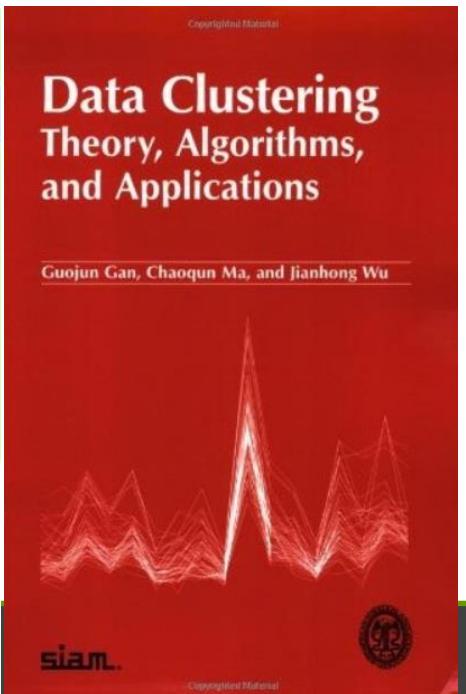
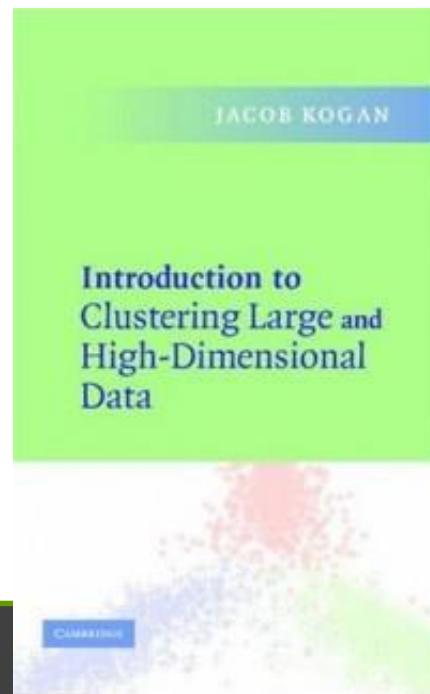
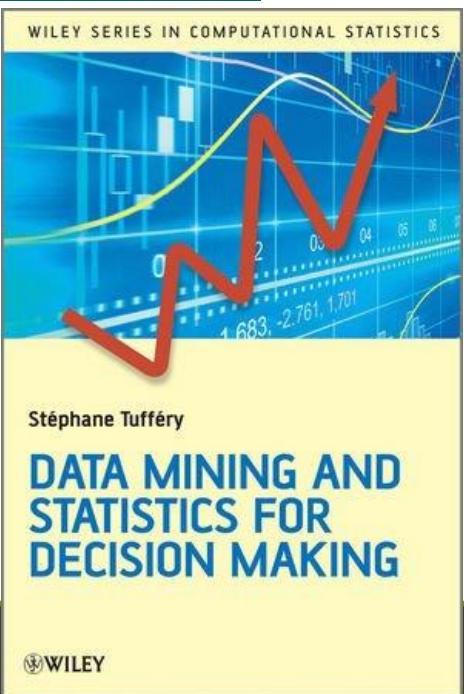
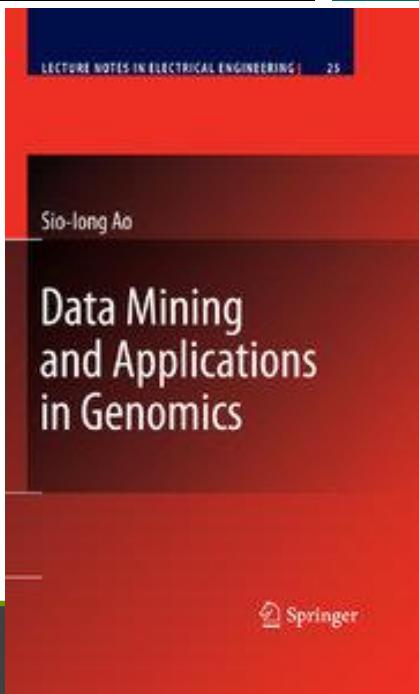
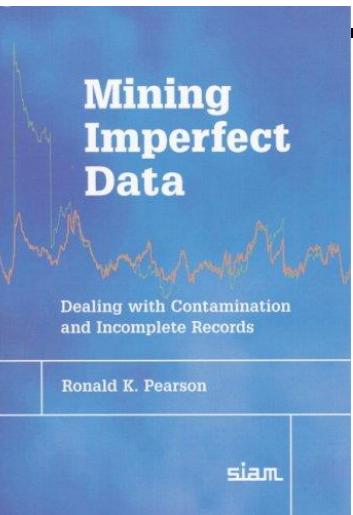
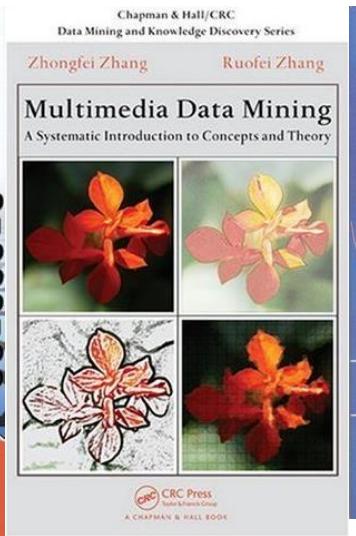
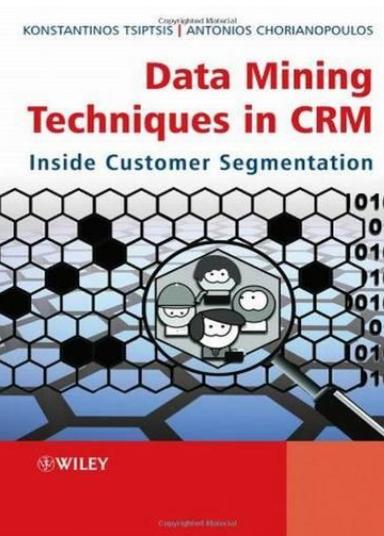
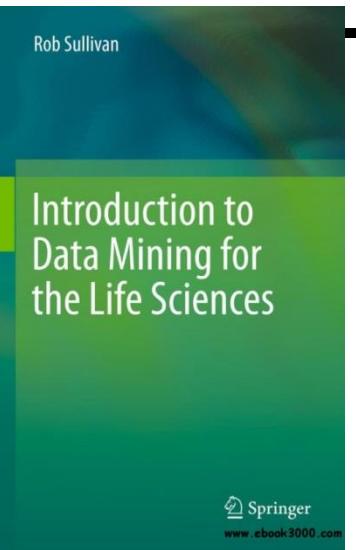
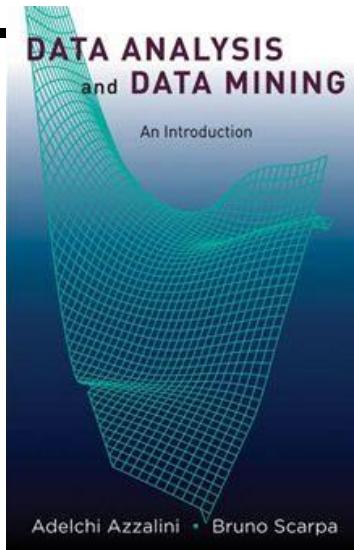
# Additional reference

<http://www.cs.uiuc.edu/homes/hanj/>

## □ Data Mining: Concepts & Techniques

- By Jiawei Han and Micheline Kamber
- 3<sup>rd</sup> edition





# Chapter 5: Go deeper and wider

## □ Overview about the evolution of Data Management & Analytics

- Evolution of DM

- Before IT
- In IT : File system, RDBMS, **ERP**, **DW**, **Big Data**

- Hints –

- Integrate DM with DA (Data Analytics)
  - ✓ Data – diverse (from structured to unstructured) and huge (from single to distributed)

## □ SQL again

## □ Final (if you want to share)

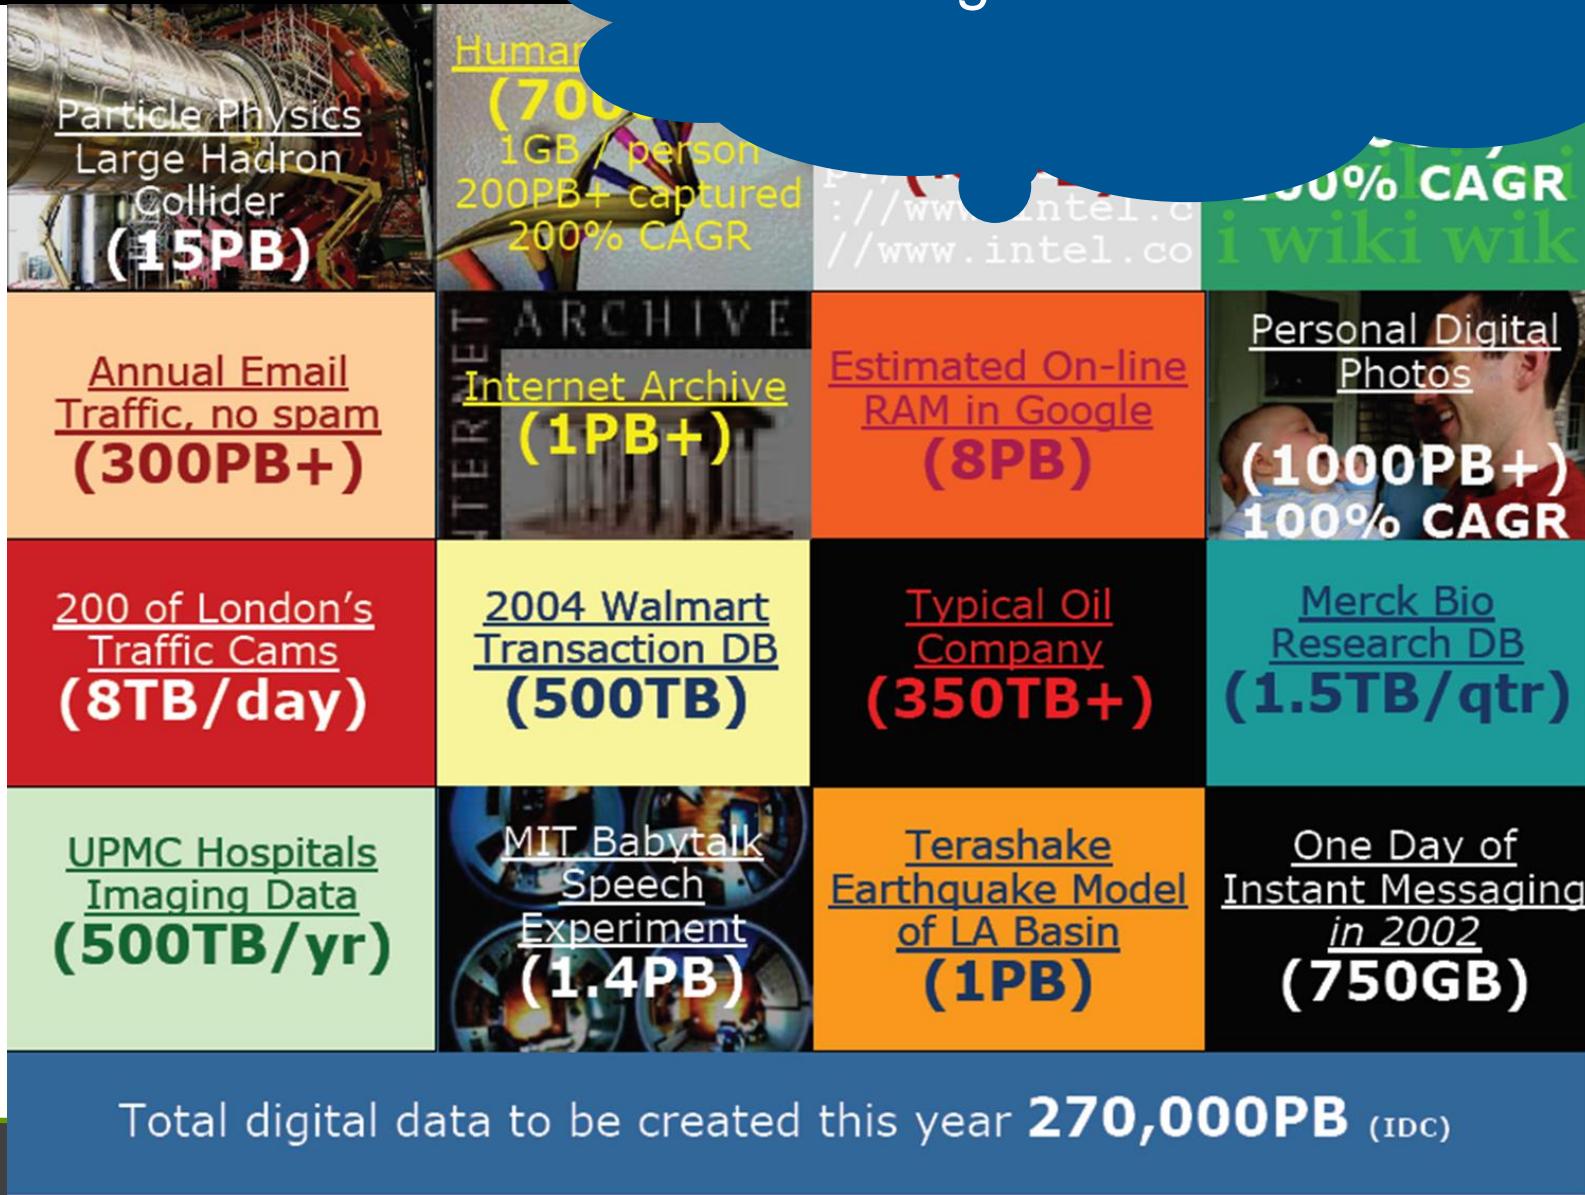
# Big Data is coming

- Ambitious to manage huge and diverse data – 3Vs: Volume, Variety, Velocity
  - 2.5 quintillion bytes of data are generated every day!
    - A quintillion is  $10^{18}$
  - coming from many quarters, like Social media sites, Sensors, Digital photos, Business transactions, Location-based data, Web data, e-commerce, Bank/Credit Card, ...
  - For information: Google processes 20 PB a day (2008)
    - <http://www.worldwidewebsize.com/>
    - ✓ “The Indexed Web contains at least **9.18 billion pages** (Sunday, 09 December, 2012).”
  - For science: NASA & Hubble scope



# Here comes Big Data

Huge and Diverse



Fundamental ideas to understand BD:  
Distributed for huge + Redundancy for data safety

## □ Making decisions!

- For governance: Edward Snowden



# The whistleblower

I can't allow the US government to destroy privacy and basic liberties

**the guardian**  
guardian.co.uk

A composite image featuring a portrait of Edward Snowden on the left and a headline from The Guardian on the right. The headline reads "The whistleblower" in large yellow letters, followed by the quote "I can't allow the US government to destroy privacy and basic liberties" in white. Below the quote is the "the guardian" logo in white on a dark blue background, with the website "guardian.co.uk" in smaller text below it.

# Are you pleased with the response speed?

- Found about 307,000,000 items (using 0.28 second)

Google search results for "google number of wen pages". The search bar shows the query. Below it, the navigation bar includes "网页" (selected), 图片, 地图, 购物, 应用, 更多, 搜索工具. A red box highlights the search result summary: "找到约 307,000,000 条结果 (用时 0.28 秒)". The main content area displays the search results:

显示以下查询字词的结果: [google number of web pages](#)  
仍然搜索: [google number of wen pages](#)

[WorldWideWebSize.com | The size of the World Wide Web \(The ...\)](#)  
[www.worldwidewebsize.com/](http://www.worldwidewebsize.com/) - 网页快照 - 翻译此页

The estimated minimal size of the indexed World Wide Web is based on the estimations of the numbers of pages indexed by Google, Bing, Yahoo Search.

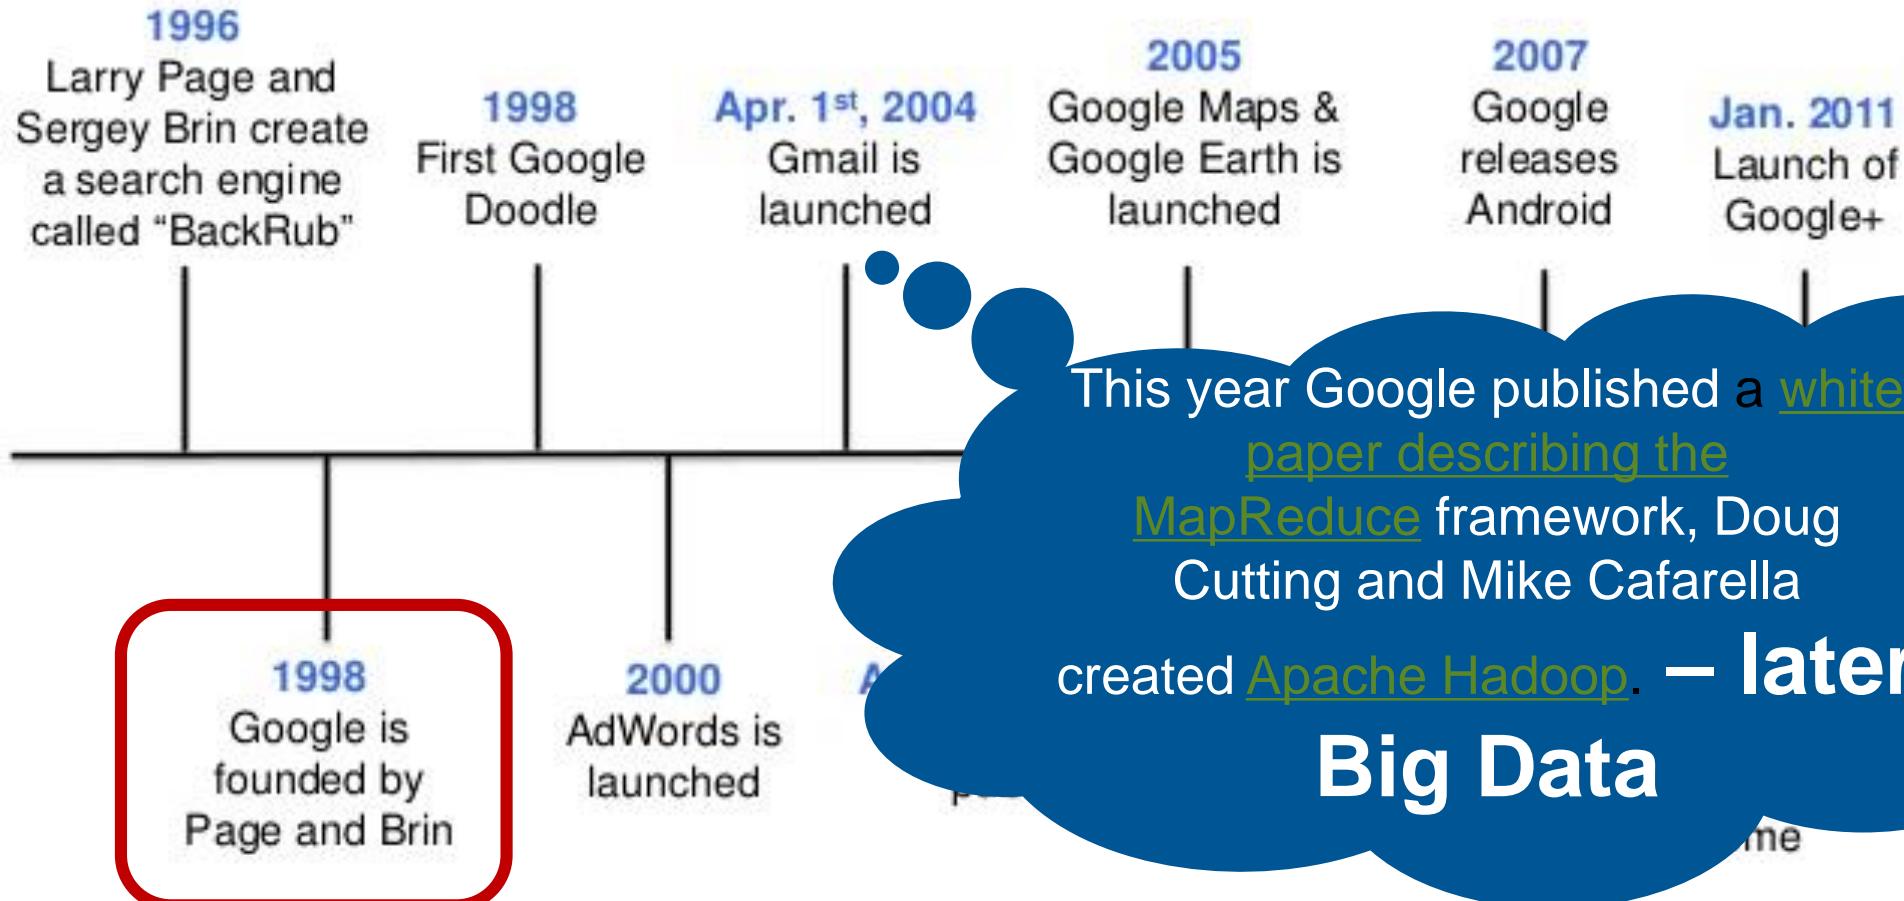
[Last Month](#) - [Last Year](#) - [Last Two Years](#) - [The Dutch Indexed Web](#)

[We knew the web was big... | Official Google Blog](#)  
[googleblog.blogspot.com/.../we-knew-web-was-bi...](http://googleblog.blogspot.com/.../we-knew-web-was-bi...) - 网页快照 - 翻译此页

25 Jul 2008 – The first Google index in 1998 already had 26 million pages, and by 2000 the ... So how many unique pages does the web really contain?

# Google triggers other searches ...

## □ CBIR, XML, ...

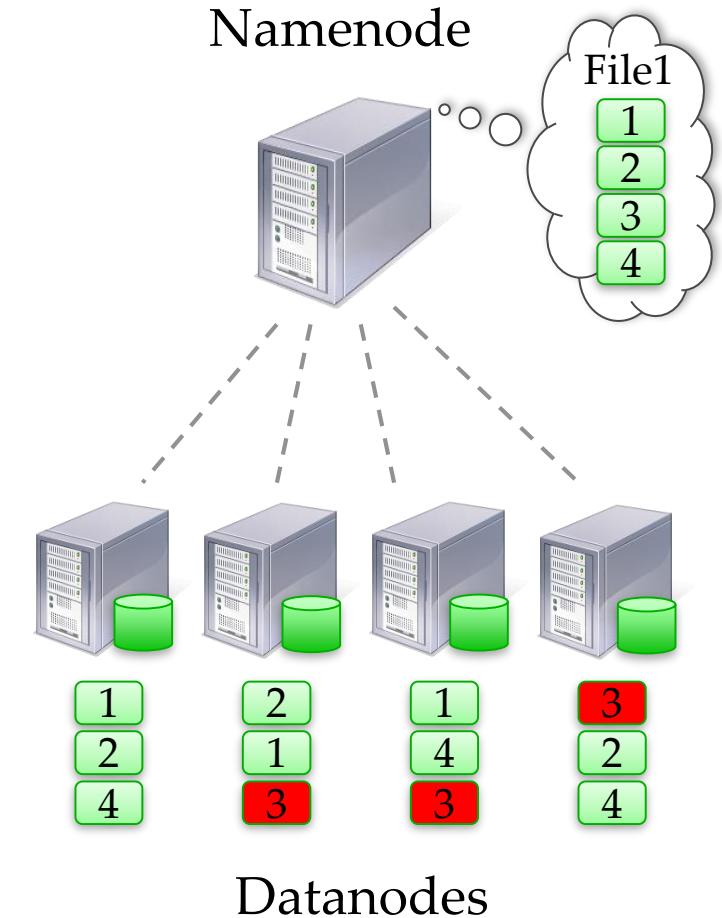


# HDFS/Hadoop with Cluster system



## □ Hadoop Distributed file system (HDFS)

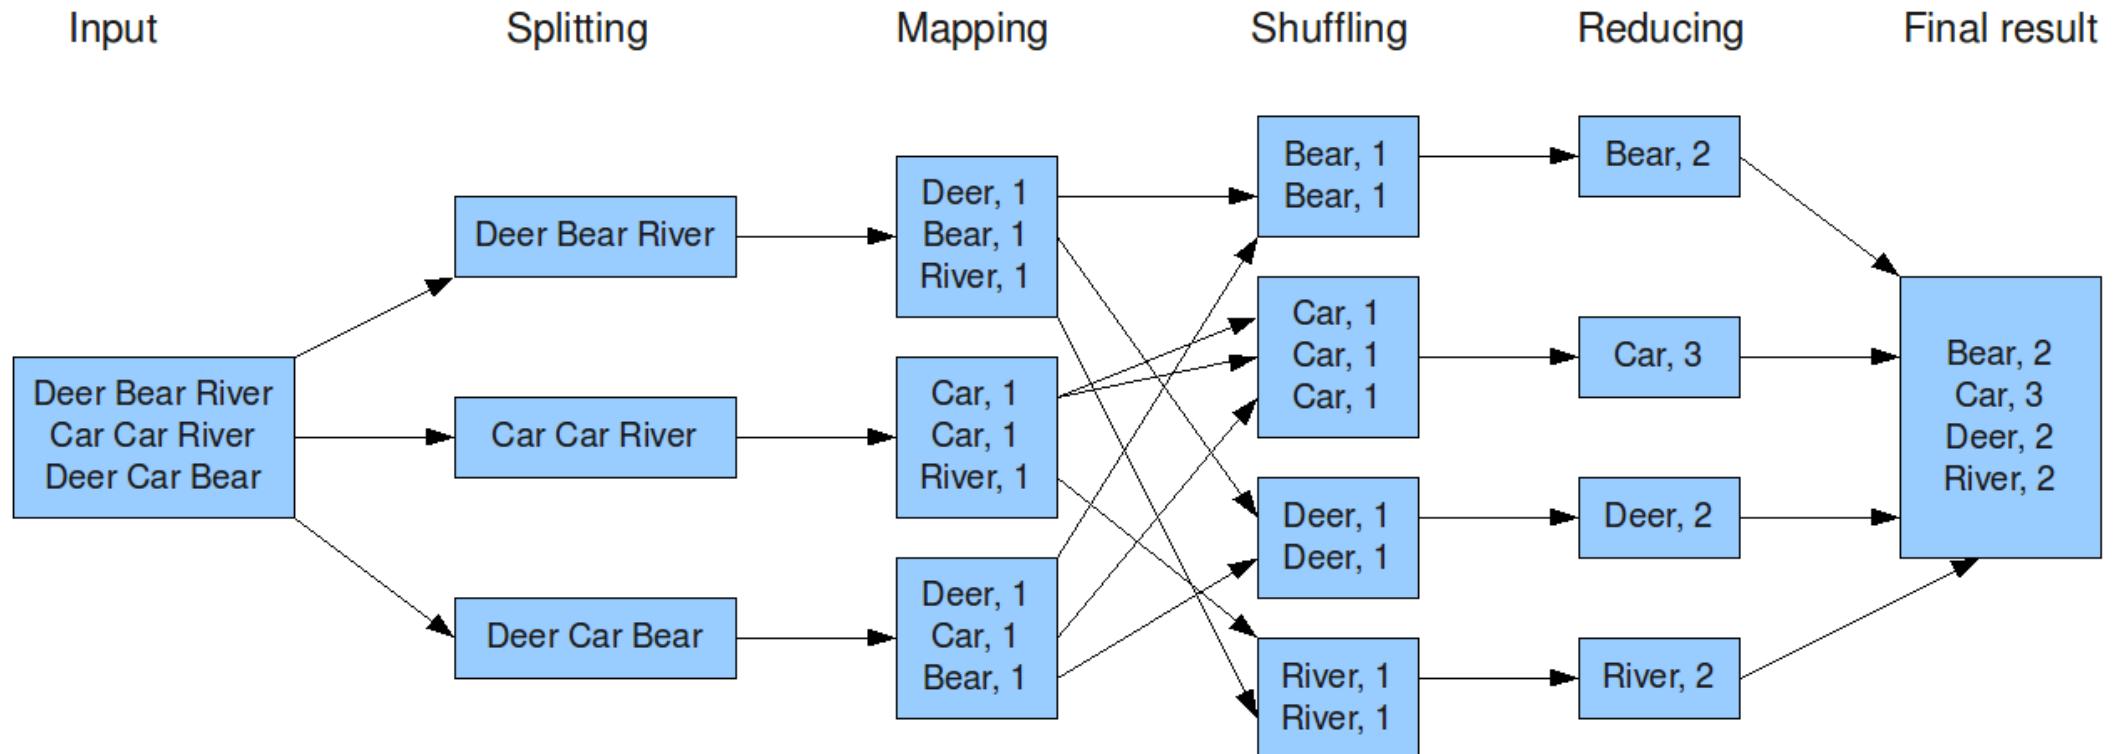
- Single namespace for entire cluster
  - Namenode stores metadata (file names, locations, etc)
- Replicates data 3x for fault-tolerance
  - Files split into 128MB blocks
  - Blocks replicated across several datanodes (often 3)
- Optimized for large files, sequential reads
- Files are append-only



# Map/Reduce - YARN

M/R is too naïve to manipulate data – word counting. We need more flexible query for analytics

The overall MapReduce word count process



# Computing model

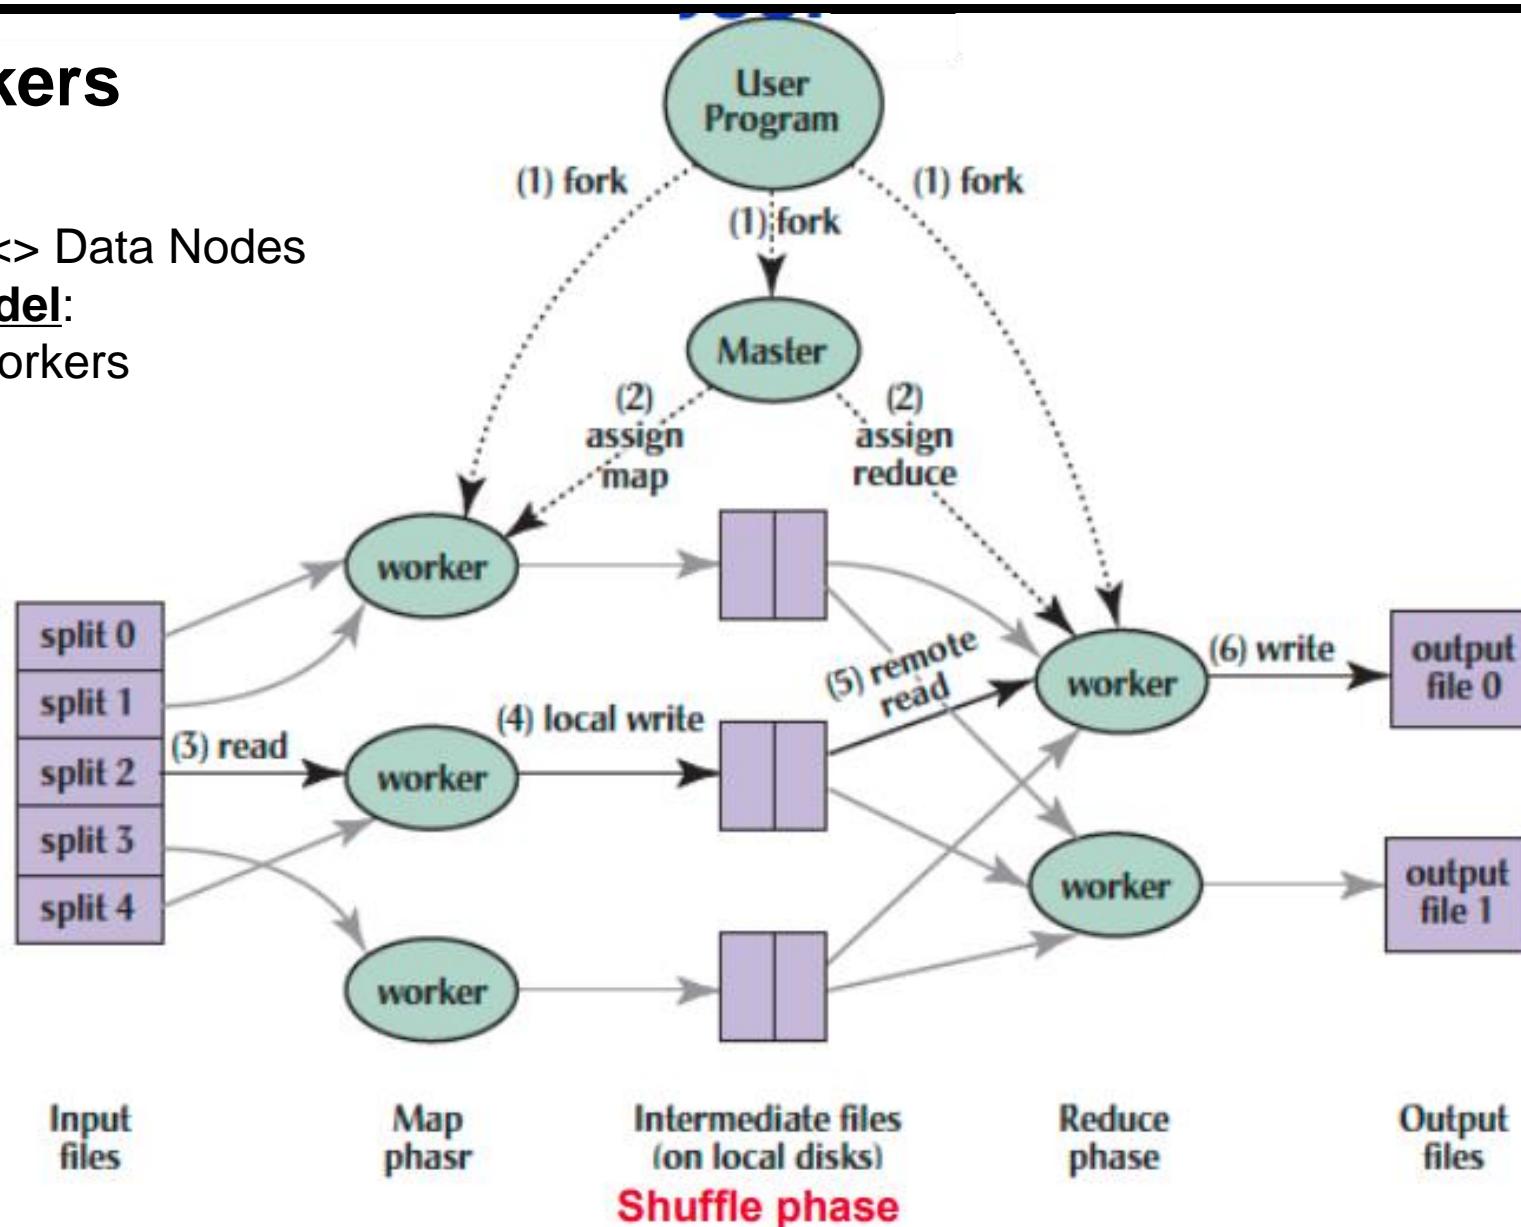
## □ Master – Workers

### Data Storage:

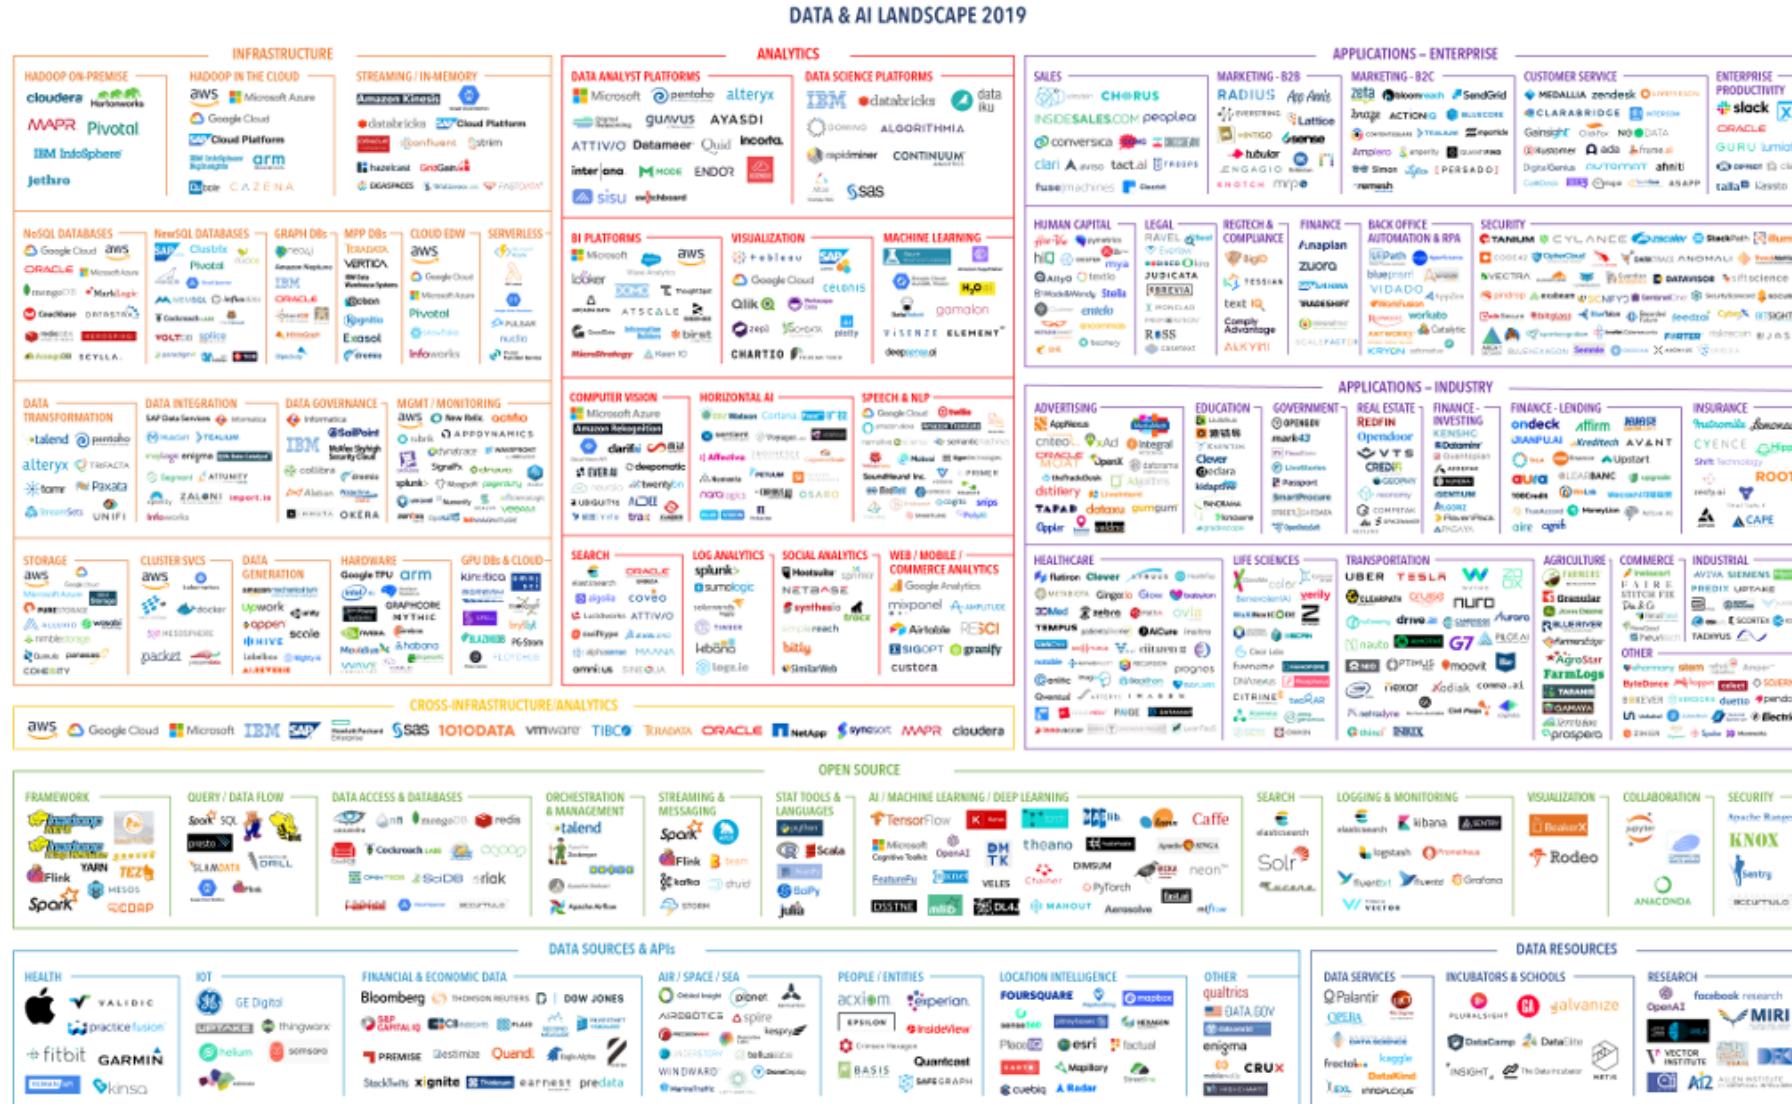
Name node <> Data Nodes

### Computing Model:

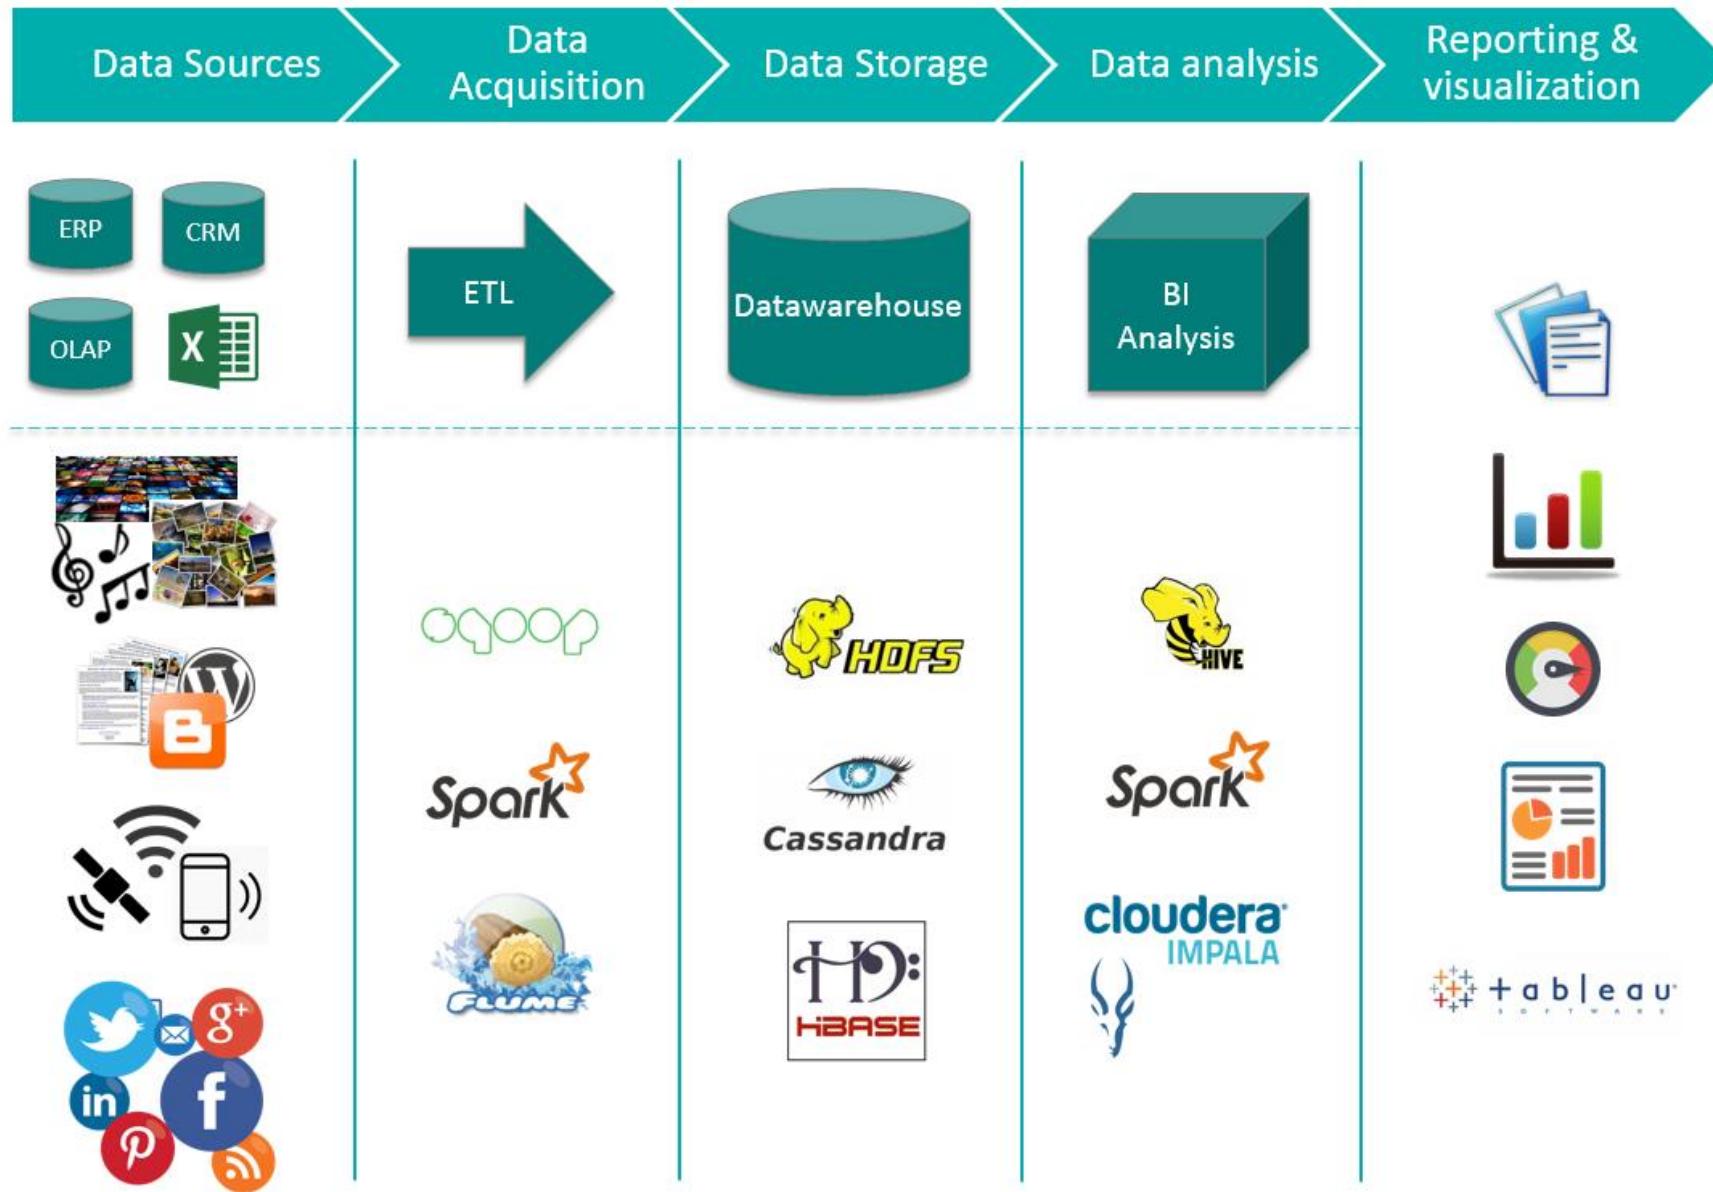
Master <> Workers



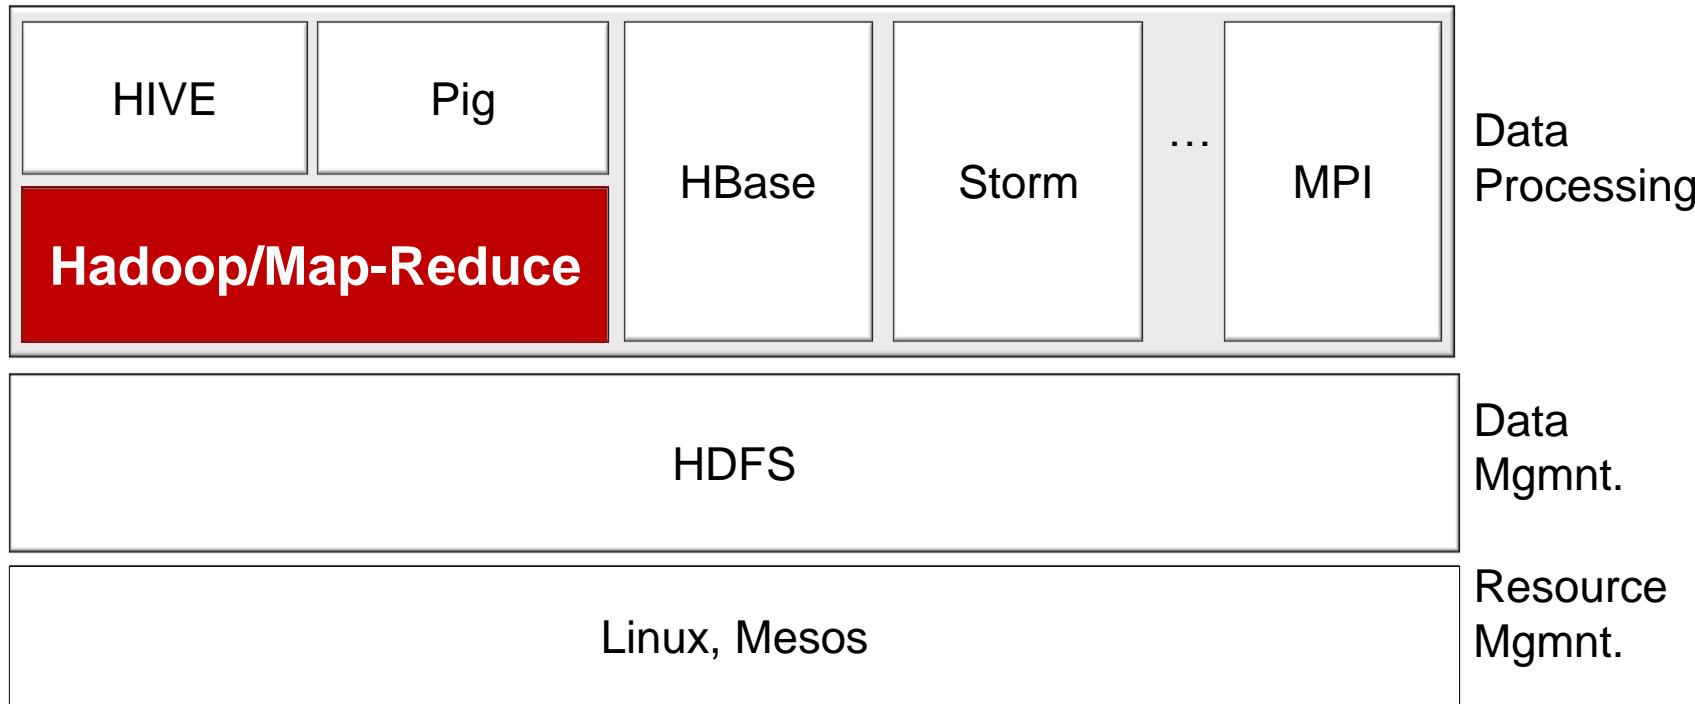
# The 2019 Big Data Landscape



# Integrating Data + Analytics is also carried in BD

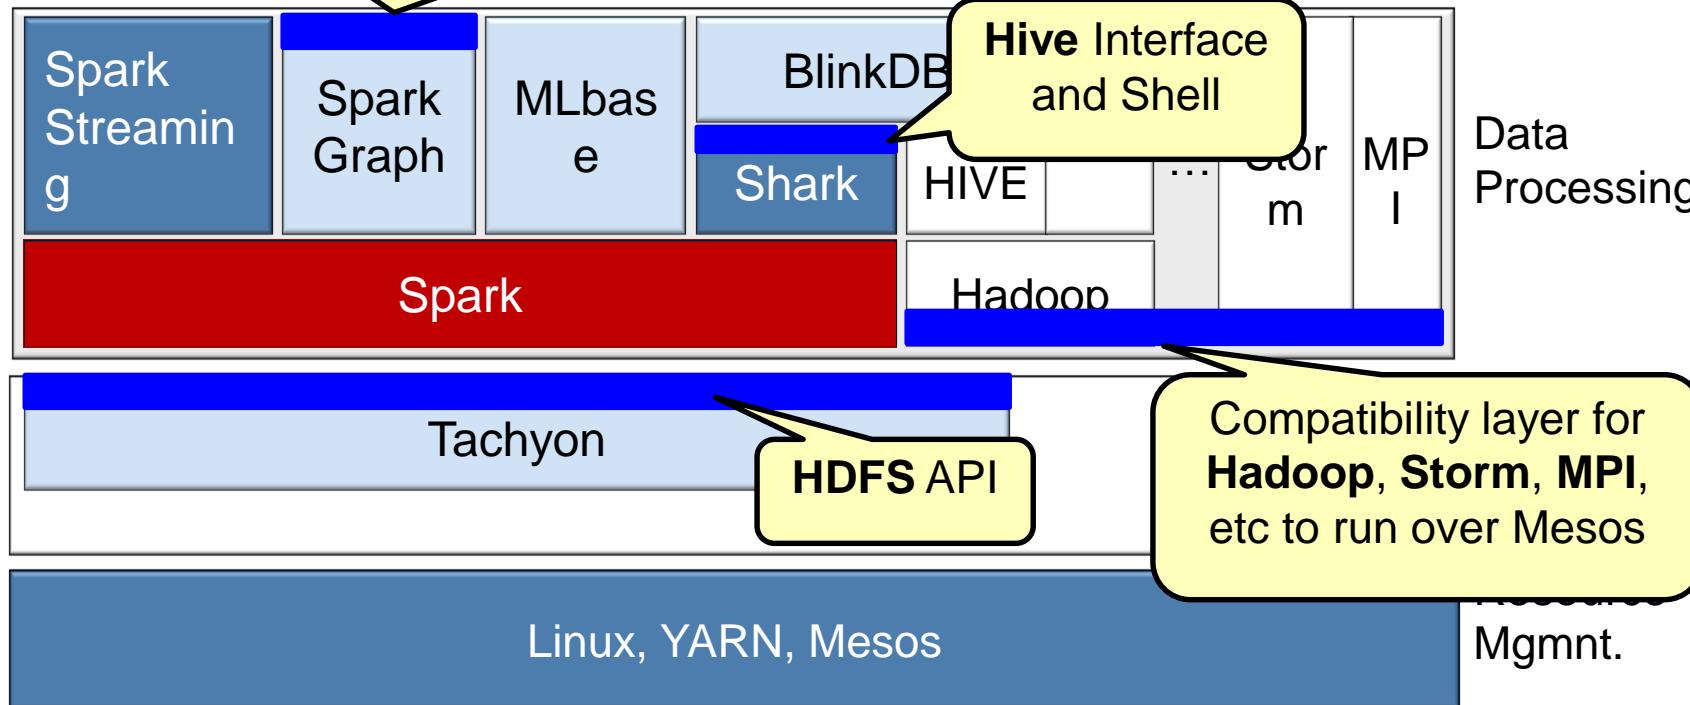


# 1<sup>st</sup> generation



- The data processing is based on M/R model, which required intermediate result **be kept into/out of disks** – low performance

# 2<sup>nd</sup> generation – Spark – In-Memory processing for higher perf.



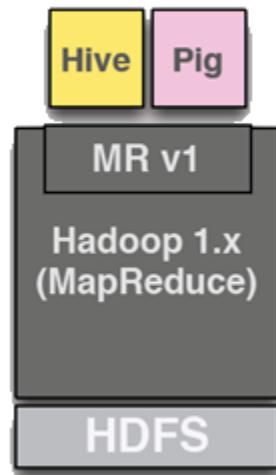
- The data processing is based on Spark model, in which intermediate result is **kept in memory** – higher performance
- **Shark → Spark SQL**, and **SQL is re-found effective to represent business logics for data analytics!**
- **Spark is further extended to Support existing interfaces now**

## □ Recently I believe there should be the 4<sup>th</sup> classic software

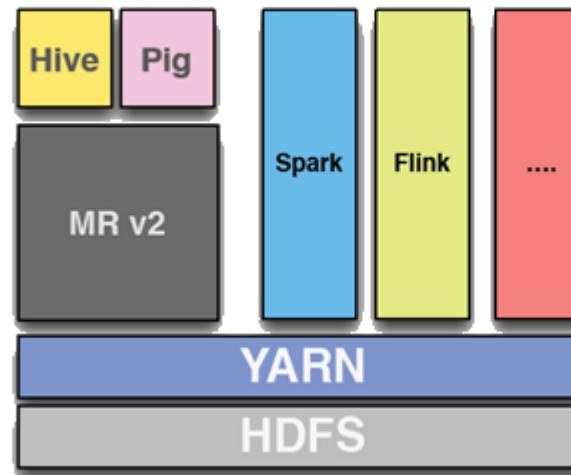
### ■ Big Data

- Hadoop/MapReduce, HIVE, Spark, HAWQ, etc.

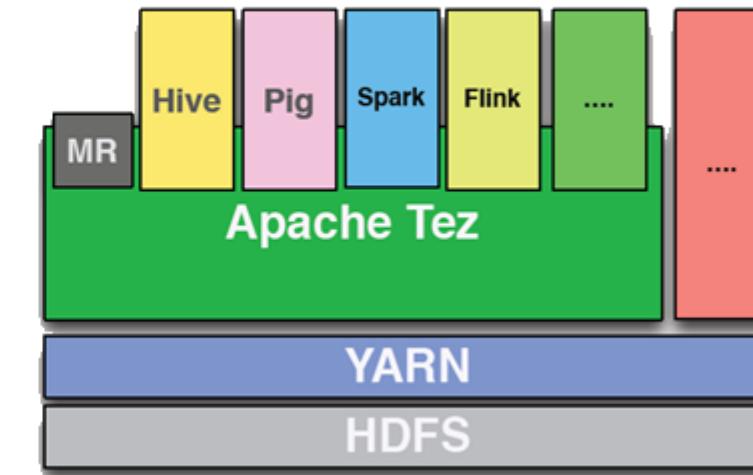
Hadoop 1



Hadoop 2

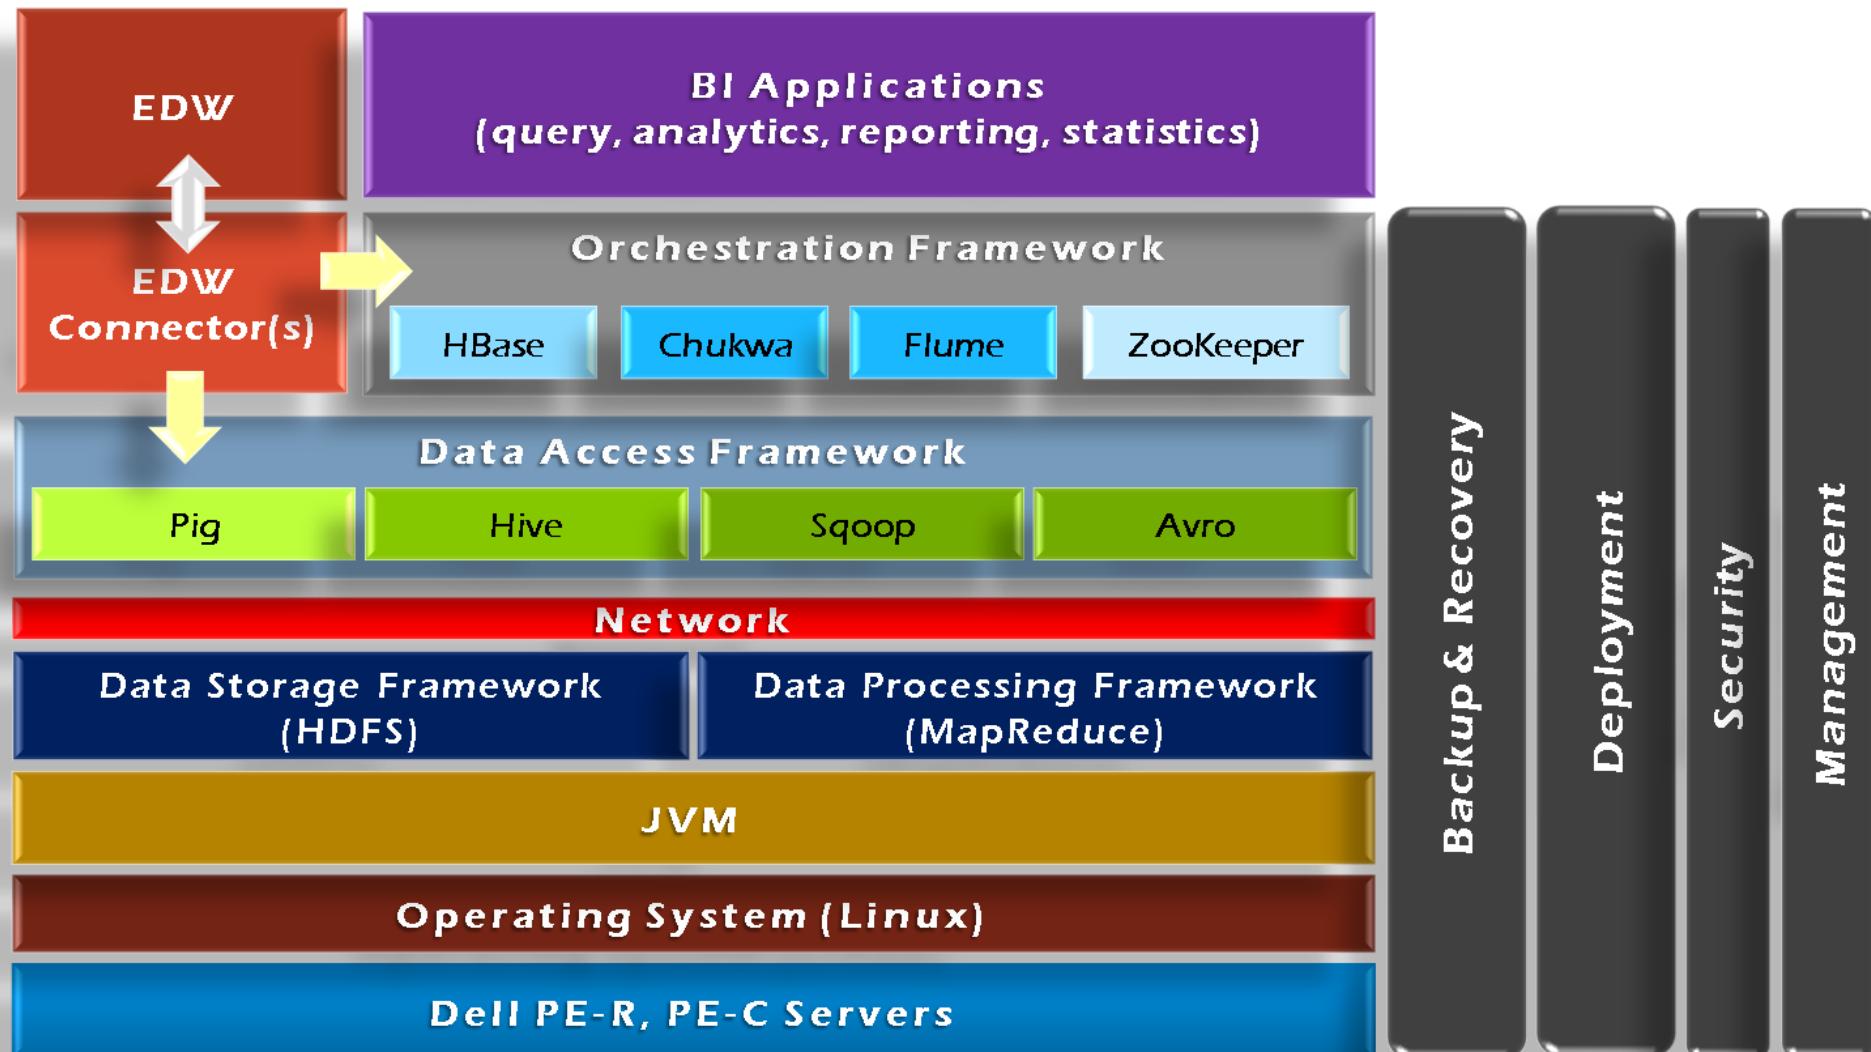


Hadoop 2 + Tez



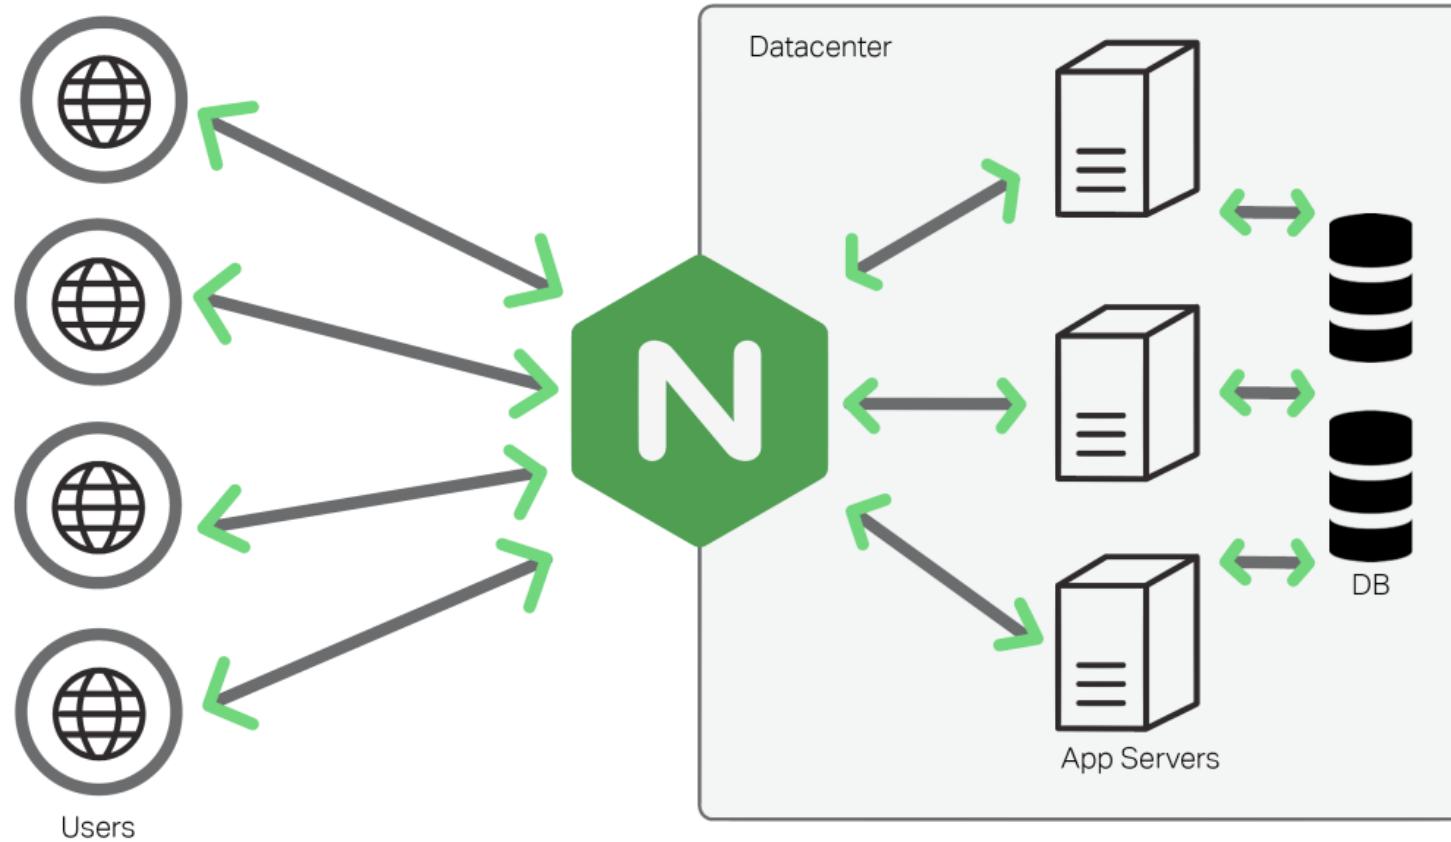
<https://stackoverflow.com/questions/39411888/why-would-someone-run-spark-flink-on-tez>

# Hadoop Framework Tools



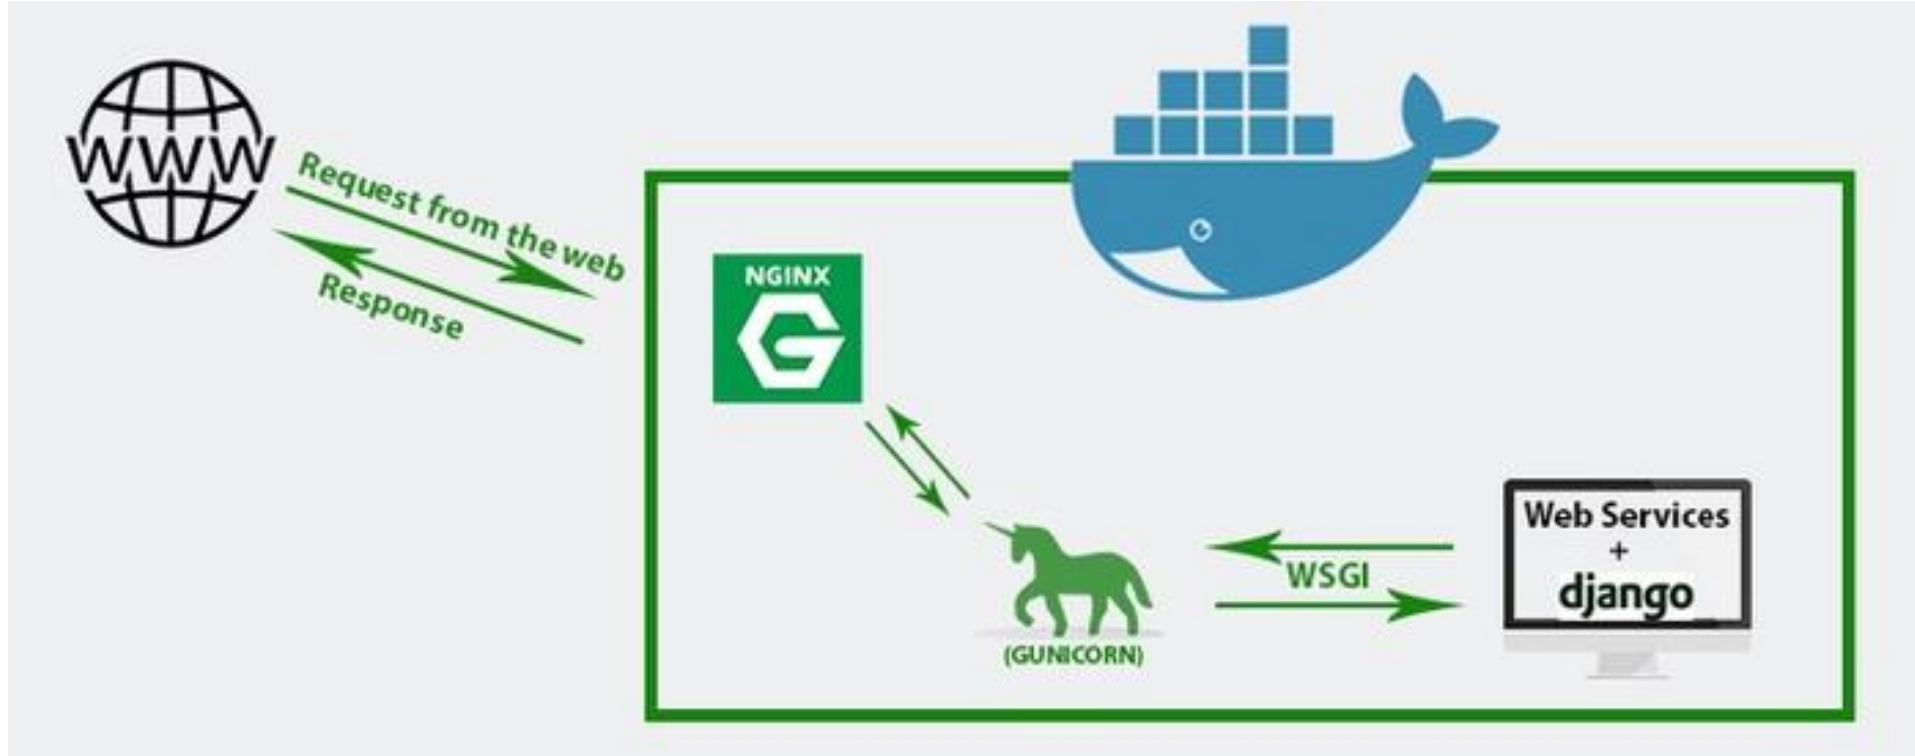
# Application architecture becomes more complex In Big Data age

□ + Nginx



# Application architecture becomes more complex In Big Data age

- Linux + Docker, Nginx, Django etc.

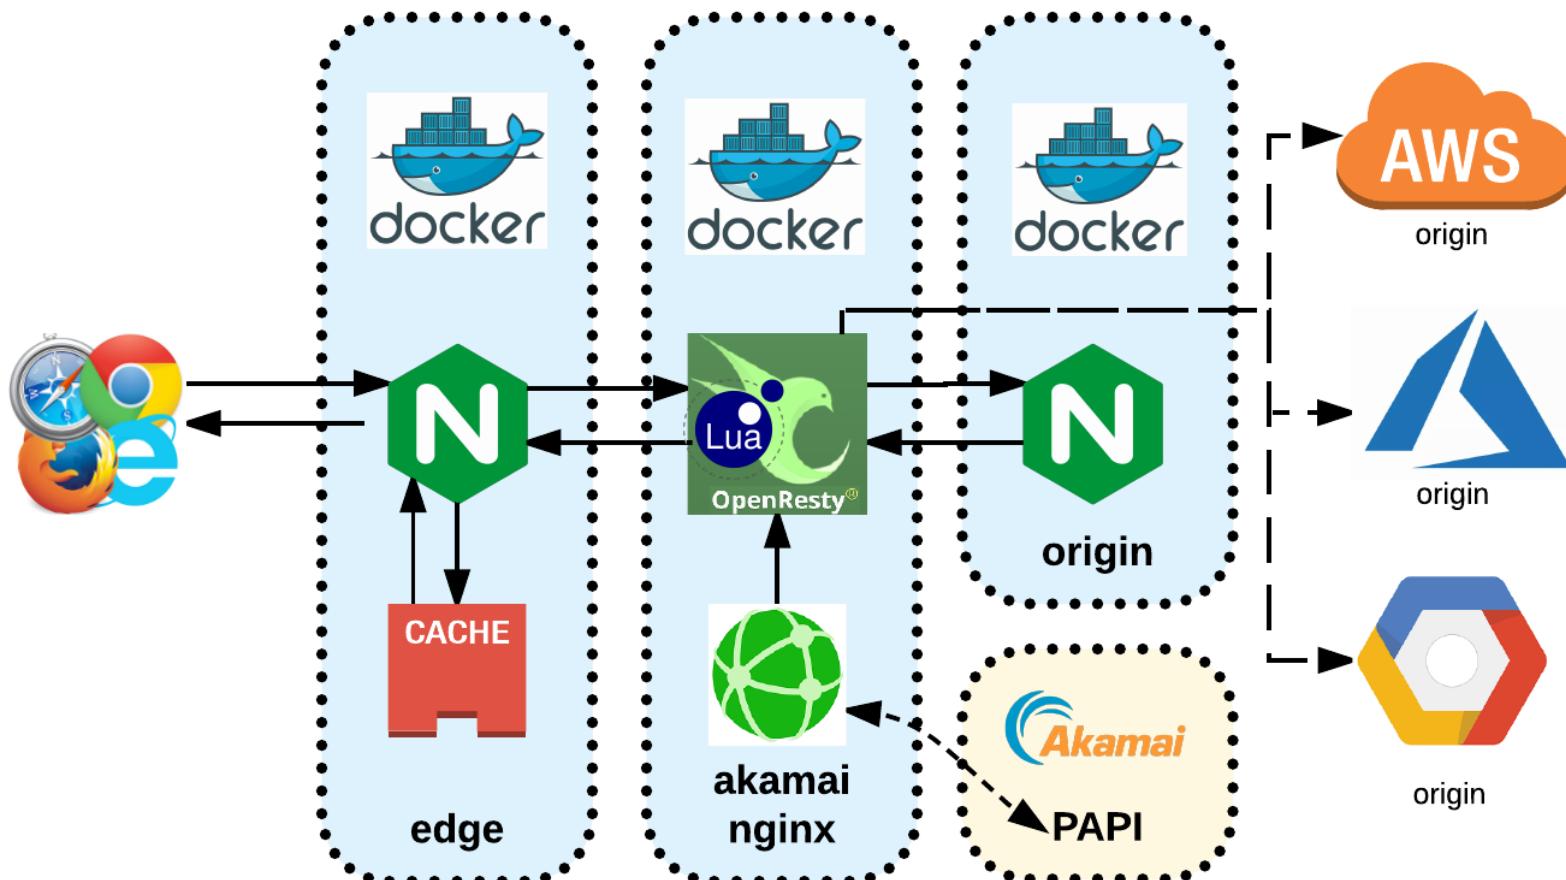


<https://www.codementor.io/samueljames/nginx-setting-up-a-simple-proxy-server-using-docker-and-python-django-f7hy4e6jv>

# Application architecture becomes more complex In Big Data age

- Linux + Docker, Nginx, Lua etc.

## akamai-nginx request flow



## □ Many Cloud platforms



# Chapter 5: Go deeper and wider

- Overview about the evolution of Data Management & Analytics
- SQL again
  - SQL → NoSQL(HIVEQL/Spark SQL) → NewSQL
- Final (if you want to share)

# You've known something about SQL, right?

## □ I hope you really are ☺

- SQL (like other computer languages) is formally defined in a notation called "Backus-Naur Form"

## □ Subset of SQL may be

```
<query specification> ::=  
    SELECT [ <set quantifier> ] <select list> <table expression>  
  
<select list> ::=  
    <asterisk>  
    | <select sublist> [ { <comma> <select sublist> }... ]  
  
<select sublist> ::= <derived column> | <qualifier> <period> <asterisk>  
  
<derived column> ::= <value expression> [ <as clause> ]  
  
<as clause> ::= [ AS ] <column name>  
  
<table expression> ::=  
    <from clause>  
    [ <where clause> ]  
    [ <group by clause> ]  
    [ <having clause> ]  
    <from clause> ::= FROM <table reference> [ { <comma> <table reference> }... ]
```



# How is SQL translated and executed?

---

## □ Math Expression → AST

- Helpful to understand SQL processing

## □ SQL Processing

- SQL → Parse Tree [[Math → AST](#)]

- Parse Tree → Canonical RA

- Canonical RA is converted into equivalent RAs (for Algebraic optimization)

- RA → PQP (Physical Query Plan) [[Math → AST](#)]

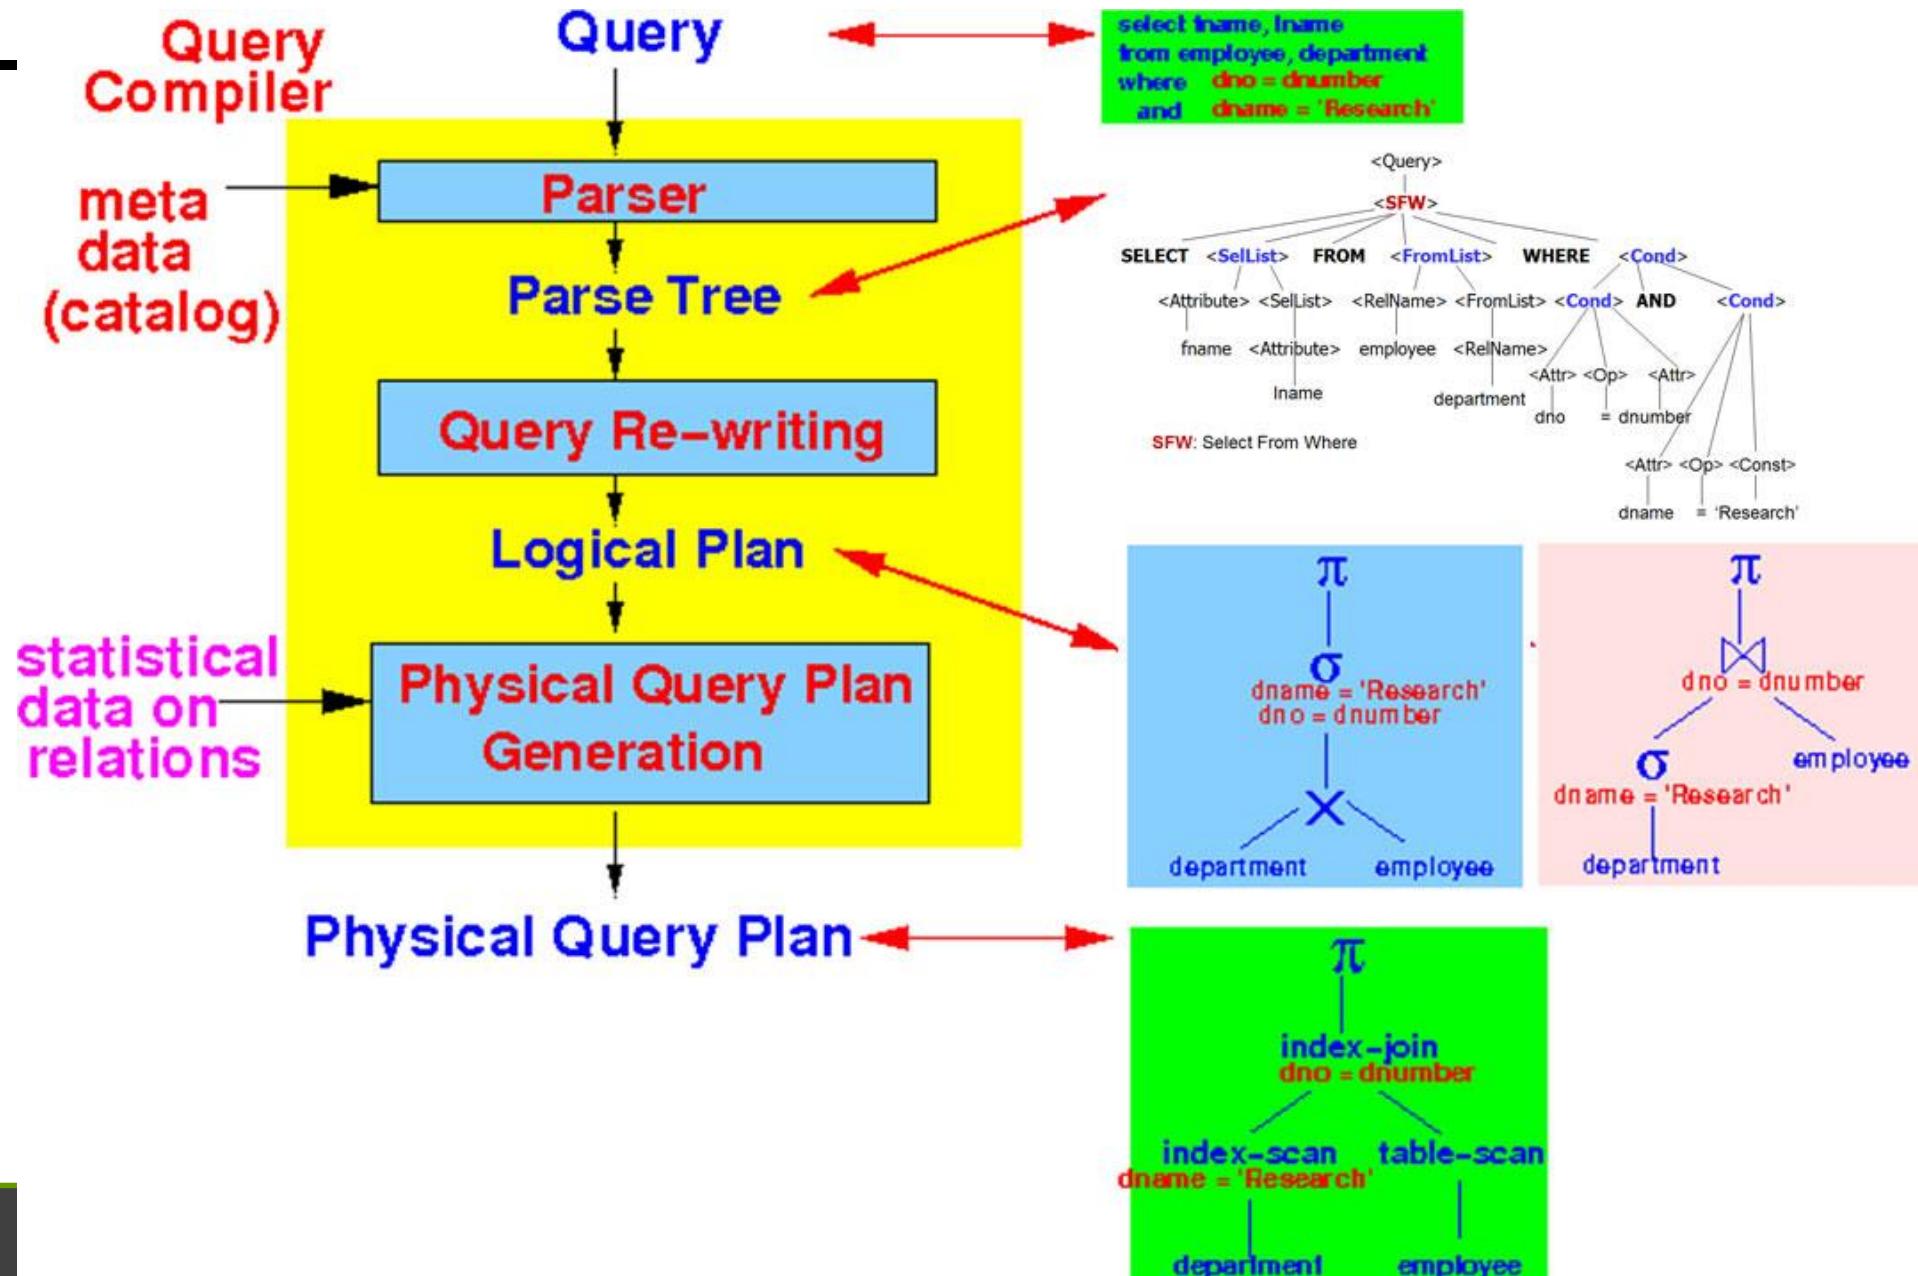
- RA could be mapped to many PQPs

- With help of catalog information, physical optimization

- The optimal PQP is executed

- Materialize or Pipeline

# In short



# NoSQL - Not Only SQL

---

## □ Stands for Not Only SQL

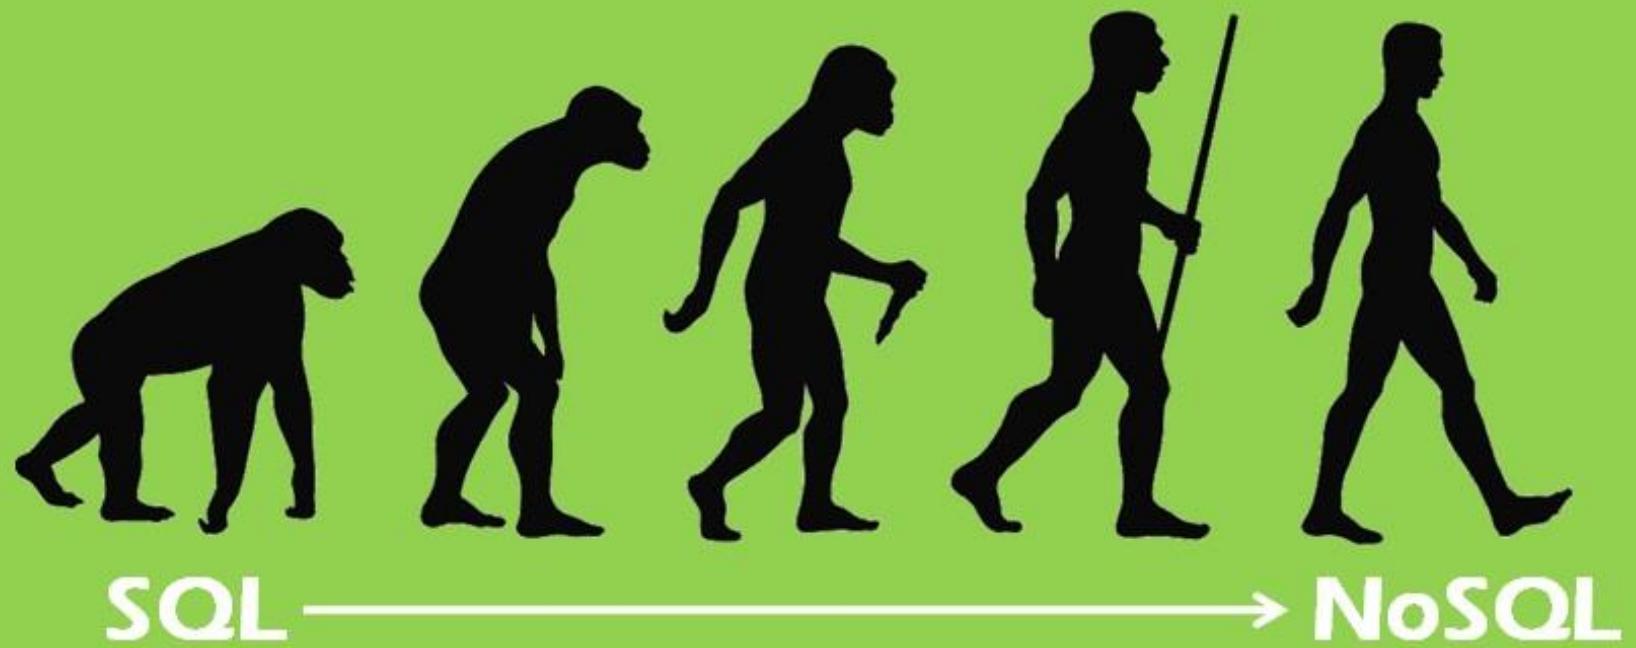
- Class of non-relational data storage systems
- Usually do not require a fixed table schema nor do they use the concept of joins

## □ All NoSQL offerings relax one or more of the ACID properties (will talk about the CAP theorem)

### ■ Key features (advantages):

- non-relational, data are replicated to multiple, nodes (so, identical & fault-tolerant) and can be partitioned, cheap, easy to implement (open-source), massive write performance, fast key-value access

## **The continued evolution of Databases**



<https://www.mytechnology.com/IT-blogs/8362/from-sql-to-nosql-the-continued-evolution-of-databases/#.WTa6NtJ9600>

---

## □ How is data stored with NoSQL?

## □ There are a variety of types:

- Column Store – Each storage block contains data from only one column
- Document Store – stores documents made up of tagged elements
- Key-Value Store – Hash table of keys

➤ **Dictionary**

## □ Other Non-SQL Databases

- XML Databases, Graph Databases, Object Oriented Databases, ...

# CouchDB JSON Example – dictionary

```
{  
  "_id": "guid goes here",  
  "_rev": "314159",  
  
  "type": "abstract",  
  
  "author": "Keith W. Hare"  
  
  "title": "SQL Standard and NoSQL Databases",  
  
  "body": "NoSQL databases (either no-SQL or Not Only SQL)  
          are currently a hot topic in some parts of  
          computing.",  
  "creation_timestamp": "2011/05/10 13:30:00 +0004"  
}
```



## MongoDB to documents (JSON):

---

```
{  
    _id: ObjectId("51156a1e056d6f966f268f81"),  
    type: "Article",  
    author: "Derick Rethans",  
    title: "Introduction to Document Databases with MongoDB",  
    date: ISODate("2013-04-24T16:26:31.911Z"),  
    body: "This arti..."  
},  
{  
    _id: ObjectId("51156a1e056d6f966f268f82"),  
    type: "Book",  
    author: "Derick Rethans",  
    title: "php|architect's Guide to Date and Time Programming with PHP",  
    isbn: "978-0-9738621-5-7"  
}
```



---

## □ Typical NoSQL API

### □ Basic API access:

- get(key) -- Extract the value given a key
- put(key, value) -- Create or update the value given its key
- delete(key) -- Remove the key and its associated value
- execute(key, operation, parameters) -- Invoke an operation to the value (given its key) which is a special data structure (e.g. List, Set, Map .... etc).

# Map Reduce is a typical NoSQL “query”

---

- Technique for indexing and searching large data volumes
- Two Phases, Map and Reduce
  - Map
    - Extract sets of Key-Value pairs from underlying data
    - Potentially in Parallel on multiple machines
  - Reduce
    - Merge and sort sets of Key-Value pairs
    - Results may be useful for other searches

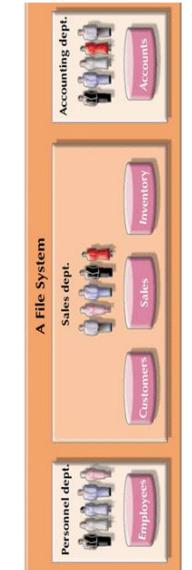




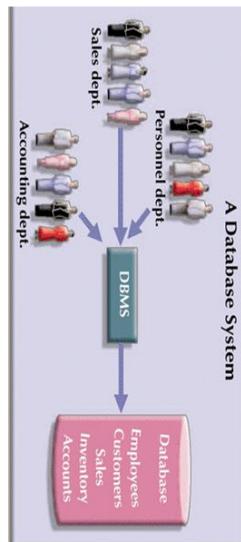
# Who is using them?



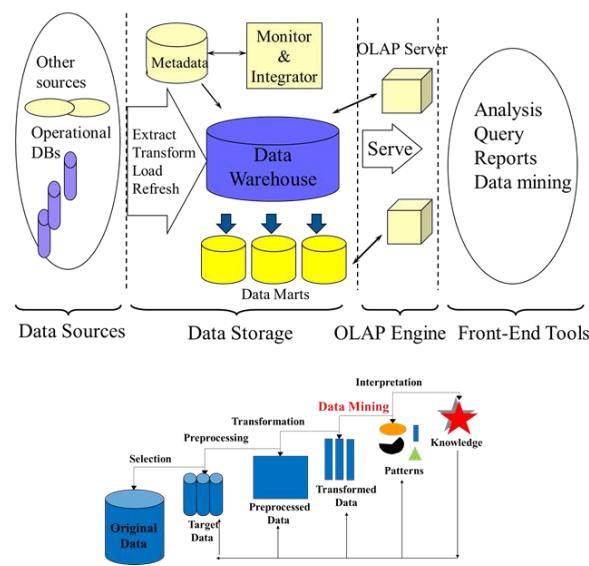
## File based DM



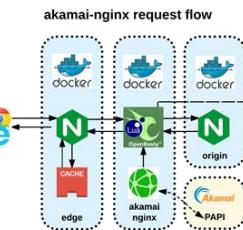
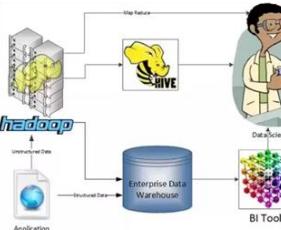
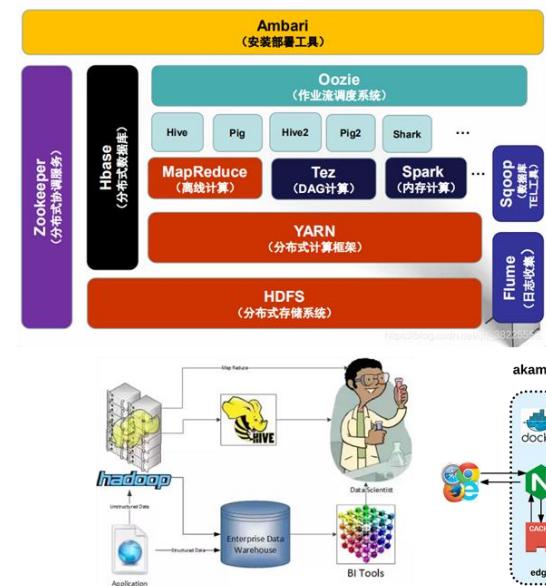
## (R)DBMS



## Data Warehouse/DM



## Big Data



## NoSQL – NewSQL



Redis - 95  
Memcached - 33  
DynamoDB - 16  
Riak - 13



MongoDB - 279  
CouchDB - 28  
Cassandra - 24  
DynamoDB - 15  
MarkLogic - 11



Cassandra - 109  
HBase - 62  
HBase - 62



neo4j - 30  
OrientDB - 4  
Titan - 3  
Giraph - 1



Solr - 81  
Elasticsearch - 70  
Splunk - 41



elasticsearch

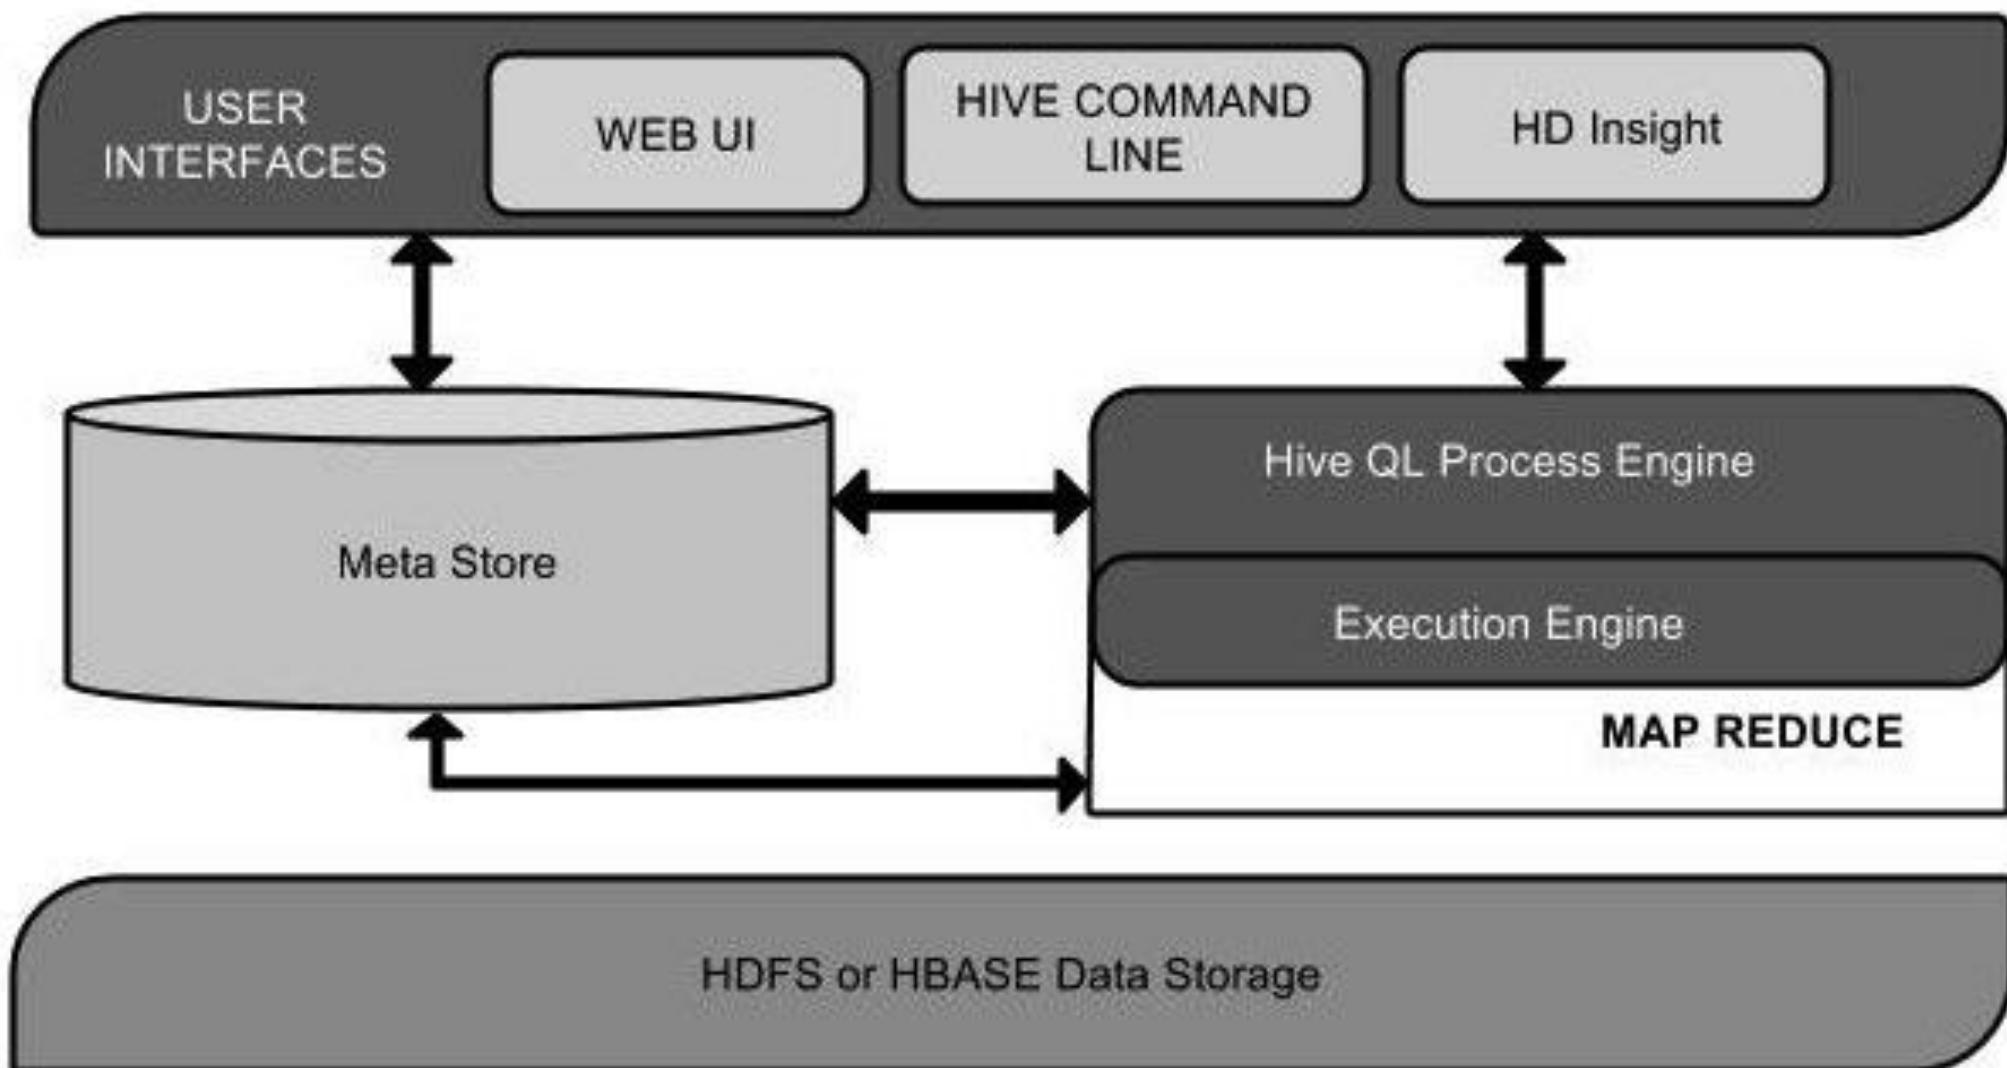
# Hive and HiveQL

## - SQL is re-found useful in Big Data



### □ Hive is a DW tool on Big Data

- Developed by [Facebook](#) and a top-level Apache project
- A [data warehousing infrastructure based on Hadoop](#)
  - like Hadoop, is designed for batch processing of large datasets
  - Immediately makes data on a cluster available to non-Java programmers via SQL like queries
- Built on [HiveQL](#) (HQL), a SQL-like query language
  - Interprets HiveQL and generates MapReduce jobs that run on the cluster
  - Enables easy data summarization, ad-hoc reporting and querying, and analysis of large volumes of data



# HiveQL

---

## □ HQL provides the basic SQL-like operations:

- Select columns using SELECT
- Filter rows using WHERE
- JOIN between tables
- Evaluate aggregates using GROUP BY
- Store query results into another table
- Download results to a local directory (i.e., export from HDFS)
- Manage tables and queries with CREATE, DROP, and ALTER



## □ Create a Complex Table

```
CREATE TABLE employees  (
    name STRING,
    salary FLOAT,
    subordinates ARRAY<STRING>,
    deductions MAP<STRING, FLOAT>,
    address STRUCT<street:STRING,
                  city:STRING,
                  state:STRING,
                  zip:INT>)

ROW FORMAT DELIMITED
FIELDS TERMINATED BY '\t'
STORED AS TEXTFILE;
```

# Sample Select Clauses

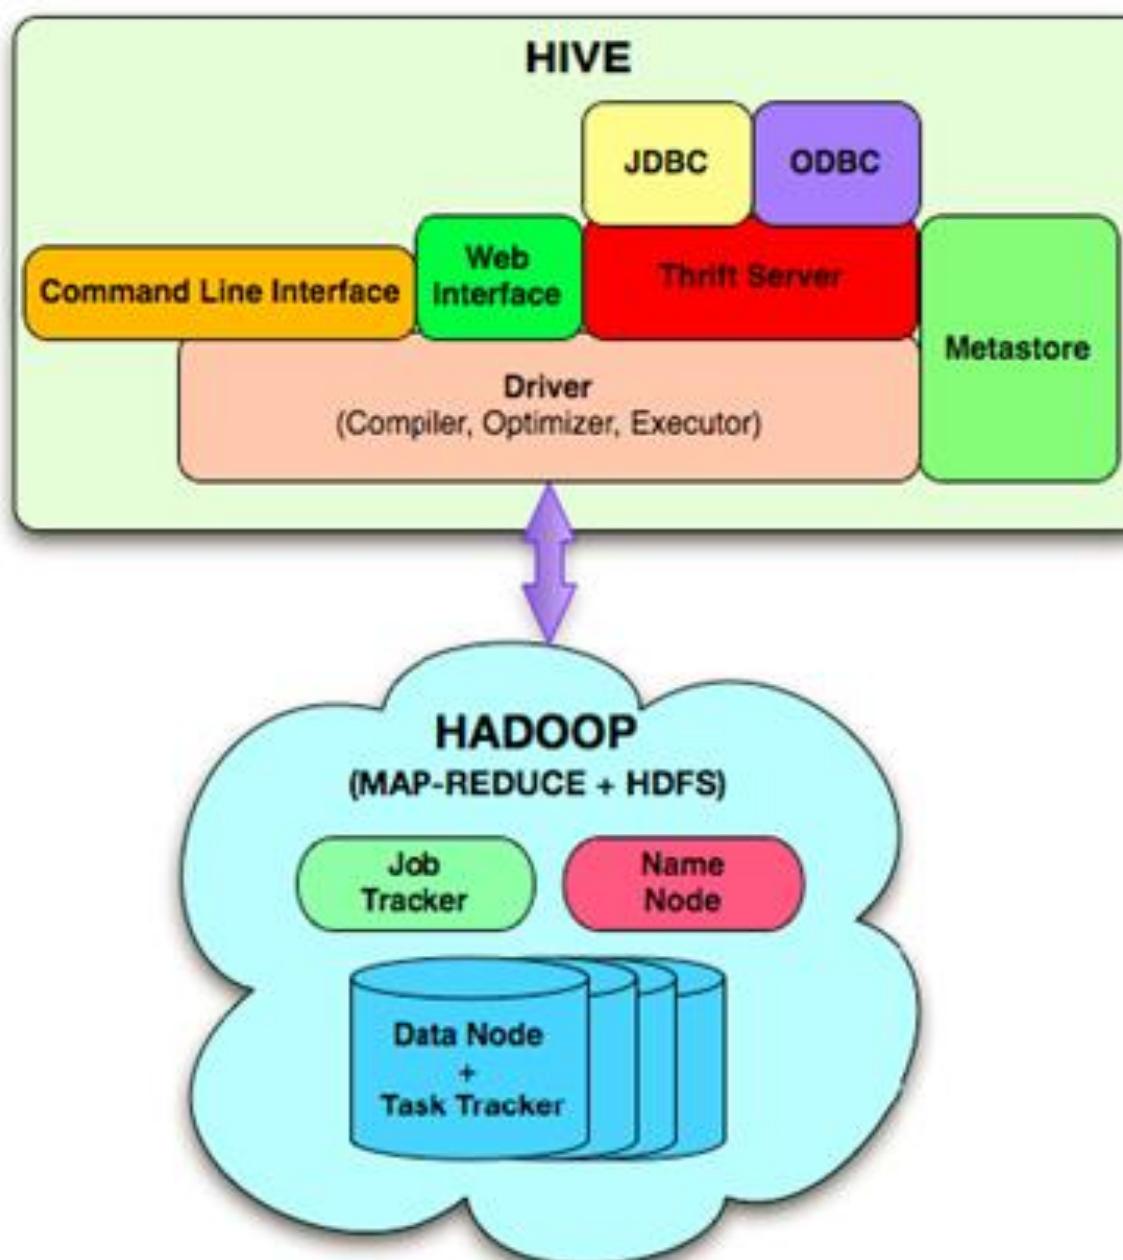
---

## □ Select from a single table

```
SELECT *
FROM sales
WHERE amount > 10 AND
      region = "US";
```

## □ Select from a partitioned table

```
SELECT page_views.*
FROM page_views
WHERE page_views.date >= '2008-03-01' AND
      page_views.date <= '2008-03-31'
```

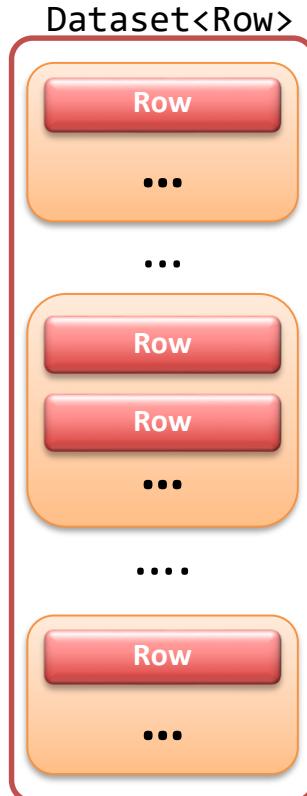


<https://jasperpeilee.wordpress.com/tag/hive/>

Fig. 1: Hive System Architecture

# SQL is re-found useful in Big Data – in Spark

## □ Return top 20 frequent words



```
case class Word(text: String, n: Long)
val wordsFreq = freq.map {
    case (text, count) => Word(text, count)
}.toDF() // Dataset<Word>
wordsFreq.createOrReplaceTempView("wordsFreq")

//Dataset<Row>
val topWords = sparkSession.sql("select text, n
                                from wordsFreq
                                order by n desc
                                limit 20")
topWords.collect().foreach(println)
```

Everyone  
could write  
SQL, right?



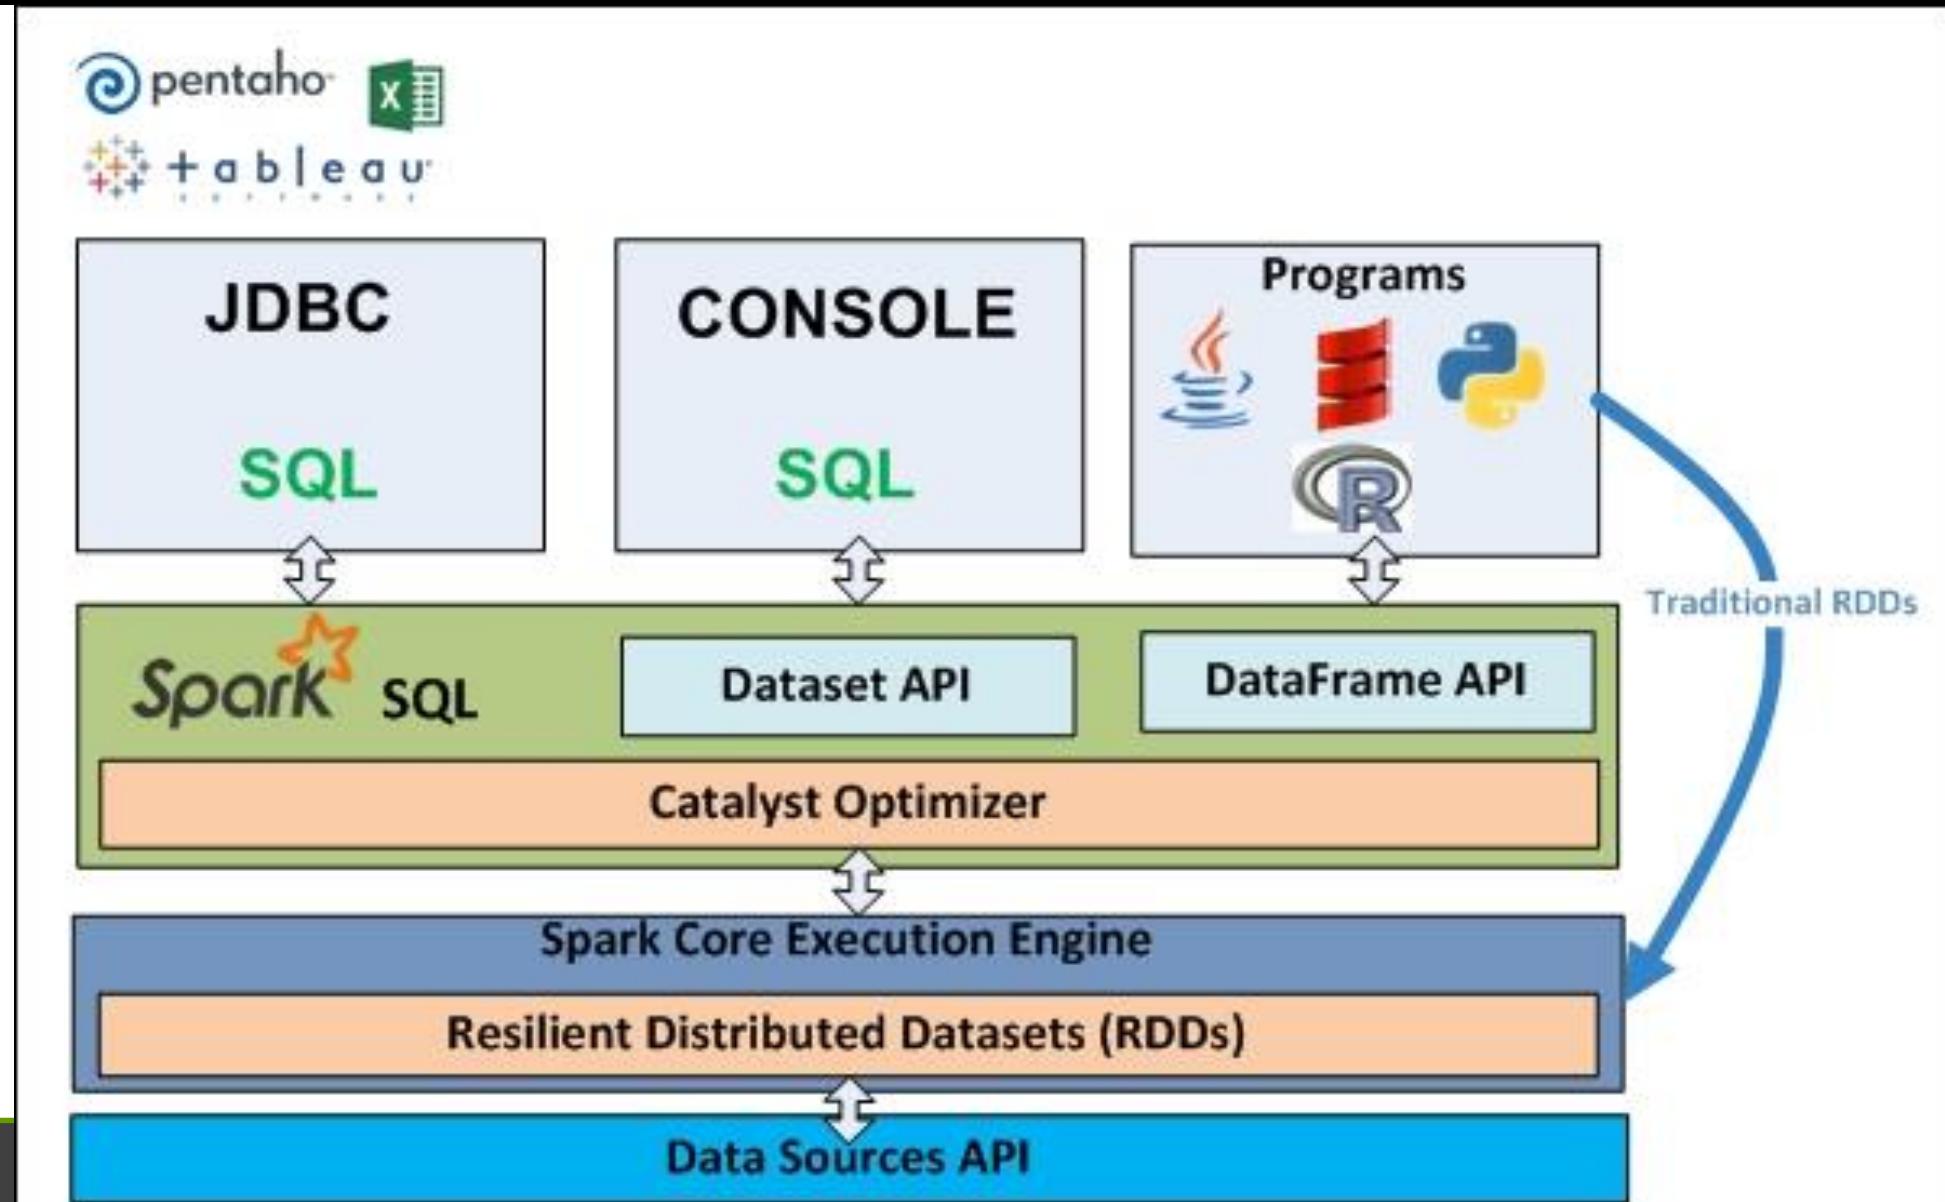
```
from pyspark.sql import Row
sc = spark.sparkContext
# Load a text file and convert each line to a Row.
lines = sc.textFile("examples/src/main/resources/people.txt")
parts = lines.map(lambda l: l.split(","))
people = parts.map(lambda p: Row(name=p[0], age=int(p[1])))
# Infer the schema, and register the DataFrame as a table.
schemaPeople = spark.createDataFrame(people)
# SQL can be run over DataFrames that have been registered as a table.
teenagers = spark.sql("SELECT name FROM people WHERE age >= 13 AND age <= 19")
# The results of SQL queries are DataFrame
# rdd returns the content as an RDD
teenNames = teenagers.rdd.map(lambda x: x['name'])
for name in teenNames:
    print(name)
```

Everyone  
could write  
SQL, right?

How to process SQL is a  
basic question to dive  
into the design and  
implementation of Big  
Data

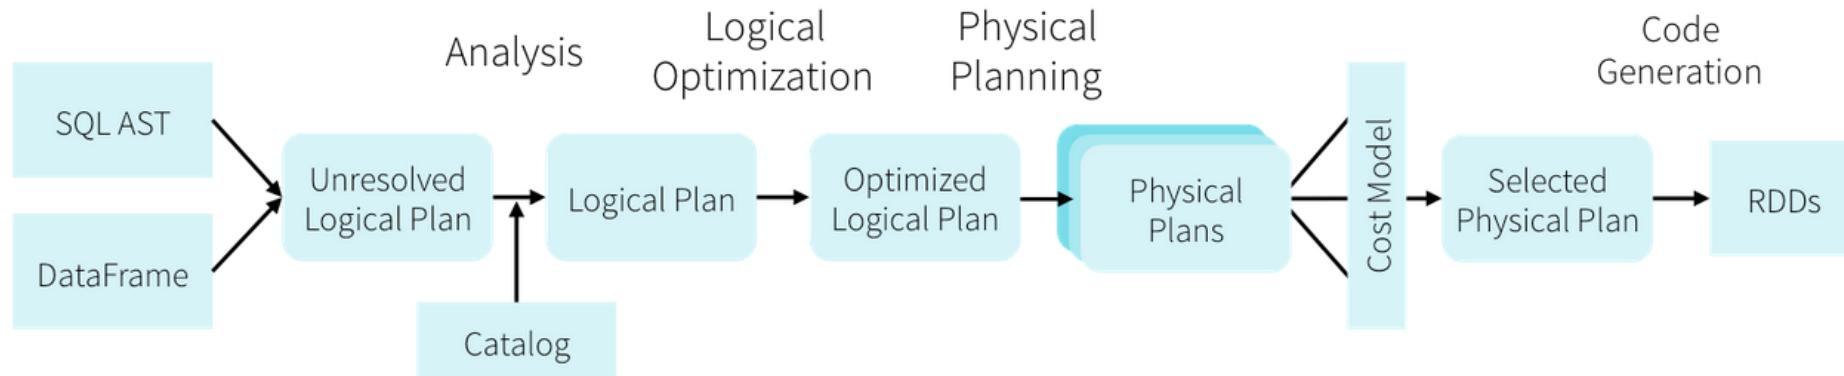
# Architecture of Spark SQL

<https://www.packtpub.com/mapt/book/big-data-and-business-intelligence/9781785884696/4/ch04lvl1sec26/Architecture+of+Spark+SQL>



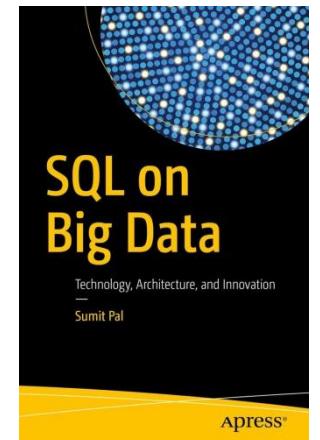
# Challenges of SQL in Big Data

## ■ Share similar flowchart



## ■ But, Data are scattered in different nodes

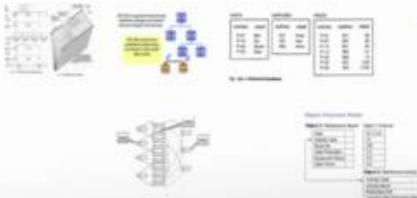
- Reducing the amount of data to be shuffled is a major challenge
- How to optimize?
- How to carry out the execution efficiently?
- How to process SQL queries on streaming data?
- ...



# □ HAWQ



## SQL (分布式数据库) 回归?



SQL回归?

Greenplum 中文社区所有

MR被其创造者Google 在2011年就放弃了

2008: MapReduce 是技术大倒退



作为一种数据处理模式，MapReduce 是一种大倒退。

- 模式 ( Schema ) 很有价值
- 实现糟糕，不支持 Access Method
- 高级访问语言很有价值

Google 内部2011年放弃MR，2015年公开放弃

Greenplum中文社区所有

Cloud Spanner 的文章！

SQL (分布式数据库) 从未离开



# 2019

## □ 王益 在蚂蚁金服 推出 SQLFlow + ElasticDL

<https://zhuanlan.zhihu.com/p/87682668>

<https://developer.aliyun.com/article/719553> [原文]



项目官网: <https://sqlflow.org>

GitHub地址: <https://github.com/sql-machine-learning/sqlflow>

您也可以使用docker, 运行文章中的汽车价  
格预测模型: docker run -p

8888:8888sqlflow/sqlflow:didi

## □ SQLFlow

<https://developer.aliyun.com/article/719553> [原文]

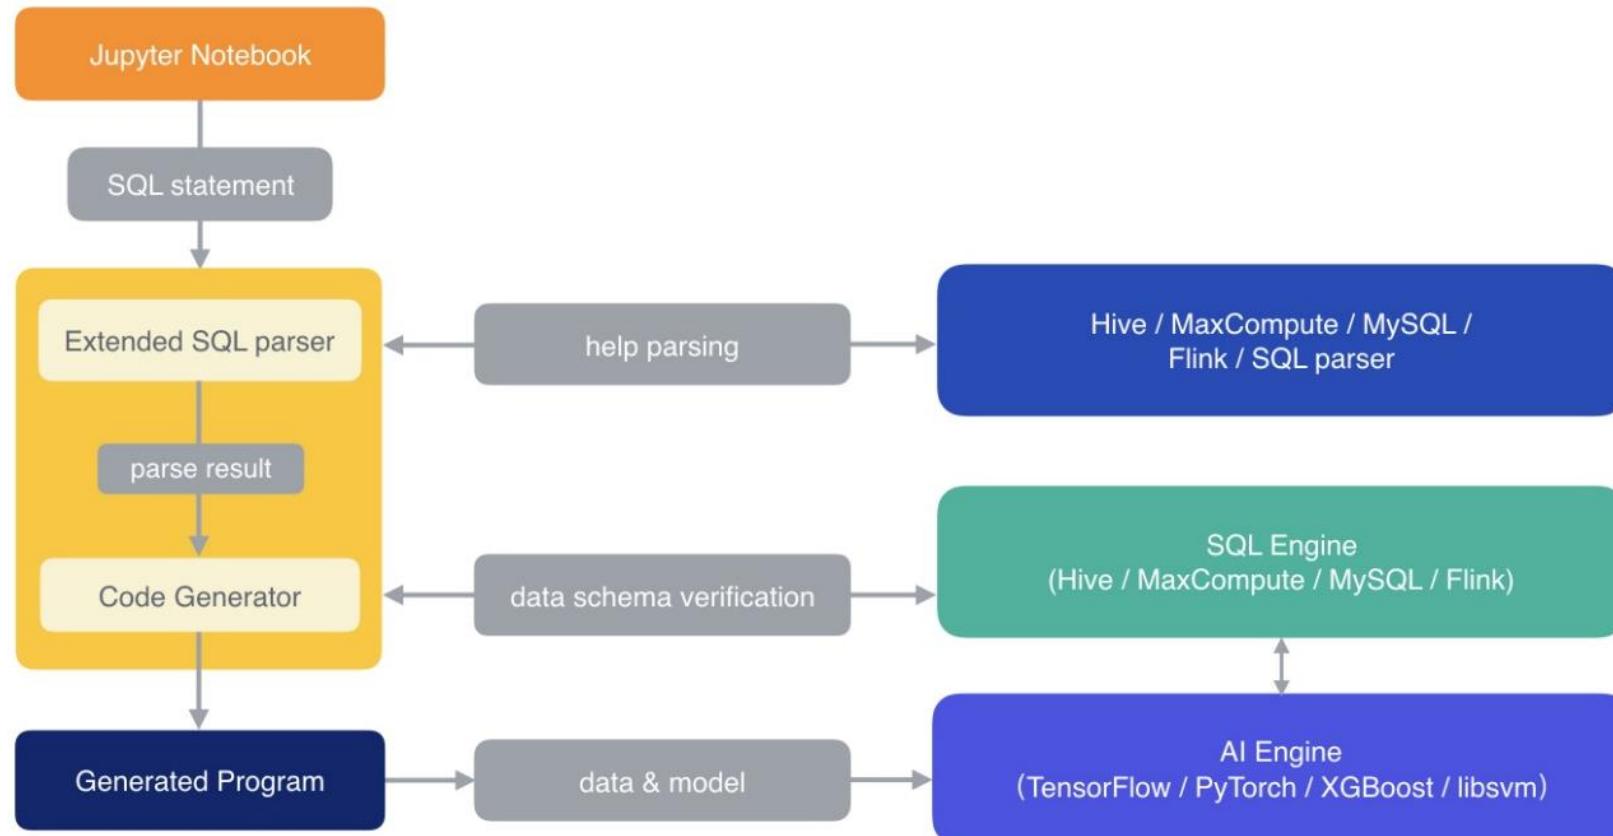
构建AI模型像SQL查询一样简单

```
01  SELECT * FROM ant_smart_promotion.train  
02  
03  SELECT * FROM ant_smart_promotion.prod  
04  
05  
06  
07  
08  
09  
10  
11
```

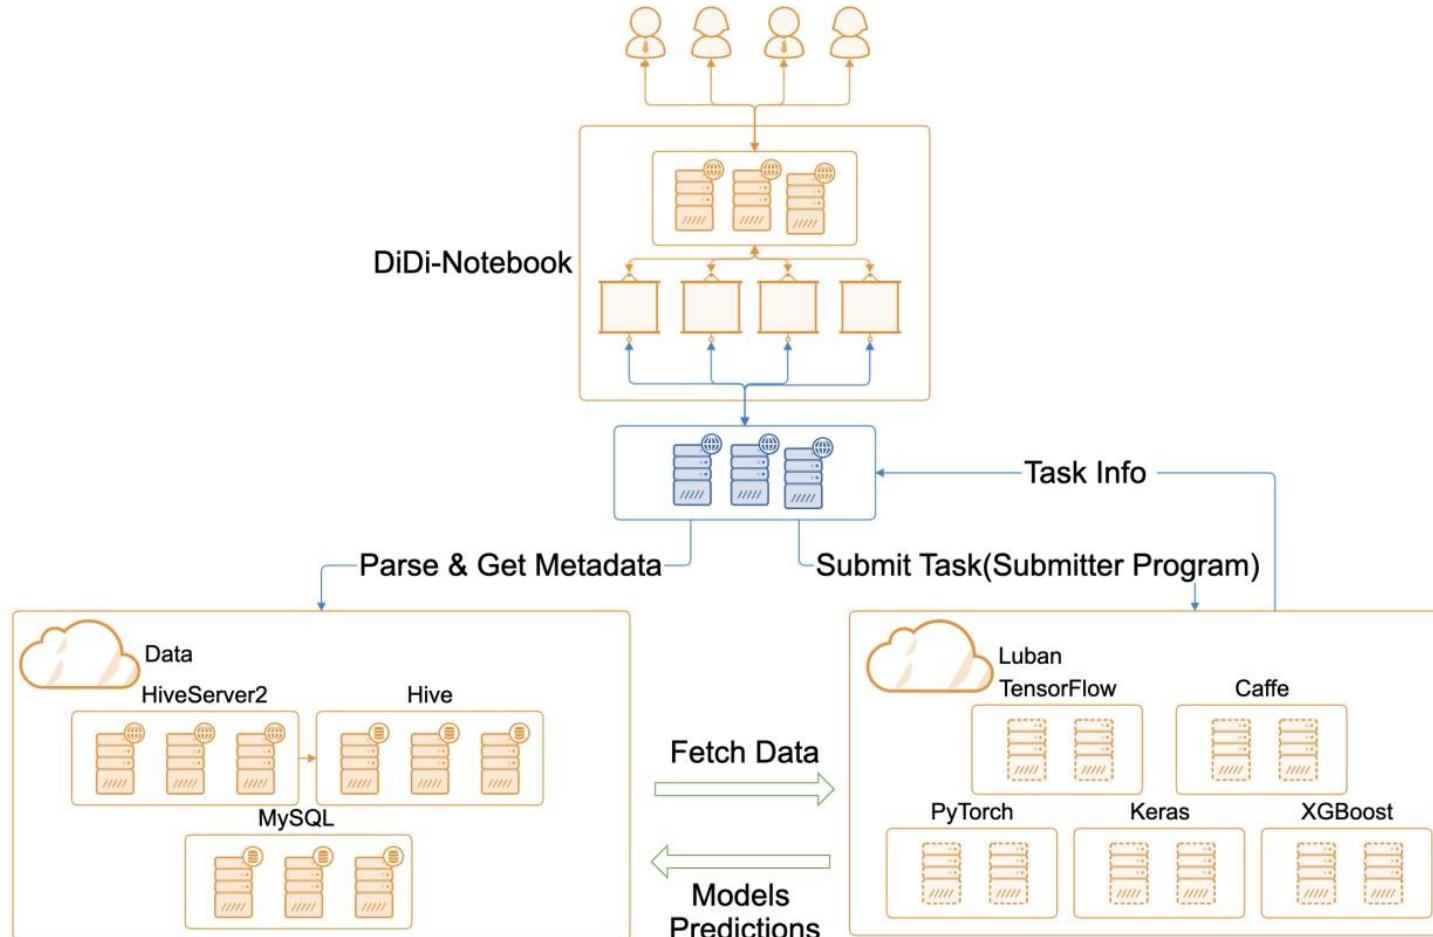
<https://github.com/sql-machine-learning>



## □ SQLFlow的技术架构



## □ SQLFlow 和滴滴数据的整合逻辑



# SQLOva: English to SQL translation

**Table:**

Player	Country	Points	Winnings (\$)
Steve Stricker	United States	9000	1260000
K.J. Choi	South Korea	5400	756000
Rory Sabbatini	South Africa	3400	4760000
Mark Calcavecchia	United States	2067	289333
Ernie Els	South Africa	2067	289333

**Question:** What is the points of South Korea player?

**SQL:** SELECT Points WHERE Country = South Korea

**Answer:** 5400

知乎 @yif

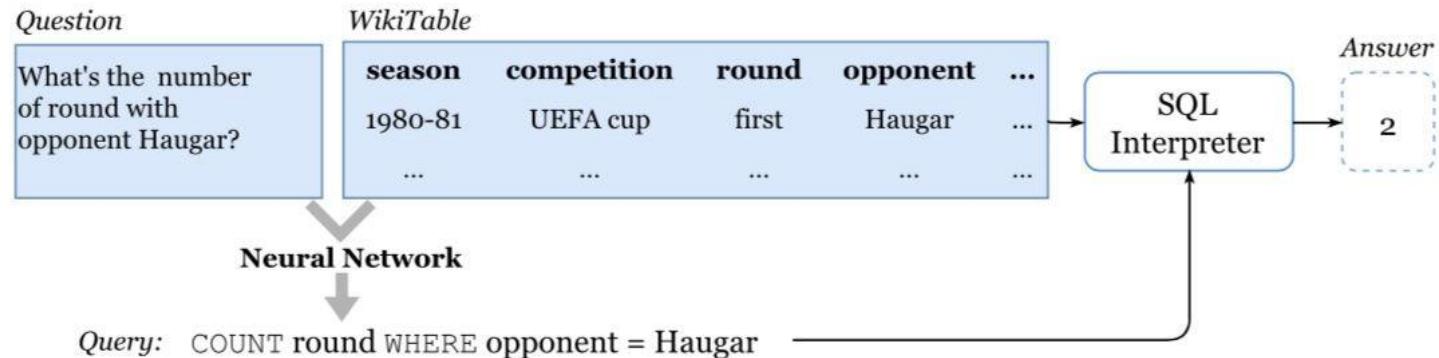


Figure 1: Answering a table question by synthesizing a query and executing it on the provided table.

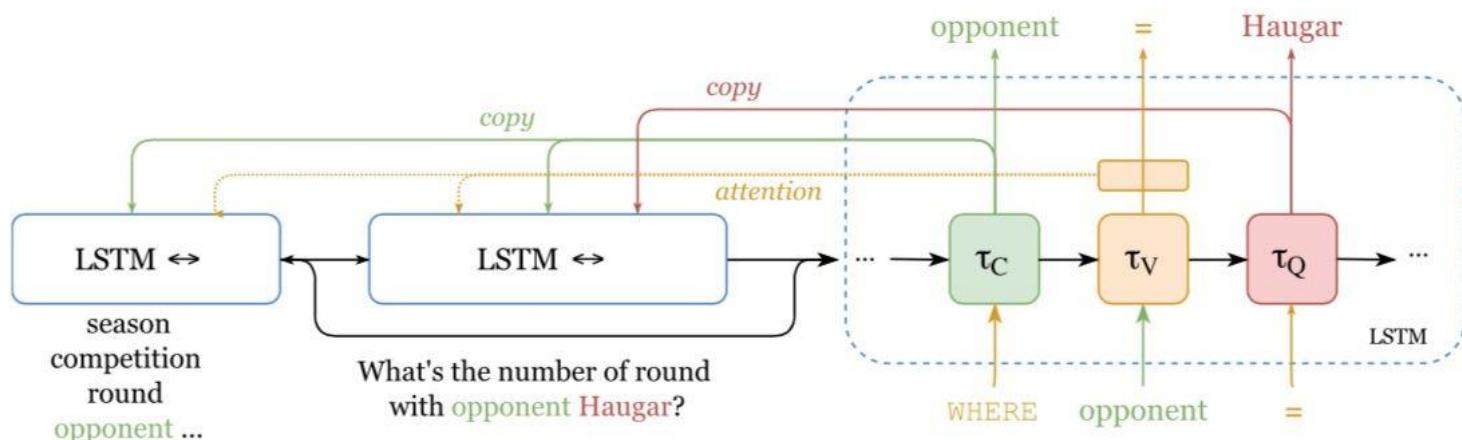


Figure 2: Model overview for the example in Figure 1. The model encodes table columns as well as the user question with a bidirectional LSTM and then decodes the hidden state with a typed LSTM, where the decoding action for each cell is statically determined.

知乎 @yif

### Question

What's the number of round with opponent Haugar in UEFA competition?

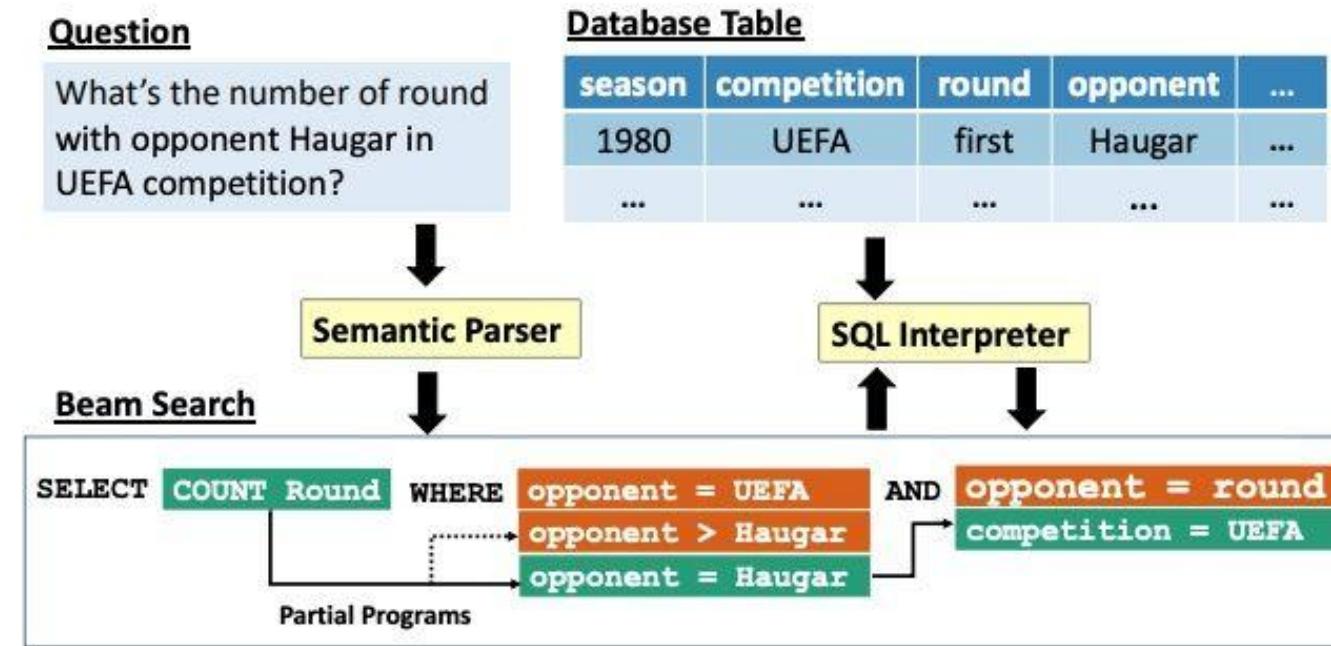


Figure 1: An execution-guided decoder evaluates partially generated queries at appropriate timesteps and then excludes those candidates that cannot be completed to a correct SQL query (red background). Here, “opponent > Haugar” would yield a runtime error, whereas “opponent = UEFA” would yield an empty result.

知乎 @yif

$[CLS], T_{n,1}, \dots T_{n,L}, [SEP], T_{h_1,1}, T_{h_1,2}, \dots, [SEP], \dots, [SEP], T_{h_{N_h},1}, \dots, T_{h_{N_h},M_{N_h}}, [SEP]$

### Table-aware BERT-Encoder

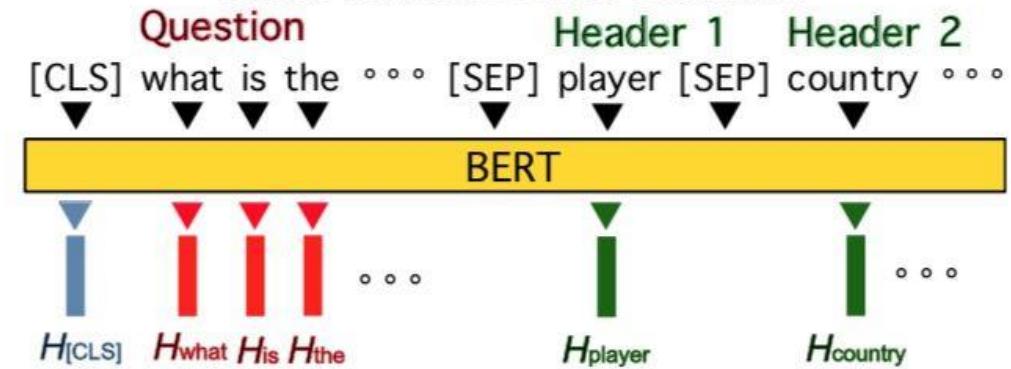


Figure 1: (A) The scheme of input encoding process by table-aware BERT. Final output vectors are represented by colored bars: light blue for [CLS] output, red for question words, and green for tokens from table headers.

知乎 @yif

github.com/naver/sqlova

Issues 33 Pull requests 2 Actions Projects Wiki Security Insights

master 1 branch 2 tags Go to file Add file Code

whhwang299 Merge pull request #56 from jinyeong/master ... fc68af6 on 9 Dec 2019 44 commits

bert	Force to read text file with utf-8	2 years ago
human_eval	Update of the human performance	13 months ago
sqlnet	Bug fix for deterministic pseudo random number generation based on a...	2 years ago
sqlova	load_jsonl added	16 months ago
wikisql	Shallow-Layer, Decoder-Layer have added	2 years ago
.gitignore	Force to read text file with utf-8	2 years ago
LICENSE.md	Initial commit	2 years ago
NOTICE	Initial commit	2 years ago
README.md	README updated to update publication information	14 months ago
add_csv.py	bug fix of considering w+d as float	16 months ago
add_question.py	add a prediction script	2 years ago
annotate_ws.py	add a prediction script	2 years ago
evaluate_ws.py	Initial commit	2 years ago
predict.py	remove unused get_opt from predict.py; revert unneeded change to tri...	2 years ago

About

No description, website, or topics provided.

Readme Apache-2.0 License

Releases 2

Trained Shallow Layer and De... Latest on 30 Sep 2019 + 1 release

Packages

No packages published

Contributors 6

paulfitz hanrelan wenfengand

	Dev			Test		
	Acc <sub>lf</sub>	Acc <sub>qm</sub>	Acc <sub>ex</sub>	Acc <sub>lf</sub>	Acc <sub>qm</sub>	Acc <sub>ex</sub>
Content Insensitive						
Dong and Lapata (2016)	23.3%	-	37.0%	23.4%	-	35.9%
Augmented Pointer Network (Zhong et al., 2017)	44.1%	-	53.8%	42.8%	-	52.8%
Seq2SQL (Zhong et al., 2017)	49.5%	-	60.8%	48.3%	-	59.4%
SQLNet (Xu et al., 2017)	-	63.2%	69.8%	-	61.3%	68.0%
TypeSQL w/o type-awareness (ours)	-	66.5%	72.8%	-	64.9%	71.7%
TypeSQL (ours)	-	<b>68.0%</b>	<b>74.5%</b>	-	<b>66.7%</b>	<b>73.5%</b>
Content Sensitive						
Wang et al. (2017a)	59.6%	-	65.2%	59.5%	-	65.1%
TypeSQL+TC (ours)	-	<b>79.2%</b>	<b>85.5%</b>	-	<b>75.4%</b>	<b>82.6%</b>

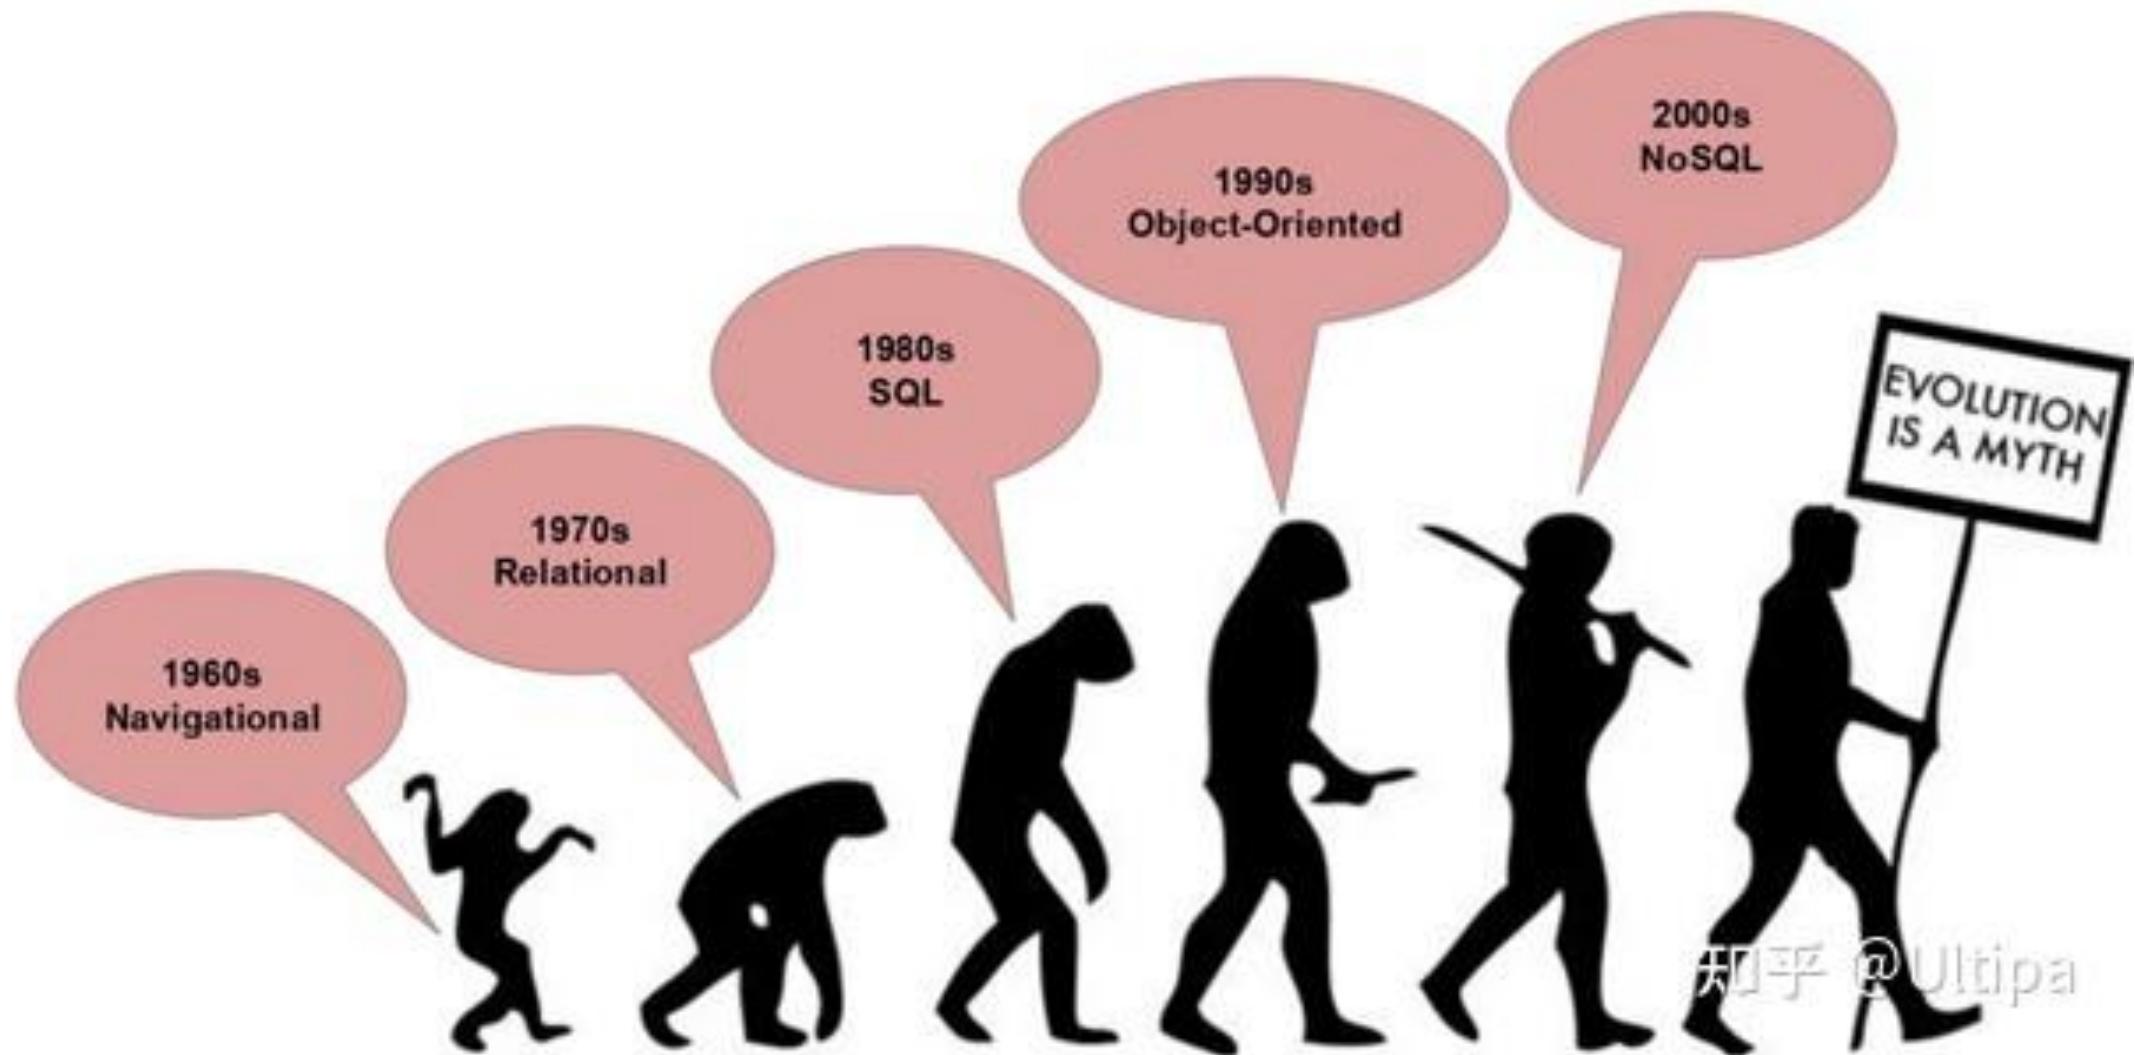
Table 1: Overall results on WikiSQL. Acc<sub>lf</sub>, Acc<sub>qm</sub>, and Acc<sub>ex</sub> denote the accuracies of exact string, canonical representation, and execute result matches between the synthesized SQL with the ground truth respectively. The top six results are content-insensitive, which means only the question and table schema are used as inputs. The bottom two are content-sensitive, where the models use the question, the table schema, and the content of databases.

	Dev			Test		
	Acc <sub>agg</sub>	Acc <sub>sel</sub>	Acc <sub>where</sub>	Acc <sub>agg</sub>	Acc <sub>sel</sub>	Acc <sub>where</sub>
Seq2SQL (Zhong et al., 2017)	90.0%	89.6%	62.1%	90.1%	88.9%	60.2%
SQLNet (Xu et al., 2017)	90.1%	91.5%	74.1%	90.3%	90.9%	71.9%
TypeSQL (ours)	90.3%	<b>93.1%</b>	<b>78.5%</b>	90.5%	<b>92.2%</b>	<b>77.8%</b>
TypeSQL+TC (ours)	90.3%	<b>93.5%</b>	<b>92.8%</b>	90.5%	<b>92.1%</b>	<b>87.9%</b>

Table 2: Breakdown results on WikiSQL. Acc<sub>agg</sub>, Acc<sub>sel</sub>, and Acc<sub>where</sub> are the accuracies of canonical representation matches on AGGREGATOR, SELECT COLUMN, and WHERE clauses between the synthesized SQL and the ground truth respectively.

知乎 @yif

# NewSQL?



知乎 @Ultipa

## □ SQL(Structured Query Language): 数据库，指关系型数据库。

- 好处来源于它的统一性和易用性，缺点是面对大量的数据时，他的性能会随着数据库的增大而急剧下降。
- 主要代表：SQL Server、Oracle、MySQL、PostgreSQL。

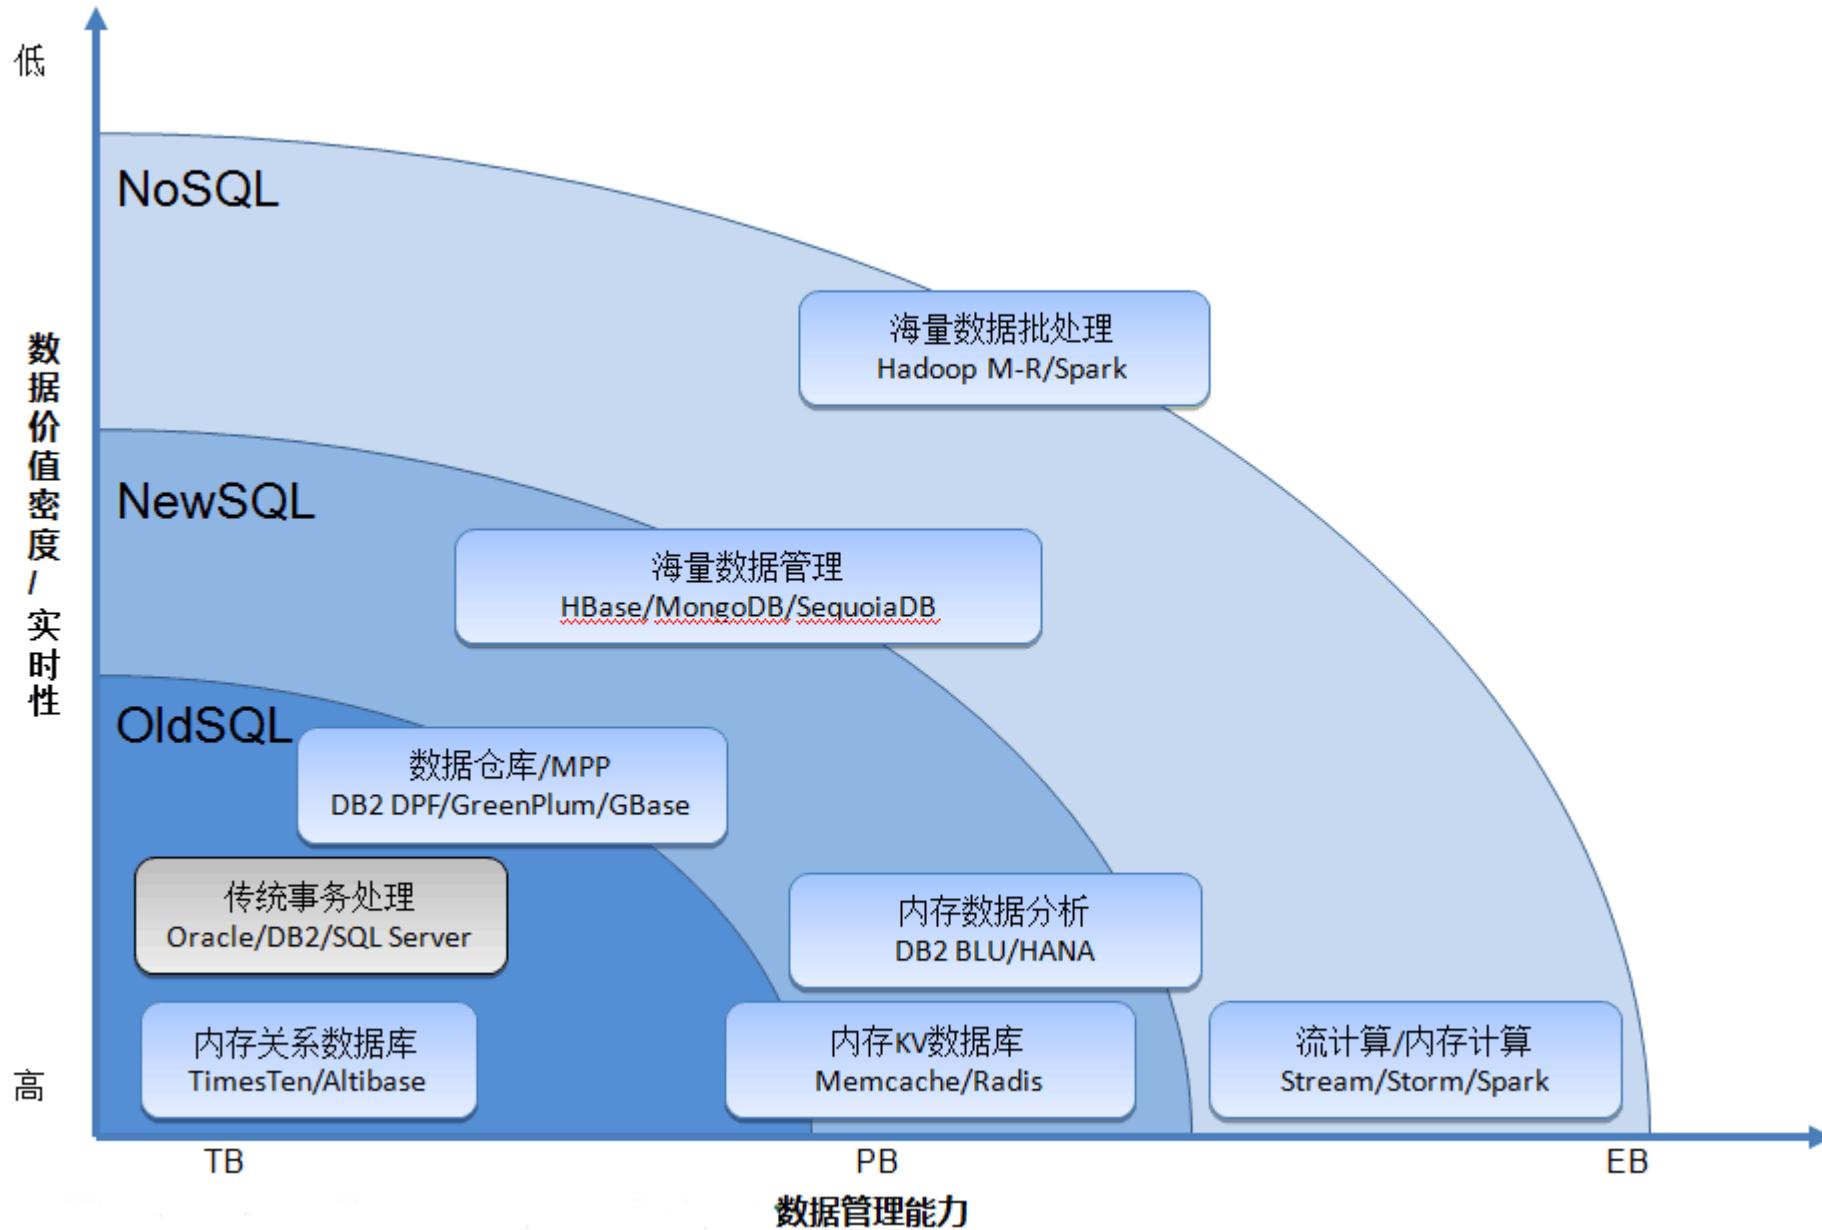
## □ NoSQL(Not Only SQL): 泛指非关系型数据库。

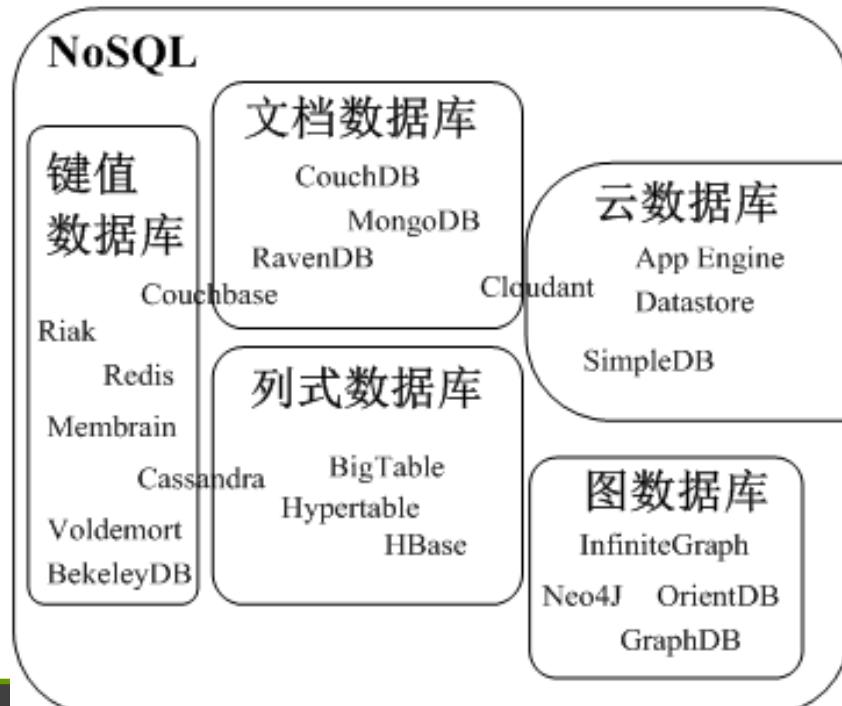
- 以放宽ACID原则为代价，NoSQL采取的是最终一致性原则，而不是像关系型数据库那样地严格遵守着ACID的原则，这意味着如果在特定时间段内没有特定数据项的更新，则最终对其所有的访问都将返回最后更新的值。这就是这样的系统通常被描述为提供基本保证的原因（基本可用，软状态，最终一致性）—而不是ACID。
- 主要代表：MongoDB、Redis、CouchDB。

## □ NewSQL：对各种新的可扩展/高性能数据库的简称。

- NewSQL选择汲取了SQL和NewSQL的优点，希望将ACID和可扩展性以及高性能结合，但是目前而言，不适用于所有的场景。
- 主要代表：Clustrix、GenieDB。







## 关系型数据库

### OldSQL

Infobright   Netezza   ParAccel   SAP Sybase IQ  
 Teradata   EMC   Calpont   IBM InfoSphere  
 Aster Data   Greenplum   VectorWise   HP Vertica  
 Oracle   IMB DB2   SQL Server   JustOne  
 MySQL   Ingres   PostgreSQL

### NewSQL

HandlerSocket Amazon RDS Database.com Xeround	Akiban MySQL Cluster Clustrix Drizzle FathomDB	HandlerSocket MySQL Cluster Clustrix GenieDB ScalArc
Schooner MySQL Tokutek Continuent	CodeFutures ScaleBase	NimbusDB VoltDB Translattice

# Chapter 5: Go deeper and wider

- Overview about the evolution of Data Management & Analytics
- SQL again
- Final (if you want share)
  - How much do you know now about D&I of DBMS and OS?

Other models like Hierarchical, Network

1961 Charles Bachman at GE – IDS (Integrated Data Store: 集成数据存储) – 1<sup>st</sup> NDBMS (Network), also 1<sup>st</sup> DBMS

1968 Vern Watts at IBM – IMS (Information Management System: 信息管理系统) – 1<sup>st</sup> HDBMS (Hierarchical)

1970 Edgar F. Codd at IBM published paper “A Relational Model of Data for Large Shared Data Banks” – Father of RDBMS

## OLTP

1970, E. F. Codd  
Relational Algebra

## OLAP/DW(DM)

- System R (1974-1978@IBM)  
prototype of RDBMS –  
1<sup>st</sup> try of SQL

- **Ingres (1973-1985@UCB)**  
INteractive Graphics  
REtrieval System

- Oracle 2.0 (1978 1<sup>st</sup> SQL based  
RDBMS)

- DB2 (1983)  
Database2 for MVS

- PostgreSQL (1996)  
1989 Postgres 1

- SQL Server (1988)  
MS+Sybase+Ashton-Tate

- MySQL 1.0 (1996)

- ...

INGRES

ORACLE



PostgreSQL

MySQL

Greenplum

2005 Pivotal

IBM Informix  
software



## Big Data

2000s Big Data

- 2004 Google  
published 3 papers

1998, Carlo Strozzi used NoSQL to name his lightweight Strozzi NoSQL open-source relational which is distinct from the more general concept of NoSQL databases.

NoSQL

**中国不能， 也不会缺席！**



2017杭州·云栖大会  
THE COMPUTING CONFERENCE

Alibaba Group  
WORLDWIDE PARTNER

# 内核专场

Session

核心技术

Global Shopping Fes



战略合作伙伴



2017杭州·云栖大会  
THE COMPUTING CONFERENCE

## 设计目标

- ✓ 高性能、跨Region部署
- ✓ 网络抖动高容忍性
- ✓ 独立基础库，可单独使用

## Various Distributed System

PutLogEntryOnCommit GetInfo AdminCommand

X-Paxos

### Consensus Protocol Stack

Replicate Log Module

Command Module

Heartbeat Module

Leader Election  
Module

Worker Pool

Async Timer Service

Net (libeasy)

飞天 智能  
FATE INTELLIGENCE





2017杭州云栖大会  
THE COMPUTING CONFERENCE

Alibaba Group  
WORLDWIDE PARTNER

# 内核专场

Session

核心技术

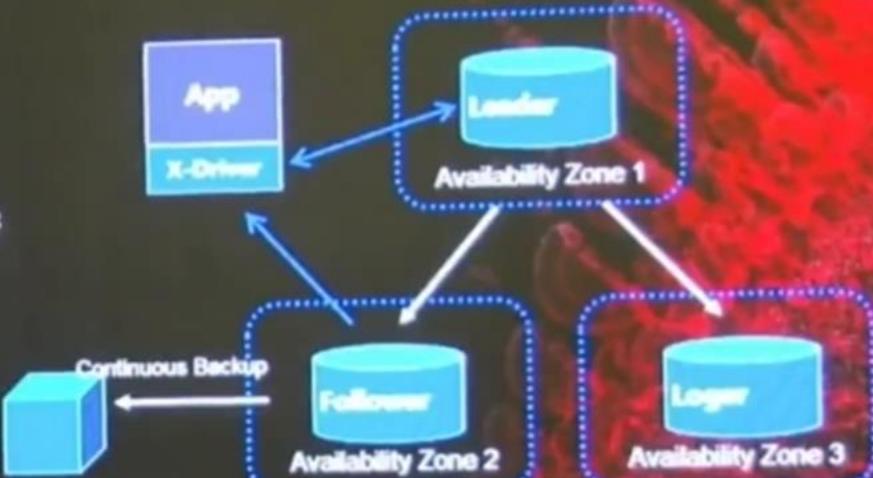
Global Shopping Festival



2017杭州·云栖大会  
THE COMPUTING CONFERENCE

## X-Cluster部署案例——Multi-AZ

- 数据强一致
  - ✓ 单AZ不可用数据0丢失
  - ✓ 单AZ不可用秒级切换
  - ✓ 切换自封闭，无第三方组件
- 0成本增加
  - ✓ 相对主备模式0成本增加
- 持续备份
  - ✓ RPO < 1s



飞天·智能  
PAAS · MIDDLEWARE



2017杭州云栖大会  
THE COMPUTING CONFERENCE

Alibaba Group  
WORLDWIDE PARTNER

# 内核专场

Session

核心技术

Global Shopping Festival



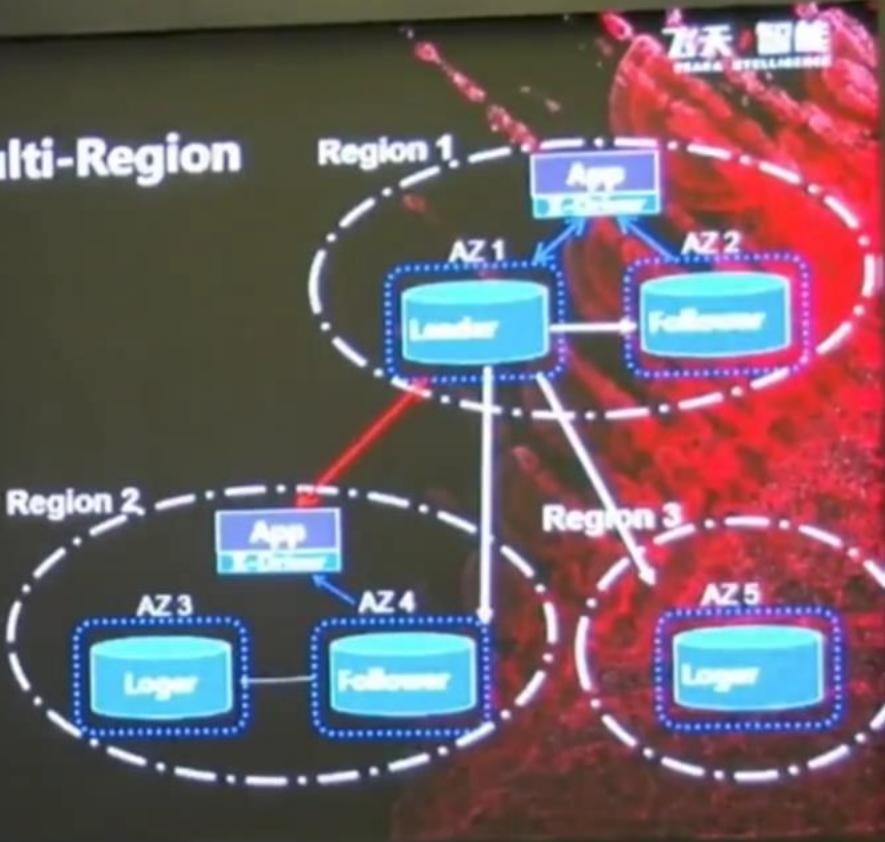
战略合作伙伴



2017杭州·云栖大会  
THE COMPUTING CONFERENCE

## X-Cluster部署案例——Multi-Region

- 真正Region级的强一致能力
  - 单个Region不可用0数据丢失
- 高性能
  - 跨Region强同步下依然保持高性能
- 灵活的切换策略
  - 优先切换同Region
  - 定制跨Region切换顺序
- 高伸缩性
  - 可无限制的扩充Region/AZ的节点
  - 可自由的调节Region/AZ内节点

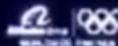




2017 杭州·云栖大会  
THE COMPUTING CONFERENCE



2017 杭州·云栖大会  
THE COMPUTING CONFERENCE



飞天·智能  
TAIEHUA INTELLIGENCE



2017 杭州·云栖大会  
THE COMPUTING CONFERENCE



# 企业高可用架构专场

High-availability Enterprise Architecture Session

让系统在双十一峰值流量下依然稳如磐石

Making the System Stable under the Peak Flow of 11.11

# 企业高可用架构专场

High-availability Enterprise Architecture Session

让系统在双十一峰值流量下依然稳如磐石

Making the System Stable under the Peak Flow of 11.11





Baidu 百度

股票代码 Stock Code: 9888.HK

热烈庆祝百度集团股份有限公司香港联合交易所成功上市  
Congratulation on the listing of Baidu Group Limited on the Hong Kong Stock Exchange

Baidu 百度

爱这个时代 星辰大海



Baidu 百度



# 全球超算500强中国上榜数量蝉联第一

全球超算500强榜单每半年发布一次

在德国法兰克福举行的国际超级计算大会当地时间6月17日发布了最新榜单



中国境内有**219**台超算上榜

中国超算上榜数量蝉联第一

这是2017年11月以来，中国超算上榜数量连续第四次位居第一



②

美国

③

日本

④

法国

⑤

英国

⑥

德国

中国企业也继续保持  
上榜数量优势

在此次  
榜单上

全球超算制造商前三位

联想  
173台

浪潮  
71台

中科曙光  
63台



外弟 却是以联想为代表的这四家中国企业的贡献。

所有系统采用Linux

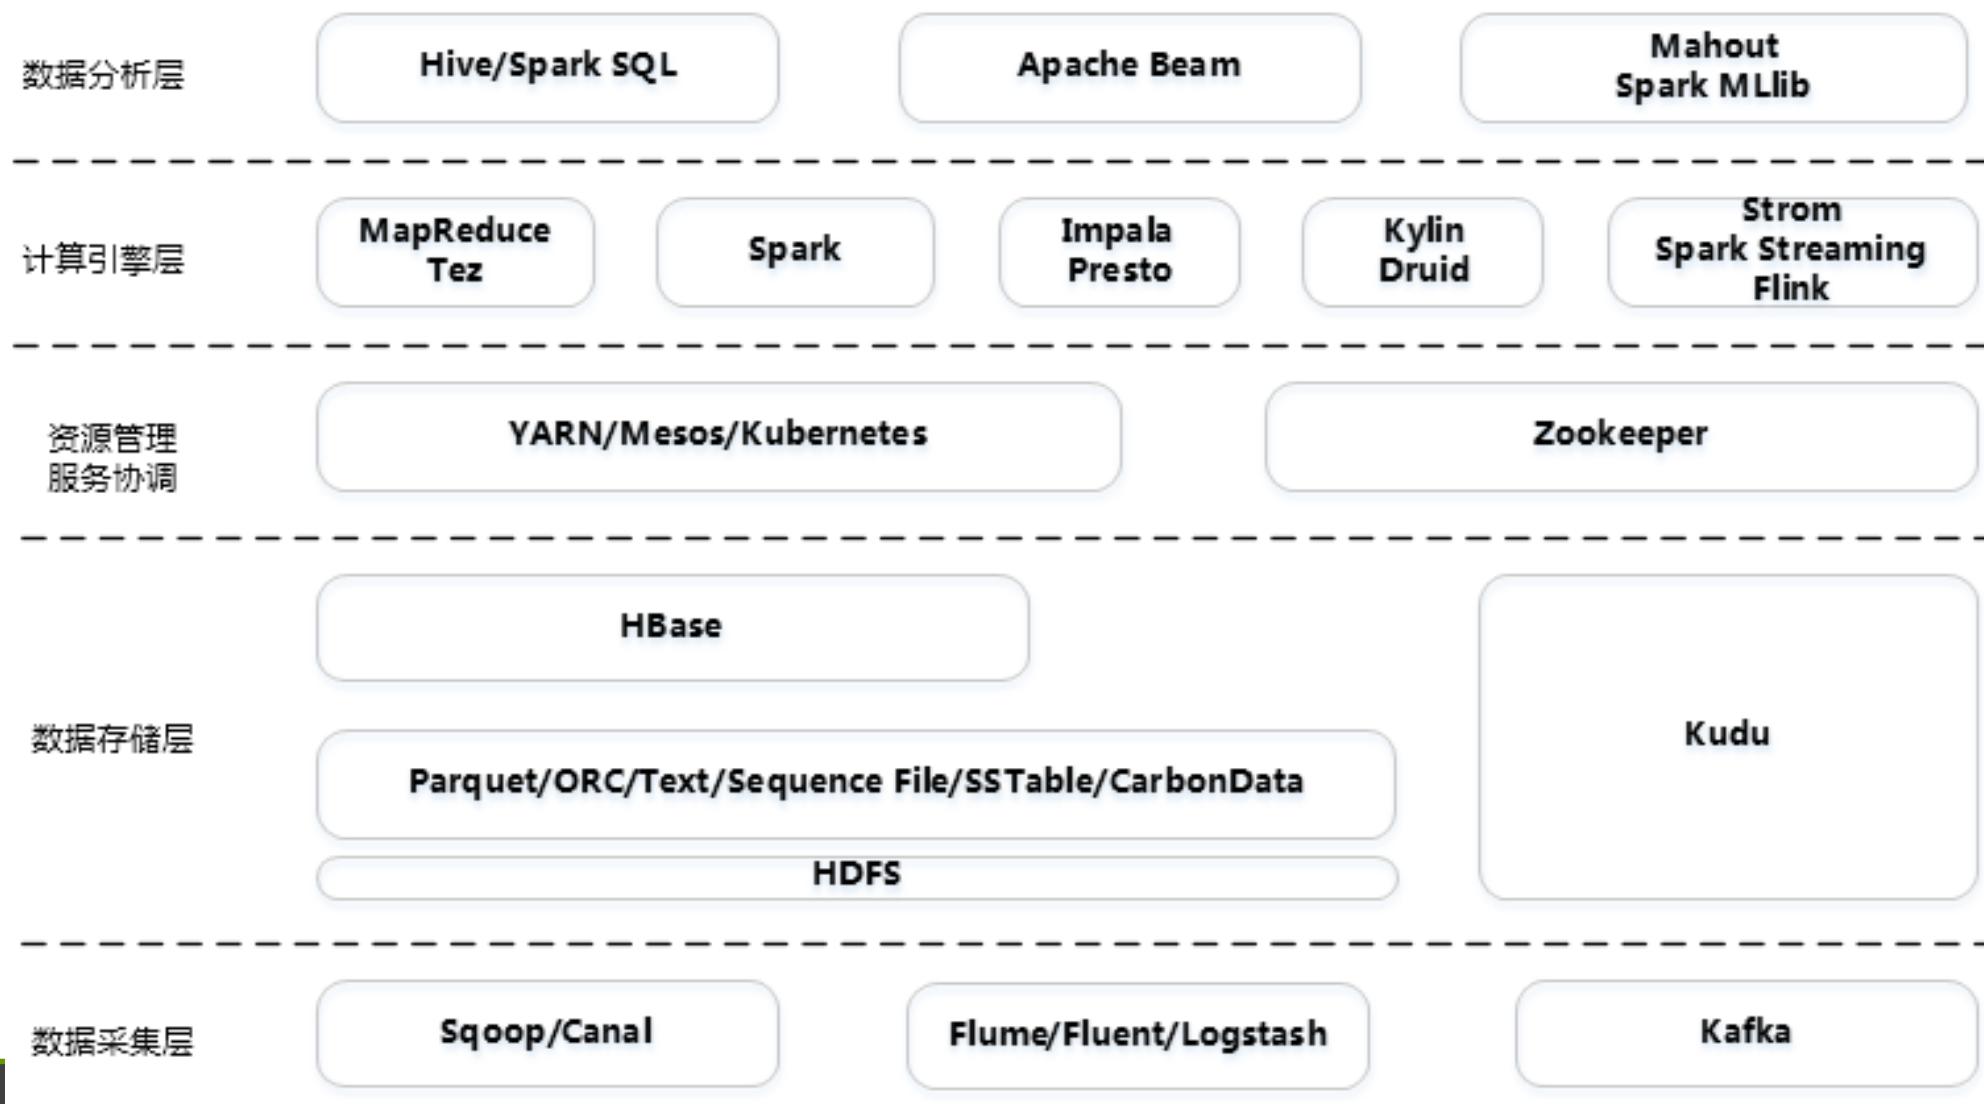


Unix、Windows、BSD Based等操作系统在2017年11月的榜单中首次  
尽数消失，TOP500所有系统都采用开源的Linux系统。

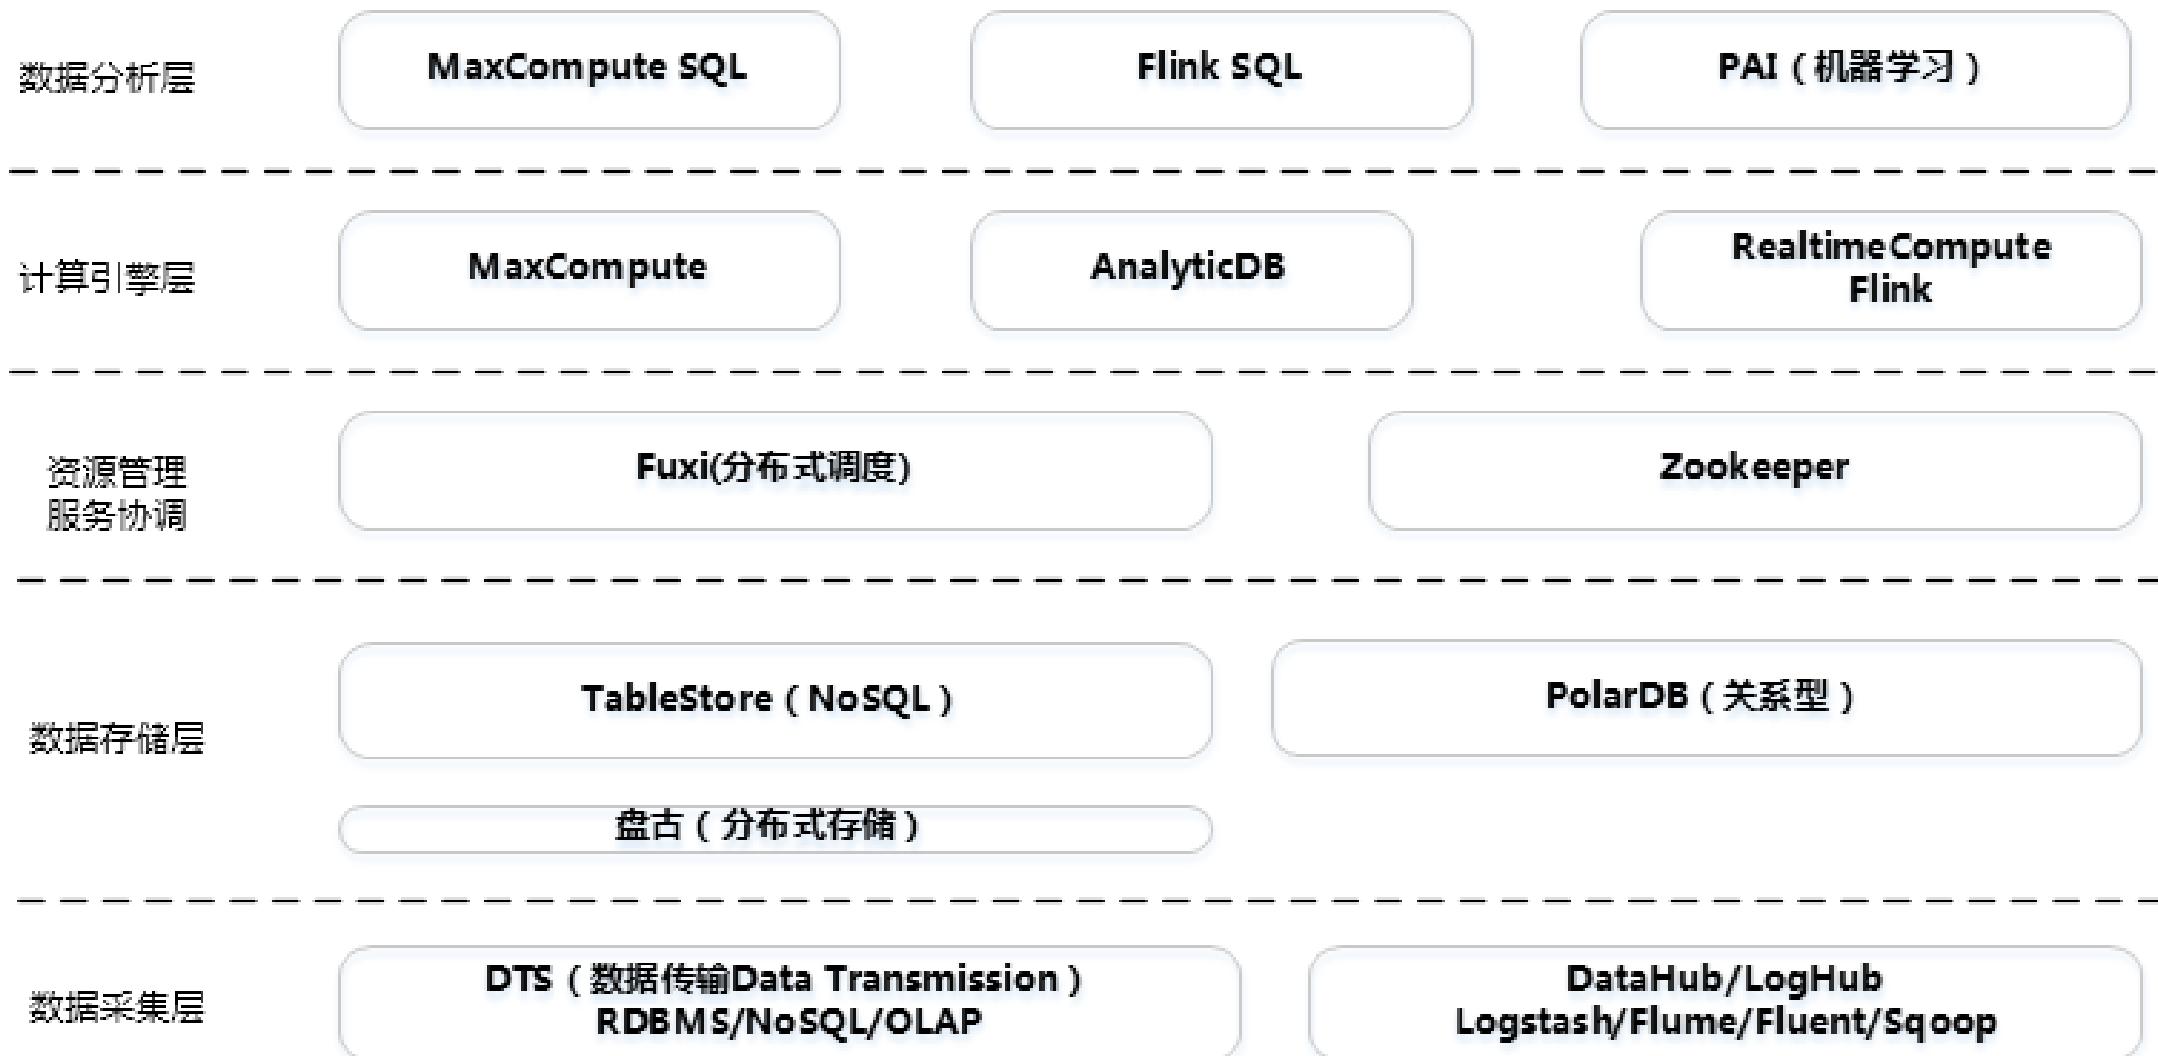
# 谷歌技术栈



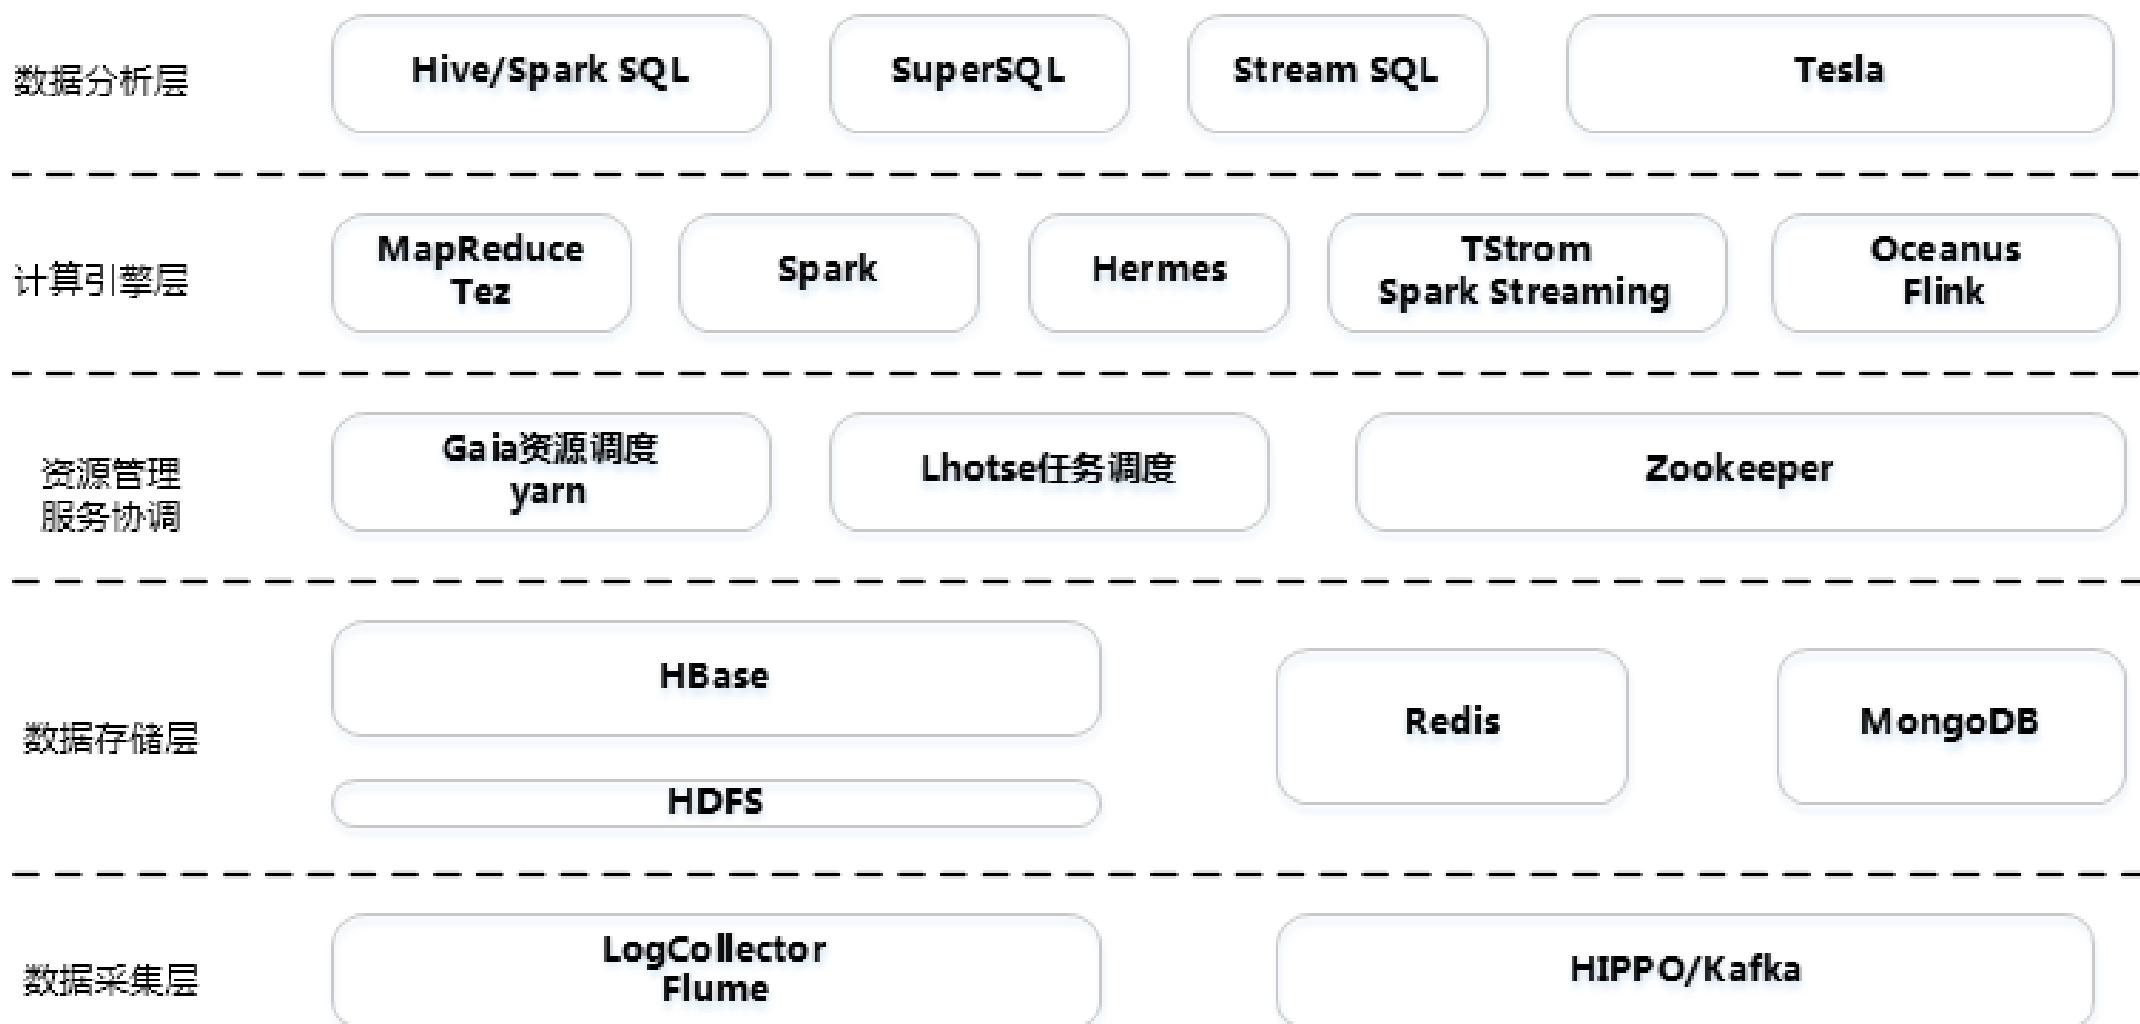
# 开源技术栈



# 阿里技术栈



# 腾讯技术栈



# 华为技术栈？呵呵

## - “为了卖硬件而不得不制作软件”，而且作业也很好！



## 鲲鹏计算产业

- 鲲鹏计算产业是基于Kunpeng处理器构建的全栈IT基础设施、行业应用及服务，包括PC、服务器、存储、操作系统、中间件、虚拟化、数据库、云服务、行业应用以及咨询管理服务等。

### 鲲鹏计算产业

#### 行业应用



政府



金融



游戏



媒体与娱乐



运营商

#### 云服务



鲲鹏ECS



鲲鹏BMS



鲲鹏容器



鲲鹏RDS



鲲鹏DWS

#### 数据库

#### 中间件

#### 操作系统



openEuler操作系统

兼容丰富的操作系统、中  
间件、数据库软件

高斯数据库

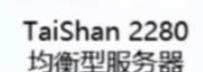
#### 虚拟化

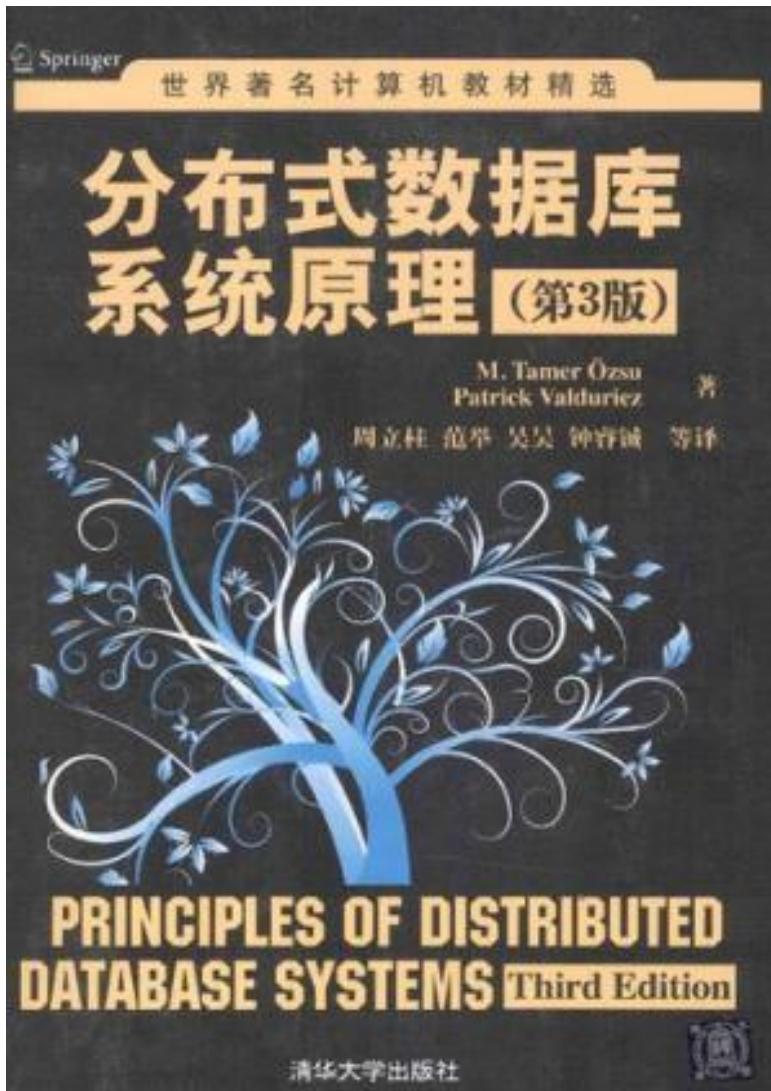
#### 存储

PC

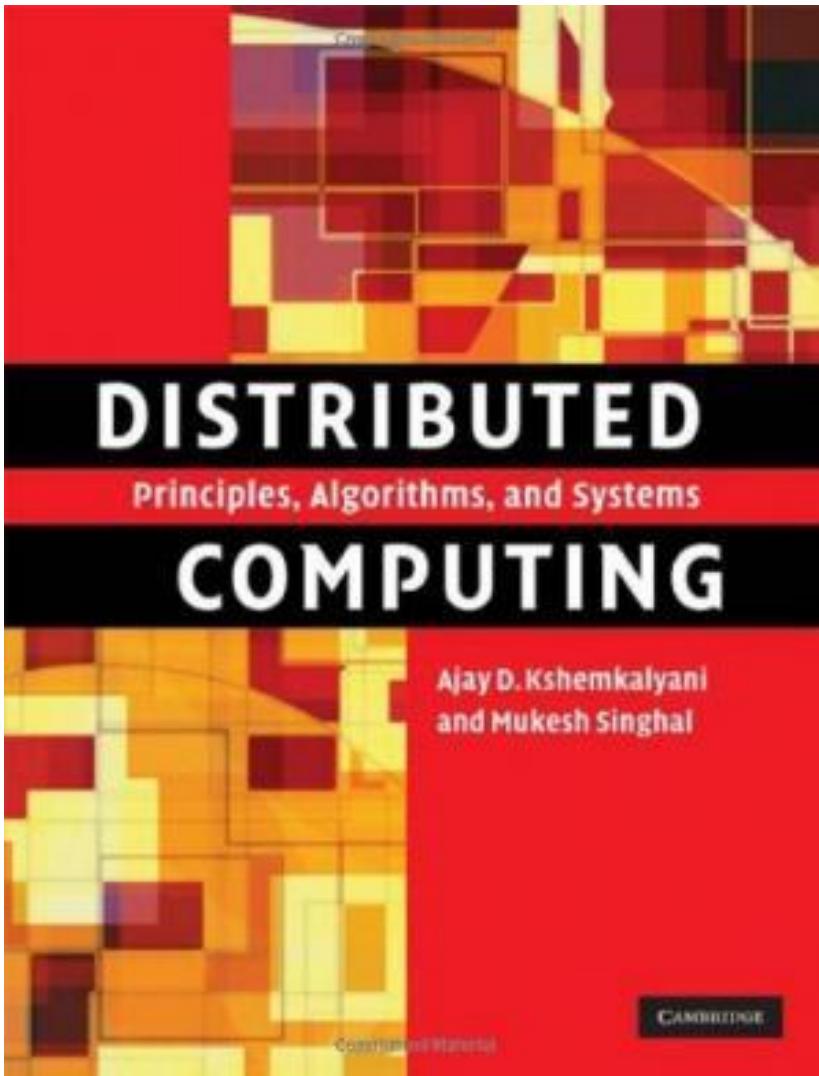
#### 服务器

#### Kunpeng 处理器

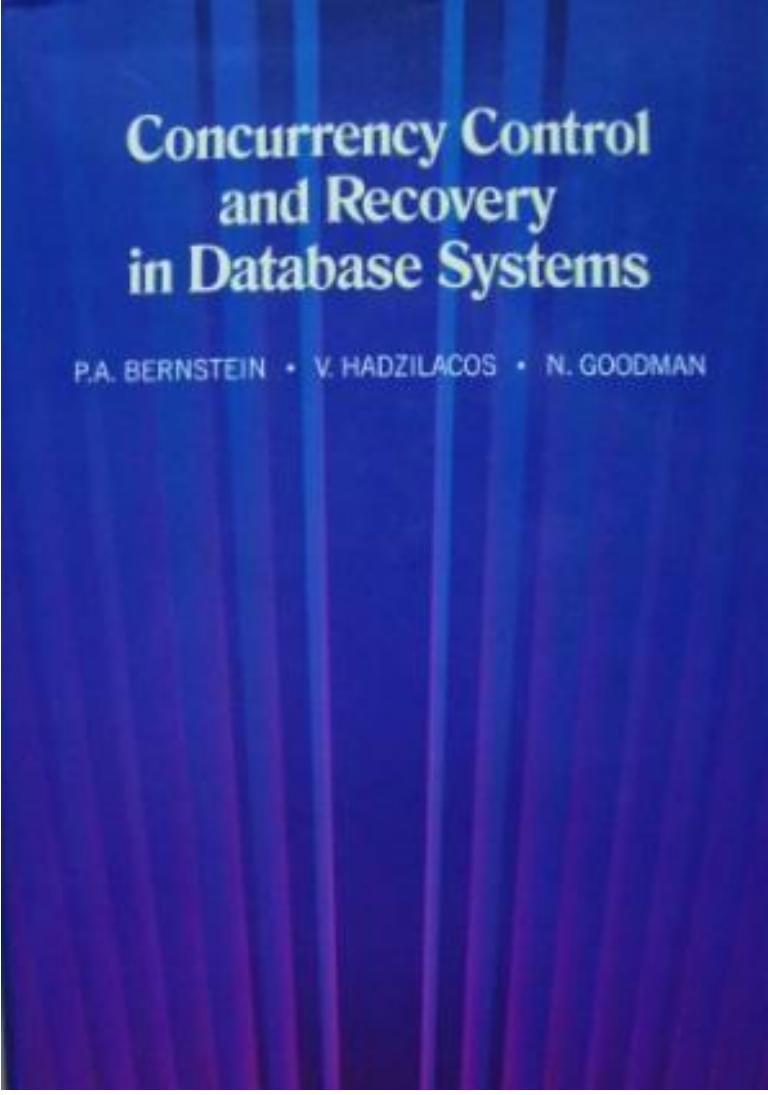
TaiShan 2280  
均衡型服务器TaiShan 5280  
存储型服务器TaiShan X6000  
高密型服务器OceanStor  
V6/F V6 存储华为鲲鹏  
处理器智能SSD  
控制器芯片智能网卡  
芯片智能管理  
芯片



- 分布式数据库系统原理
- 厄兹叙 (M.Tamer Özsu) / Patrick Valduriez 著  
, 周立柱 / 范举 译
- 2014
- 清华大学出版社



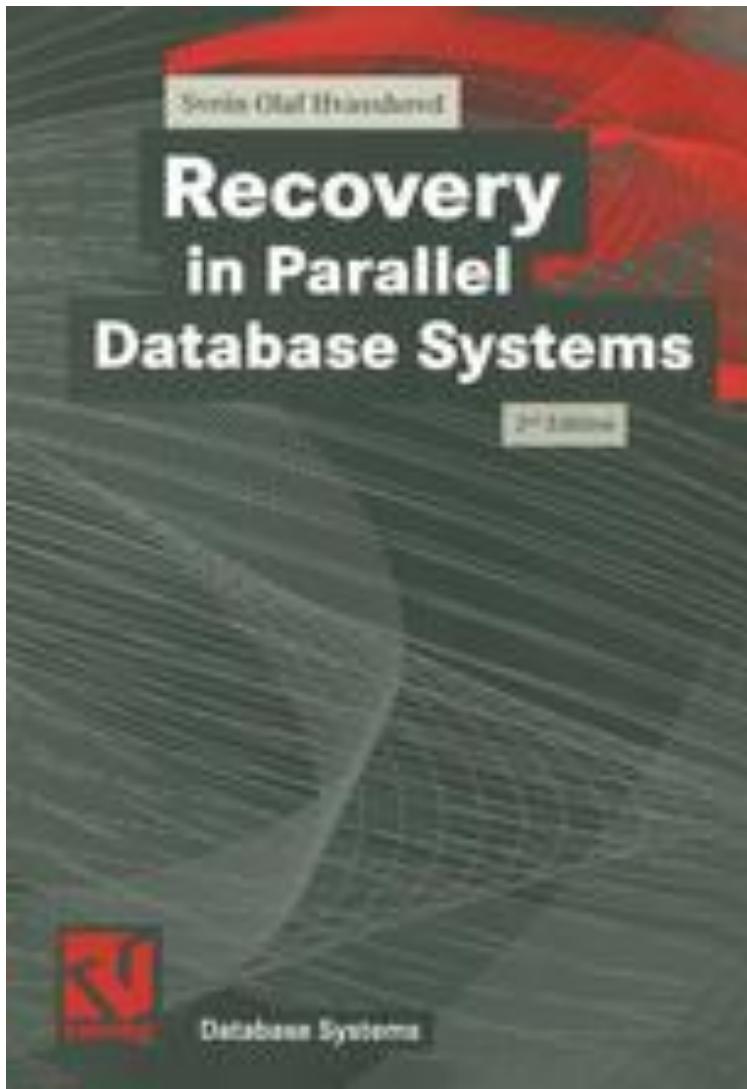
- Distributed Computing: Principles, Algorithms, and Systems
- Ajay D. Kshemkalyani, Mukesh Singhal
- 2008
- 中文版 2012



# Concurrency Control and Recovery in Database Systems

P.A. BERNSTEIN • V. HADZILACOS • N. GOODMAN

- Concurrency Control and Recovery in Database Systems
- Philip A. Bernstein, Vassos Hadzilacos, Nathan Goodman
- 1987



- Recovery in Parallel Database Systems
- Svein-Olaf Hvasshovd (auth.)
- 1999